

Least Squares

Geometric Algorithms

Lecture 23

CAS CS 132

Recap Problem

Project \vec{y} onto $\text{span}\{\vec{u}\}$

$\vec{y} = \alpha \vec{u}$

$\alpha = \frac{\vec{u}^T \vec{y}}{\vec{u}^T \vec{u}}$

$$\mathbf{u} = \begin{bmatrix} 1 \\ 3 \\ -2 \\ -1 \end{bmatrix} \quad \mathbf{v} = \begin{bmatrix} 0 \\ 1 \\ -1 \\ 0 \end{bmatrix}$$

Find the orthogonal projection of $\underline{\mathbf{u}}$ onto the span of \mathbf{v}

$$\alpha = \frac{\vec{u}^T \vec{v}}{\vec{v}^T \vec{v}} = \frac{3+2}{1+1} = 5$$

$$\hat{\mathbf{u}} = 5\vec{v} = \begin{bmatrix} 0 \\ 5 \\ -5 \\ 0 \end{bmatrix}$$

Answer

$$\hat{\mathbf{u}} = \begin{bmatrix} 0 \\ 5\cancel{12} \\ -5\cancel{12} \\ 0 \end{bmatrix}$$

Objectives

1. Introduce the least squares problem as a method of *approximating* solutions to matrix equations
2. Learn how to solve the least squares problems
3. Connect least squares solutions to projections

Keywords

general least squares problem

sum of squares error (ℓ_2 -error)

least squares solutions

orthogonal projections

normal equations

Orthogonal Matrices

Orthonormal Matrices

Definition. A matrix is **orthonormal** if its columns form an orthonormal set

The notes call a square orthonormal matrix an **orthogonal** matrix

Orthonormal Matrices

Definition. A matrix is **orthonormal** if its columns form an orthonormal set

The notes call a square orthonormal matrix an **orthogonal** matrix

This is incredibly confusing, but we'll try to be consistent and clear

Inverses of Orthogonal Matrices

Theorem. If an $n \times n$ matrix U is orthogonal (square orthonormal) then it is invertible and

$$U^{-1} = U^T$$

Orthonormal Matrices and Inner Products

Theorem. For a $m \times n$ orthonormal matrix U , and any vectors x and y in R^n

$$\langle Ux, Uy \rangle = \langle x, y \rangle$$

*Orthonormal matrices preserve inner products
and thus lengths & angles*

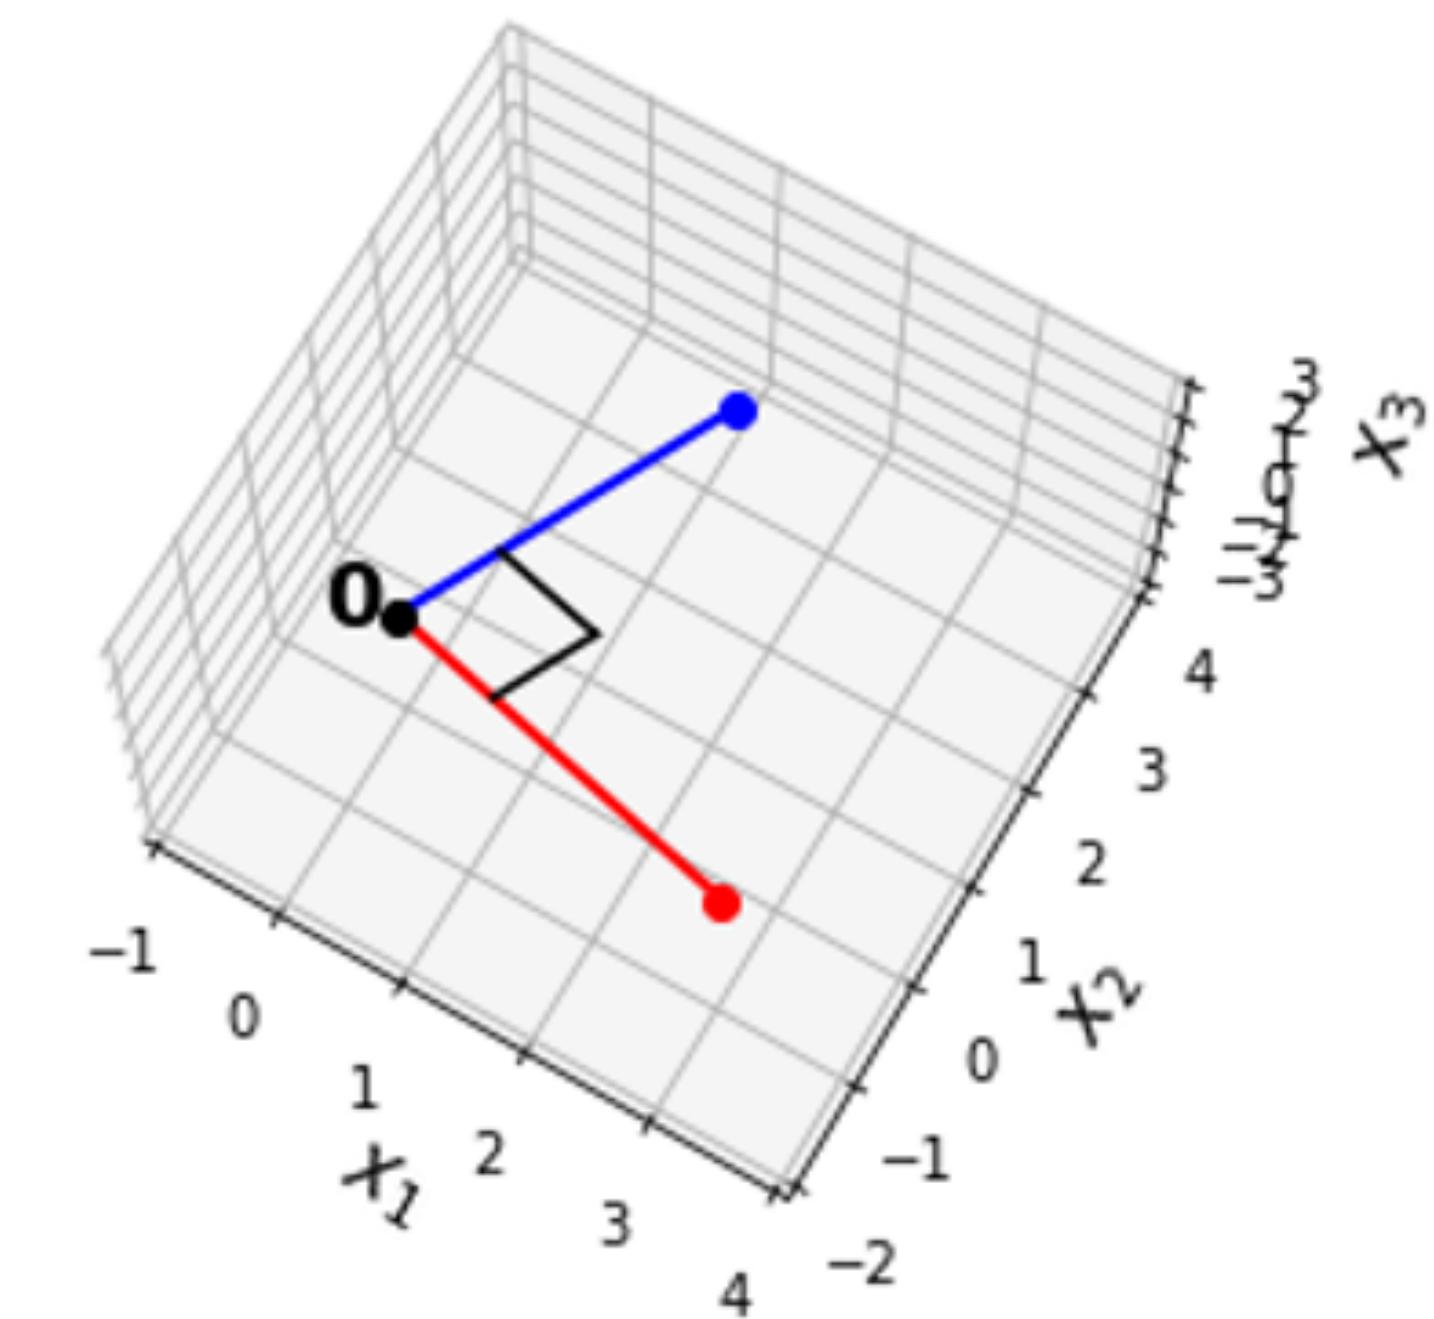
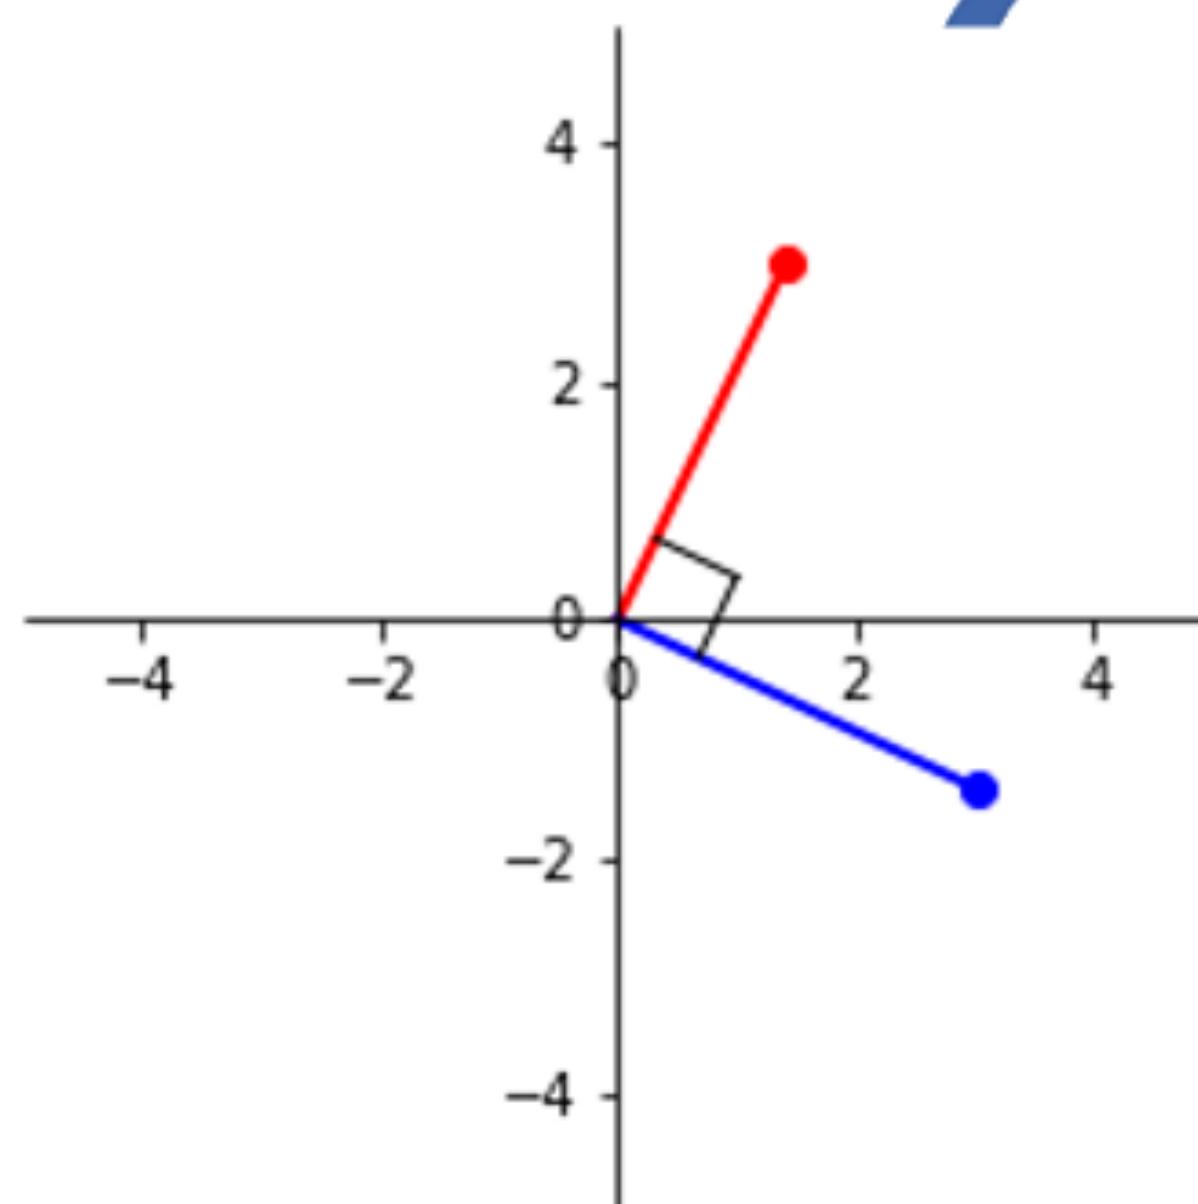
Verify:

Length, Angle, Orthogonality Preservation

Since lengths and angles are defined in terms of inner products, they are also preserved by orthonormal matrices:

The Picture

Orthonormal U



Example

$$U = \begin{bmatrix} 1/\sqrt{2} & 2/3 \\ 1/\sqrt{2} & -2/3 \\ 0 & 1/3 \end{bmatrix} \quad x = \begin{bmatrix} \sqrt{2} \\ 3 \end{bmatrix}$$

moving on. . .

Motivation

The story of an enterprising CS132 student

The story of an enterprising CS132 student

Problem. Solve the equation $Ax = b$

The story of an enterprising CS132 student

Problem. Solve the equation $Ax = b$

Answer. Use `np.linalg.solve(A, b)`

The story of an enterprising CS132 student

Problem. Solve the equation $Ax = b$

Answer. Use `np.linalg.solve(A, b)`

```
>>> A = np.array([
...     [1., 0, 5],
...     [1, -1, 4],
...     [0, 2, 2]])
>>> b = np.array([-1, 2, 3])
>>> np.linalg.solve(A, b)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/opt/homebrew/lib/python3.11/site-packages/numpy/linalg/linalg.py", line 409, in solve
    r = gufunc(a, b, signature=signature, extobj=extobj)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/opt/homebrew/lib/python3.11/site-packages/numpy/linalg/linalg.py", line 112, in _raise_linalgerror_singular
    raise LinAlgError("Singular matrix")
numpy.linalg.LinAlgError: Singular matrix
```

The story of an enterprising CS132 student

Problem. Solve the equation $Ax = b$

Answer. Use `np.linalg.solve(A, b)`

```
>>> A = np.array([
...     [1., 0, 5],
...     [1, -1, 4],
...     [0, 2, 2]])
>>> b = np.array([-1, 2, 3])
>>> np.linalg.solve(A, b)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "/opt/homebrew/lib/python3.11/site-packages/numpy/linalg/linalg.py", line 409, in solve
    r = gufunc(a, b, signature=signature, extobj=extobj)
                                ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/opt/homebrew/lib/python3.11/site-packages/numpy/linalg/linalg.py", line 112, in _raise_linalgerror_singular
    raise LinAlgError("Singular matrix")
numpy.linalg.LinAlgError: Singular matrix
```

This doesn't always work

Reads the docs...

numpy.linalg.solve

`linalg.solve(a, b)`

[\[source\]](#)

Solve a linear matrix equation, or system of linear scalar equations.

Computes the “exact” solution, x , of the well-determined, i.e., full rank, linear matrix equation $ax = b$.

Parameters: `a` : $(..., M, M)$ *array_like*

Coefficient matrix.

`b` : $\{(..., M,), (... , M, K)\}$, *array_like*

Ordinate or “dependent variable” values.

Returns: `x` : $\{(..., M,), (... , M, K)\}$ *ndarray*

Solution to the system $a x = b$. Returned shape is identical to b .

Raises: `LinAlgError`

If a is singular or not square.

 See also

[scipy.linalg.solve](#)

Similar function in SciPy

Reads the docs...

numpy.linalg.solve

`linalg.solve(a, b)`

[\[source\]](#)

Solve a linear matrix equation, or system of linear scalar equations.

Computes the “exact” solution, x , of the well-determined, i.e., full rank, linear matrix equation $ax = b$.

Parameters: `a` : $(..., M, M)$ *array_like*

Coefficient matrix.

`b` : $\{(..., M,), (... , M, K)\}$, *array_like*

Ordinate or “dependent variable” values.

Returns: `x` : $\{(..., M,), (... , M, K)\}$ *ndarray*

Solution to the system $a x = b$. Returned shape is identical to b .

Raises: `LinAlgError`

If a is singular or not square.

$$\det(a) = 0$$

 See also

[scipy.linalg.solve](#)

Similar function in SciPy

`b` : {..., M, J , ..., M, K }, `array_like`

Ordinate or “dependent variable” values.

Reads the docs...

Returns: `x` : {..., M, J , ..., M, K } `ndarray`

Solution to the system $a x = b$. Returned shape is identical to b .

Raises: `LinAlgError`

If a is singular or not square.

i See also

[scipy.linalg.solve](#)

Similar function in SciPy.

Notes

! New in version 1.8.0.

Broadcasting rules apply, see the [numpy.linalg](#) documentation for details.

The solutions are computed using LAPACK routine `_gesv`.

a must be square and of full-rank, i.e., all rows (or, equivalently, columns) must be linearly independent; if either is not true, use [lstsq](#) for the least-squares best “solution” of the system/equation.

`b` : {..., M, J , ..., M, K }, `array_like`

Ordinate or “dependent variable” values.

Reads the docs...

Returns: `x` : {..., M, J , ..., M, K } `ndarray`

Solution to the system $a x = b$. Returned shape is identical to b .

Raises: `LinAlgError`

If a is singular or not square.

i See also

[scipy.linalg.solve](#)

Similar function in SciPy.

Notes

! New in version 1.8.0.

Broadcasting rules apply, see the [numpy.linalg](#) documentation for details.

The solutions are computed using LAPACK routine `_gesv`.

a must be square and of full-rank, i.e., all rows (or, equivalently, columns) must be linearly independent; if either is not true, use [`lstsq`](#) for the least-squares best “solution” of the system/equation.

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b) ←
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

$$\begin{matrix} 0.1, & -0.0000..1, & 2 \\ & \uparrow \\ [0.1, 0, 2]^T \end{matrix}$$

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

uh... probably numerical errors...

Answer: $x = \begin{bmatrix} -1/9 \\ 7/9 \\ 2/9 \end{bmatrix}$

np.linalg.lstsq

```
>>> np.linalg.lstsq(A, b)
<stdin>:1: FutureWarning: `rcond` parameter will change to the default of machine precision times ``max(M, N)``
where M and N are the input matrix dimensions.
To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old,
explicitly pass `rcond=-1`.
(array([-0.11111111,  0.77777778,  0.22222222]), array([], dtype=float64), 2, array([6.84168488e+00,
2.27845297e+00, 6.13801942e-17]))
>>> x = np.array([-0.11111111,  0.77777778,  0.22222222])
>>> A @ x
array([ 9.9999990e-01, -9.9999994e-09,  2.0000000e+00])
>>>
```

uh...probably numerical errors...

Answer: $x = \begin{bmatrix} -1/9 \\ 7/9 \\ 2/9 \end{bmatrix}$

This is not correct

This System is Inconsistent

$$\left[\begin{array}{cccc} 1 & 0 & 5 & -1 \\ 1 & -1 & 4 & 2 \\ 0 & 2 & 2 & 3 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & 0 & 5 & -1 \\ 0 & -1 & -1 & 3 \\ 0 & 2 & 2 & 3 \end{array} \right] \sim \left[\begin{array}{cccc|c} 1 & 0 & 5 & -1 & \\ 0 & -1 & -1 & 3 & \\ 0 & 0 & 0 & 9 & \end{array} \right]$$

↗

The "correct" answer: There is no solution

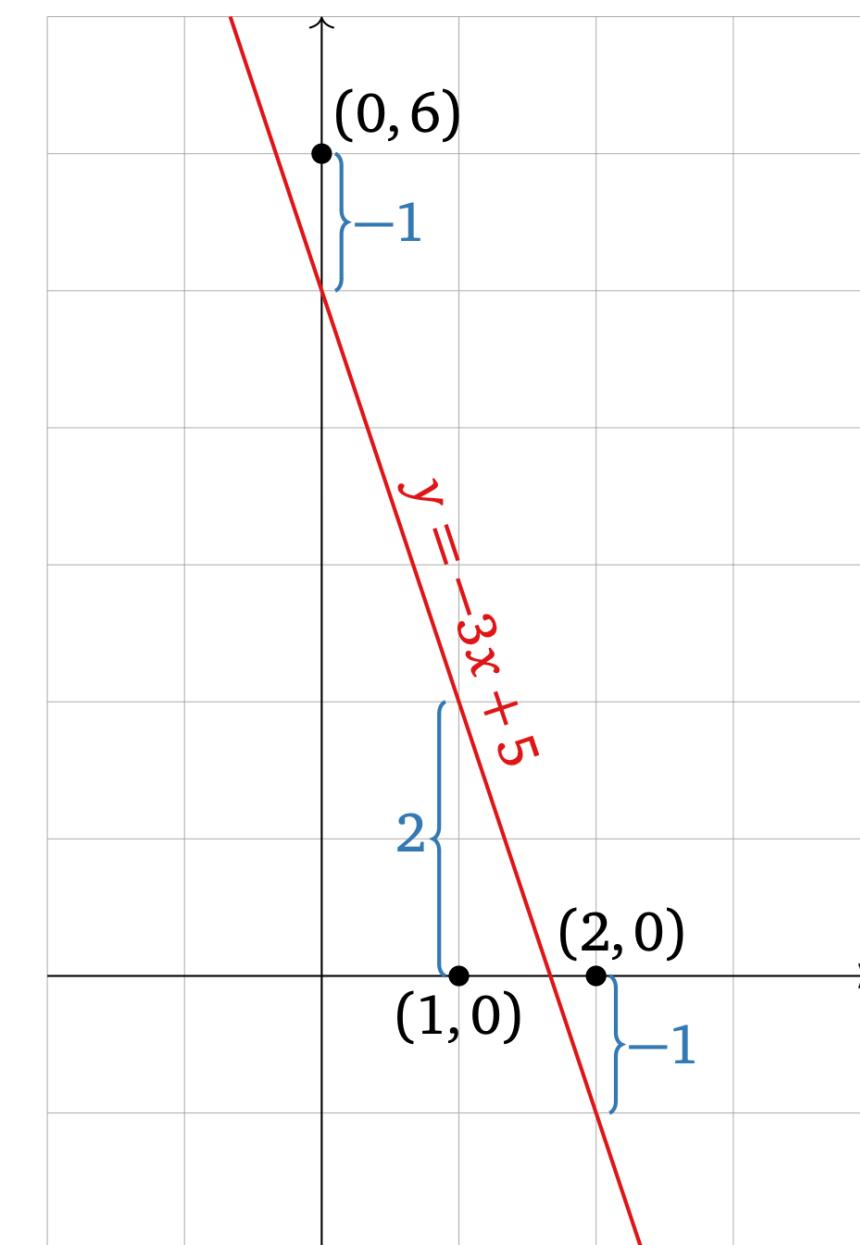
This System is Inconsistent

$$\left[\begin{array}{cccc} 1 & 0 & 5 & -1 \\ 1 & -1 & 4 & 2 \\ 0 & 2 & 2 & 3 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & 0 & 5 & -1 \\ 0 & -1 & -1 & 3 \\ 0 & 2 & 2 & 3 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & 0 & 5 & -1 \\ 0 & -1 & -1 & 3 \\ 0 & 0 & 0 & 9 \end{array} \right]$$

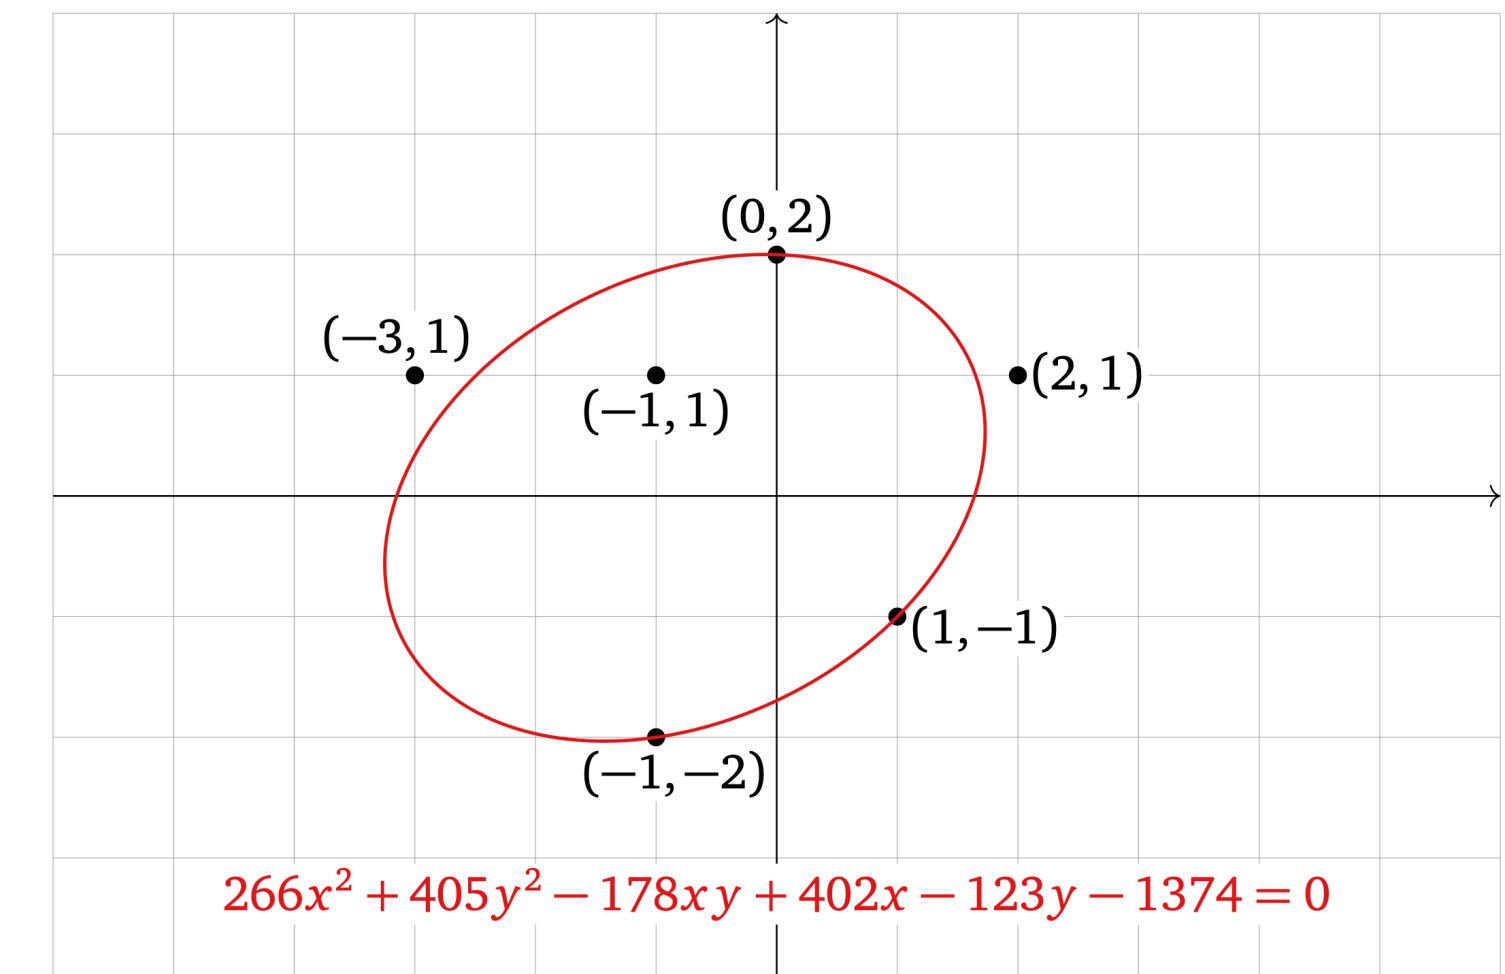
The "correct" answer: There is no solution

What's going on here?

Non-Linearity

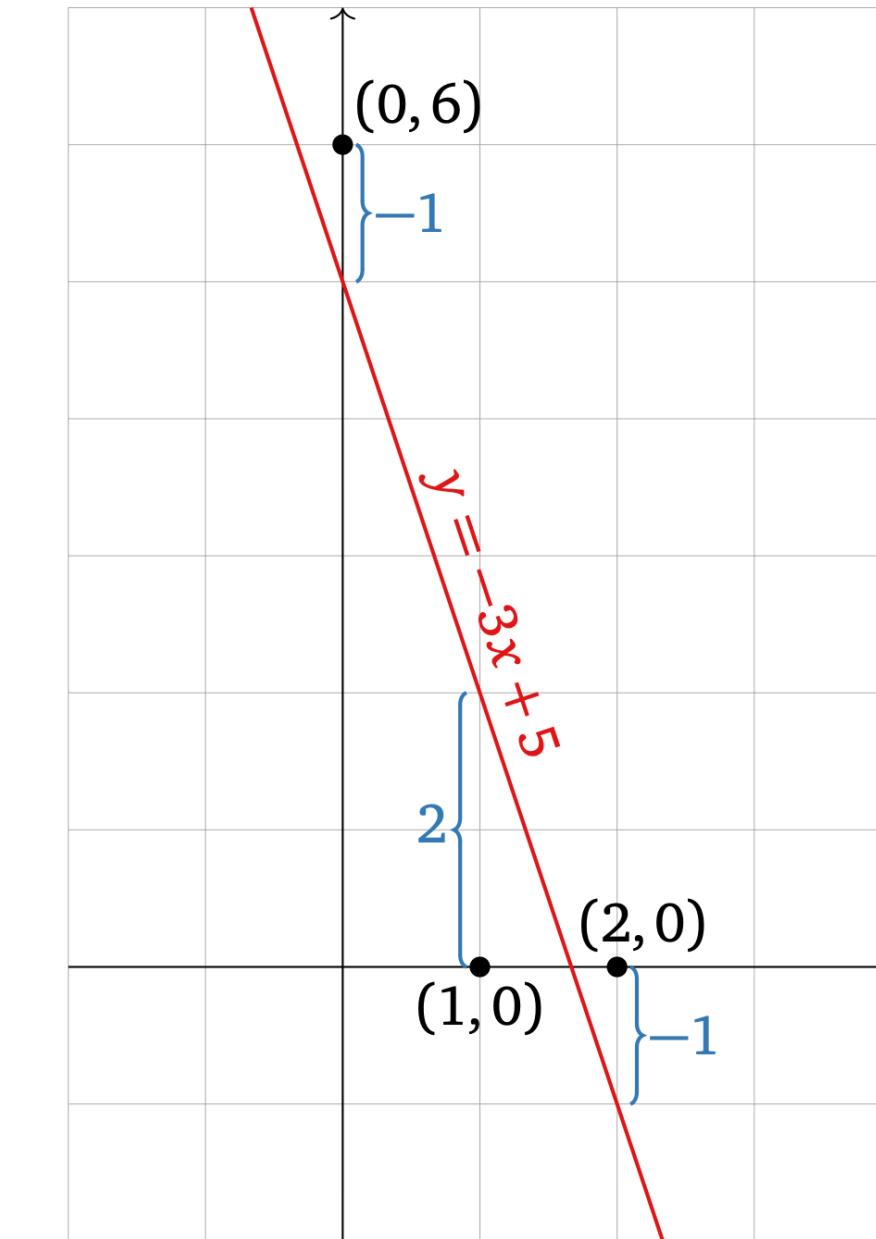


$$b - A\hat{x} = \begin{pmatrix} 6 \\ 0 \\ 0 \end{pmatrix} - A \begin{pmatrix} -3 \\ 5 \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ -1 \end{pmatrix}$$

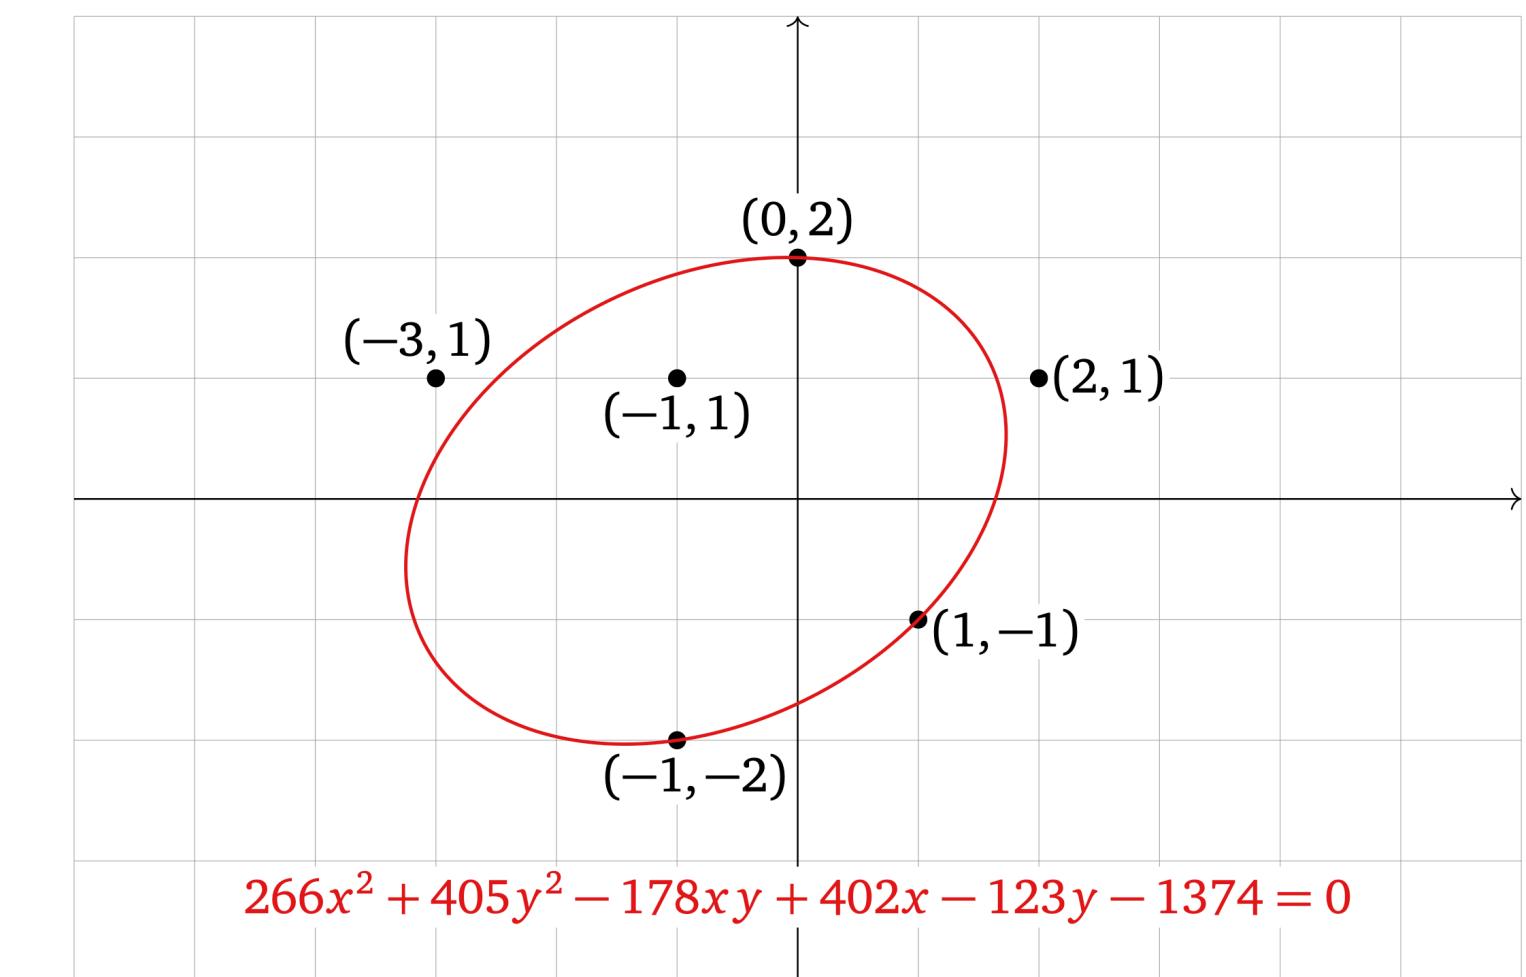


Non-Linearity

Linear algebra is very powerful and very clean, but **the world isn't linear**. There are non-linear relationships and sources of *noise*



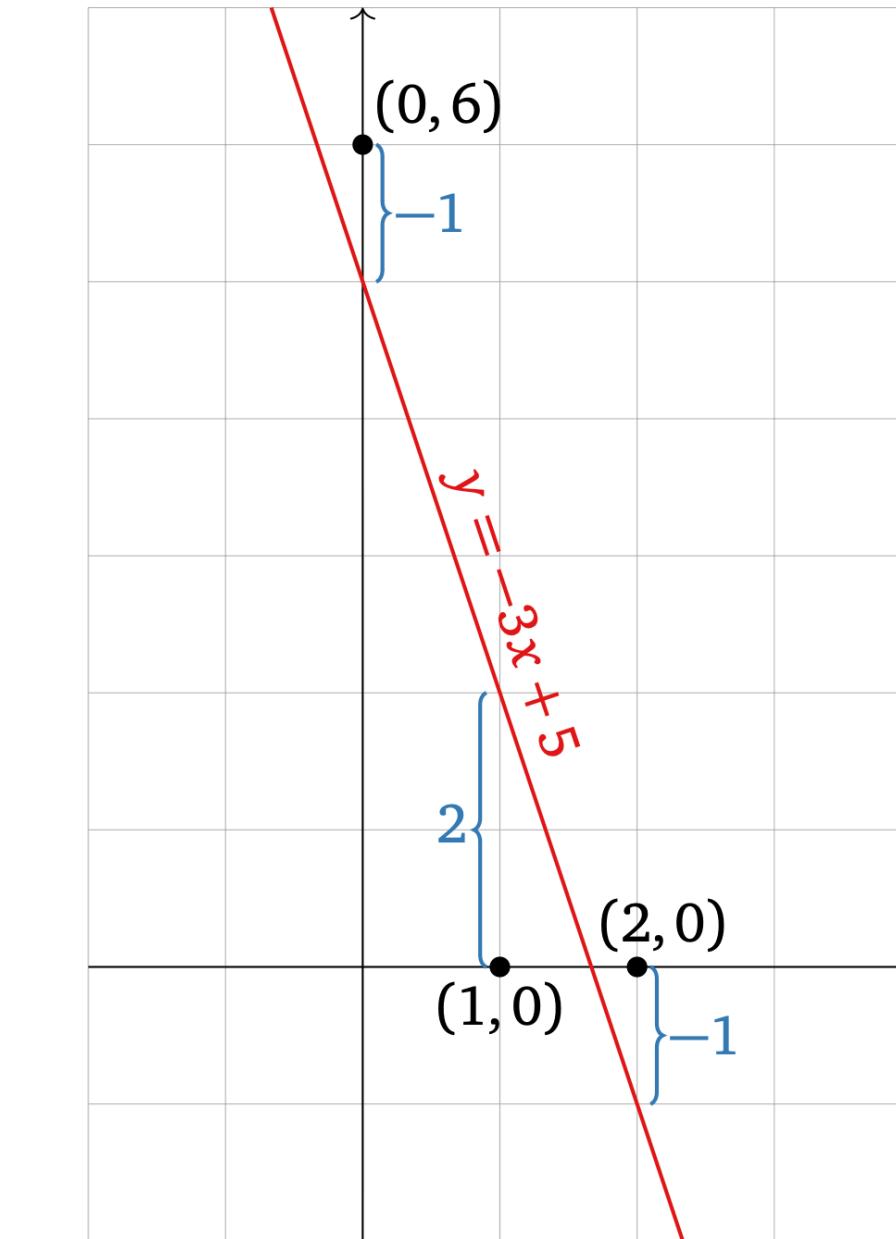
$$b - A\hat{x} = \begin{pmatrix} 6 \\ 0 \\ 0 \end{pmatrix} - A \begin{pmatrix} -3 \\ 5 \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ -1 \end{pmatrix}$$



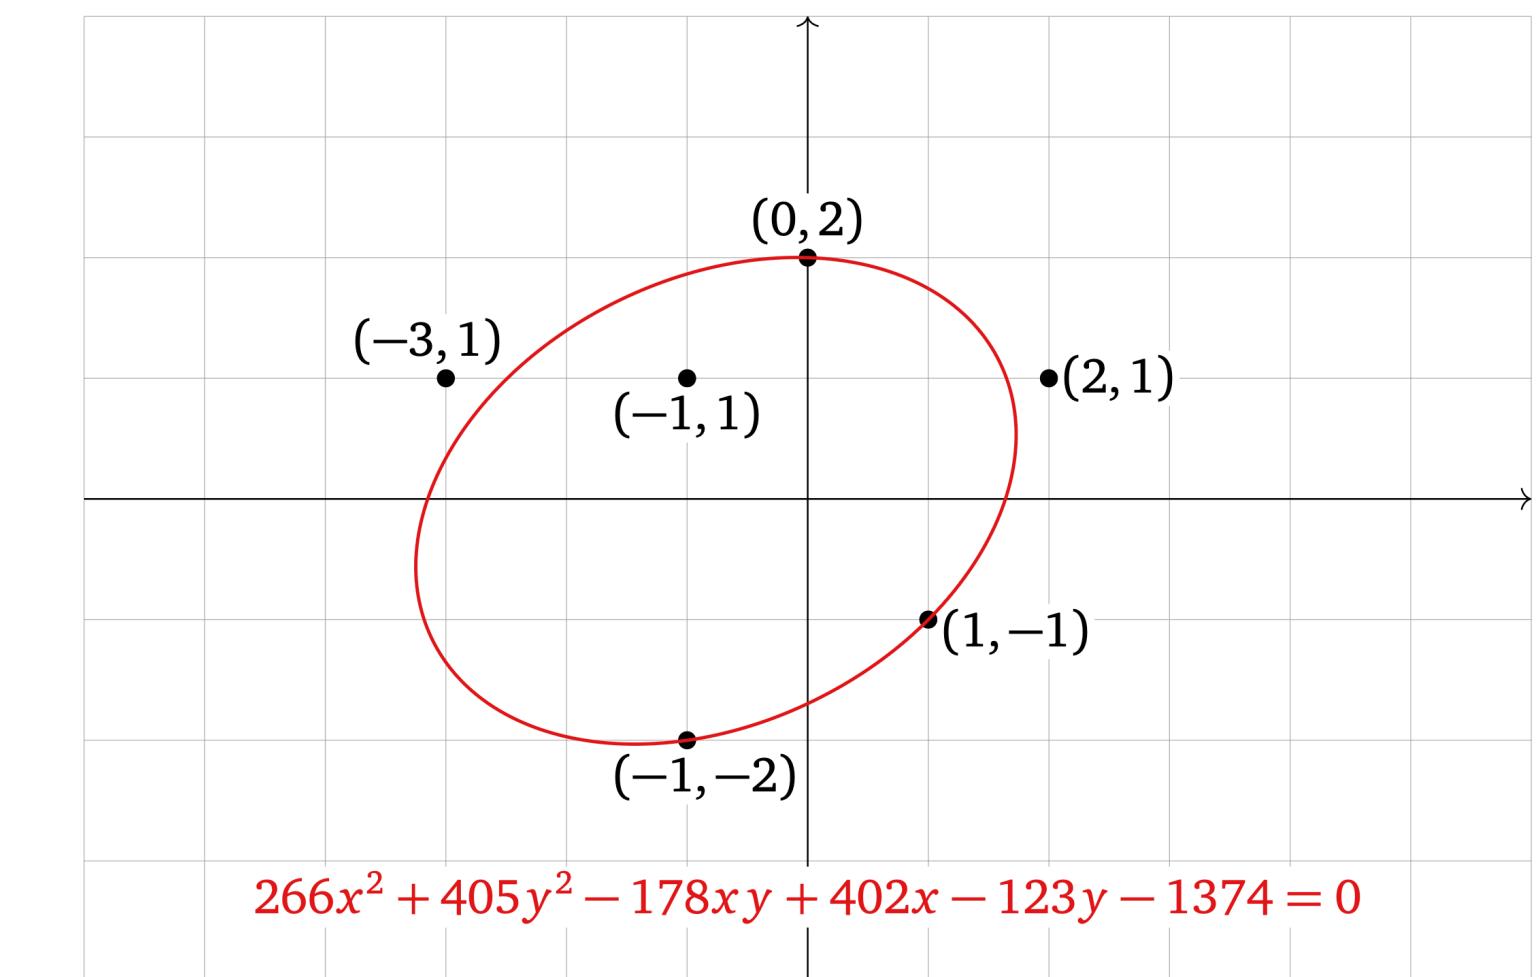
Non-Linearity

Linear algebra is very powerful and very clean, but **the world isn't linear**. There are non-linear relationships and sources of *noise*

We can't force the world to be linear



$$b - A\hat{x} = \begin{pmatrix} 6 \\ 0 \\ 0 \end{pmatrix} - A \begin{pmatrix} -3 \\ 5 \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ -1 \end{pmatrix}$$

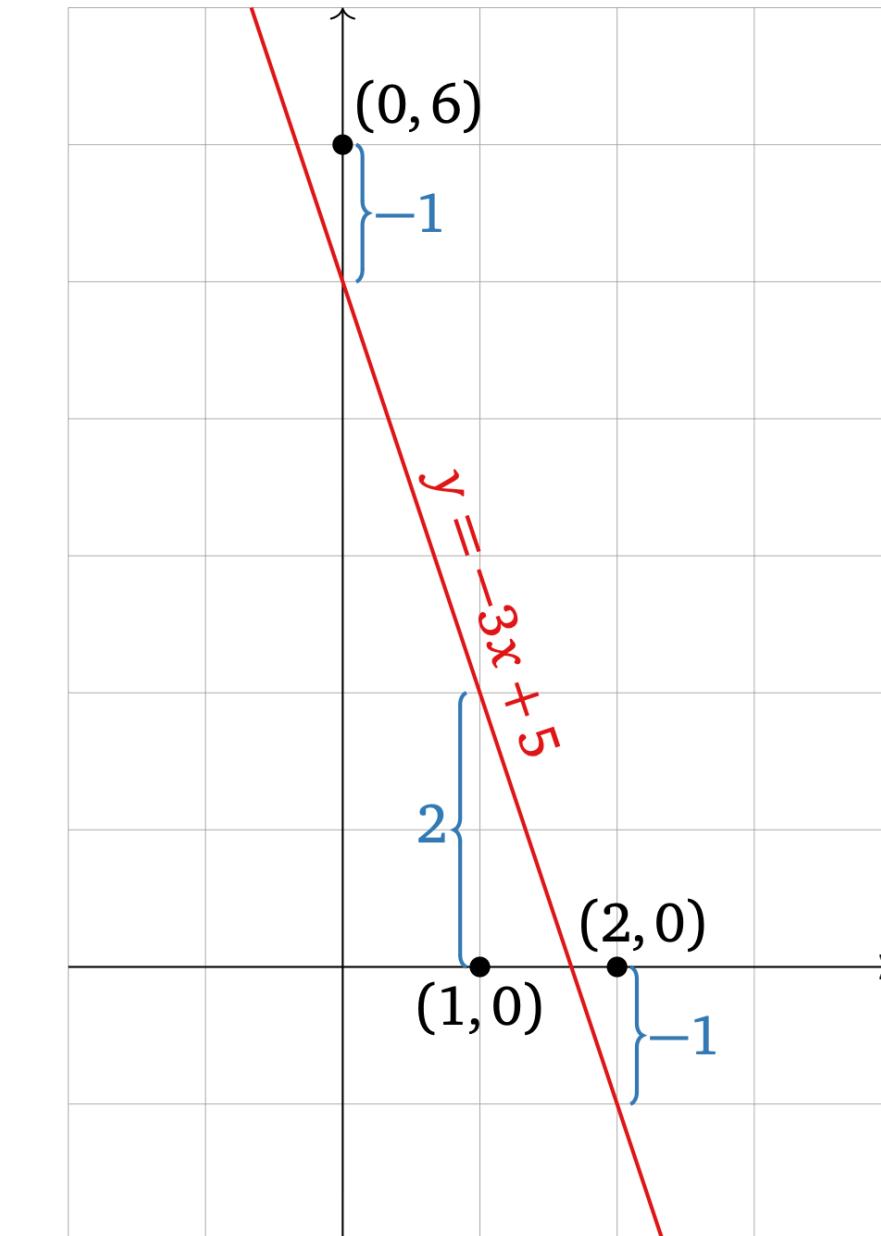


Non-Linearity

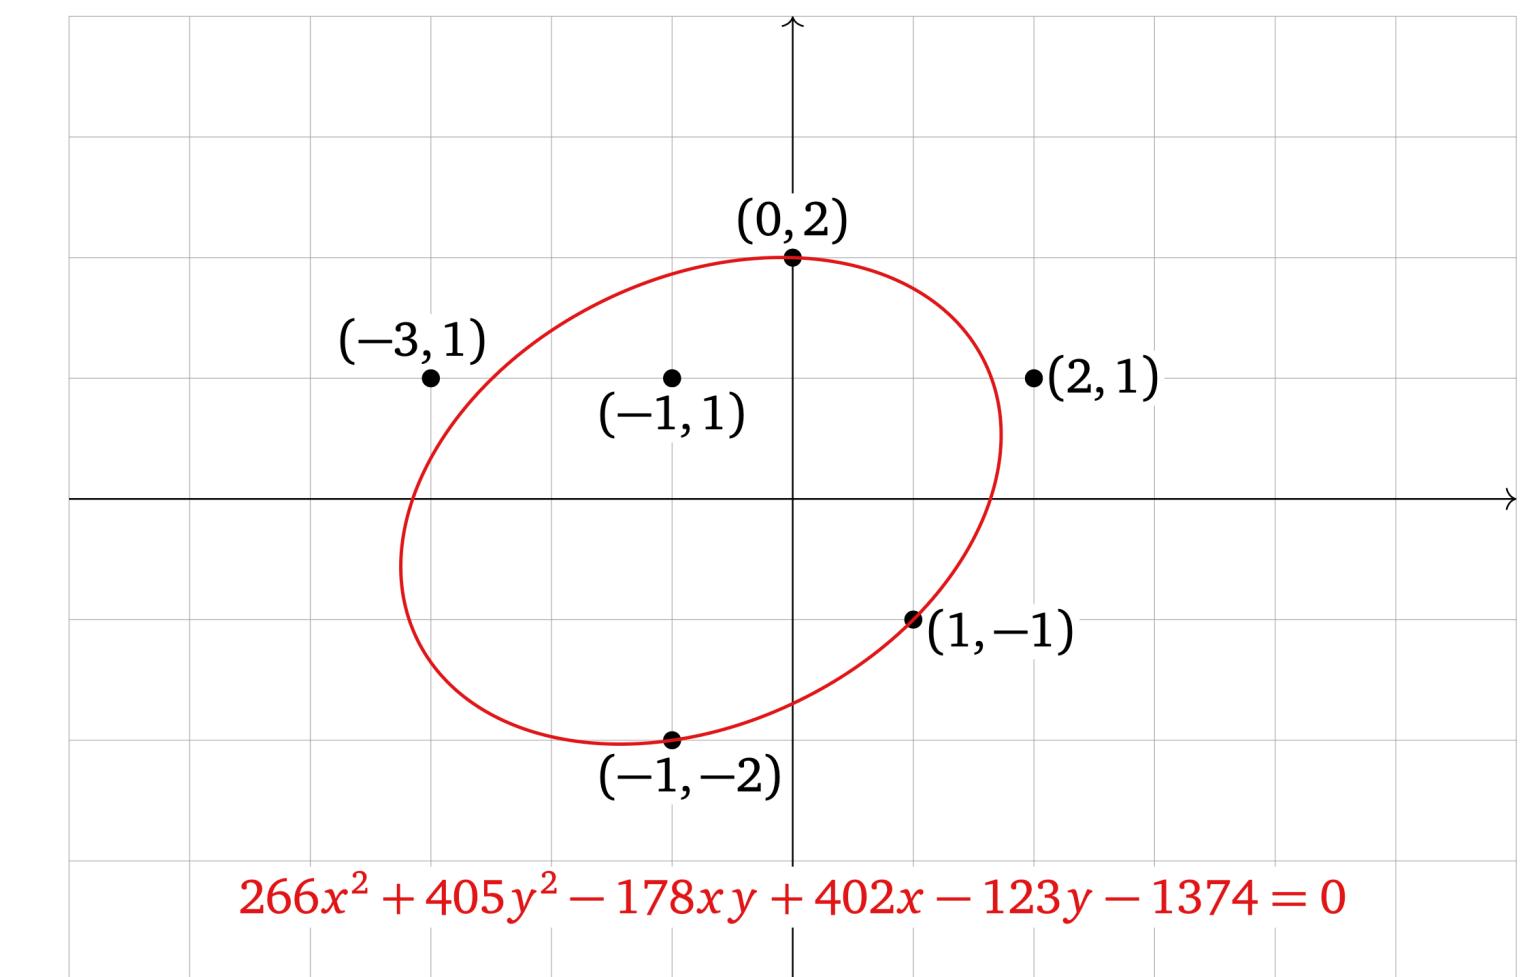
Linear algebra is very powerful and very clean, but **the world isn't linear**. There are non-linear relationships and sources of *noise*

We can't force the world to be linear

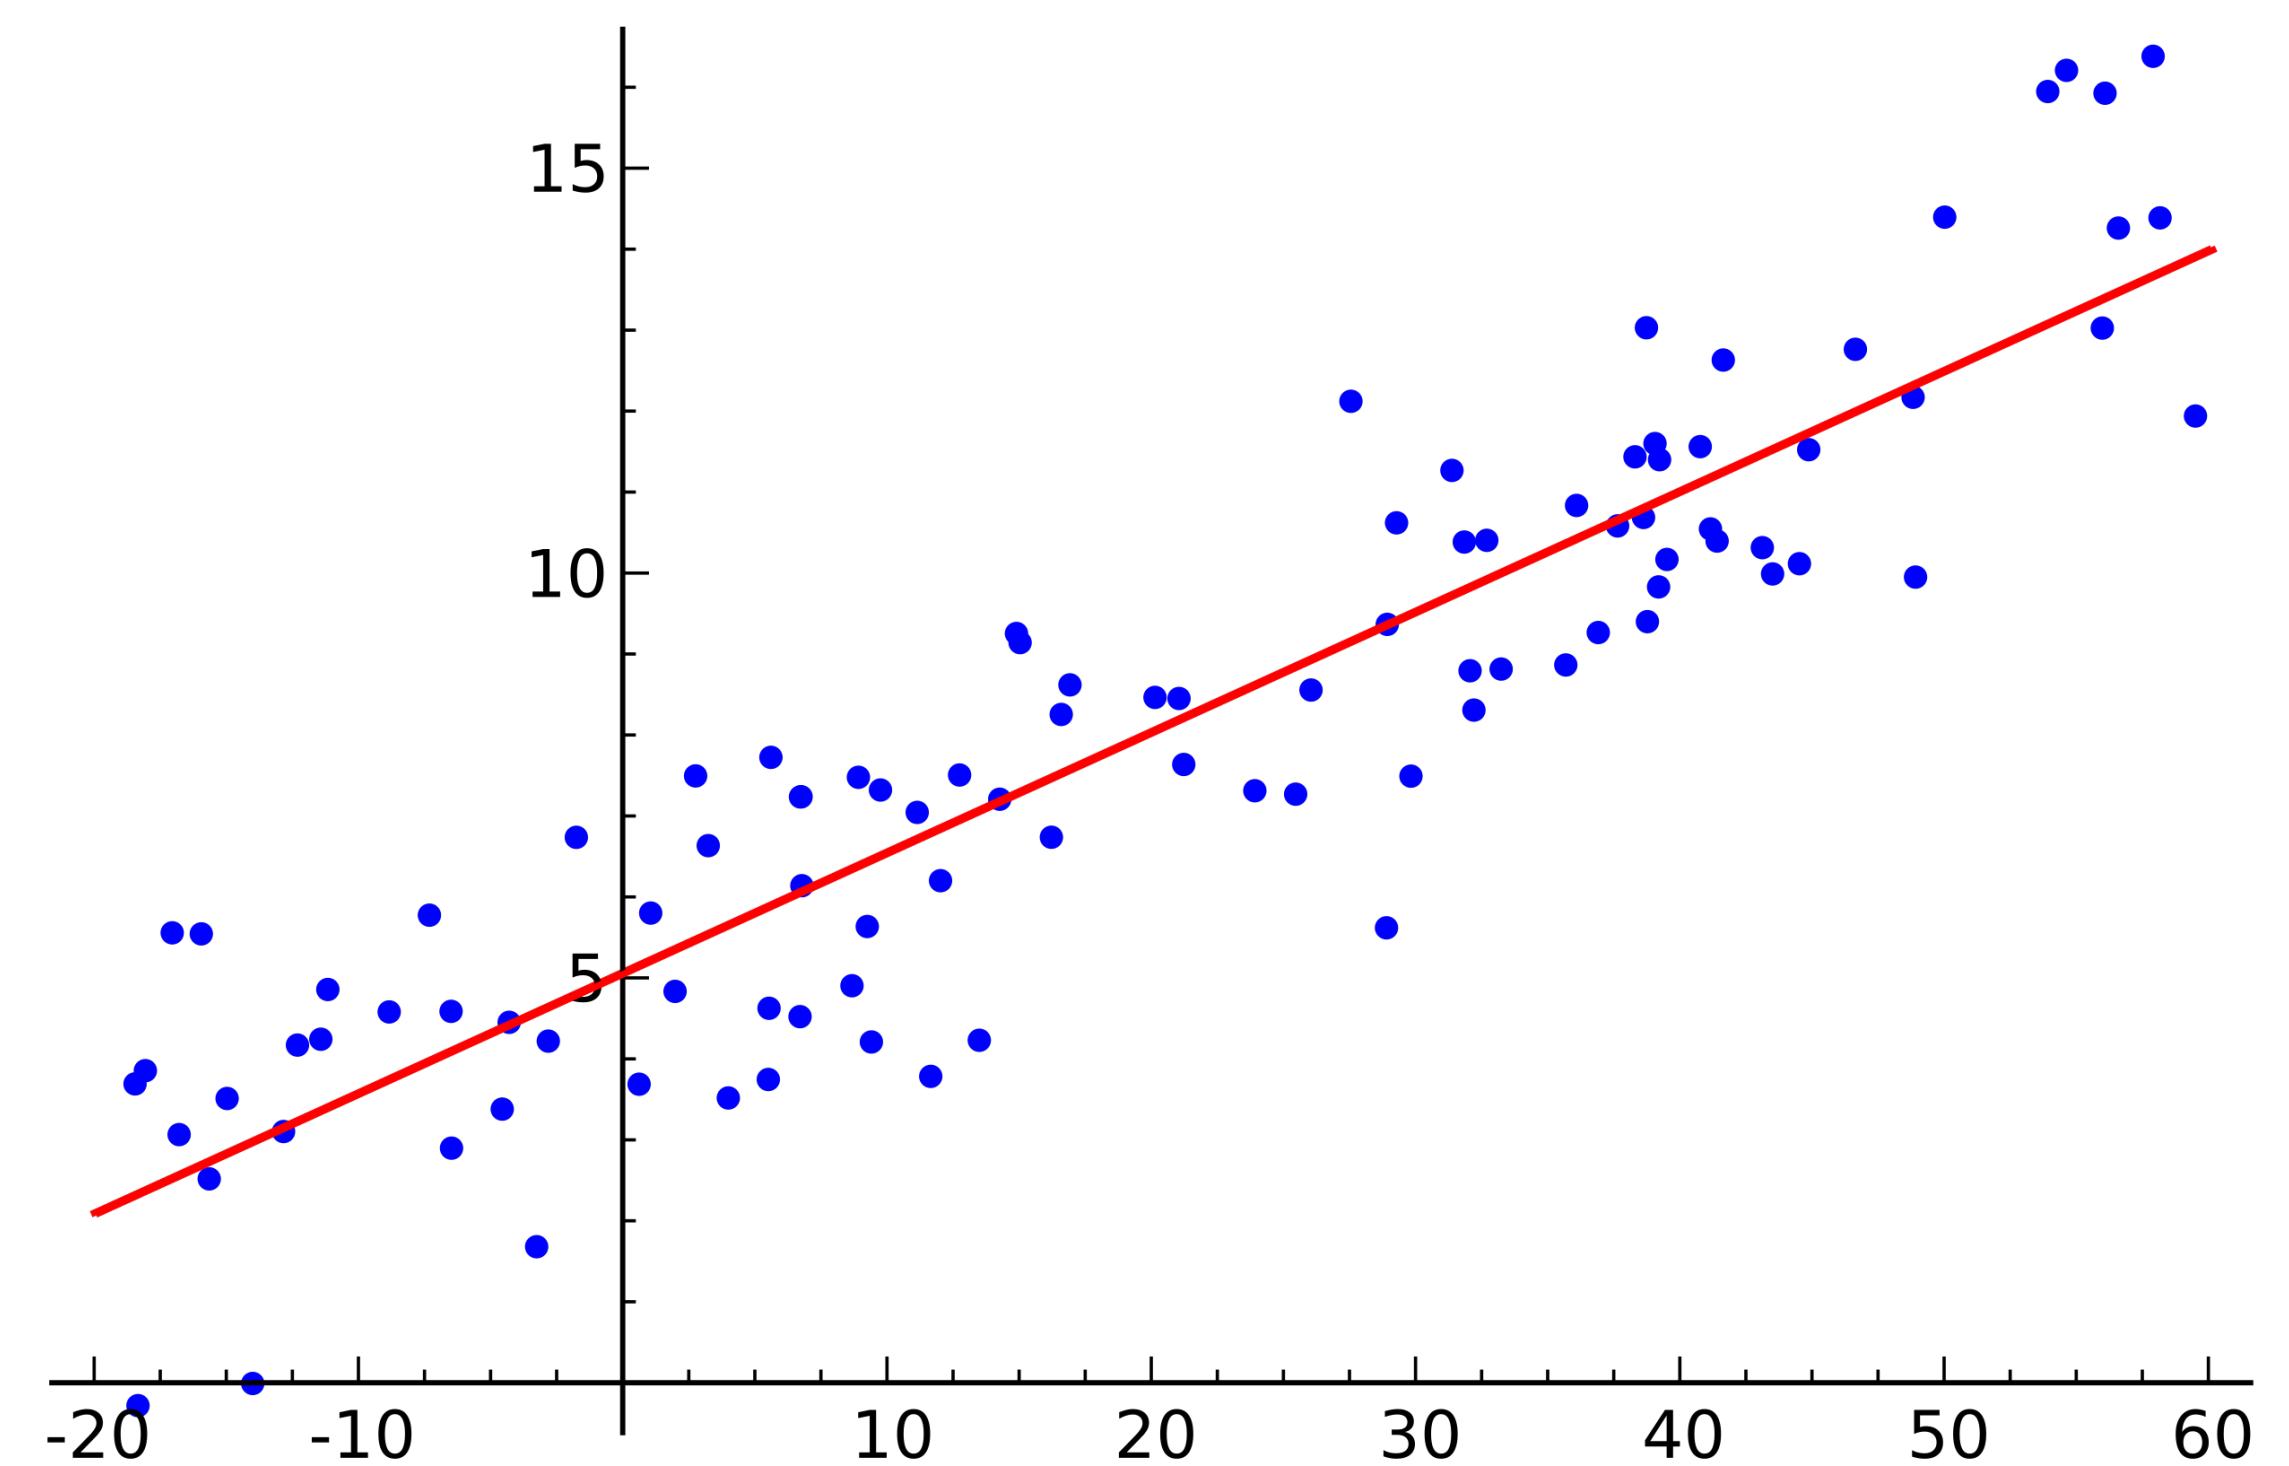
But we can try...



$$b - A\hat{x} = \begin{pmatrix} 6 \\ 0 \\ 0 \end{pmatrix} - A \begin{pmatrix} -3 \\ 5 \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ -1 \end{pmatrix}$$

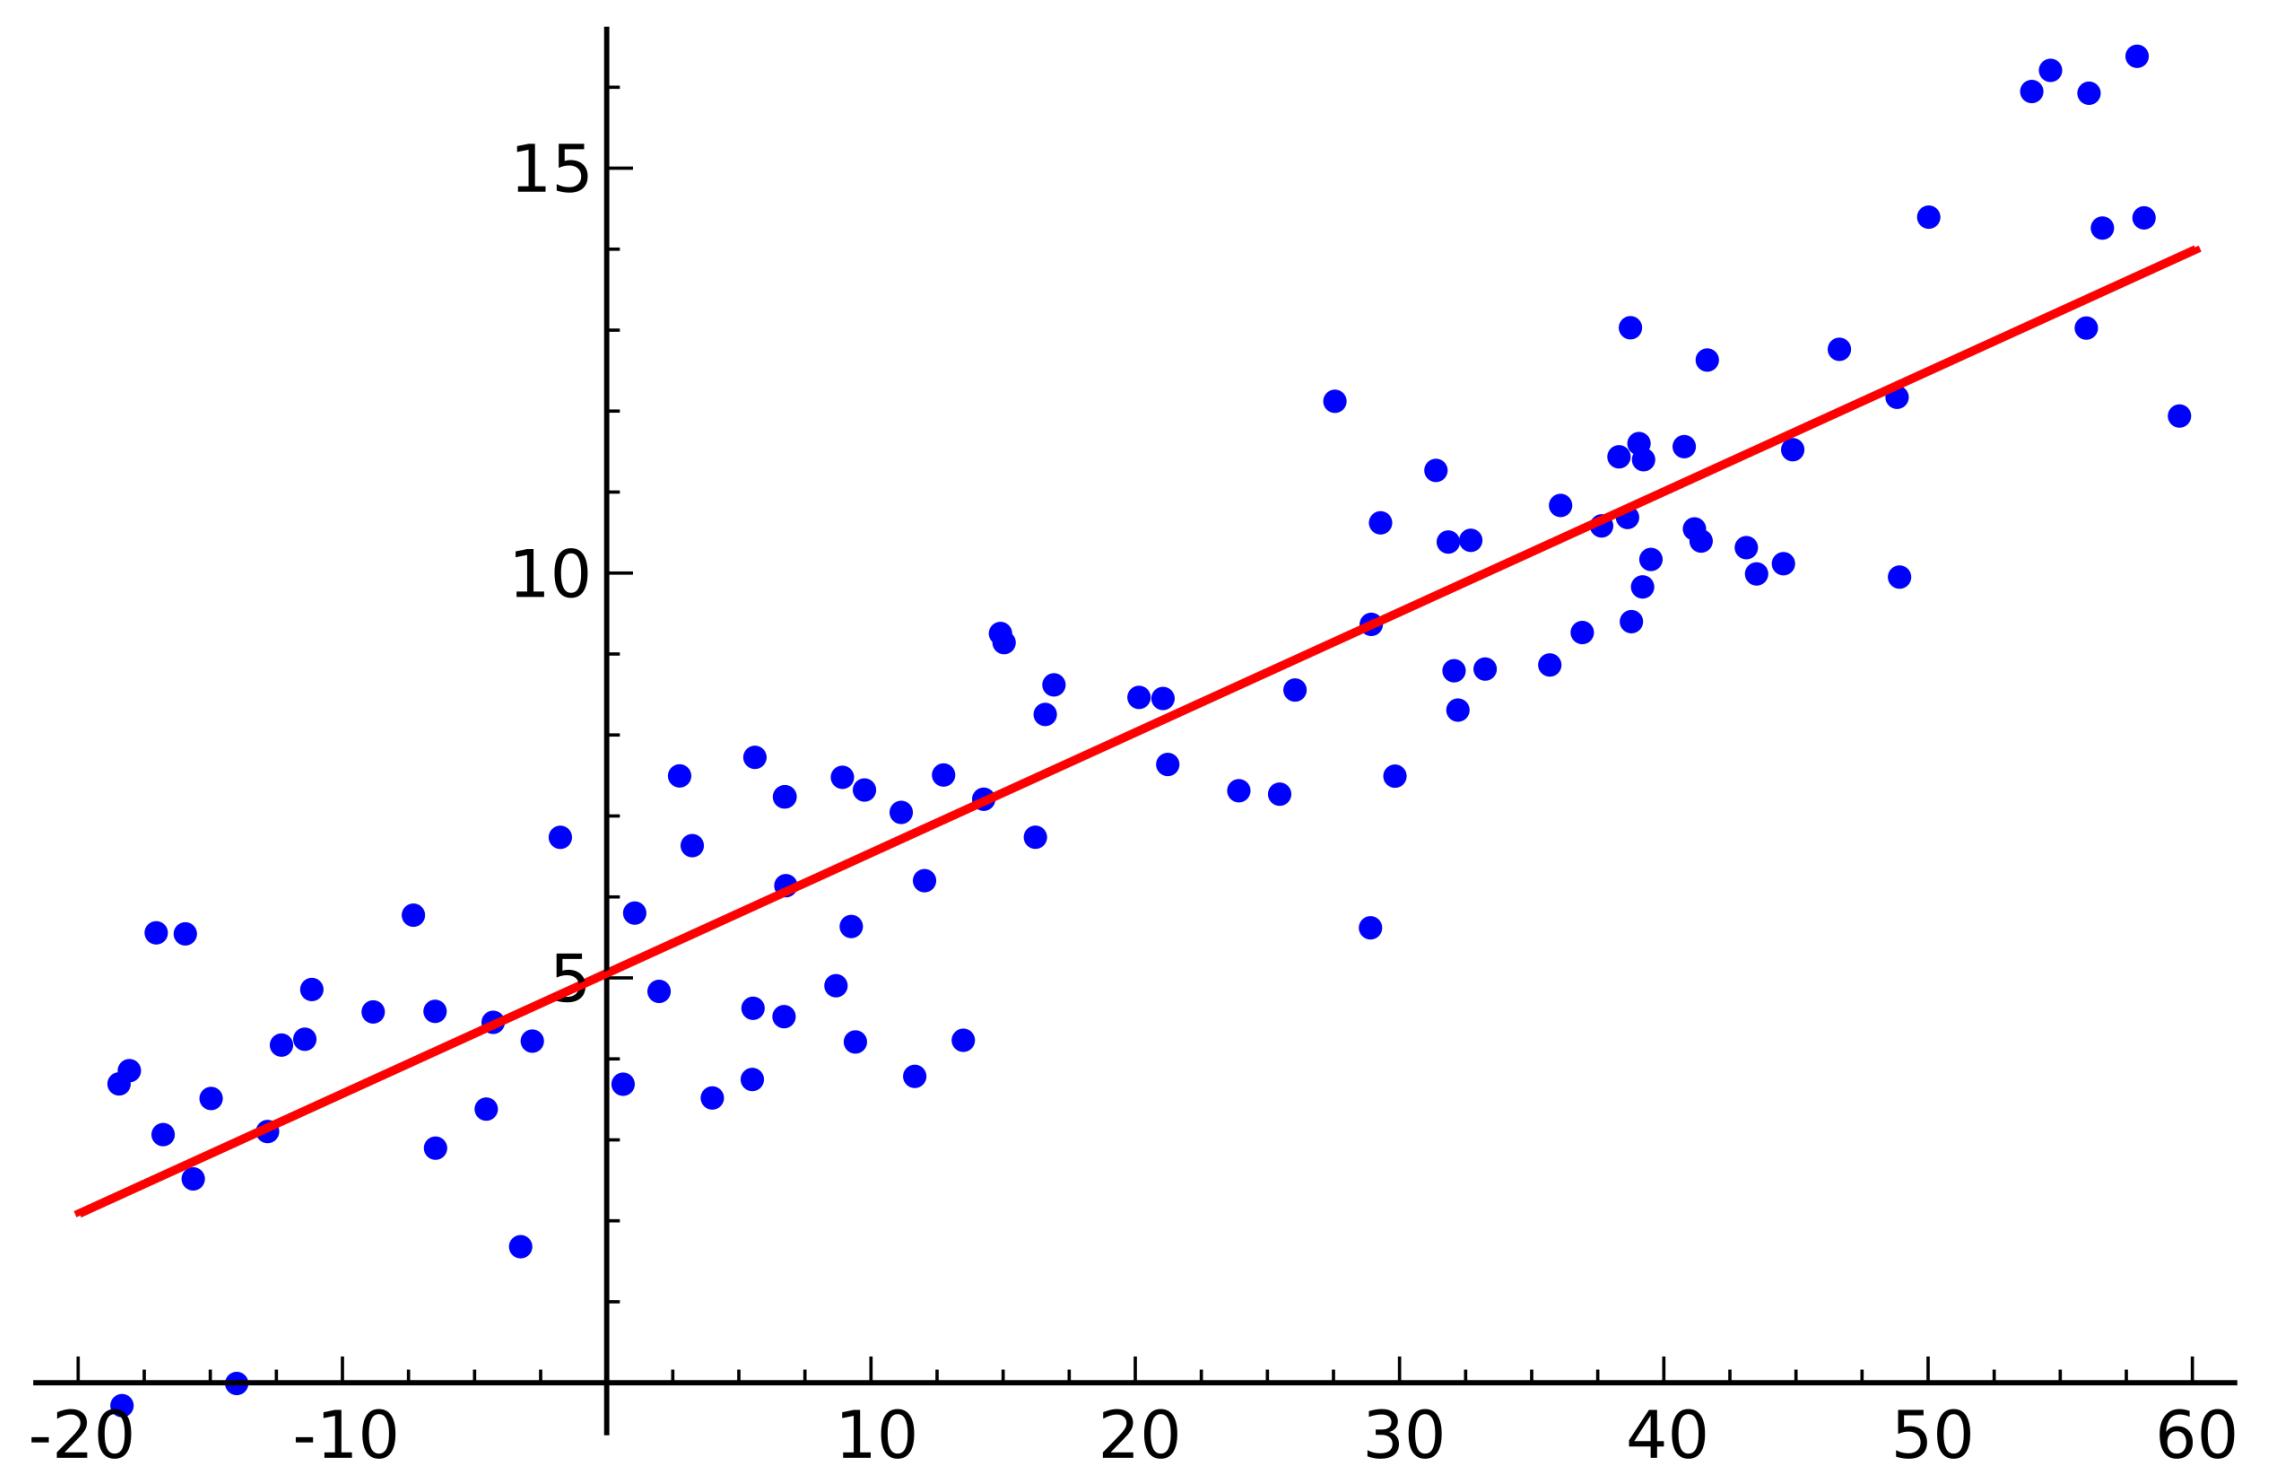


The Idea



The Idea

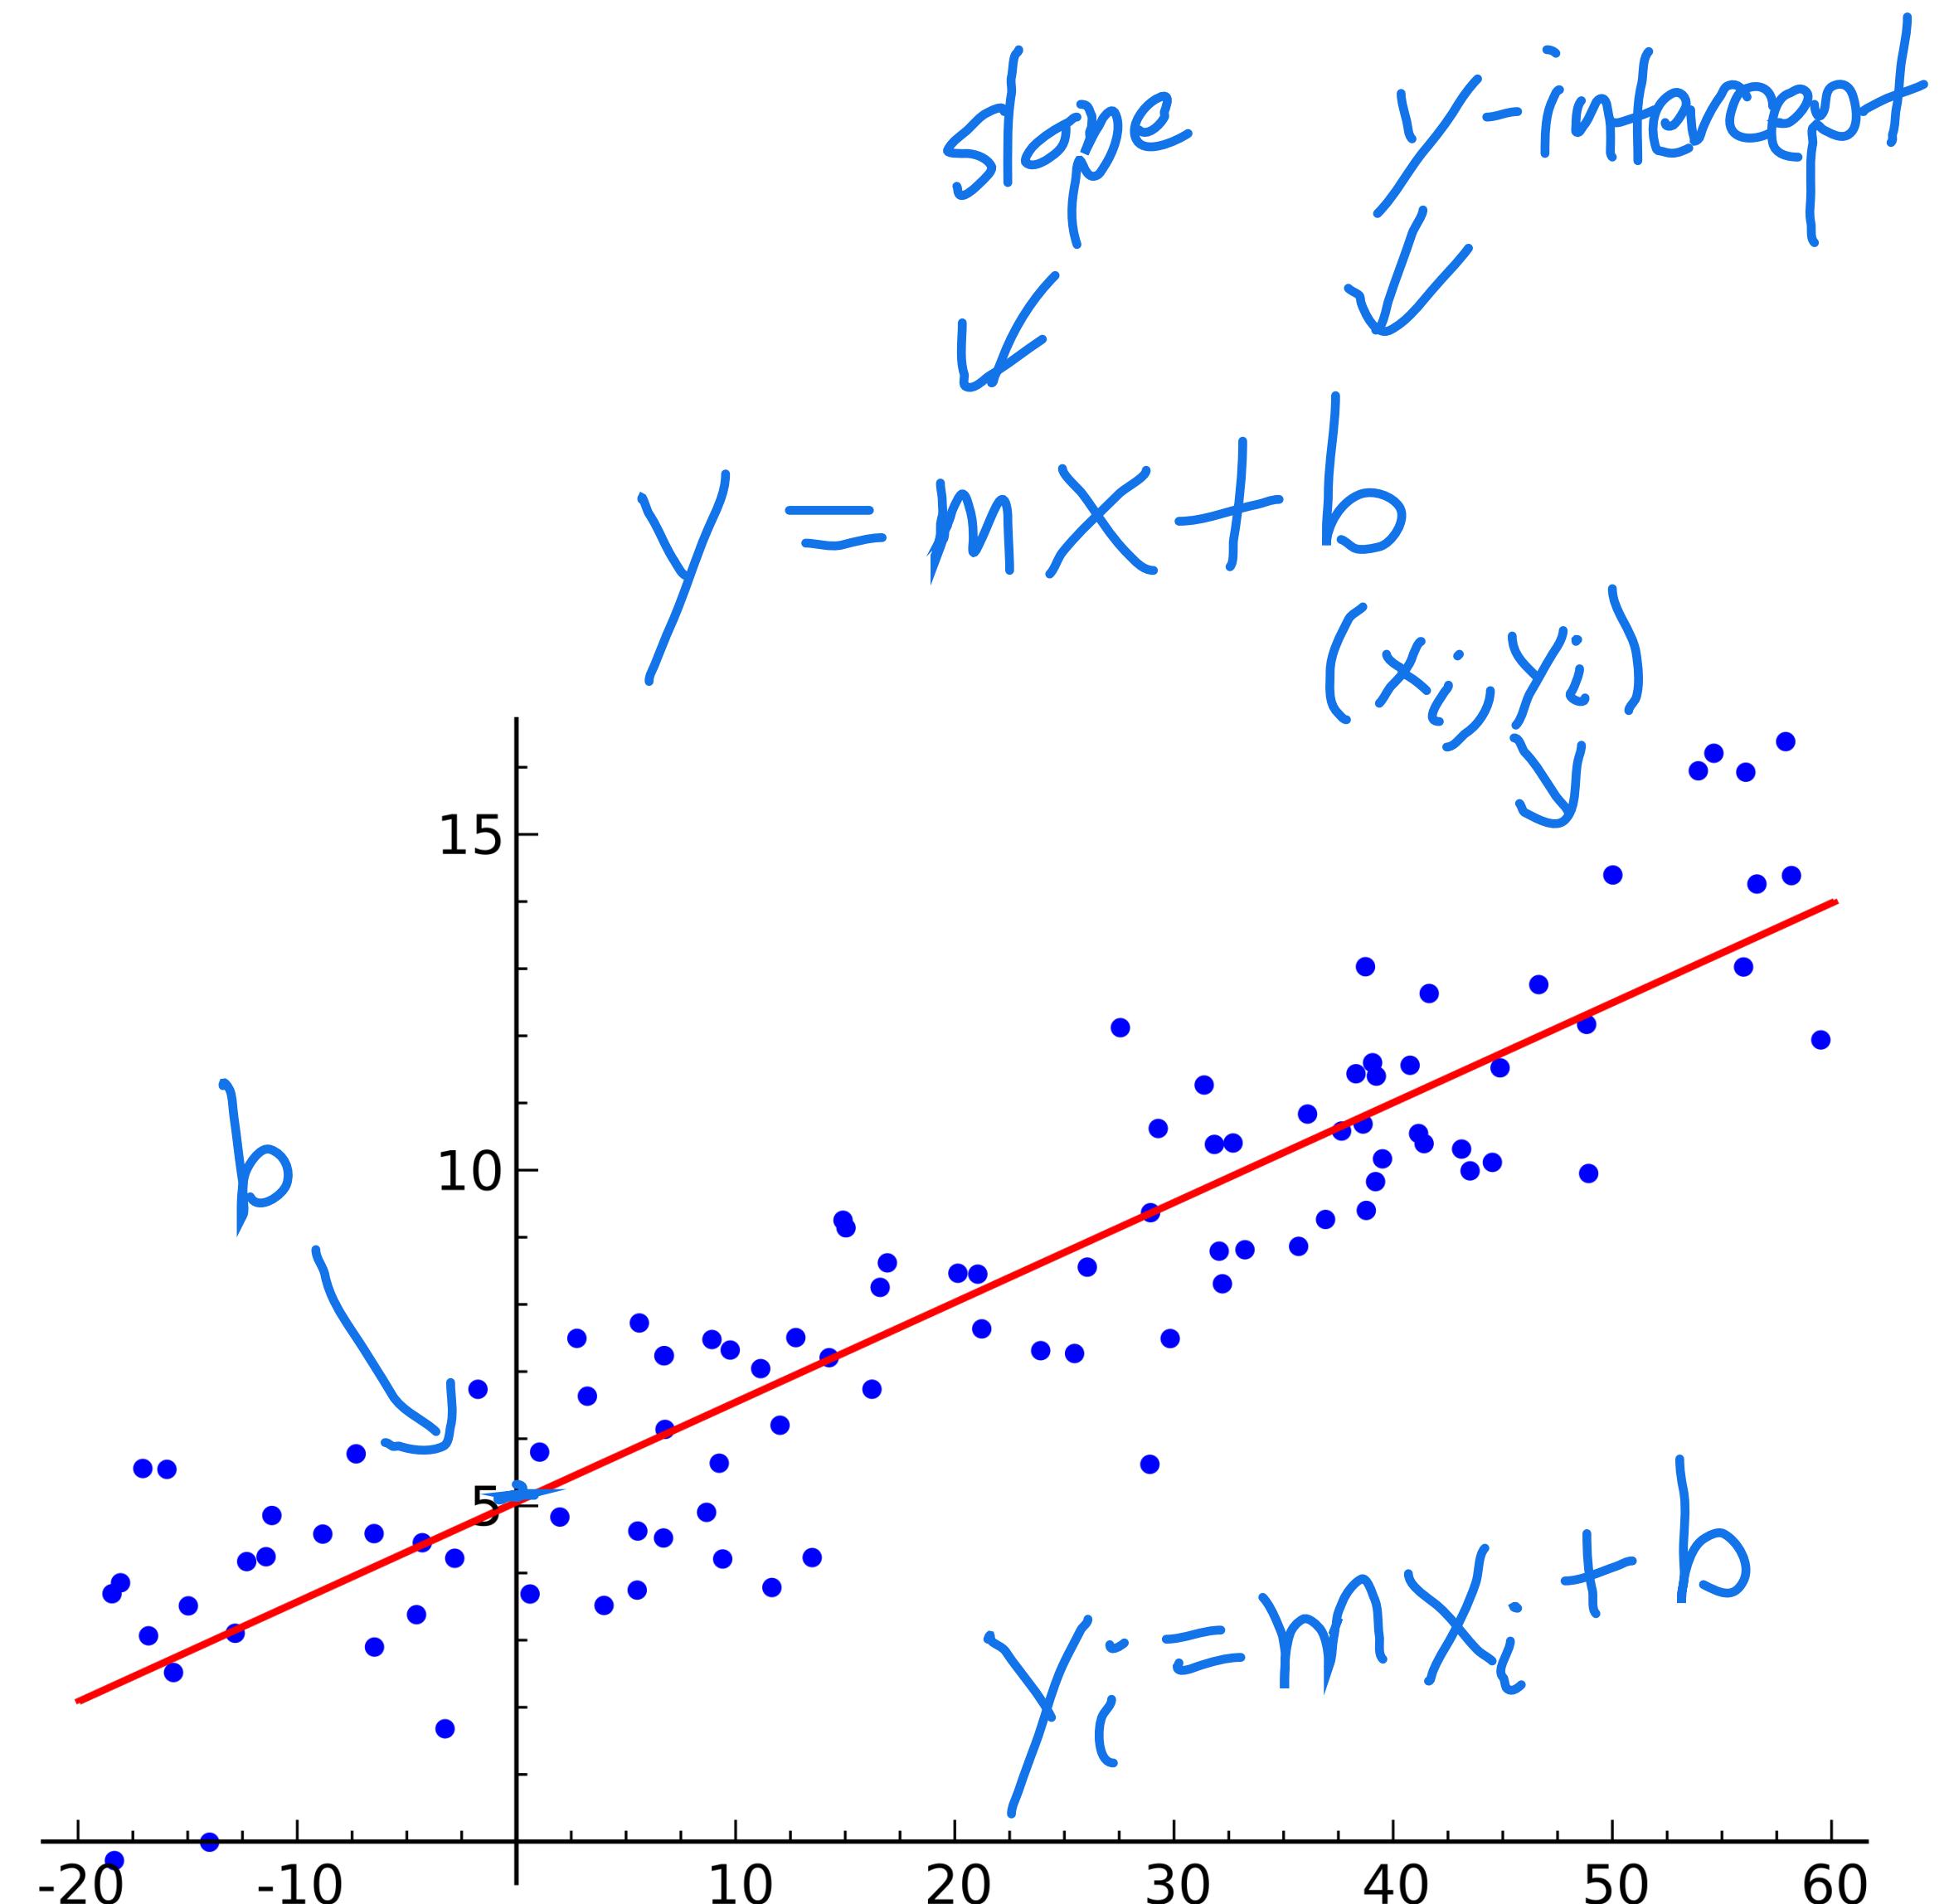
Least Squares is a method
for finding *approximate*
solutions to systems of
linear equations



The Idea

Least Squares is a method for finding *approximate* solutions to systems of linear equations

This is a **lot more useful in practice** than exact solutions

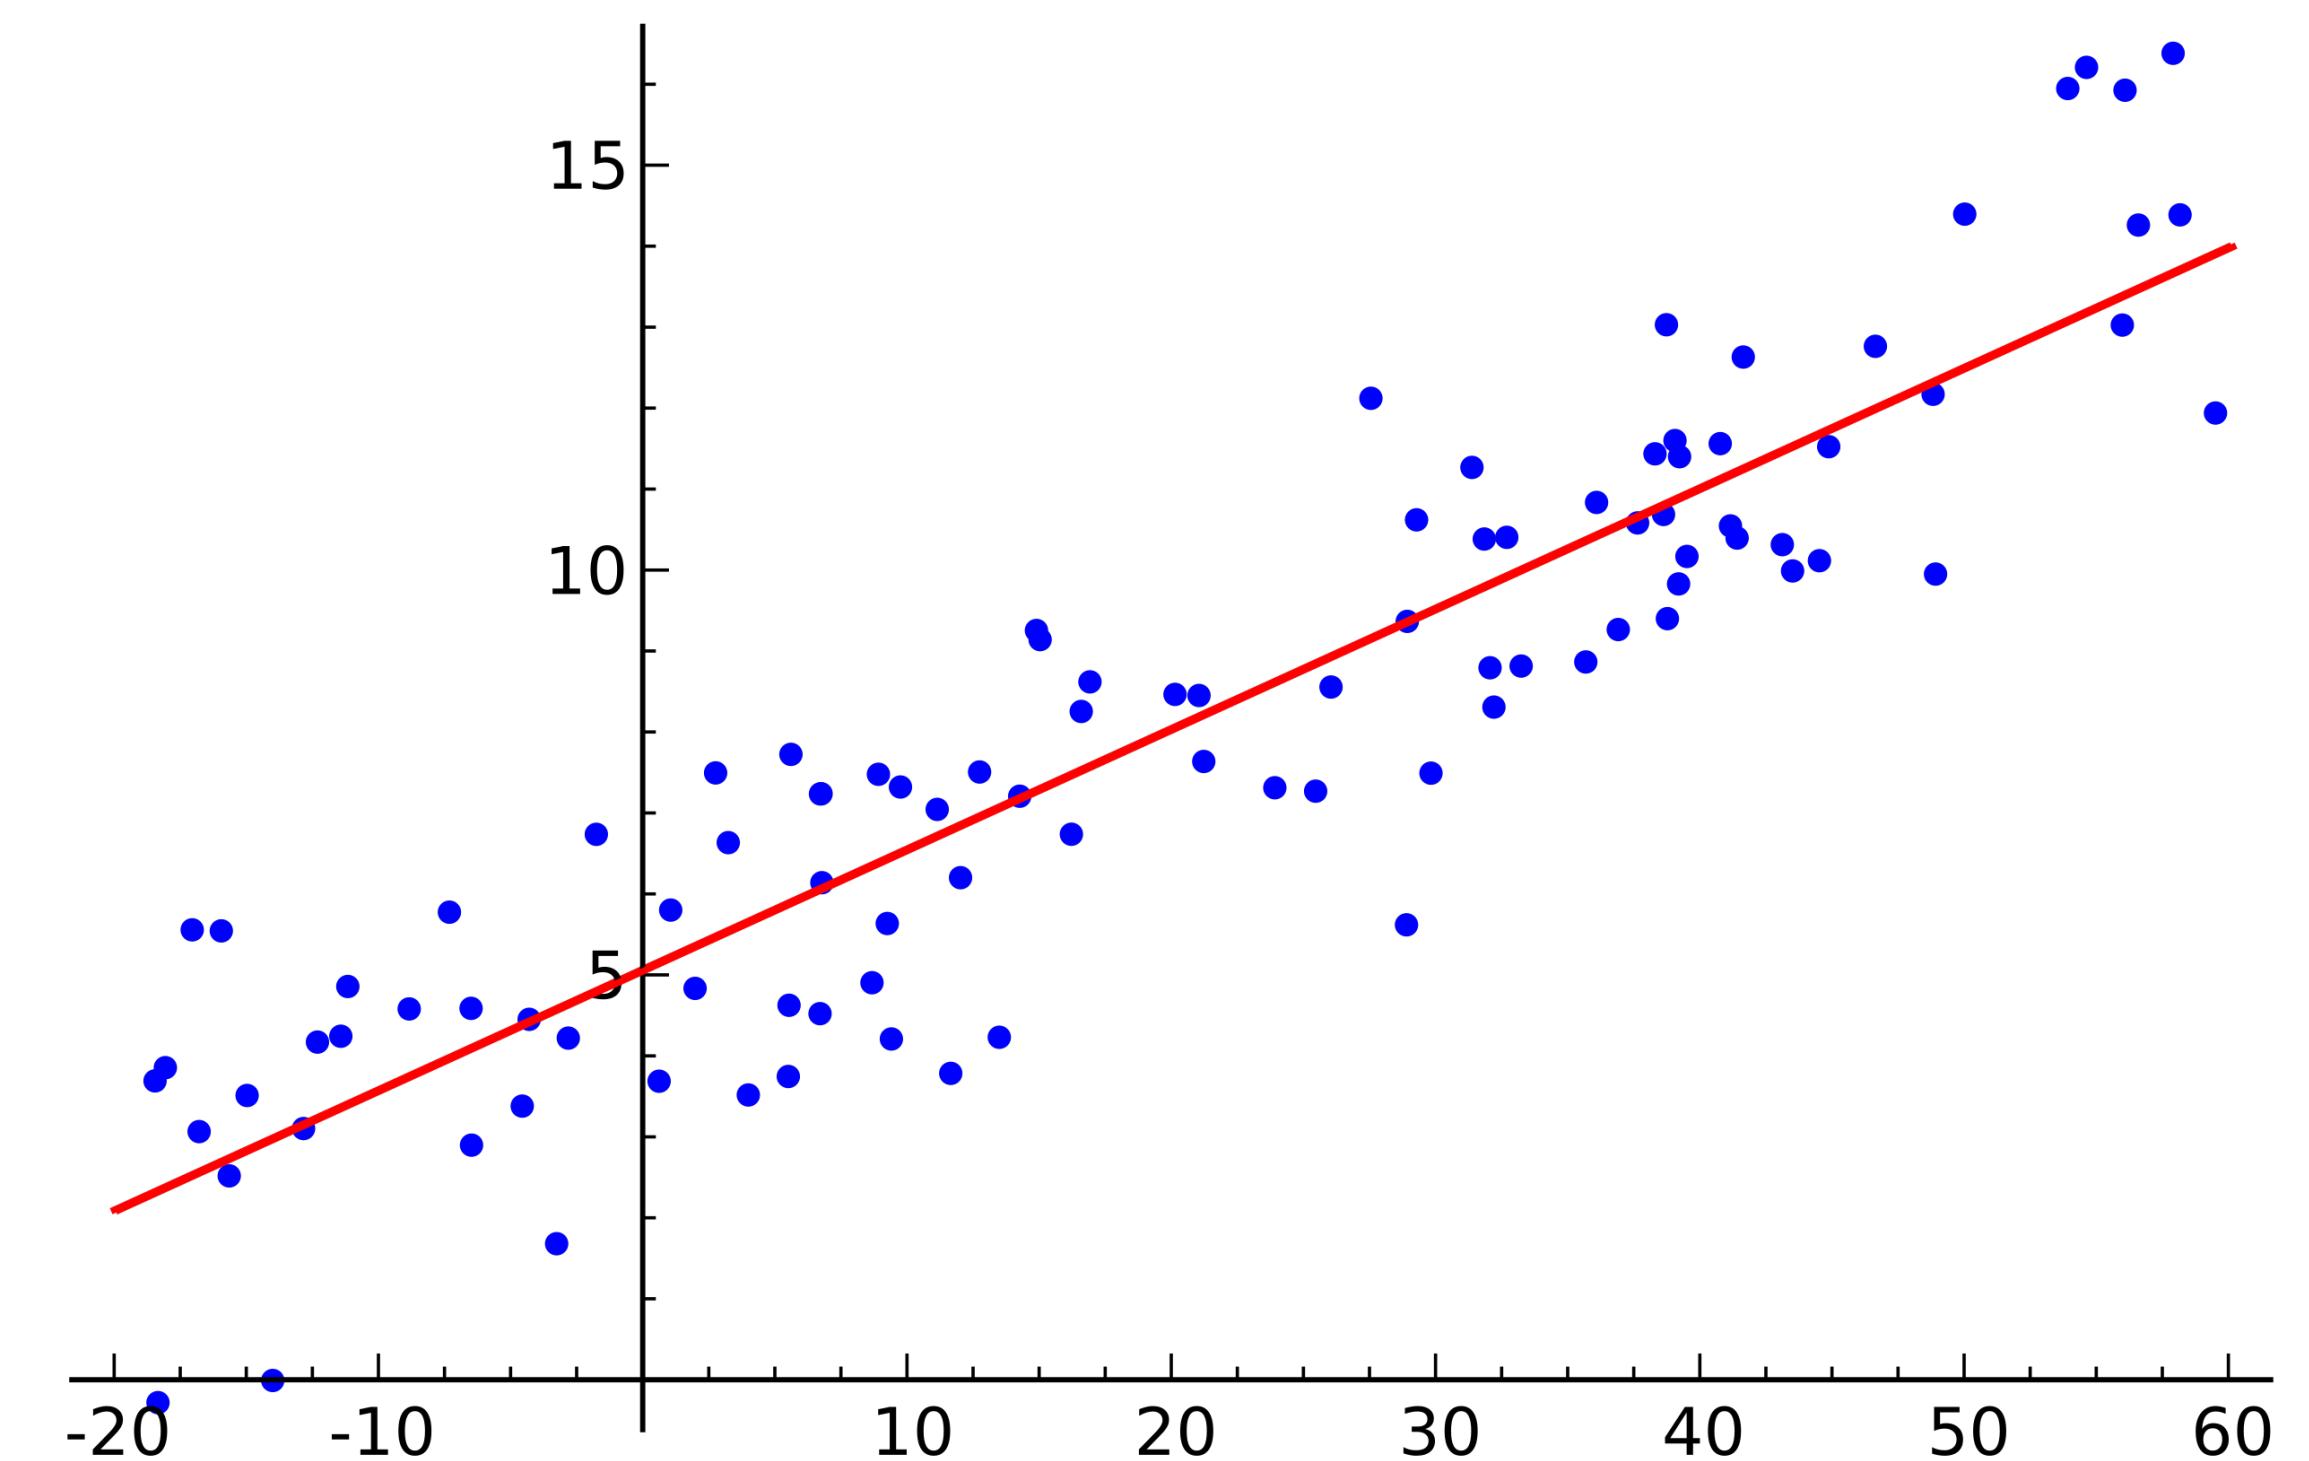


The Idea

Least Squares is a method for finding *approximate* solutions to systems of linear equations

This is a **lot more useful in practice** than exact solutions

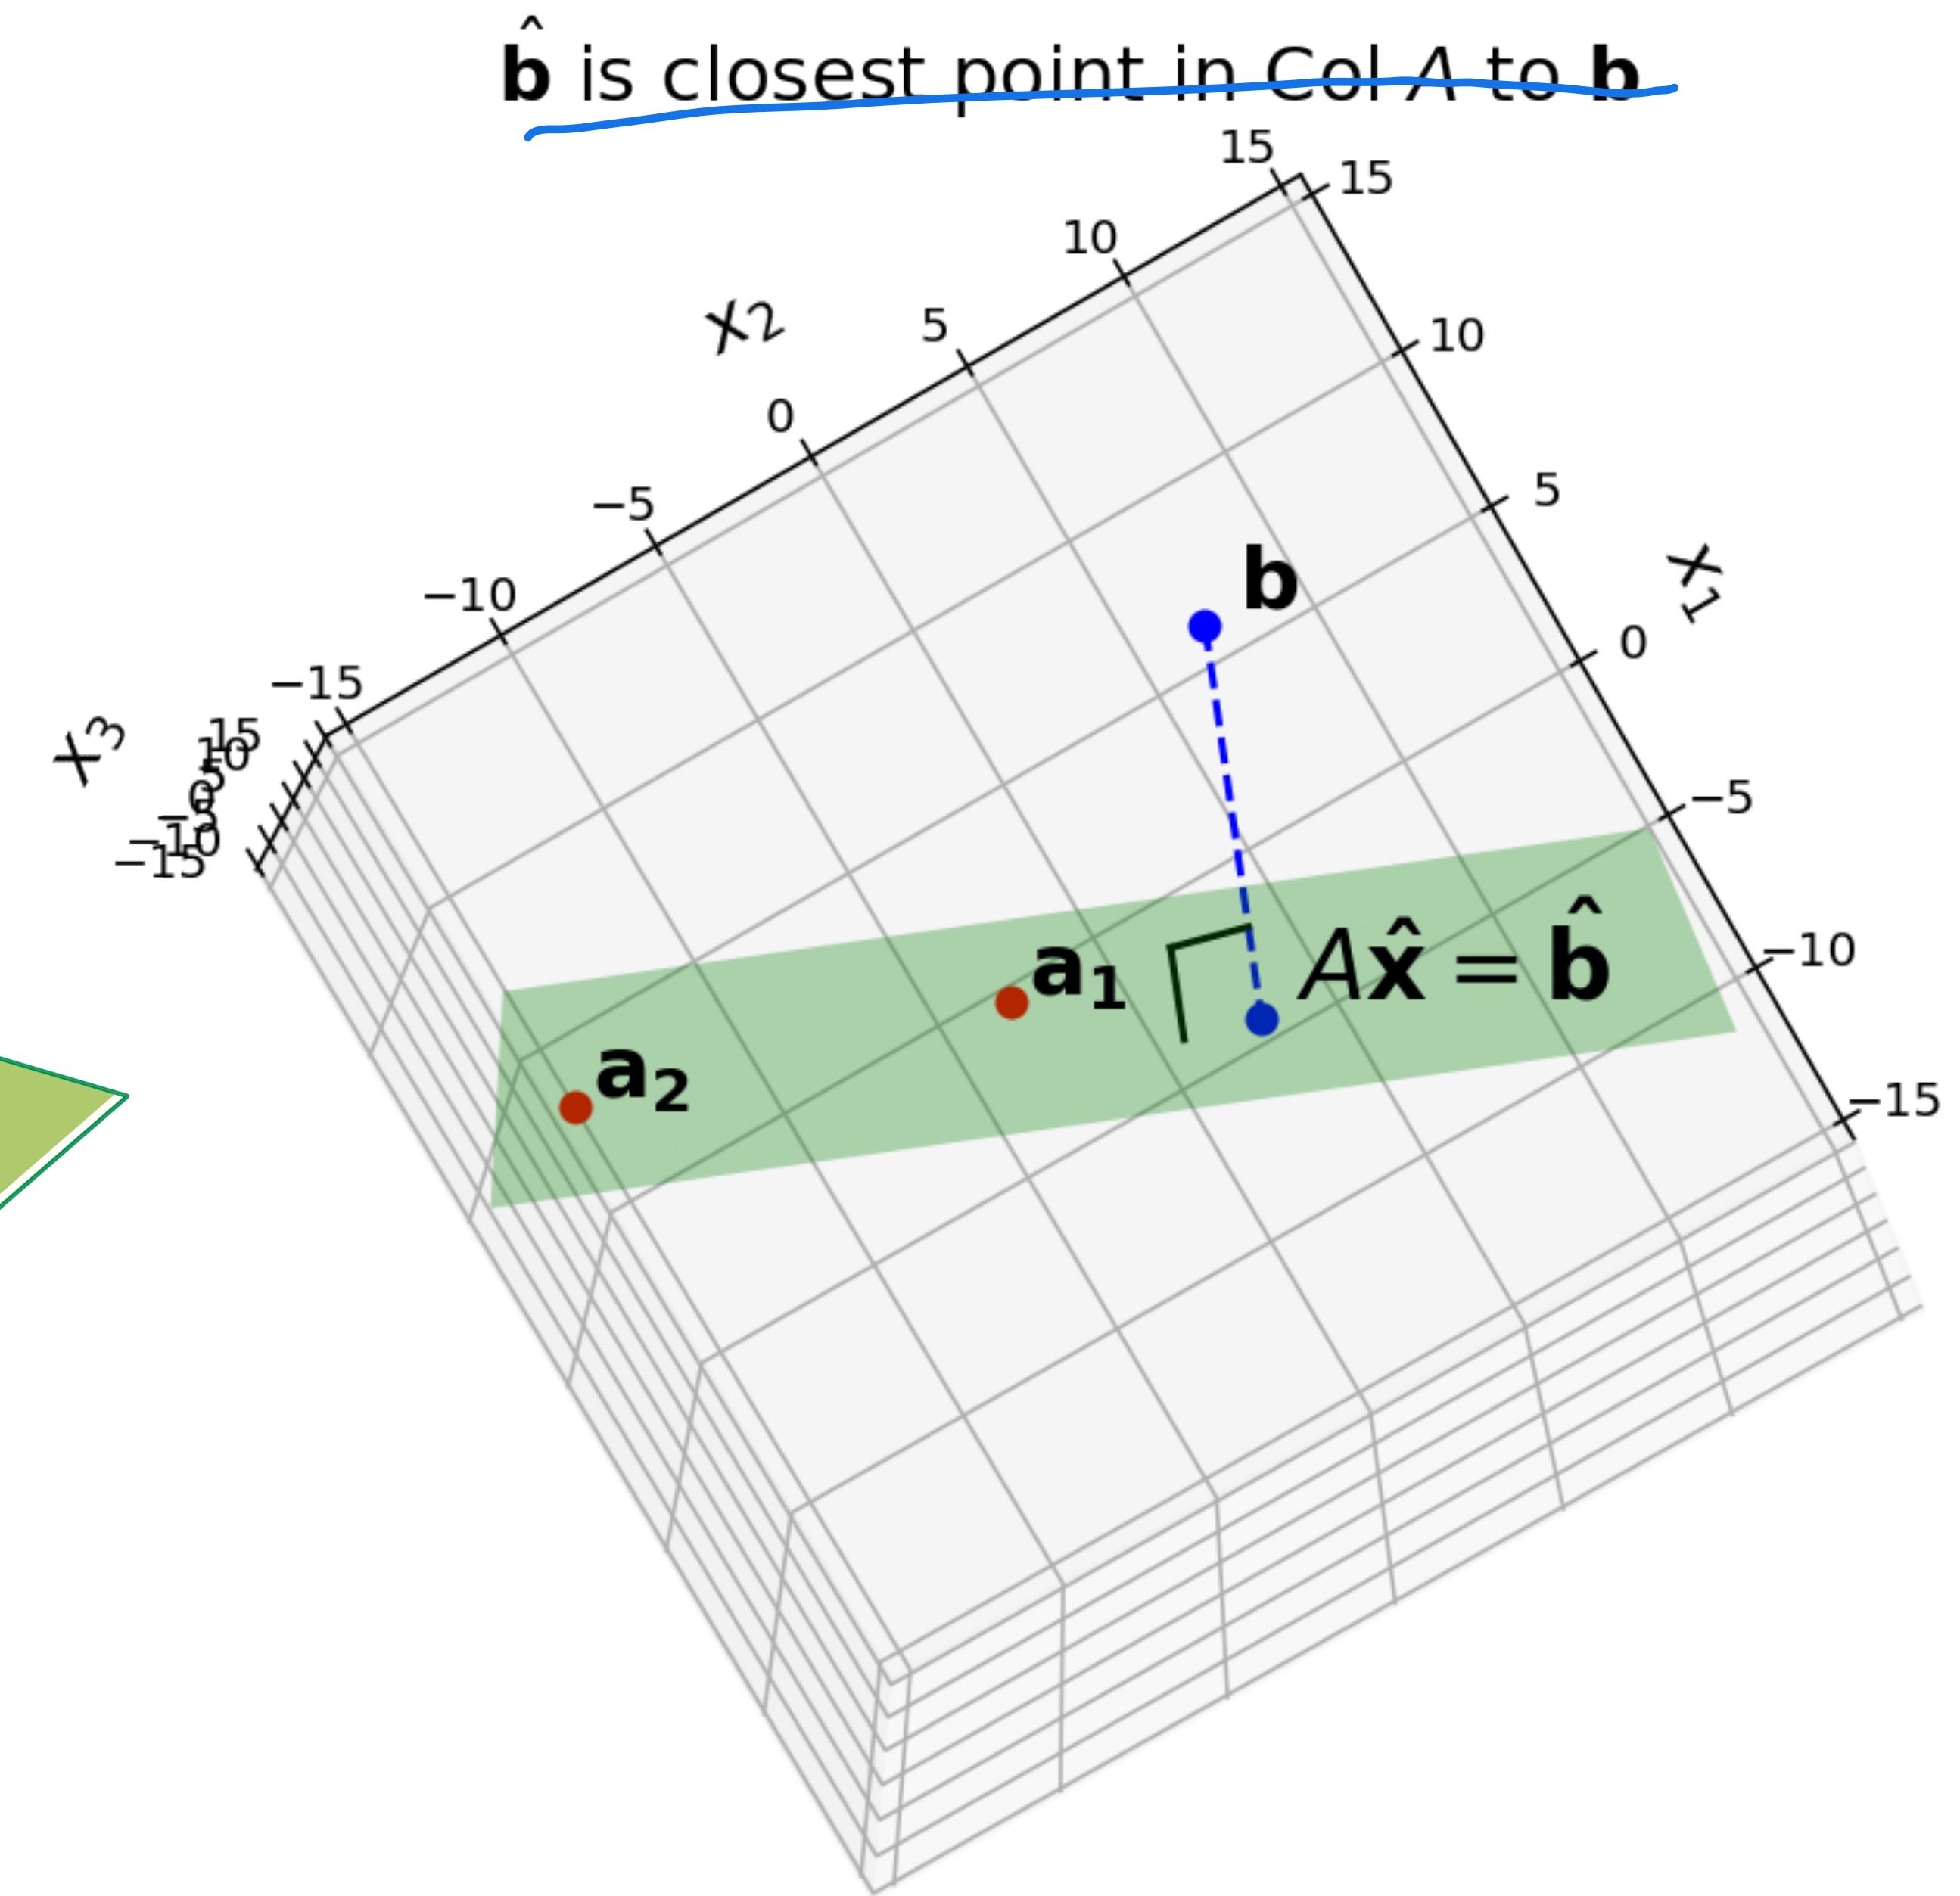
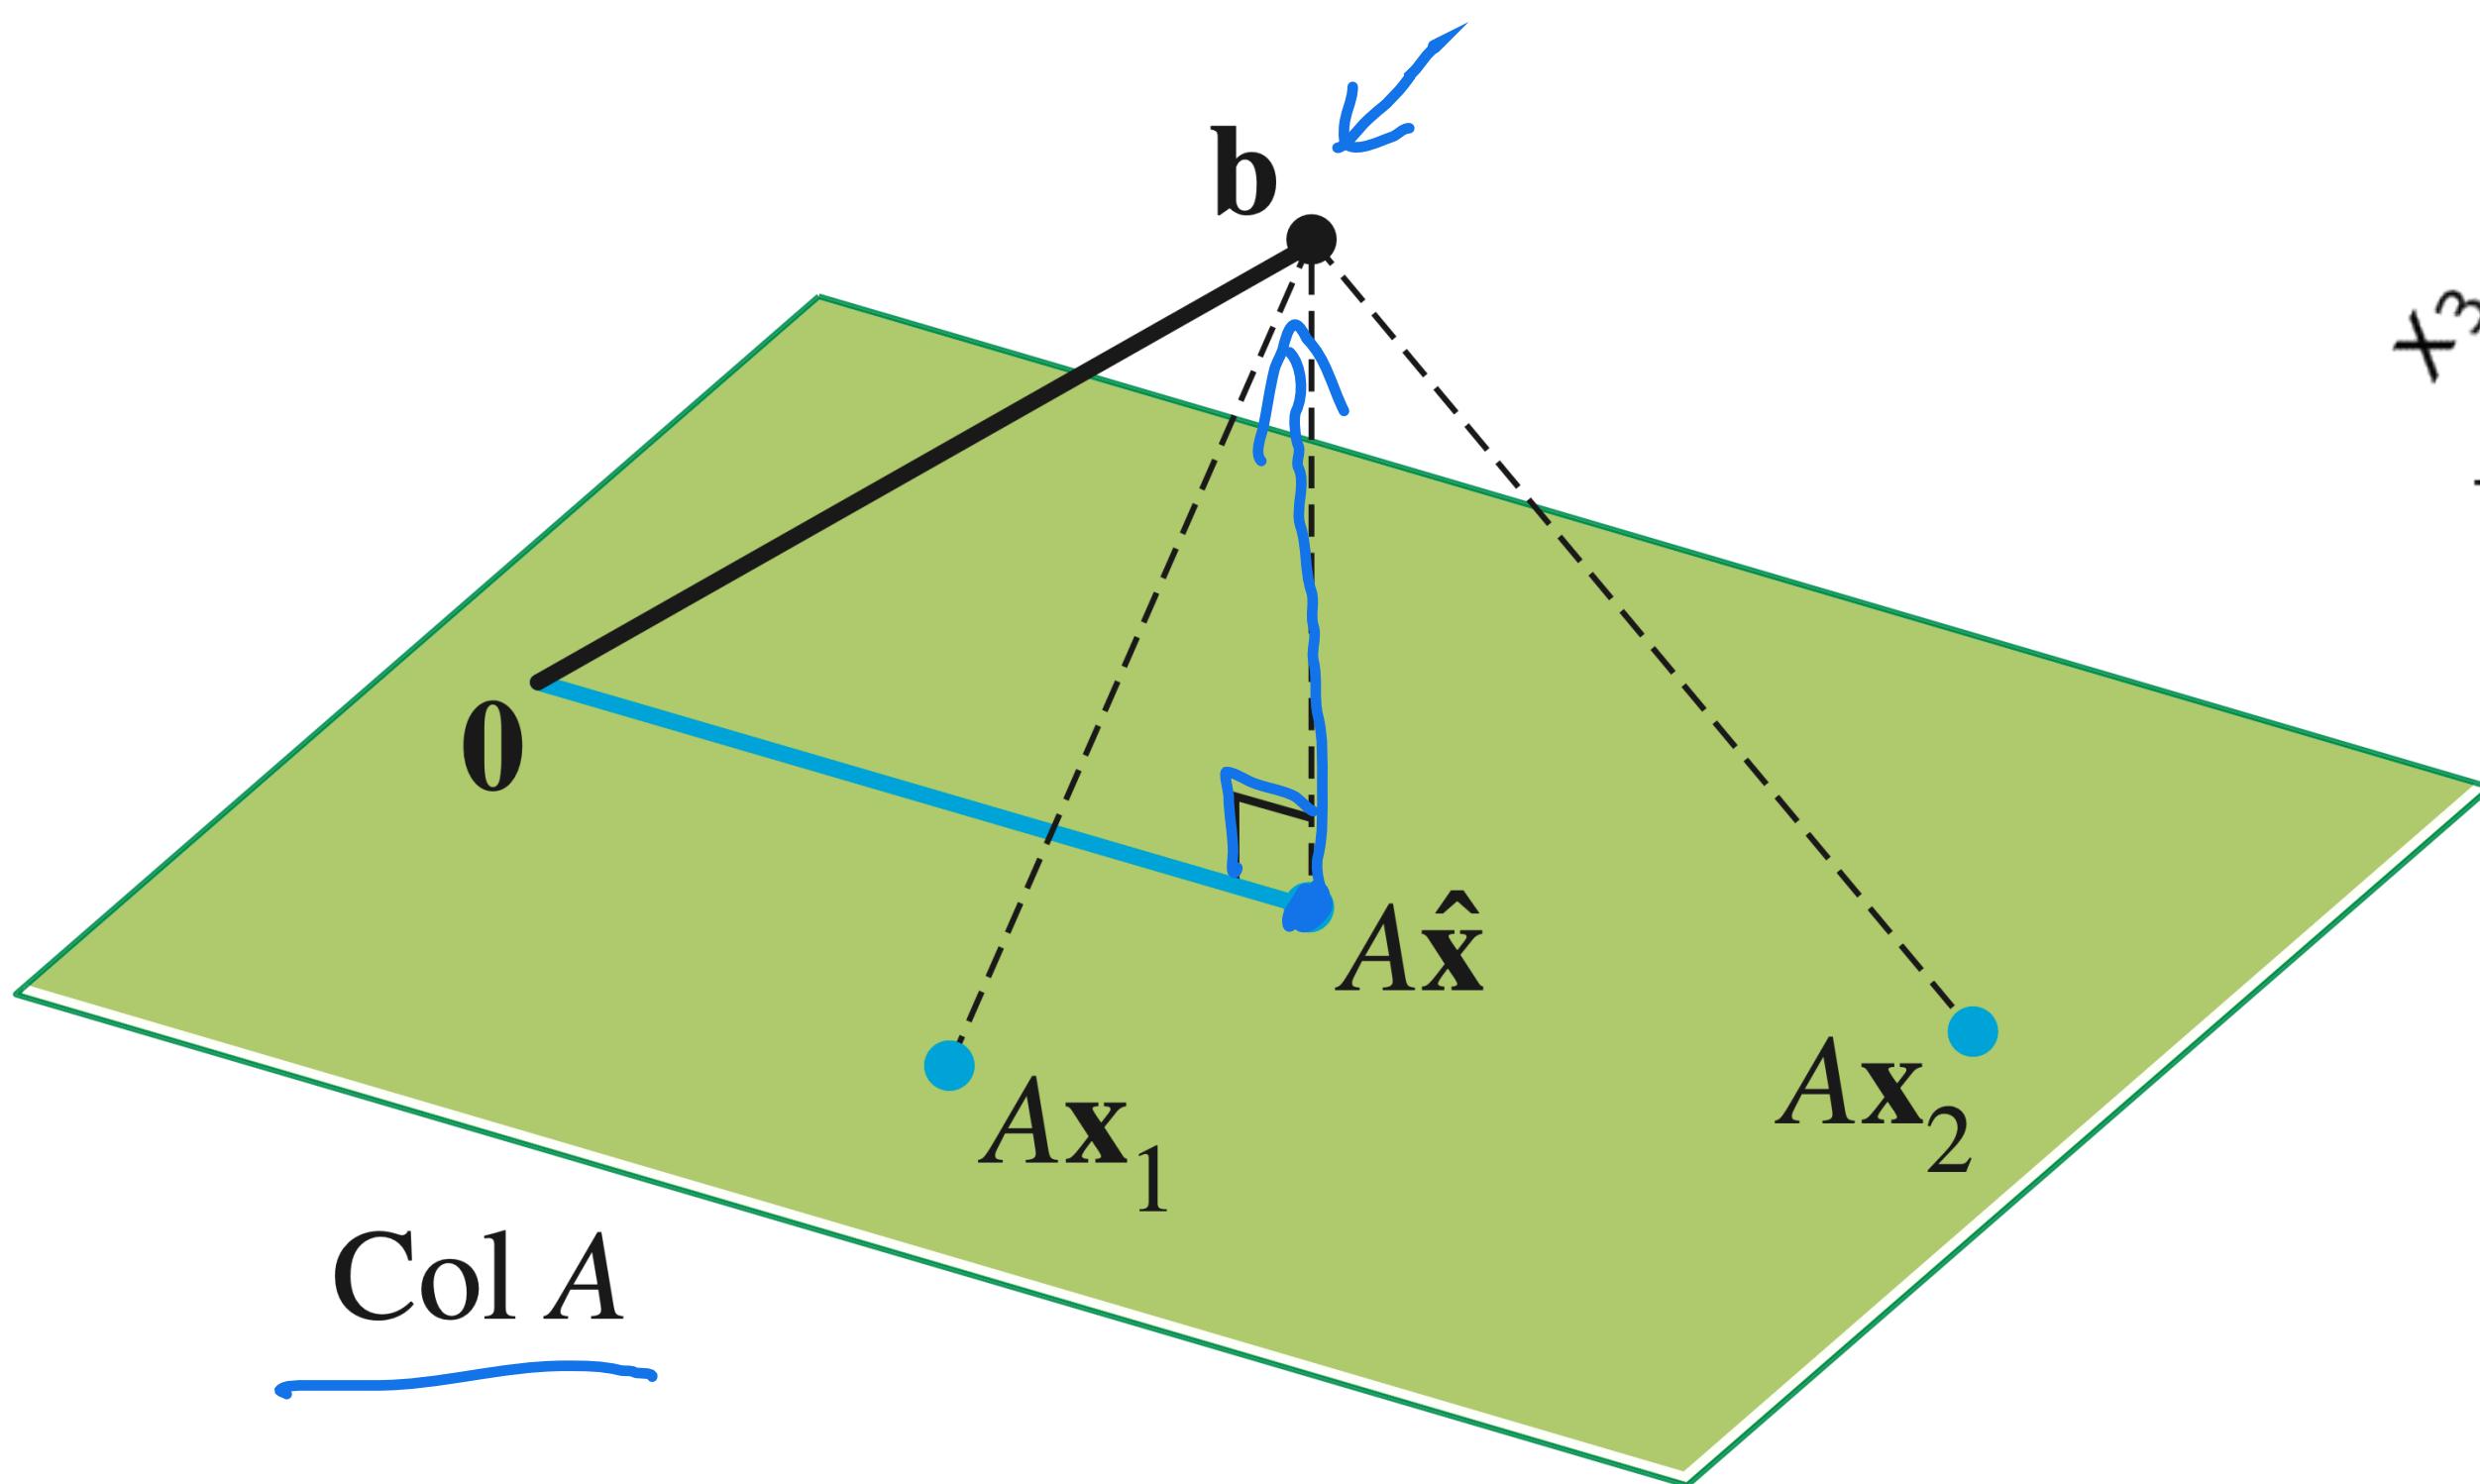
It can be used to do **linear regression** from stats class



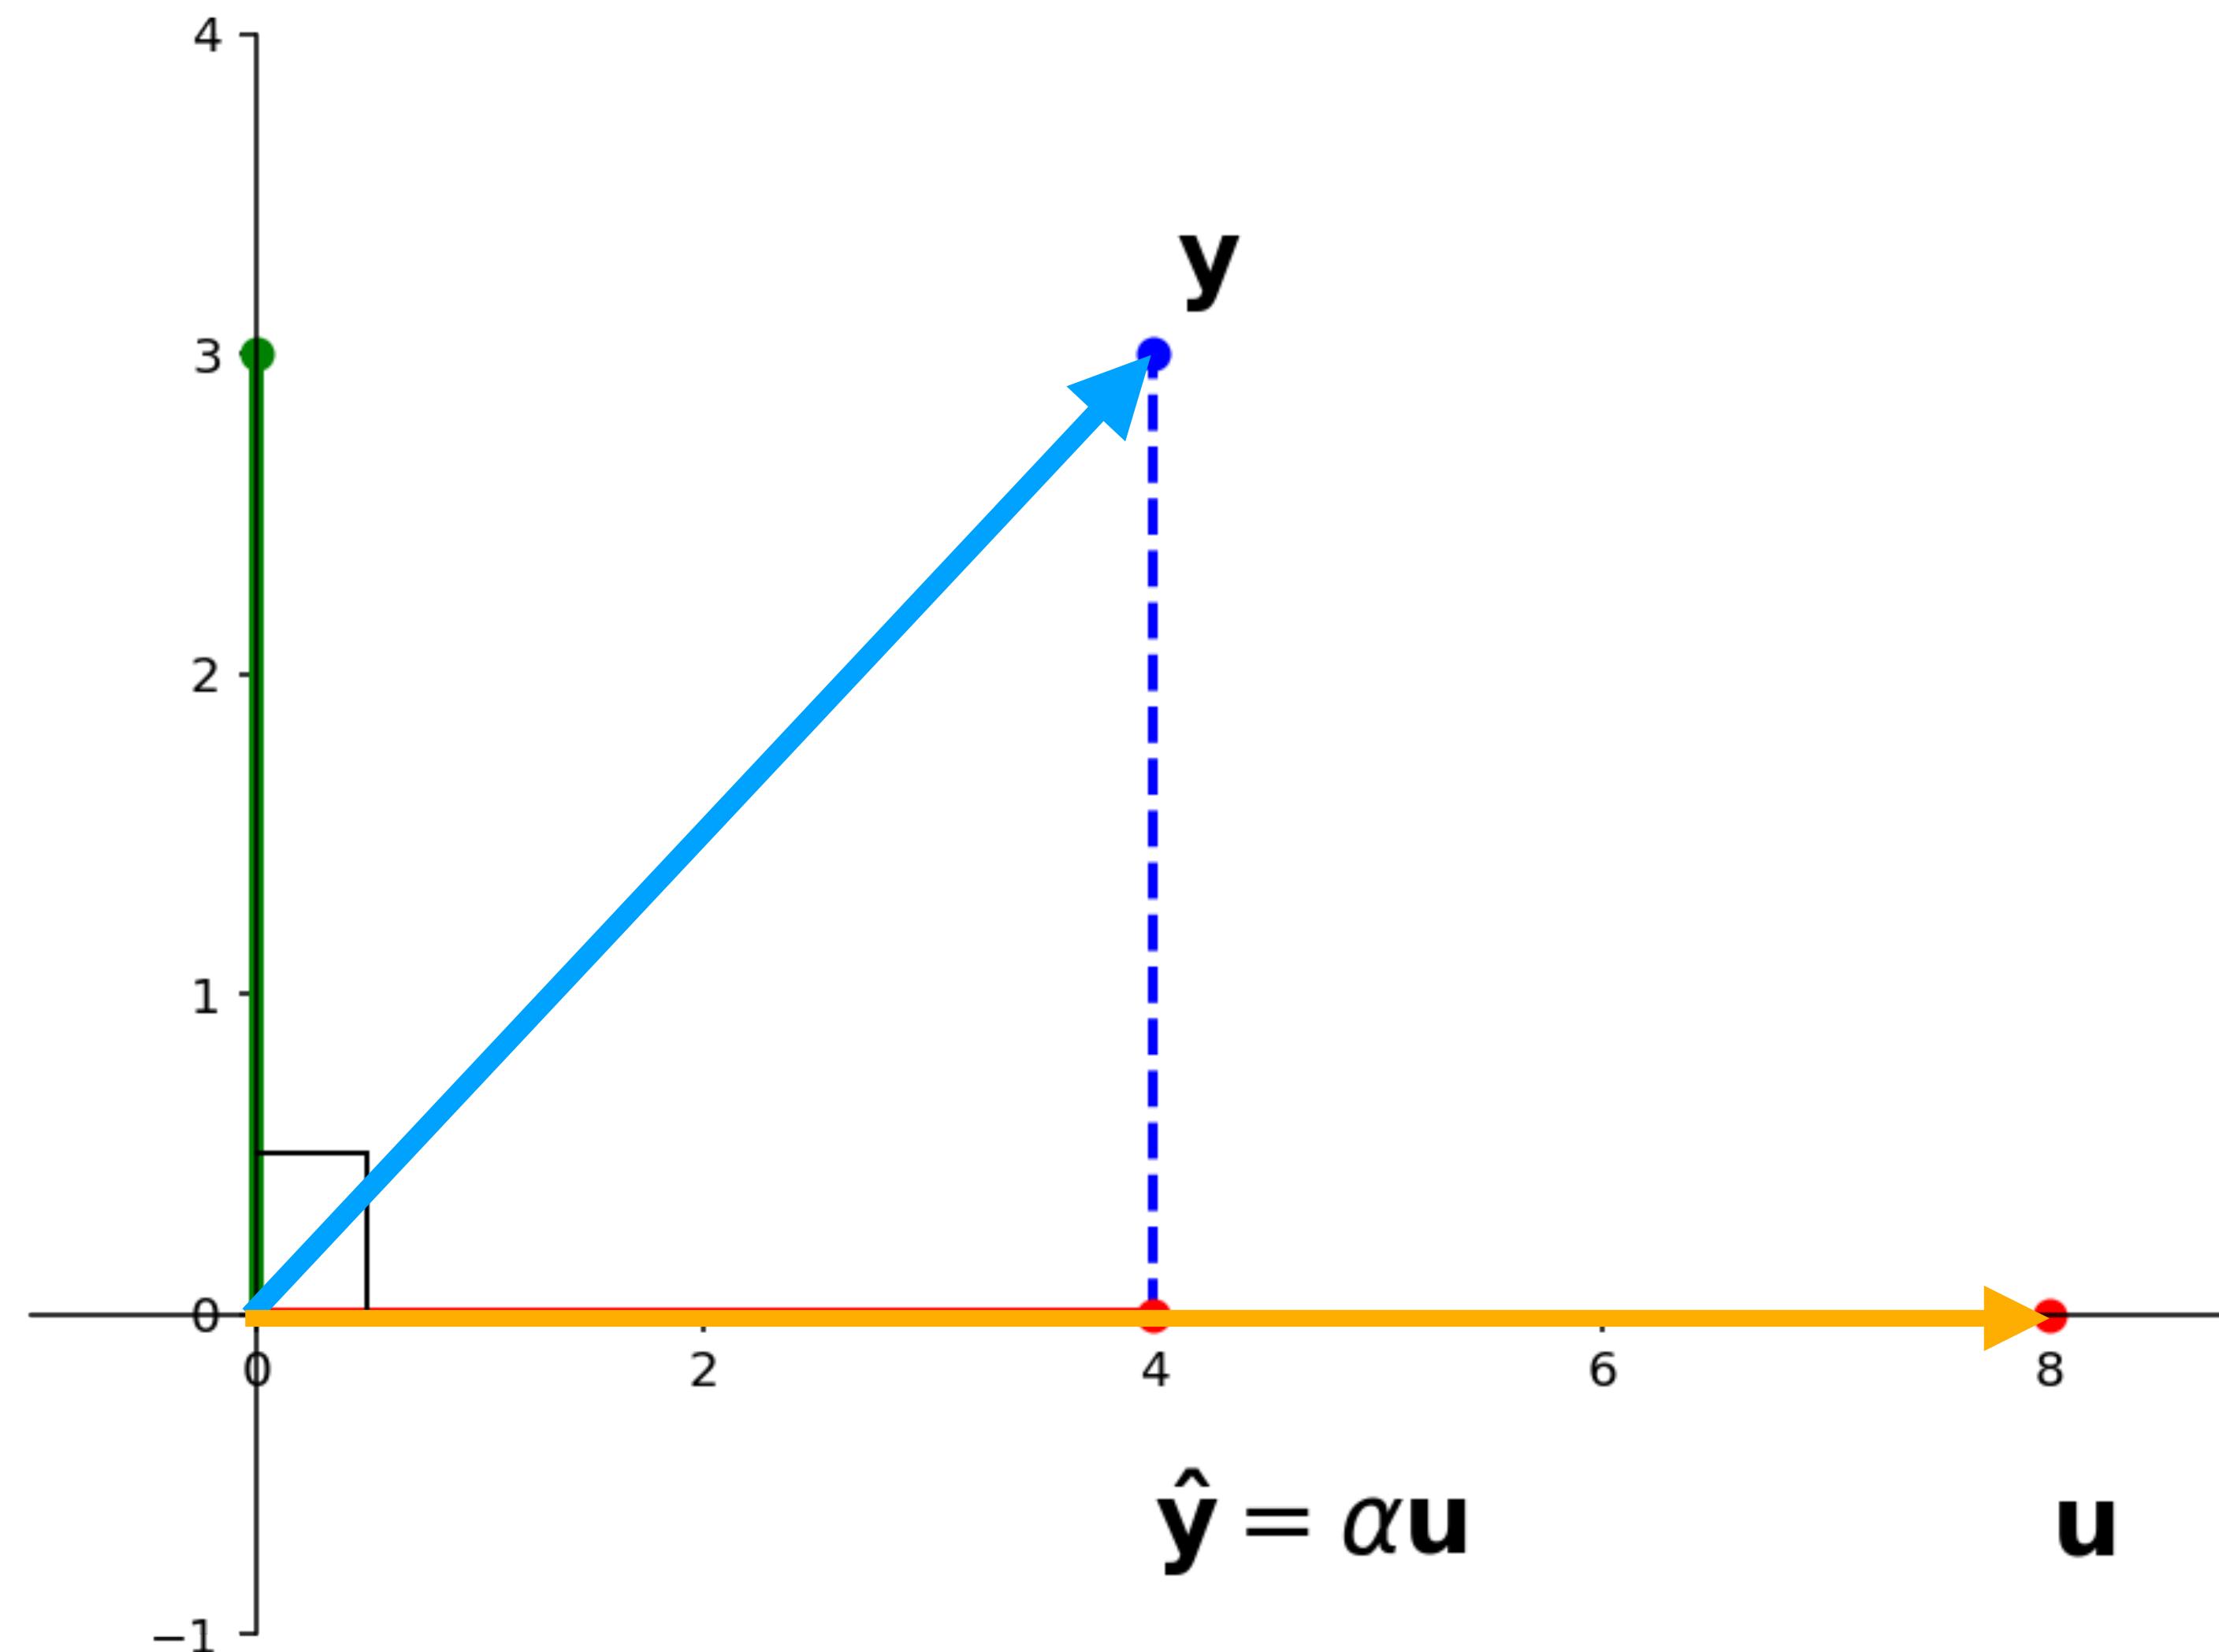
General Least Squares Problem

Figure 22.8

The Picture

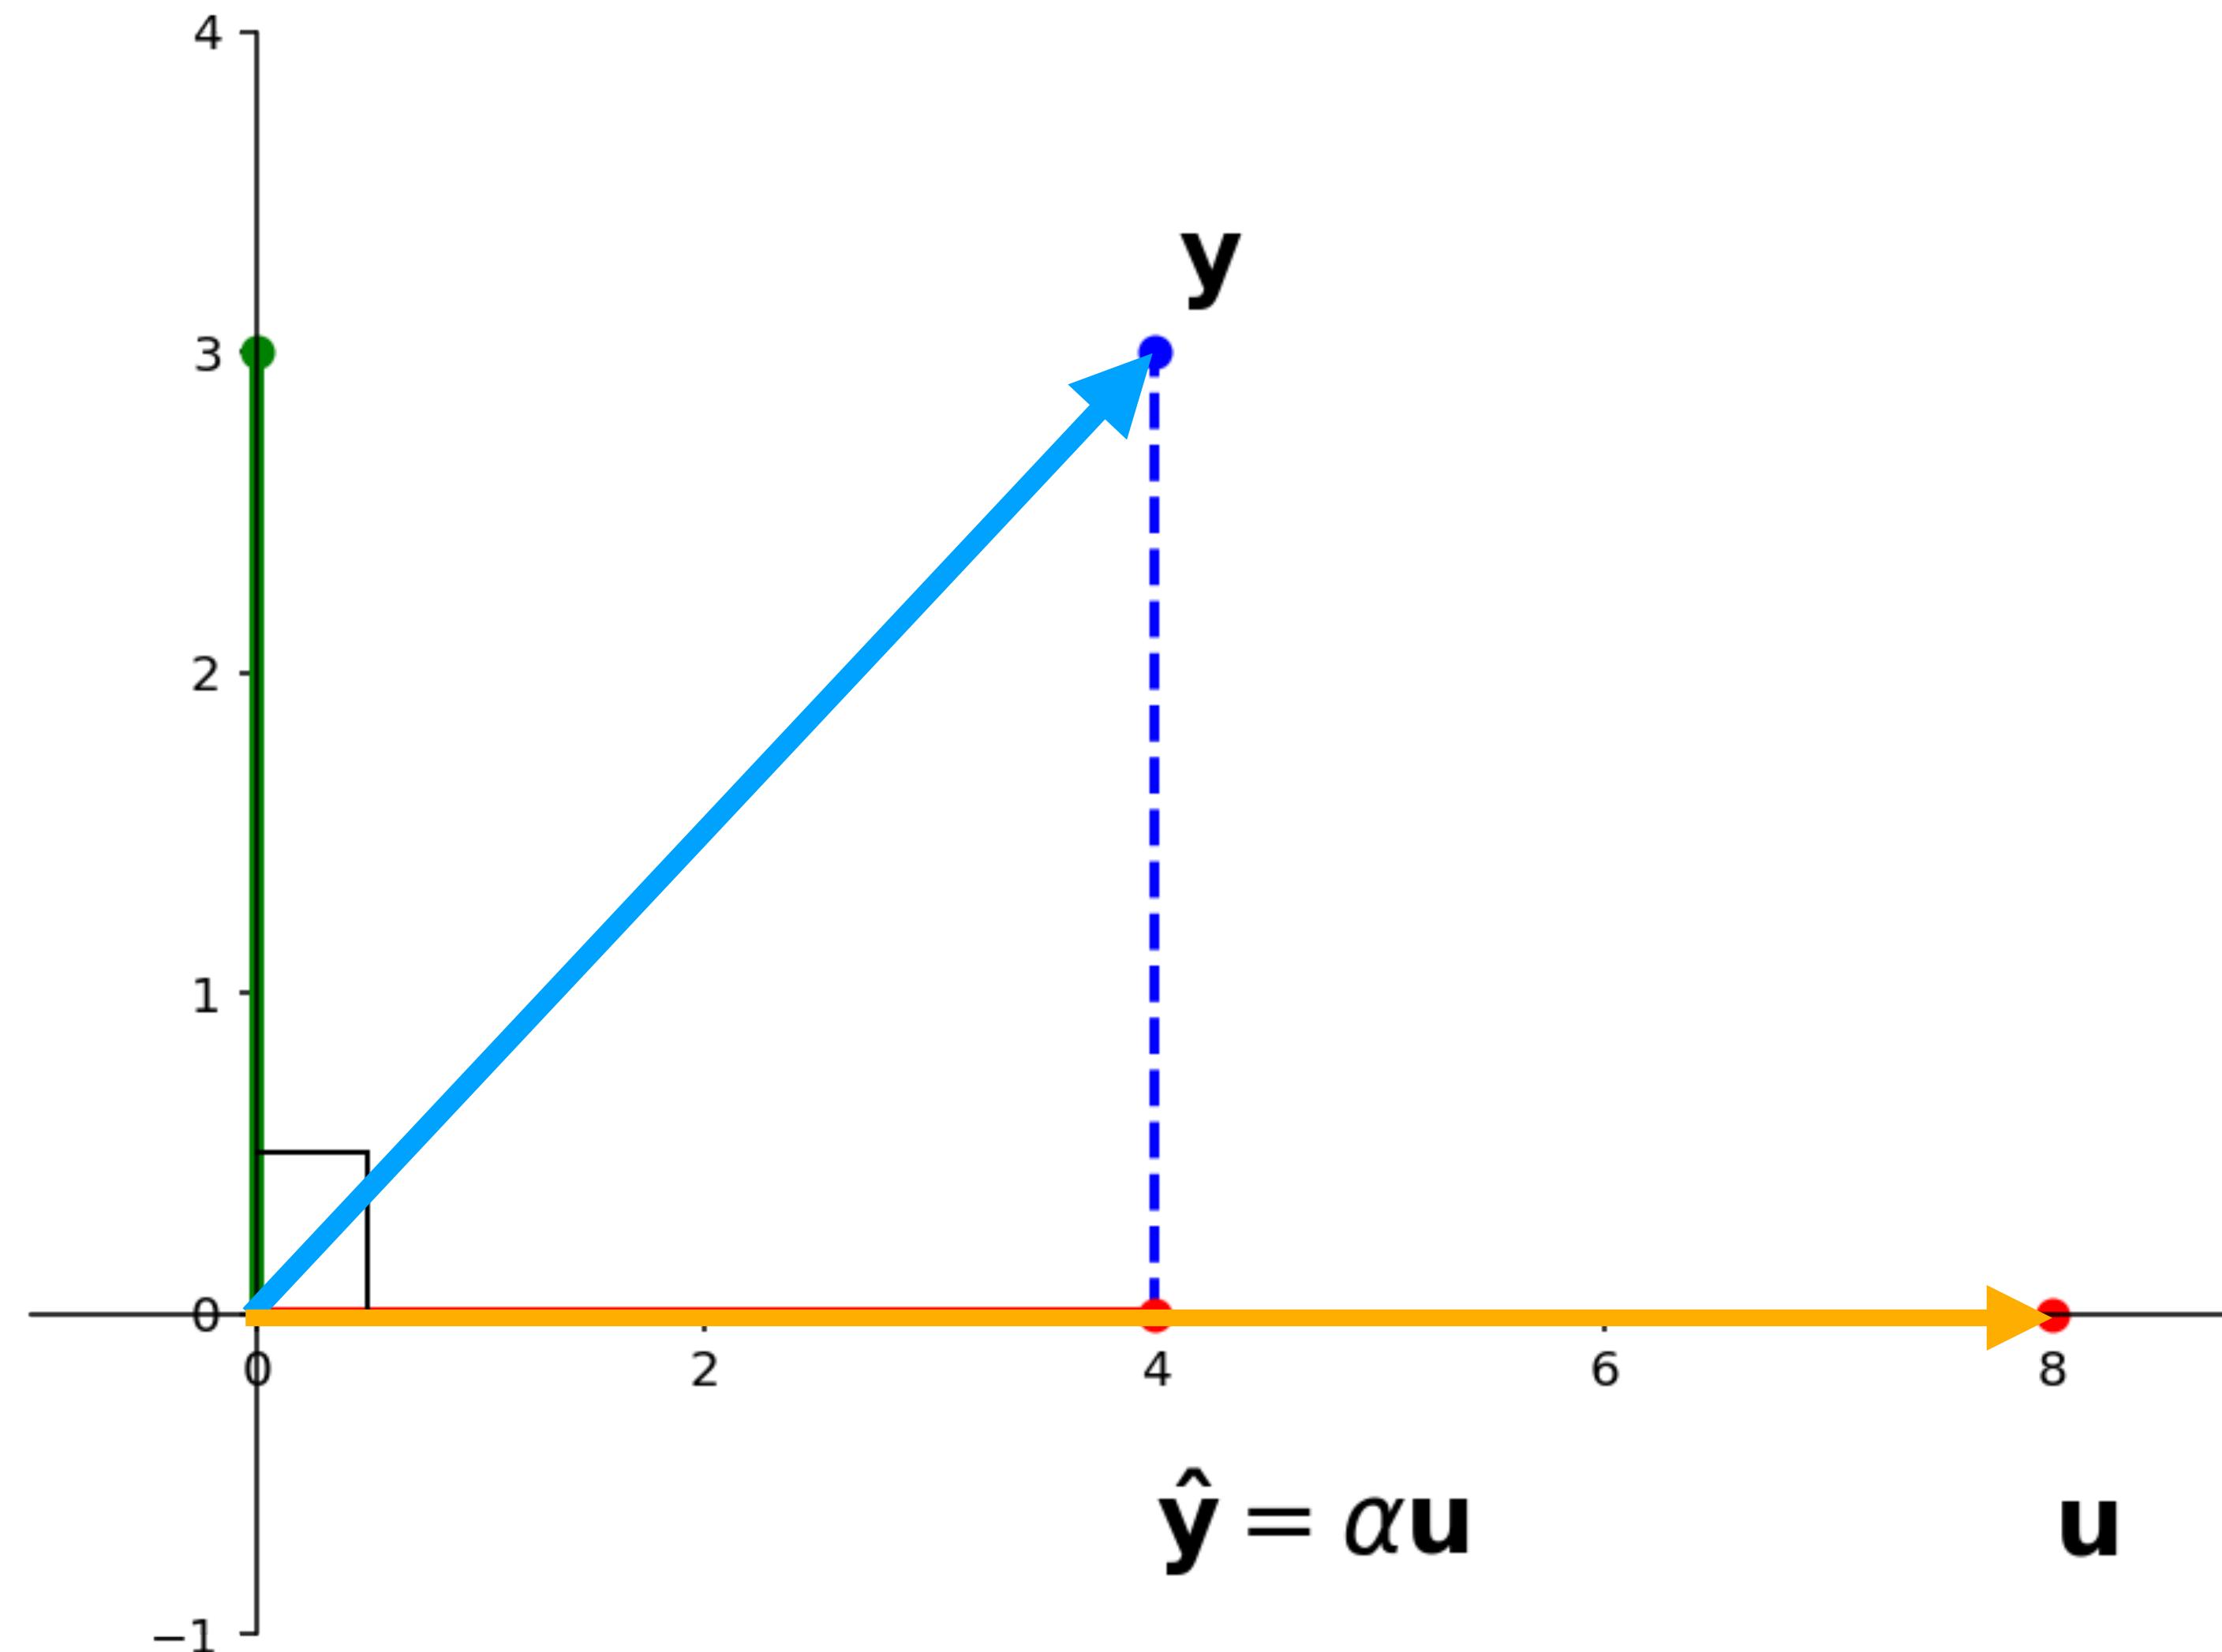


Recall: Orthogonal Projection



Recall: Orthogonal Projection

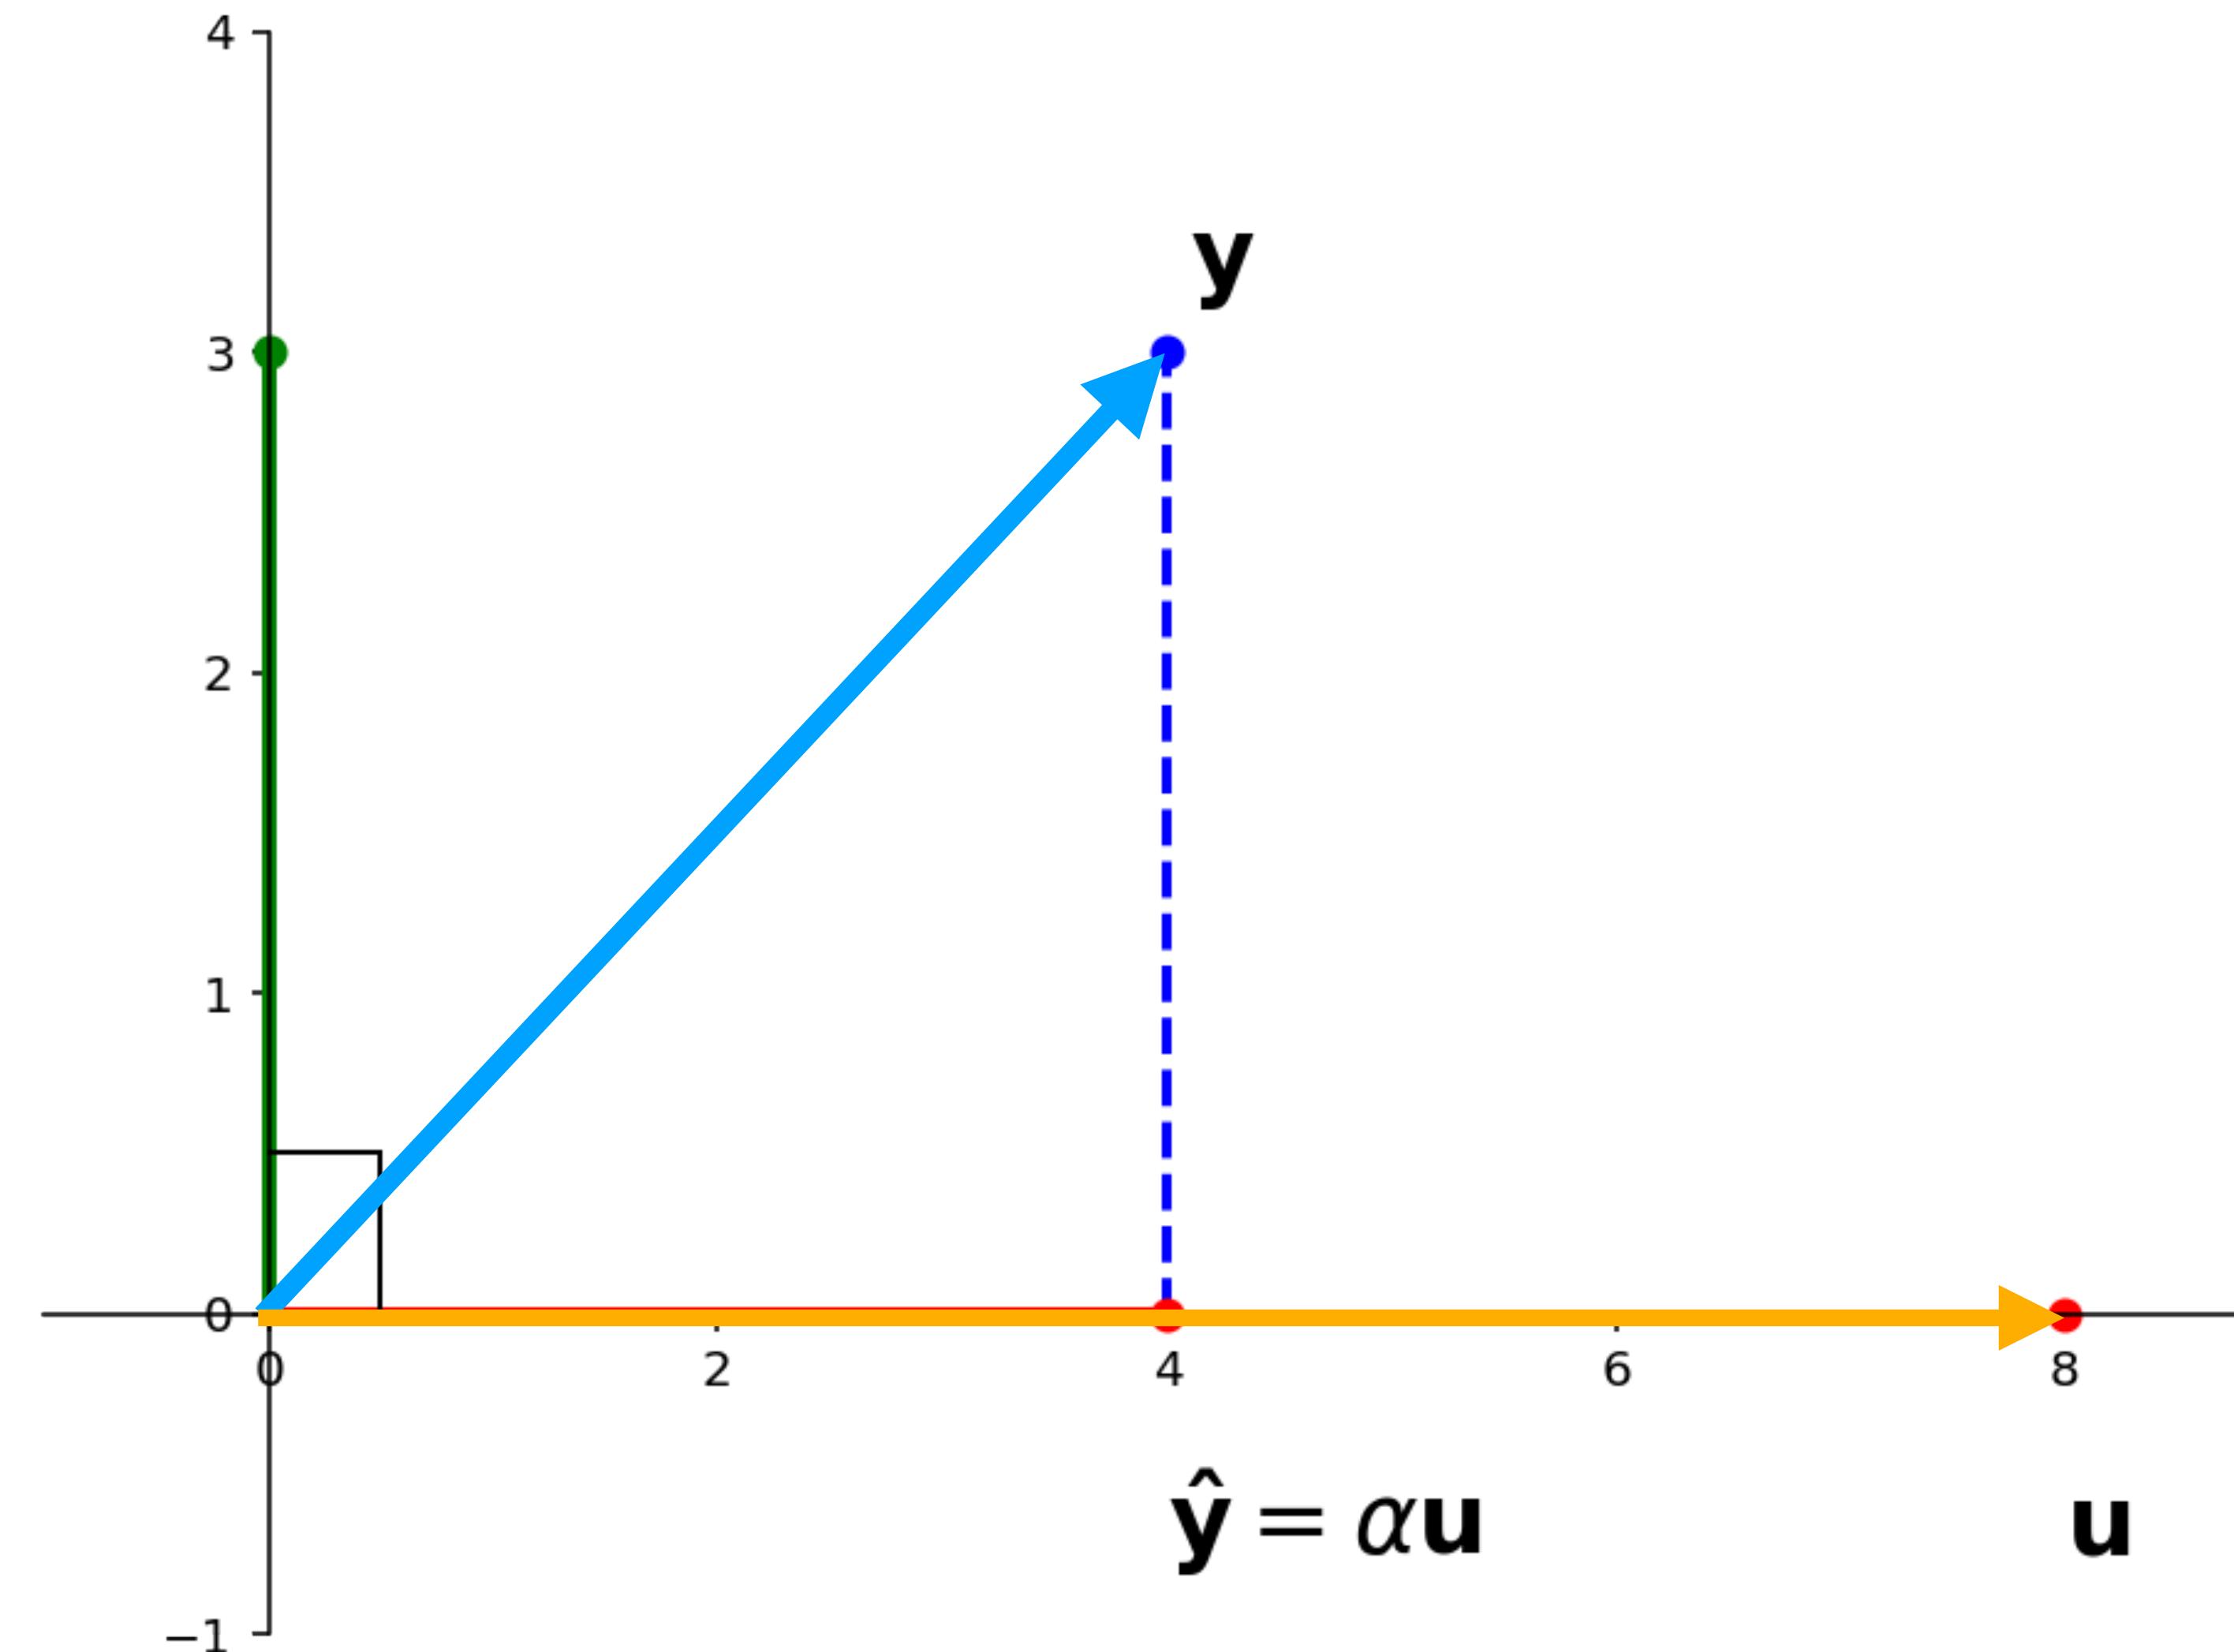
Question. Given vectors \mathbf{y} and \mathbf{u} in R^n , find vectors $\hat{\mathbf{y}}$ and \mathbf{z} such that



Recall: Orthogonal Projection

Question. Given vectors \mathbf{y} and \mathbf{u} in R^n , find vectors $\hat{\mathbf{y}}$ and \mathbf{z} such that

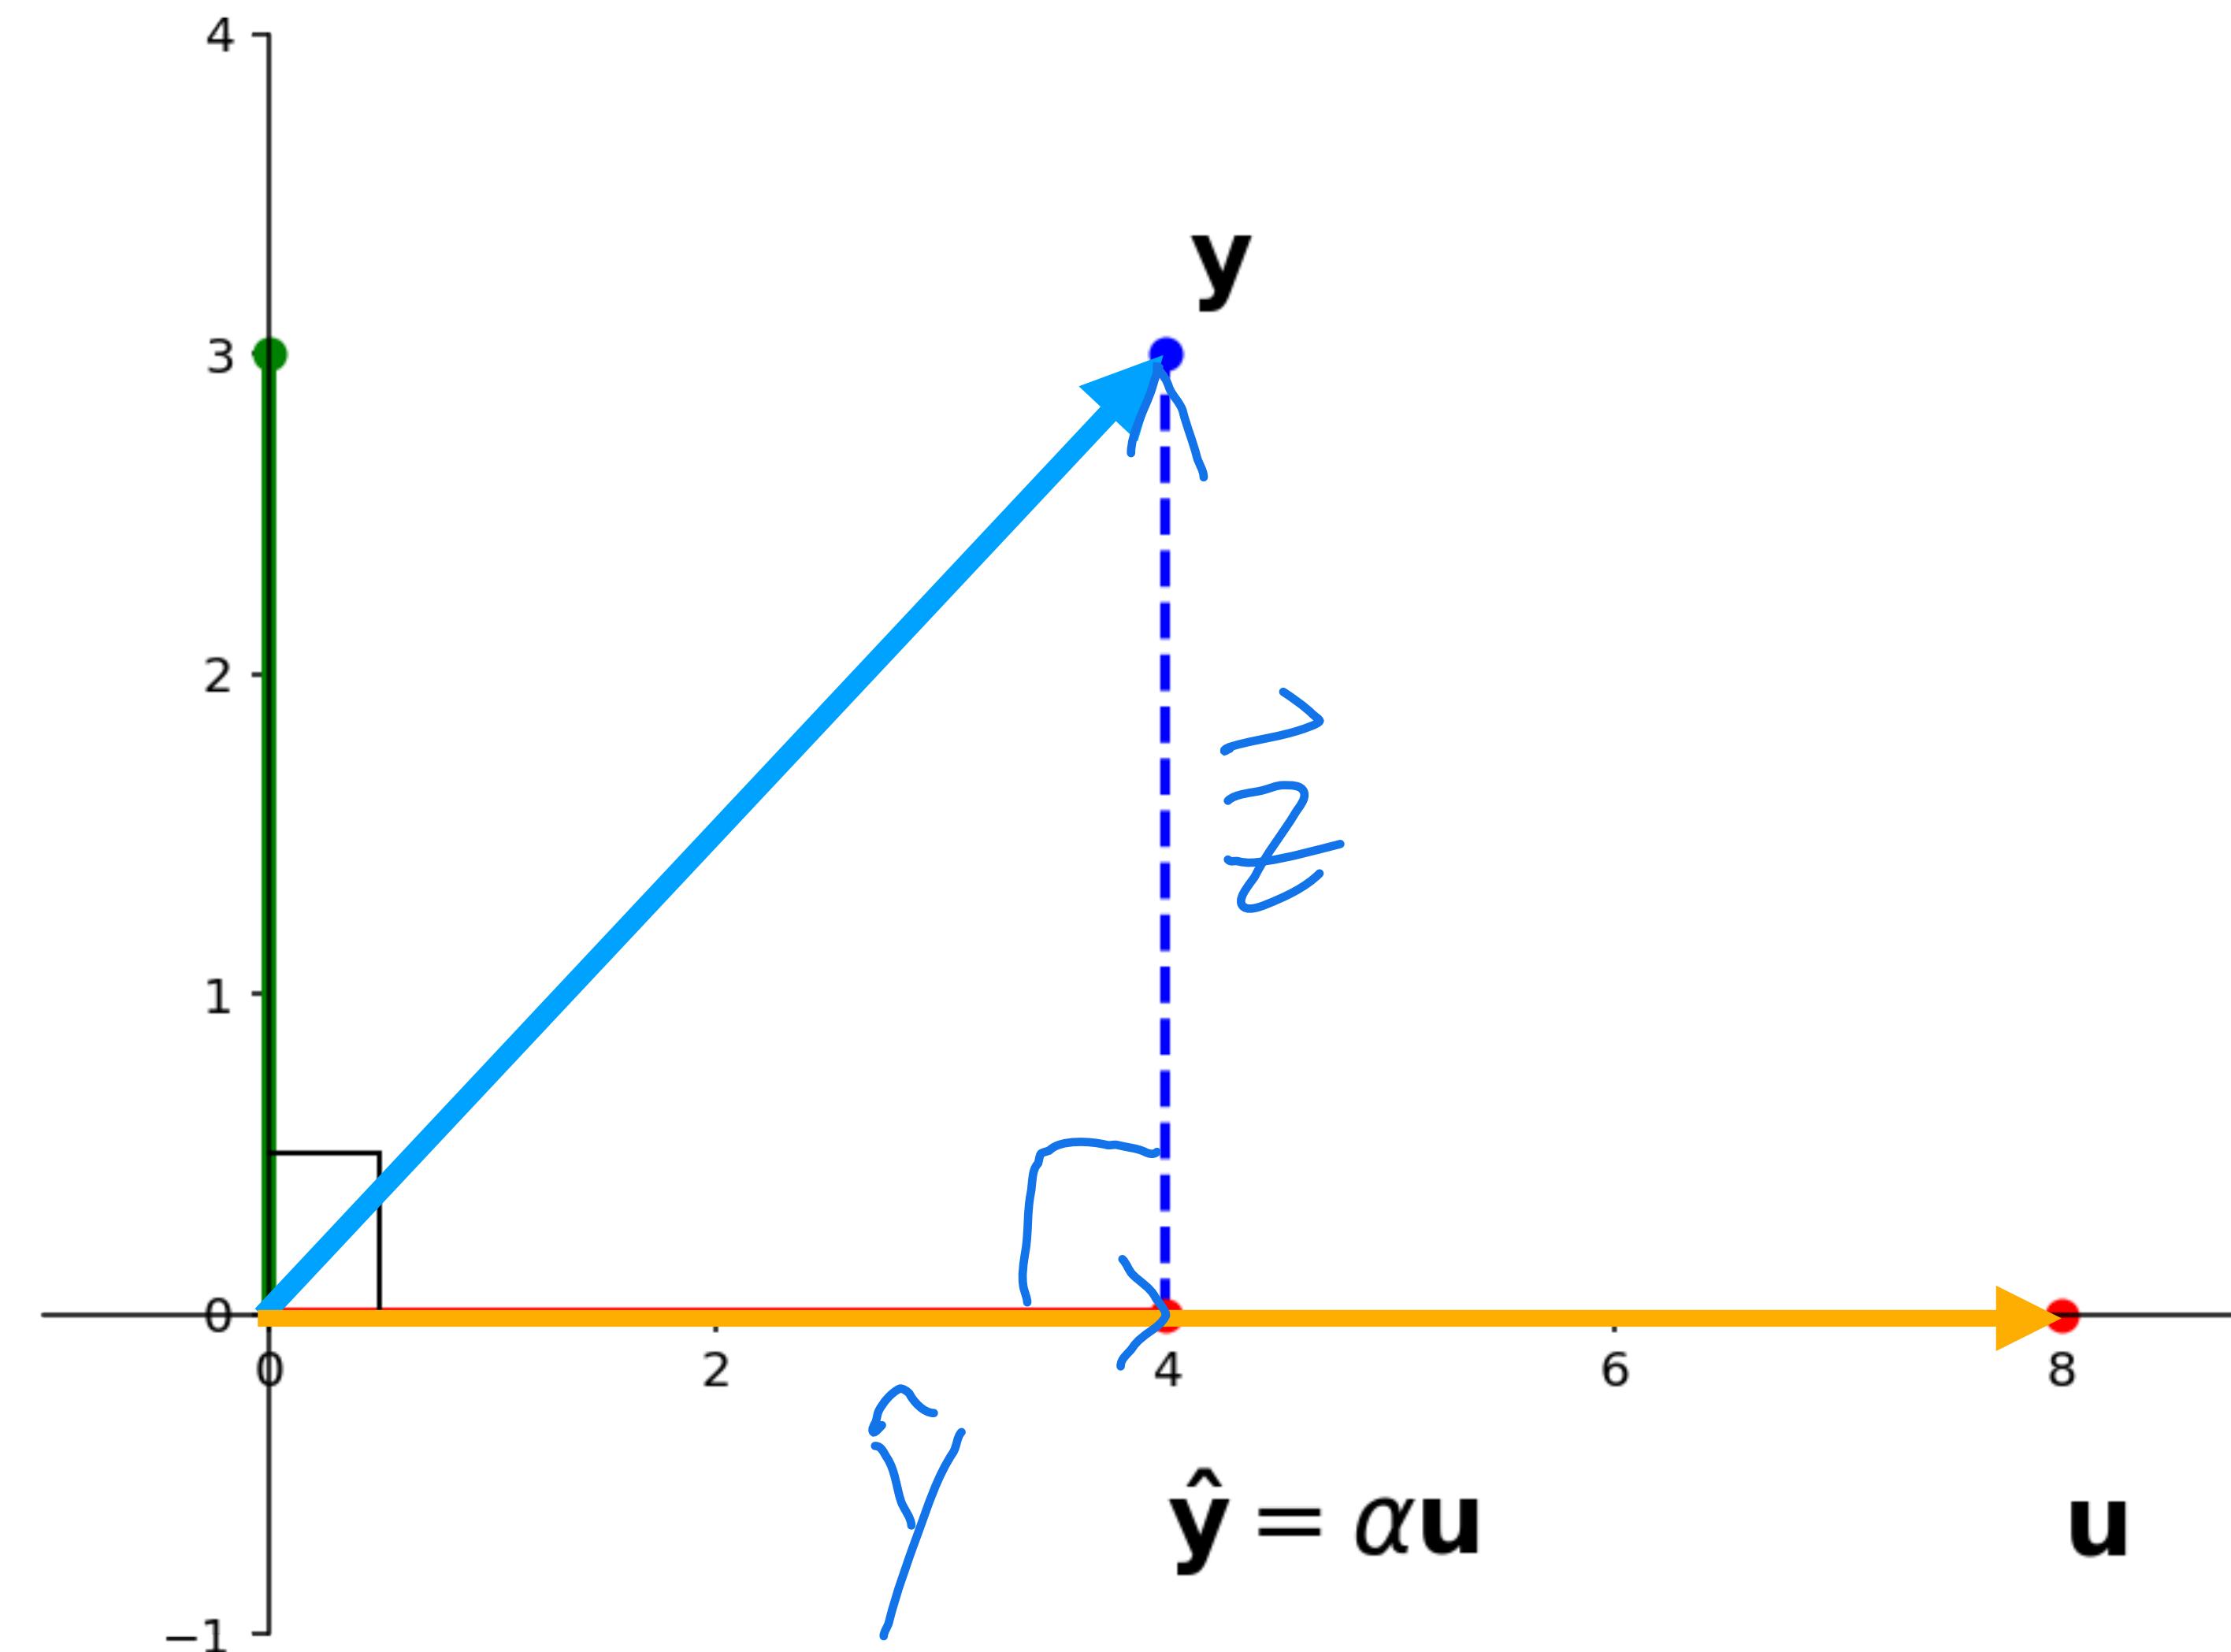
- » \mathbf{z} is orthogonal to \mathbf{u} (i.e., $\mathbf{z} \cdot \mathbf{u} = 0$)



Recall: Orthogonal Projection

Question. Given vectors \mathbf{y} and \mathbf{u} in R^n , find vectors $\hat{\mathbf{y}}$ and \mathbf{z} such that

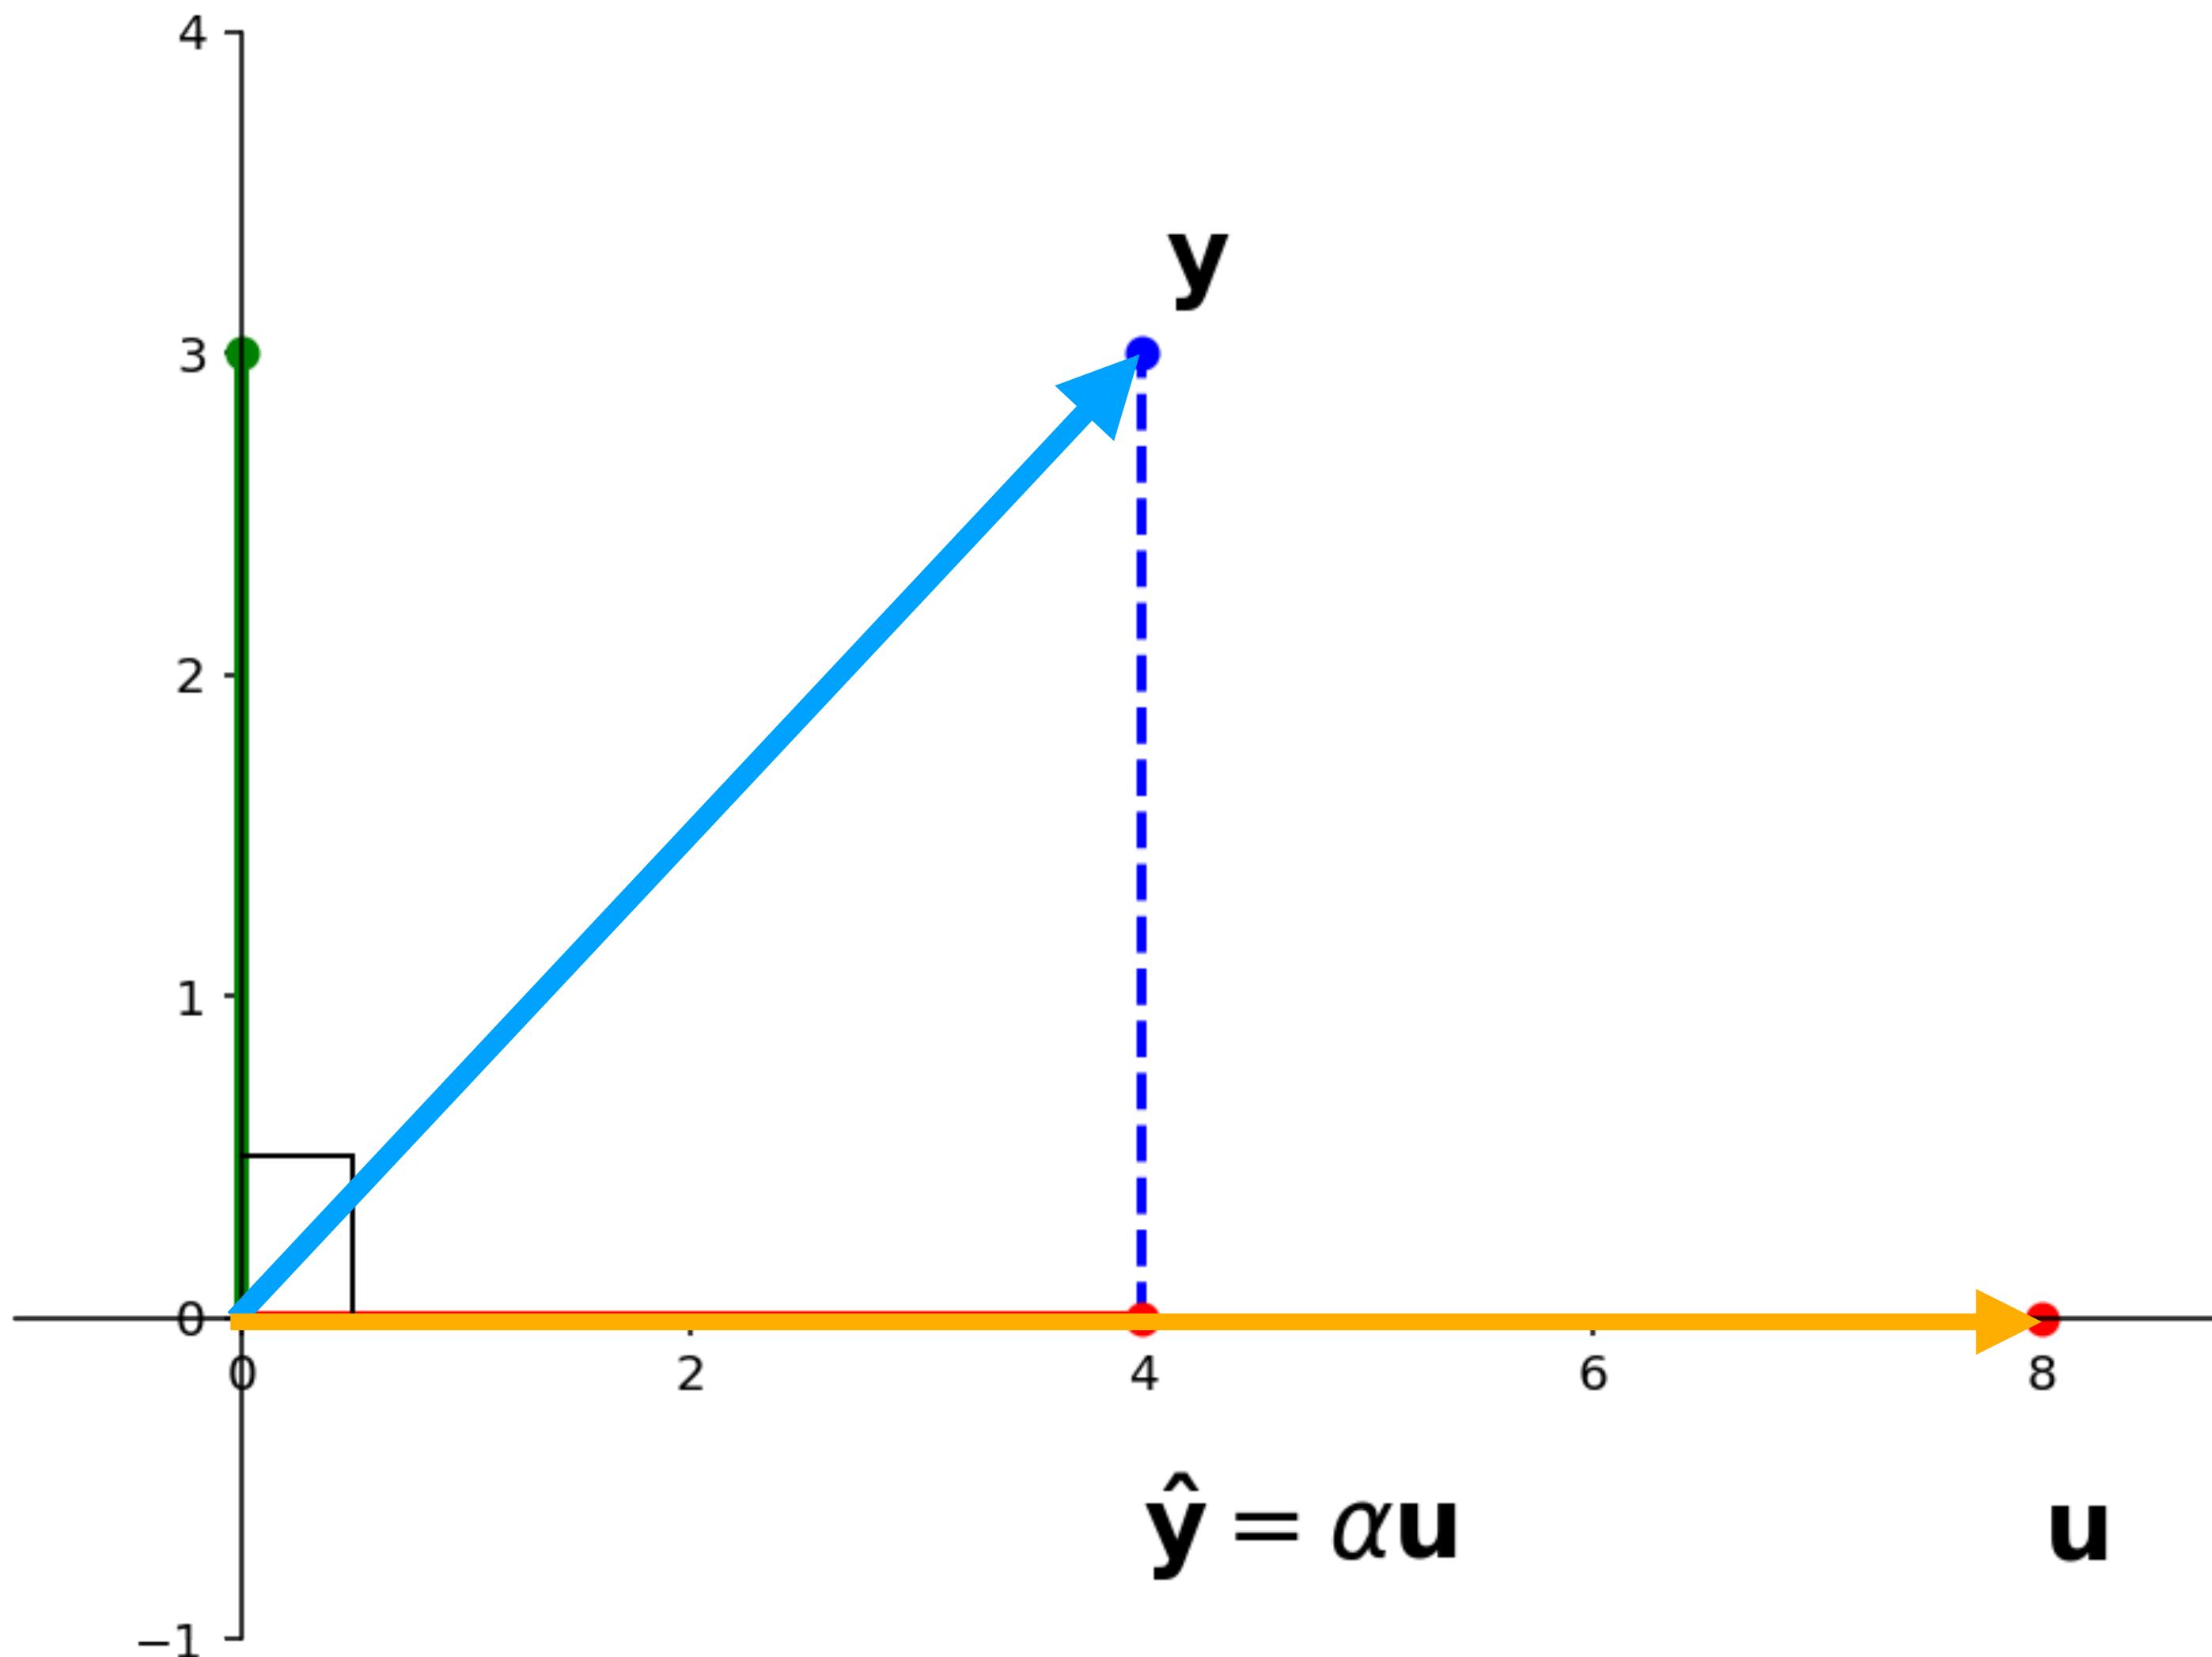
- » \mathbf{z} is orthogonal to \mathbf{u}
(i.e., $\mathbf{z} \cdot \mathbf{u} = 0$)
- » $\hat{\mathbf{y}} \in span\{\mathbf{u}\}$



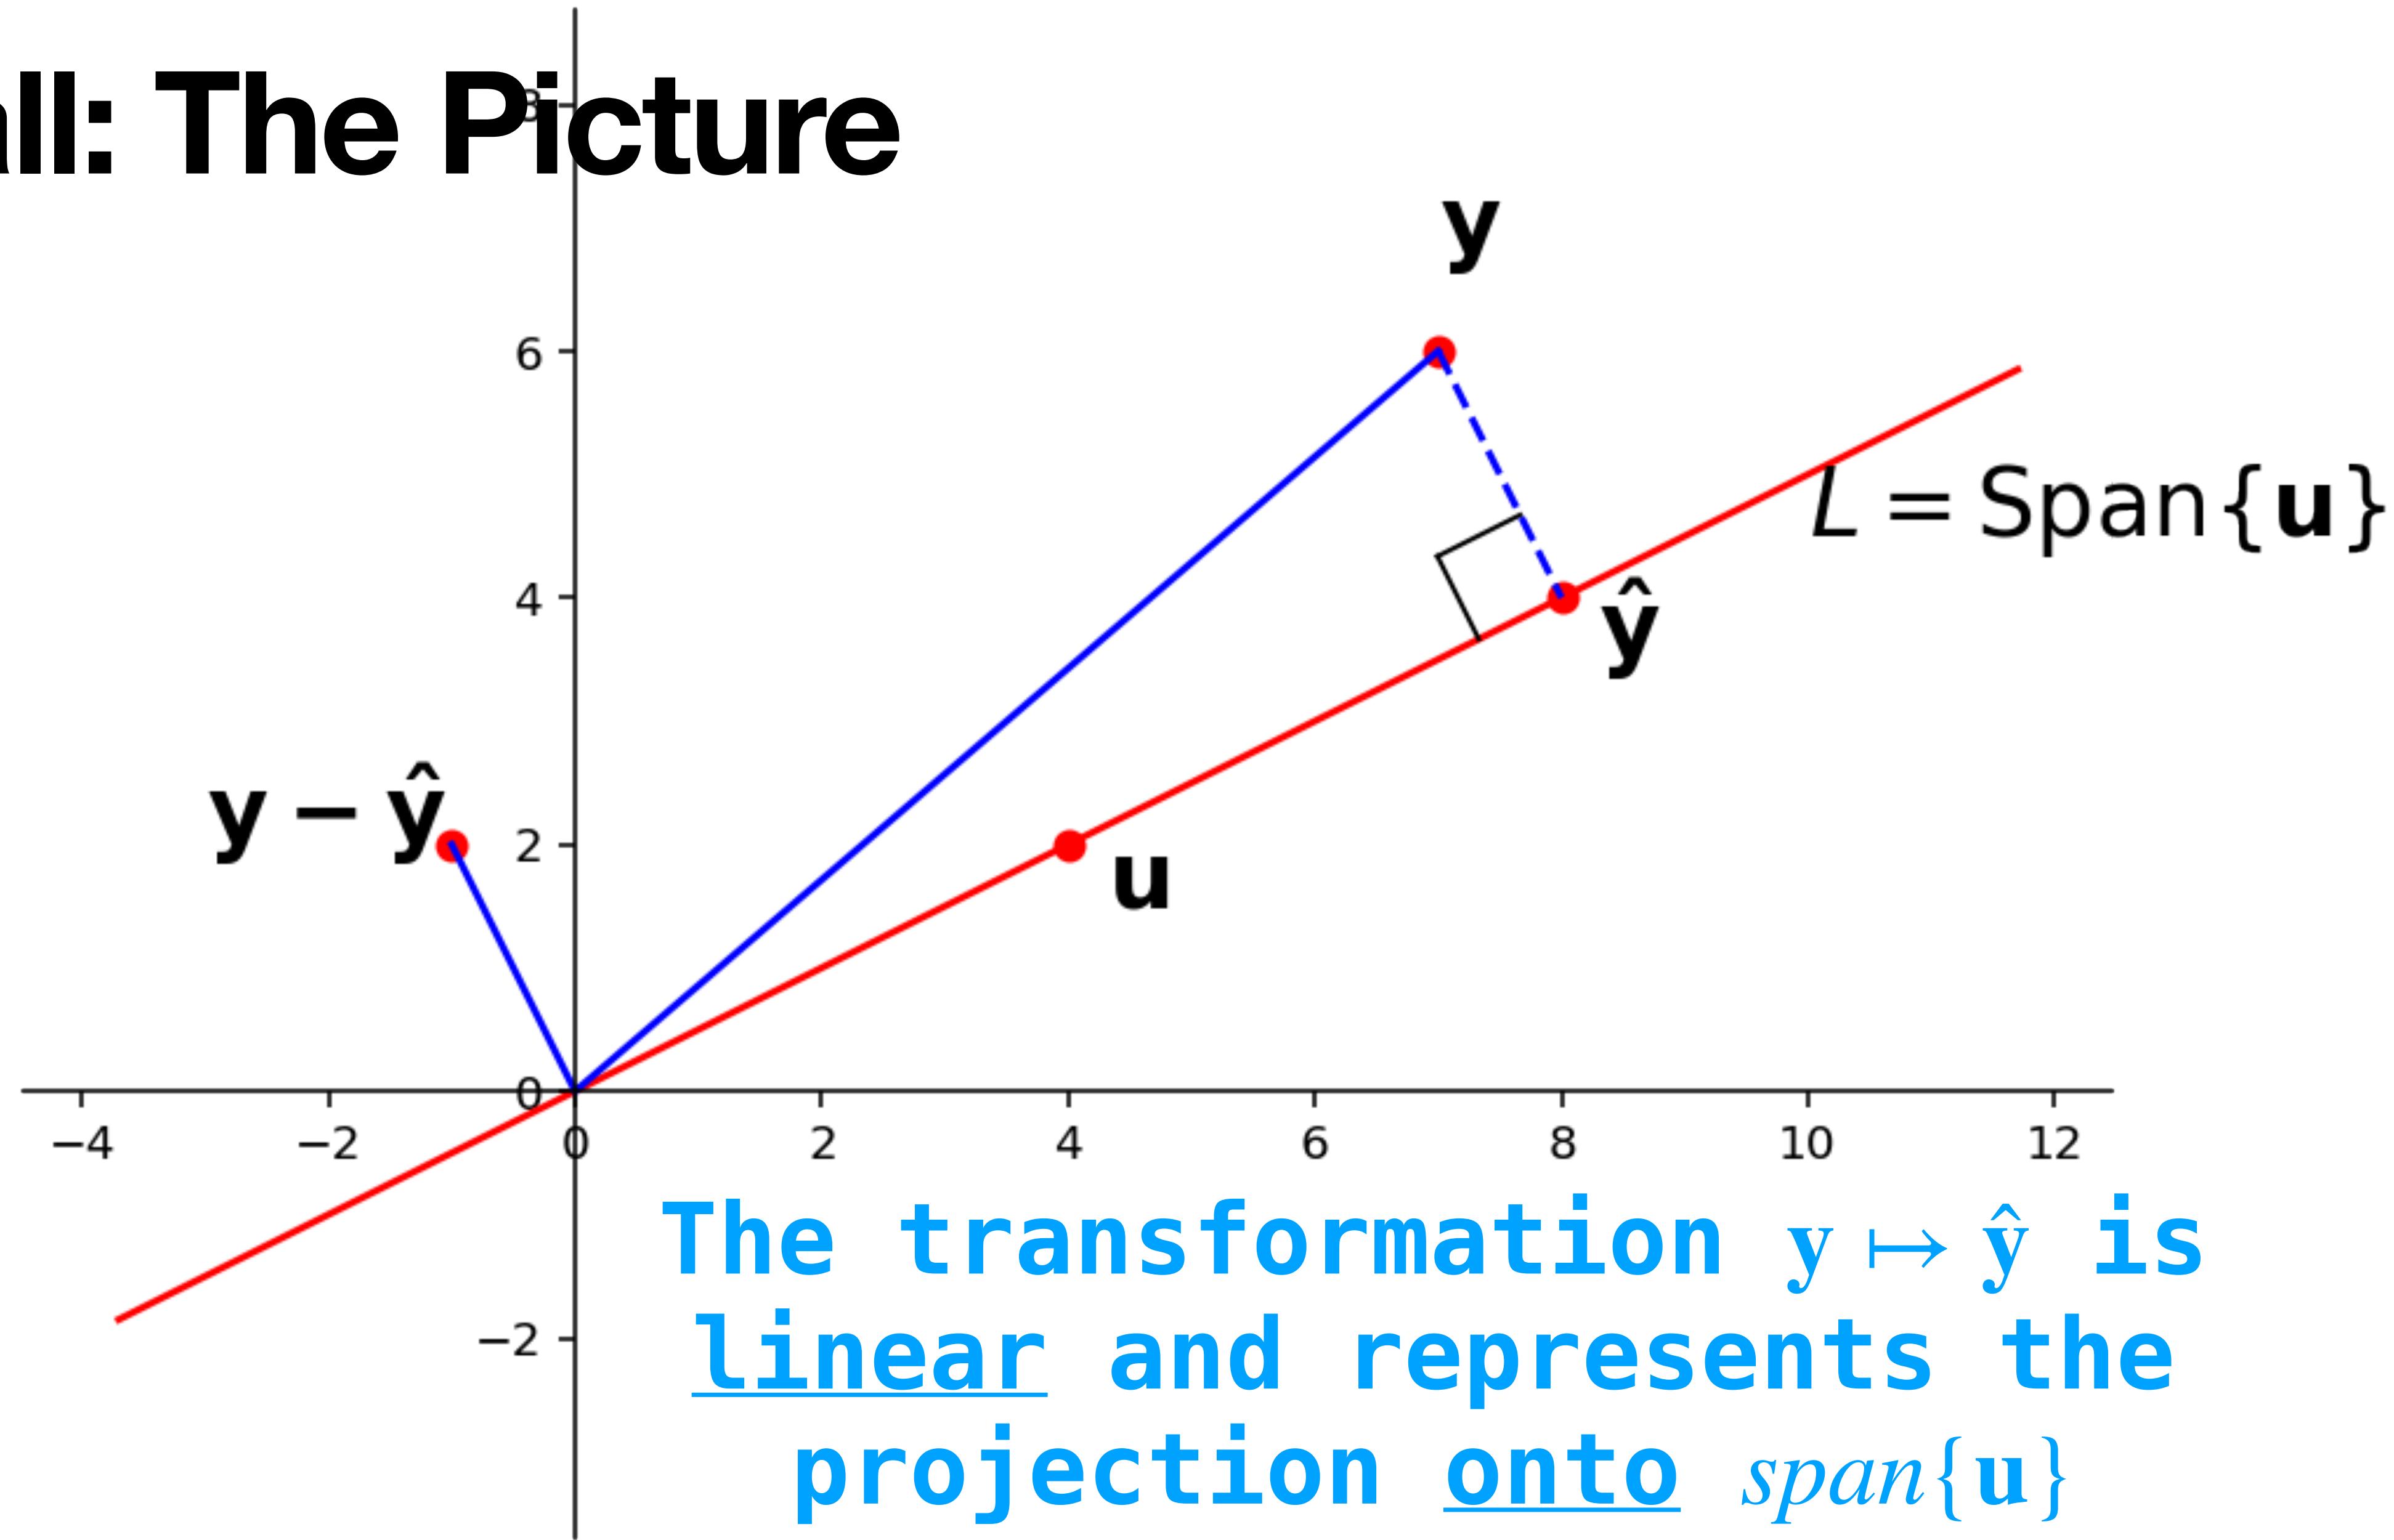
Recall: Orthogonal Projection

Question. Given vectors \mathbf{y} and \mathbf{u} in R^n , find vectors $\hat{\mathbf{y}}$ and \mathbf{z} such that

- » \mathbf{z} is orthogonal to \mathbf{u} (i.e., $\mathbf{z} \cdot \mathbf{u} = 0$)
- » $\hat{\mathbf{y}} \in span\{\mathbf{u}\}$
- » $\mathbf{y} = \hat{\mathbf{y}} + \mathbf{z}$



Recall: The Picture

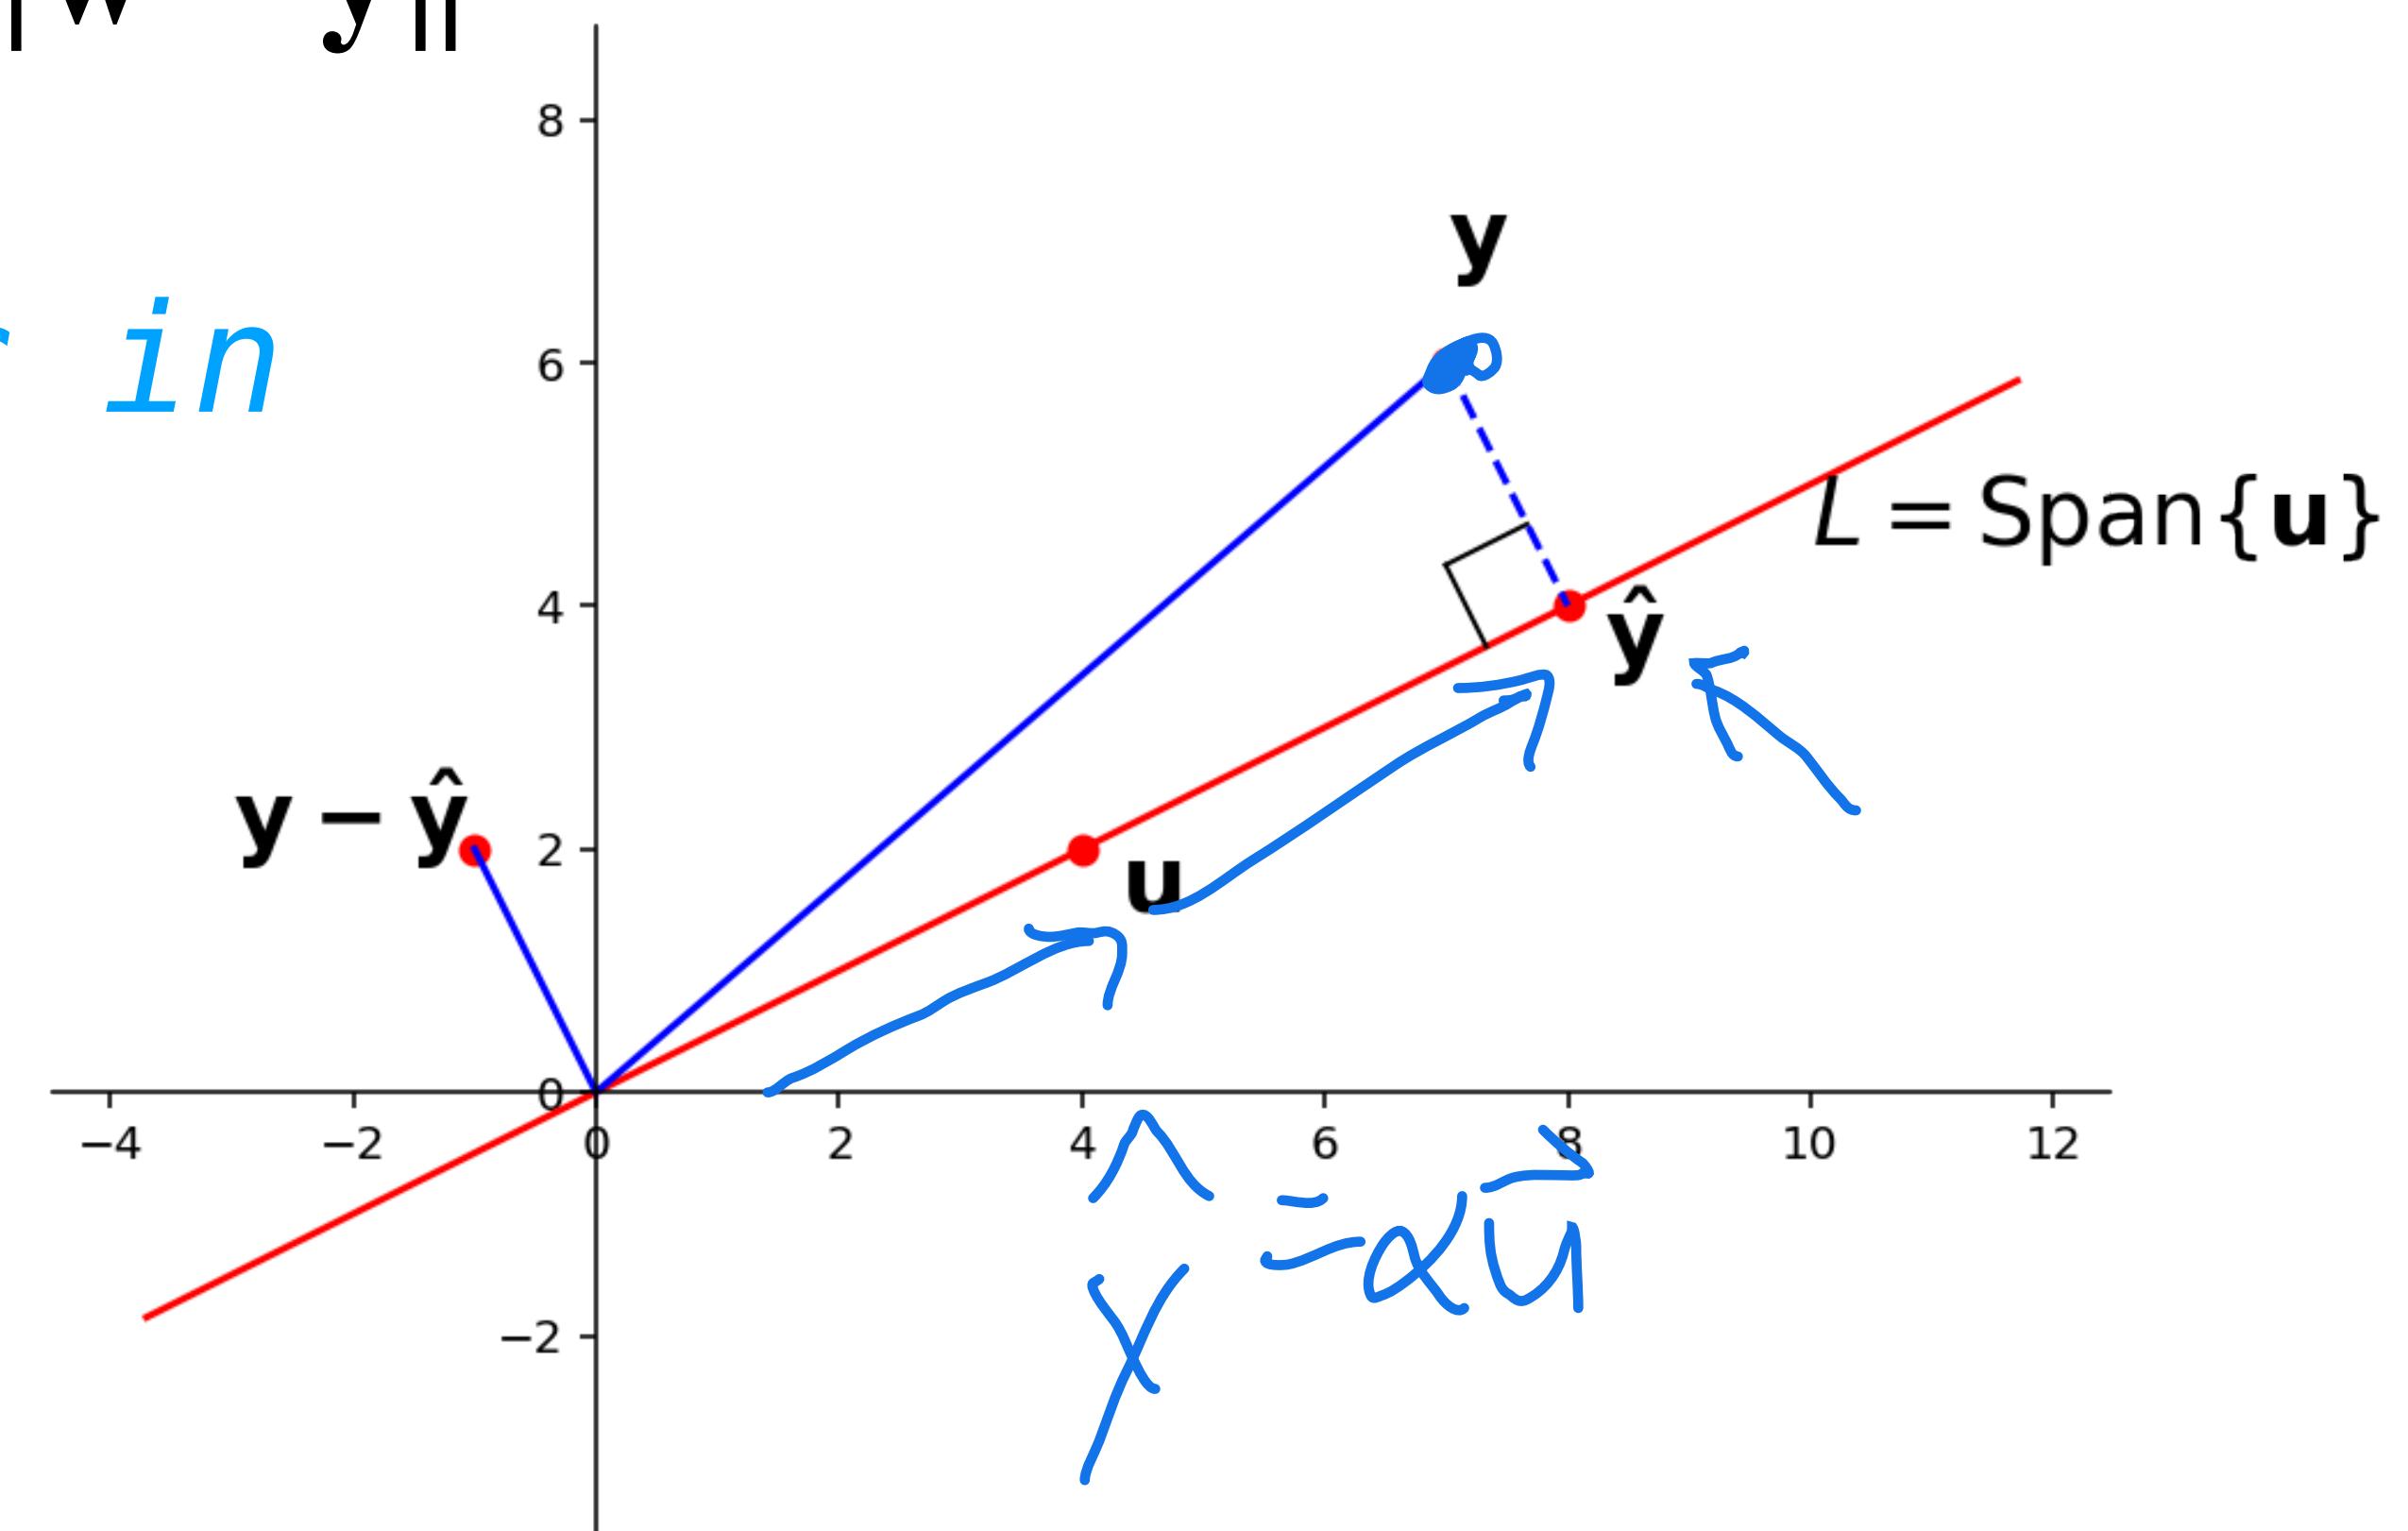


Recall: \hat{y} and Distance

Theorem. $\|\hat{y} - y\| = \min_{w \in \text{span}\{\mathbf{u}\}} \|w - y\|$

\hat{y} is the closest vector in $\text{span}\{\mathbf{u}\}$ to y

"Proof" by inspection:



The Equational Perspective

The Equational Perspective

We know the equation $x\mathbf{u} = \mathbf{y}$ may have no solution

The Equational Perspective

We know the equation $\underline{x}\mathbf{u} = \mathbf{y}$ may have no solution

Question. Find a value α such that $\alpha\mathbf{u}$ is as close as possible to \mathbf{y}

The Equational Perspective

We know the equation $x\mathbf{u} = \mathbf{y}$ may have no solution

Question. Find a value α such that $\alpha\mathbf{u}$ is as close as possible to \mathbf{y}

That is, the distance $dist(\mathbf{y}, \alpha\mathbf{u}) = \underline{\|\mathbf{y} - \alpha\mathbf{u}\|}$ is as small as possible

The Equational Perspective

We know the equation $\hat{A}xu = y$ may have no solution

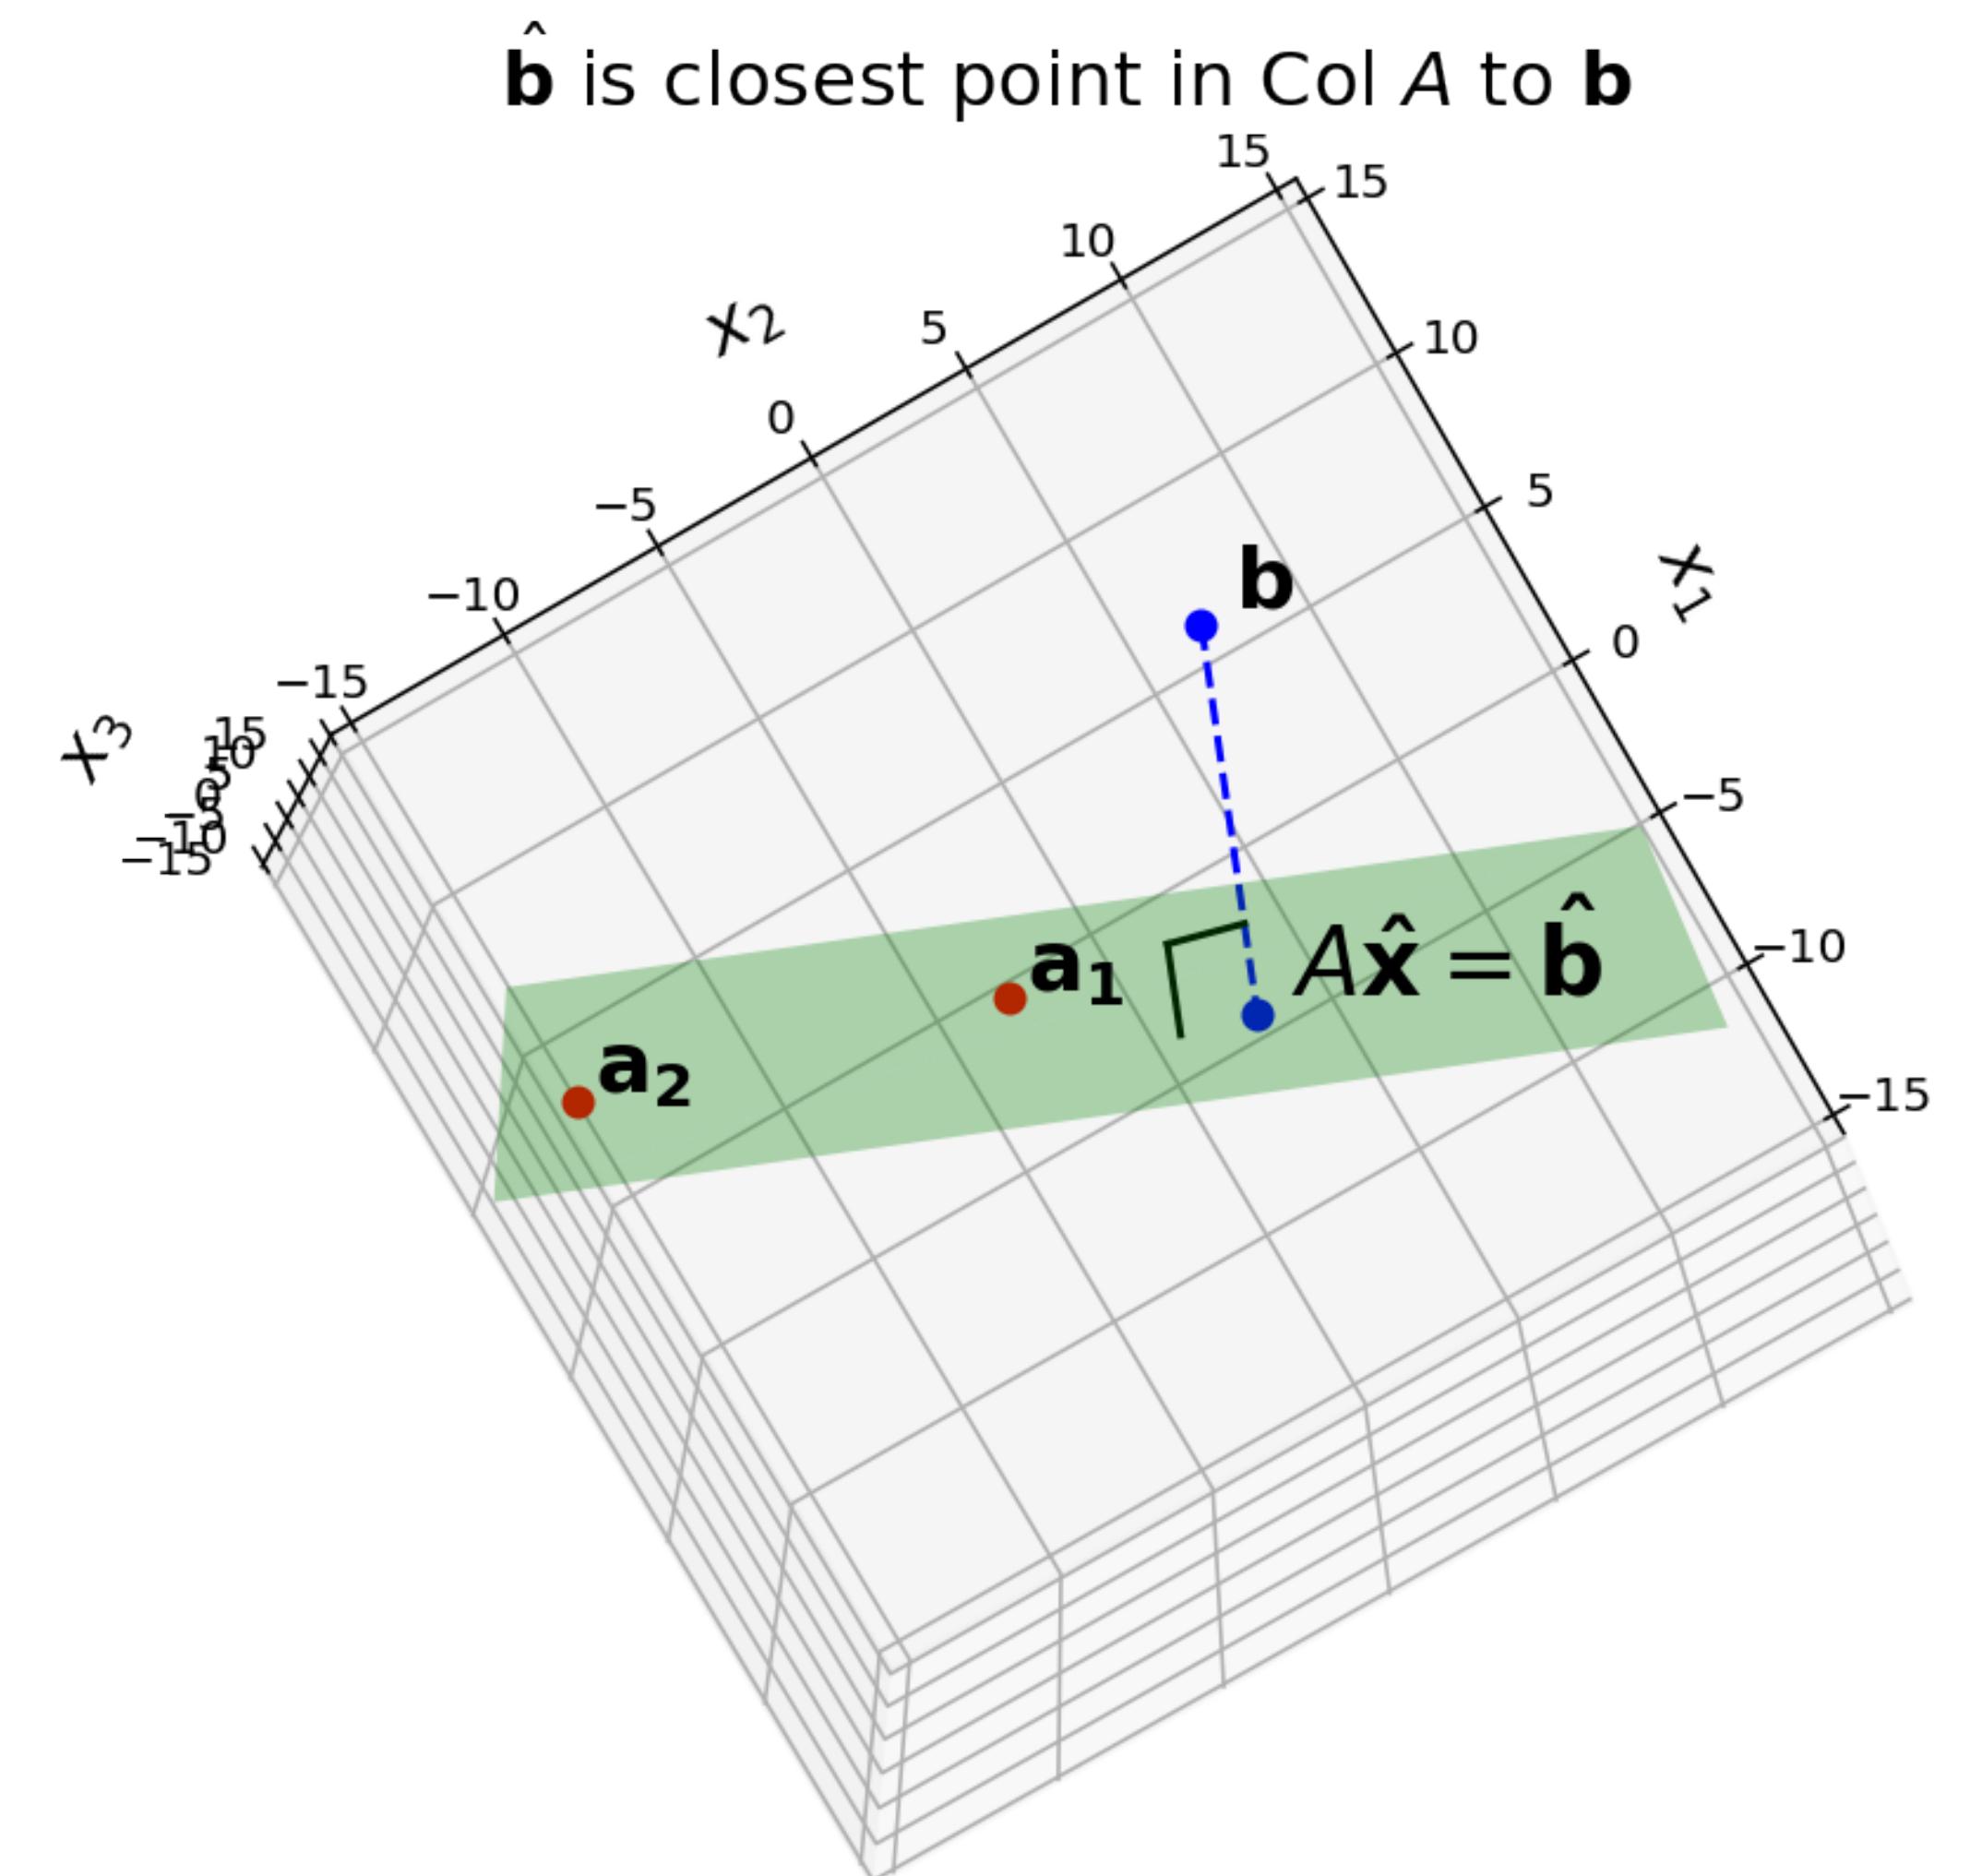
Question. Find a value α such that αu is as close as possible to y

That is, the distance $dist(y, \alpha u) = \|y - \alpha u\|$ is as small as possible

We need to generalize this to arbitrary matrix equations

The General Least Squares Problem

Figure 22.8

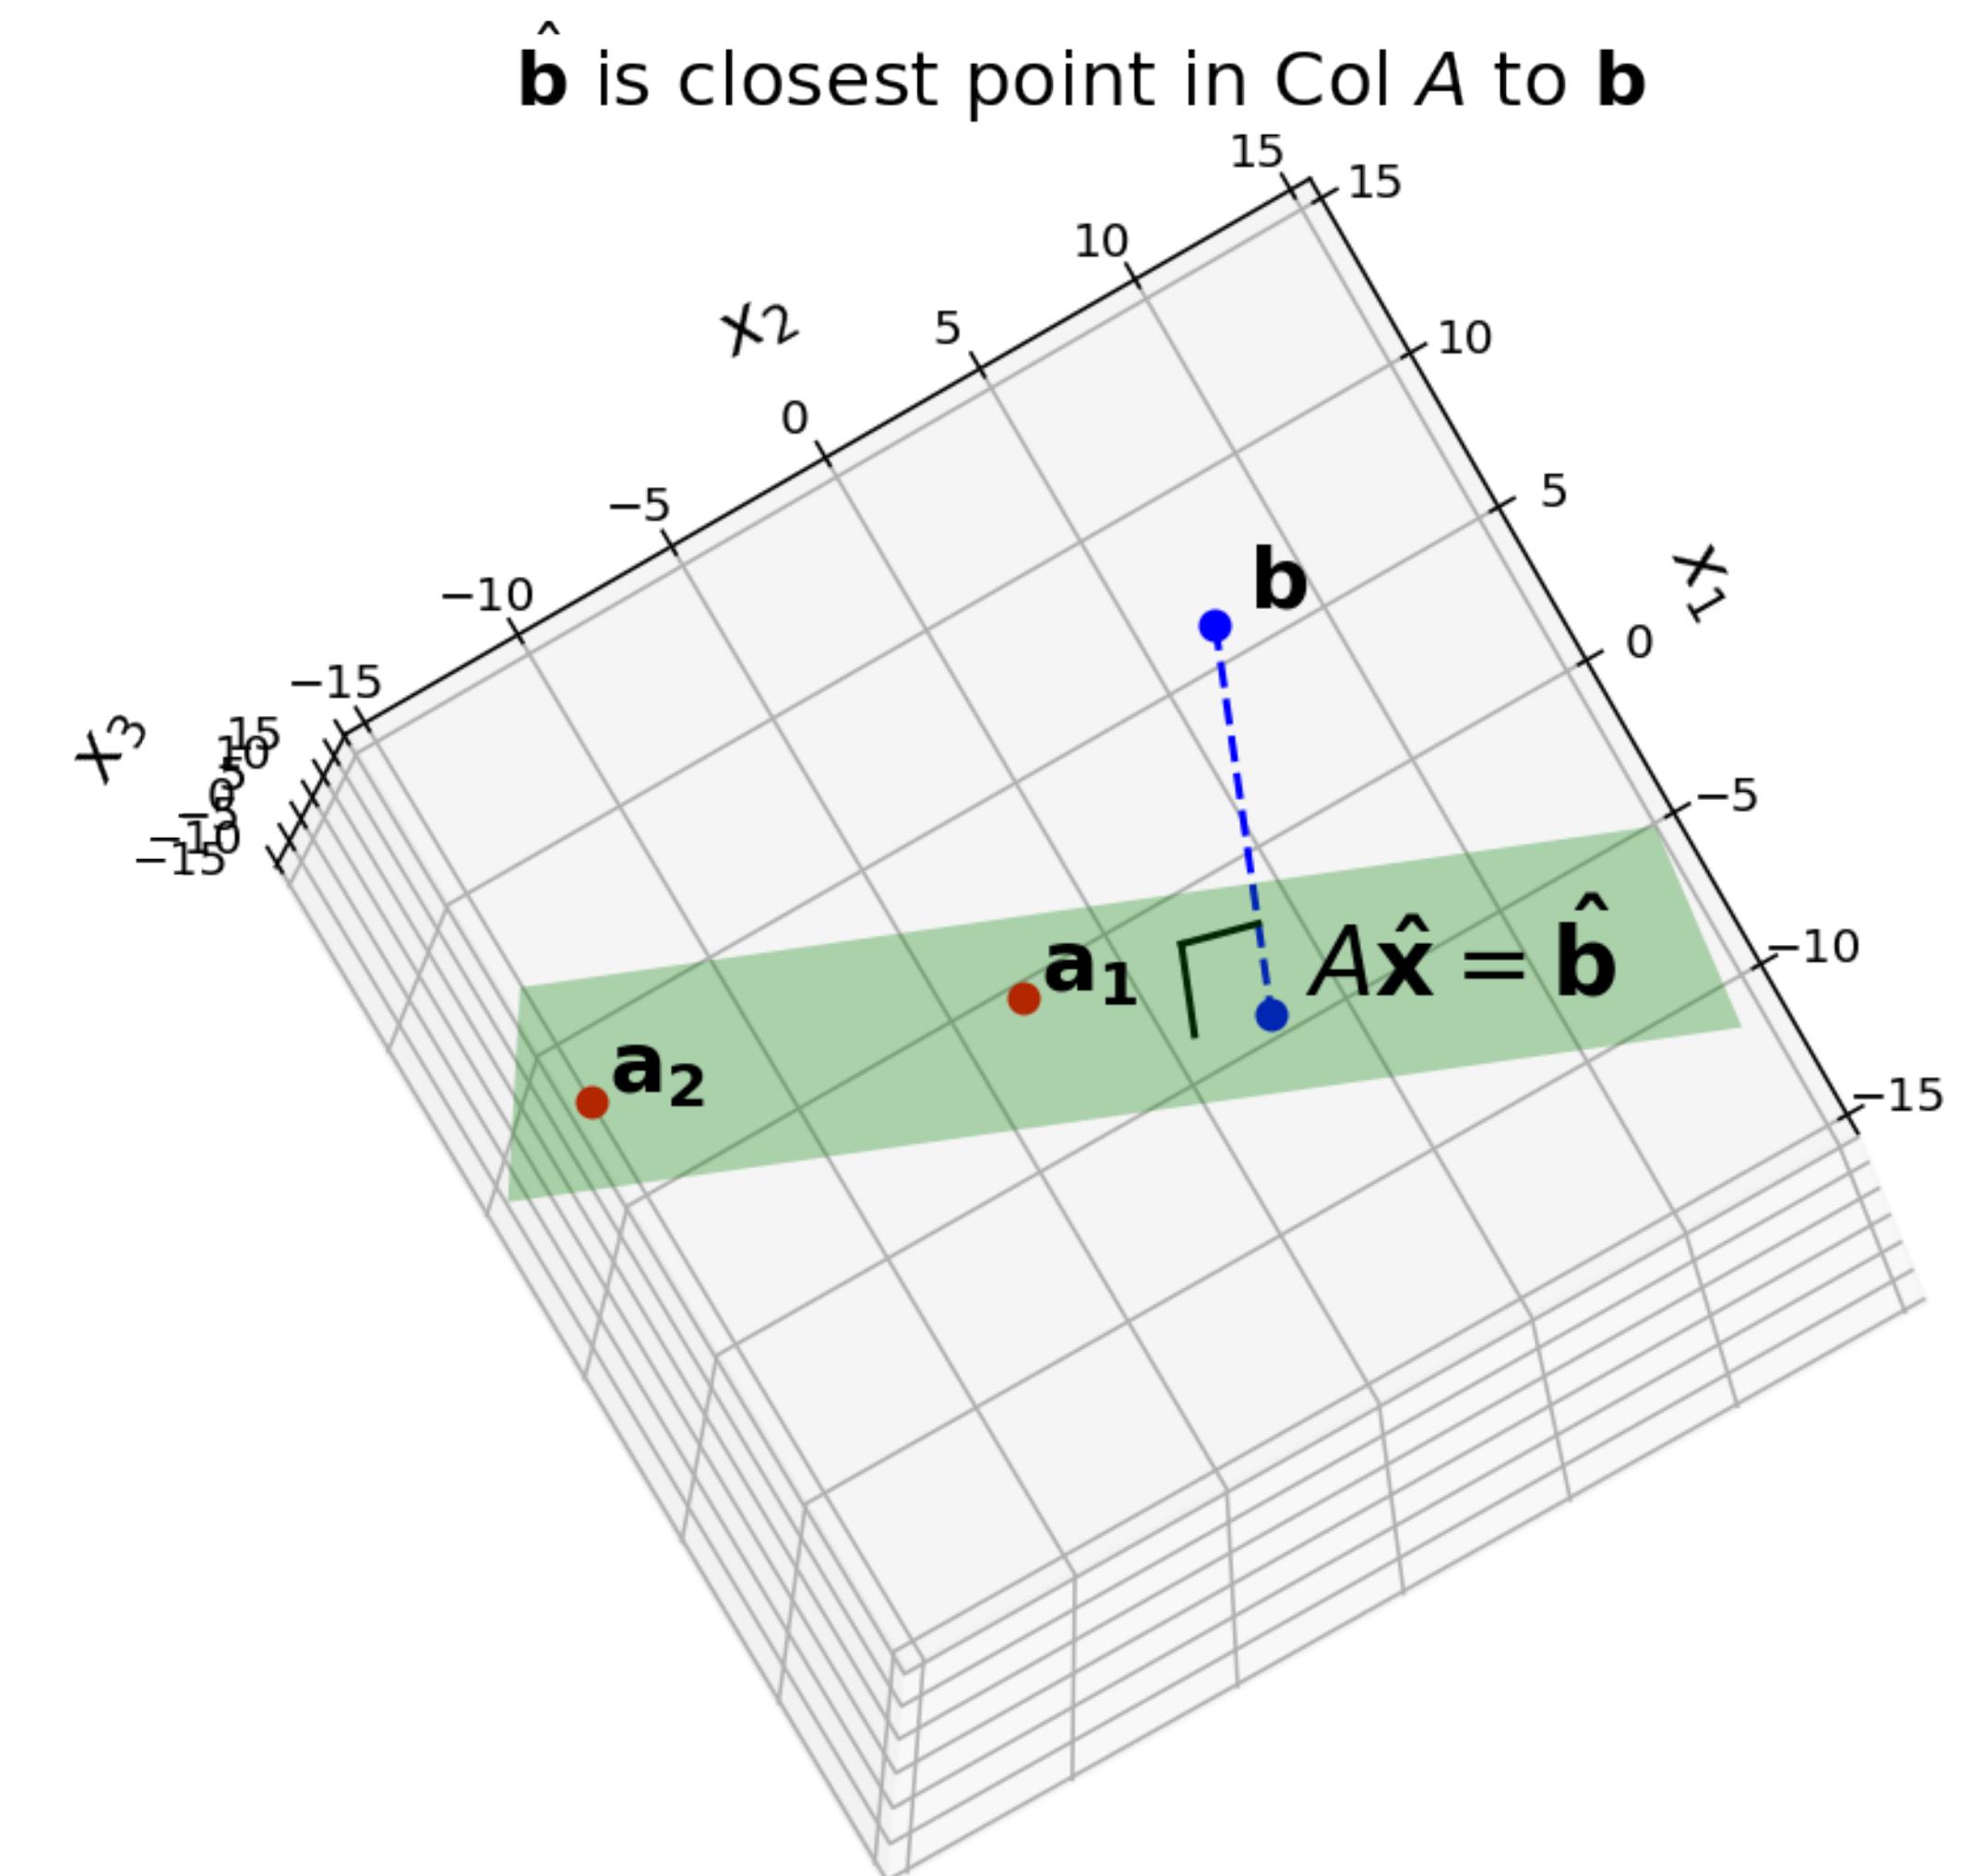


The General Least Squares Problem

Figure 22.8

Problem. Given a $m \times n$ matrix A and a vector b from \mathbb{R}^m , find a vector x in \mathbb{R}^n which minimizes

$$dist(Ax, b) = \|Ax - b\|$$



The General Least Squares Problem

Figure 22.8

Problem. Given a $m \times n$ matrix A and a vector b from \mathbb{R}^m , find a vector x in \mathbb{R}^n which minimizes

$$dist(Ax, b) = \|Ax - b\|$$

Find a vector x which makes $\|Ax - b\|$ as small as possible

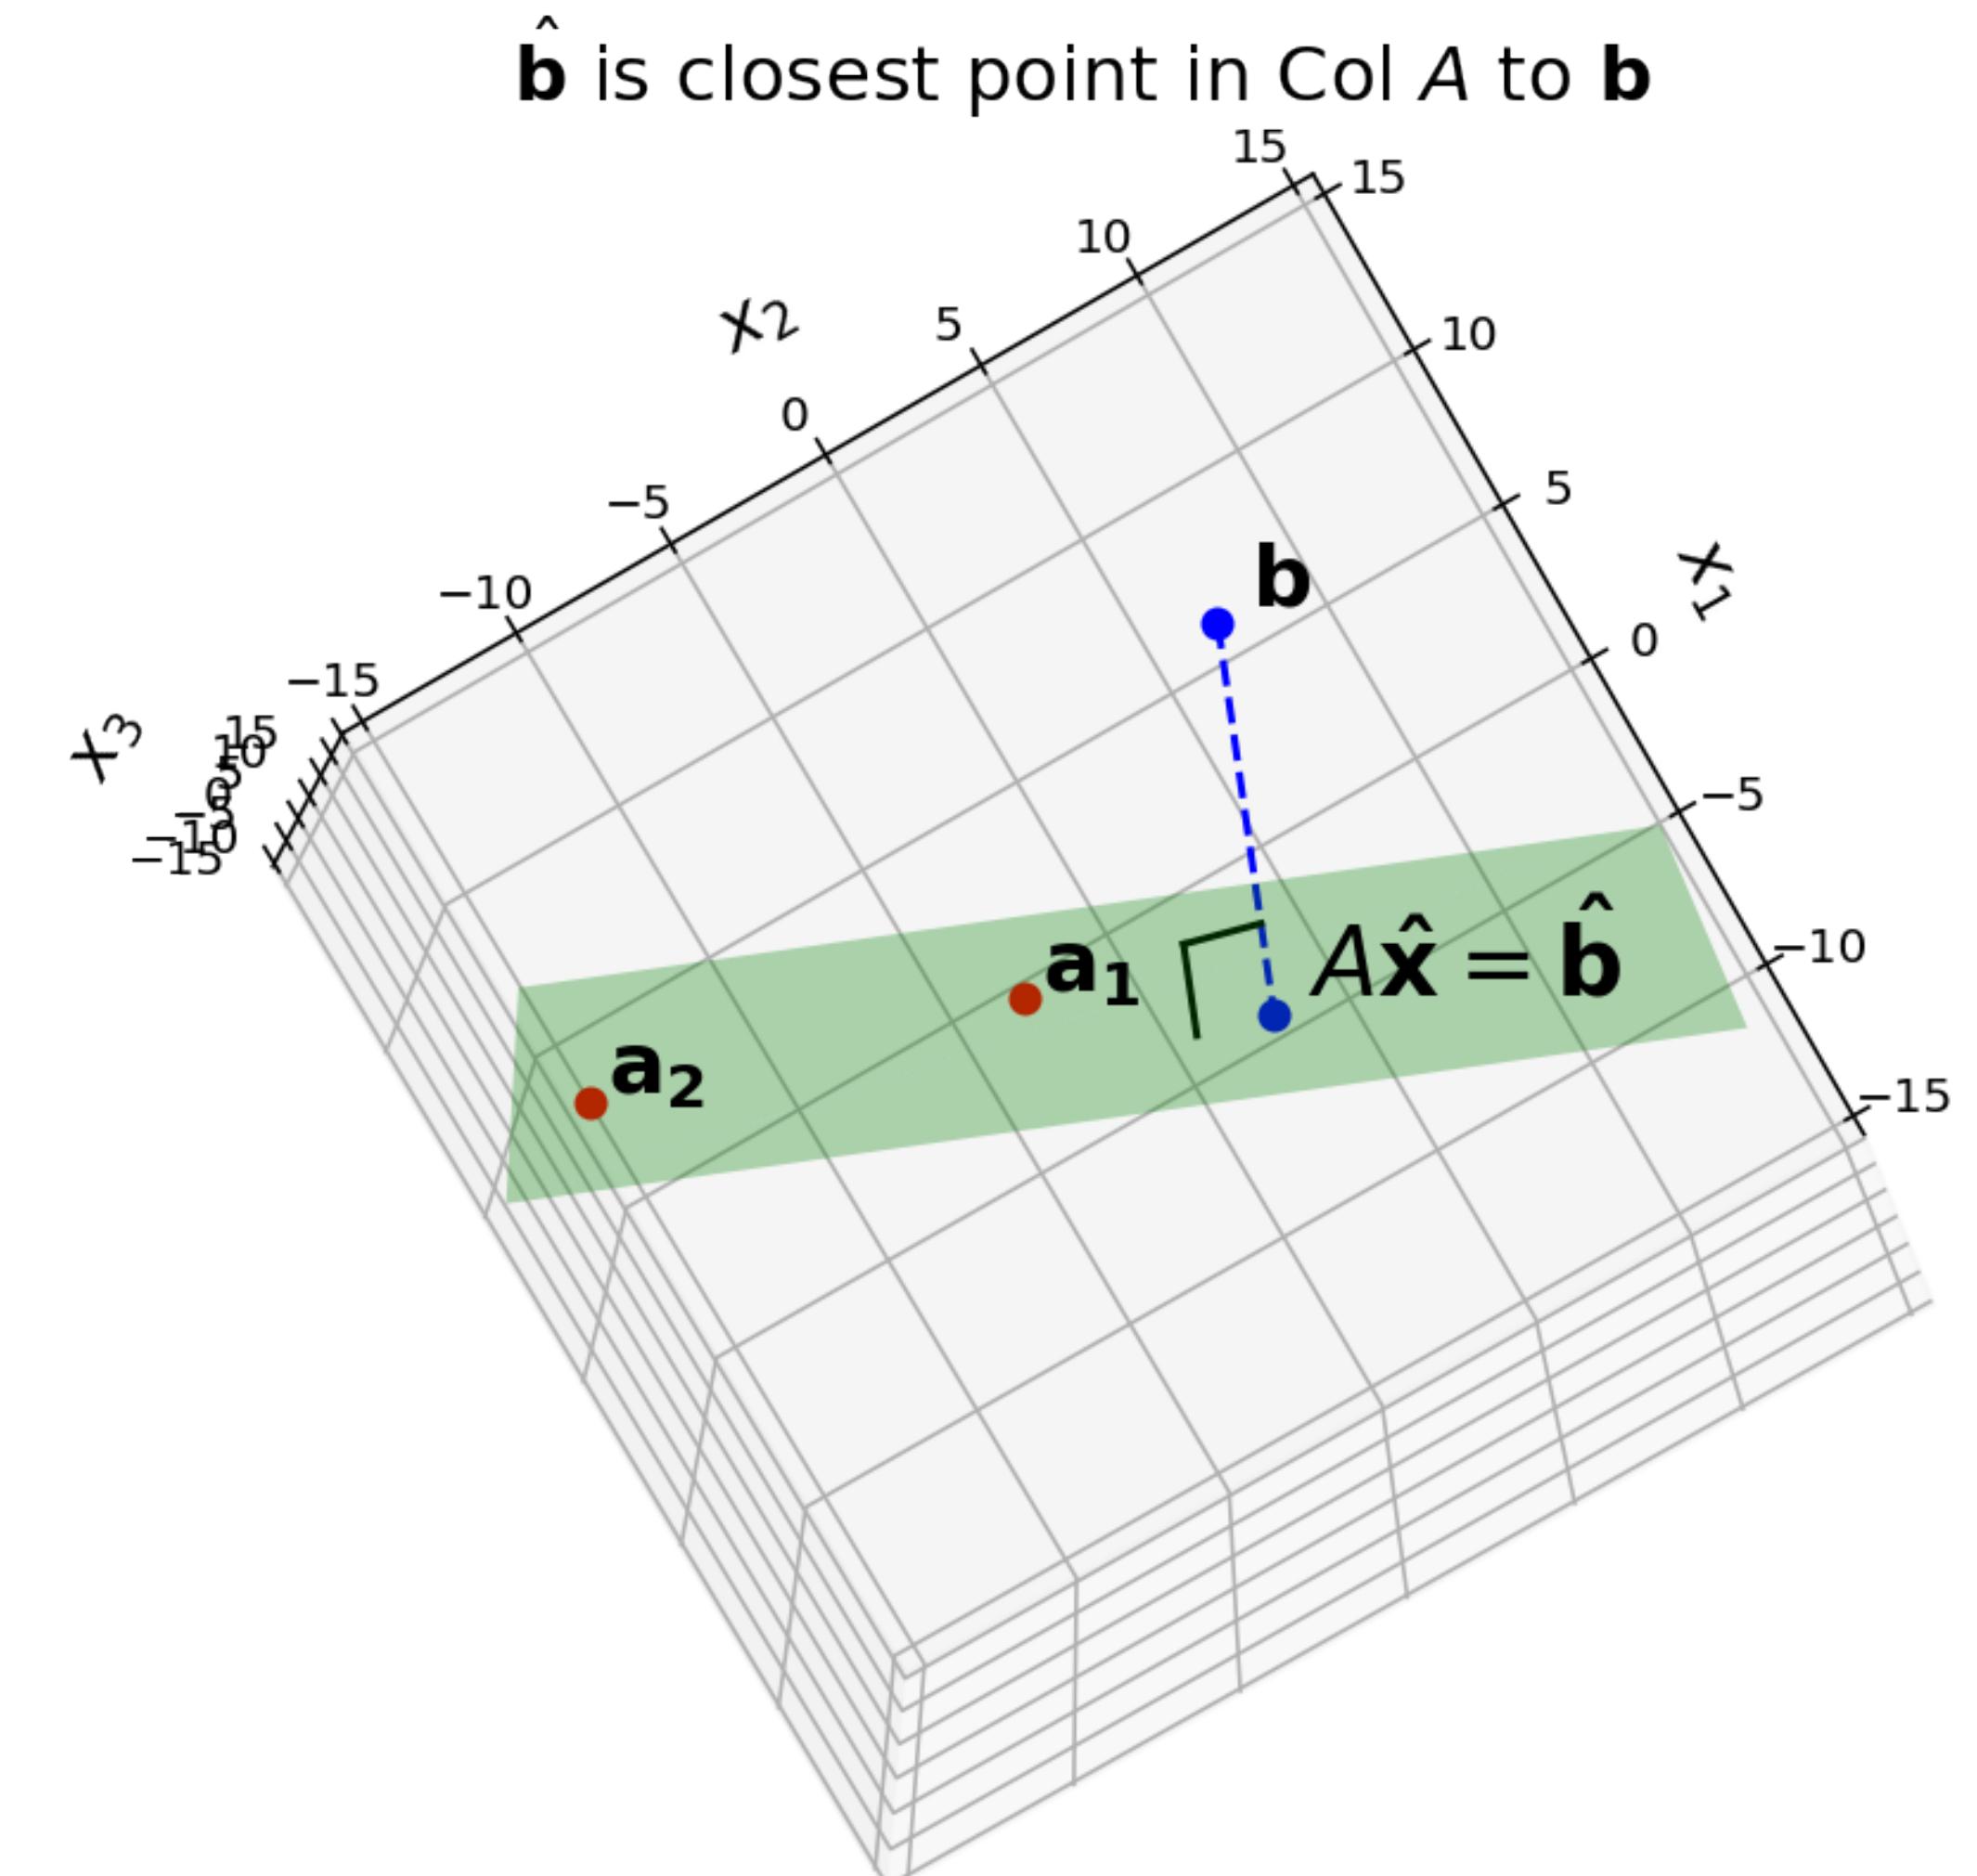
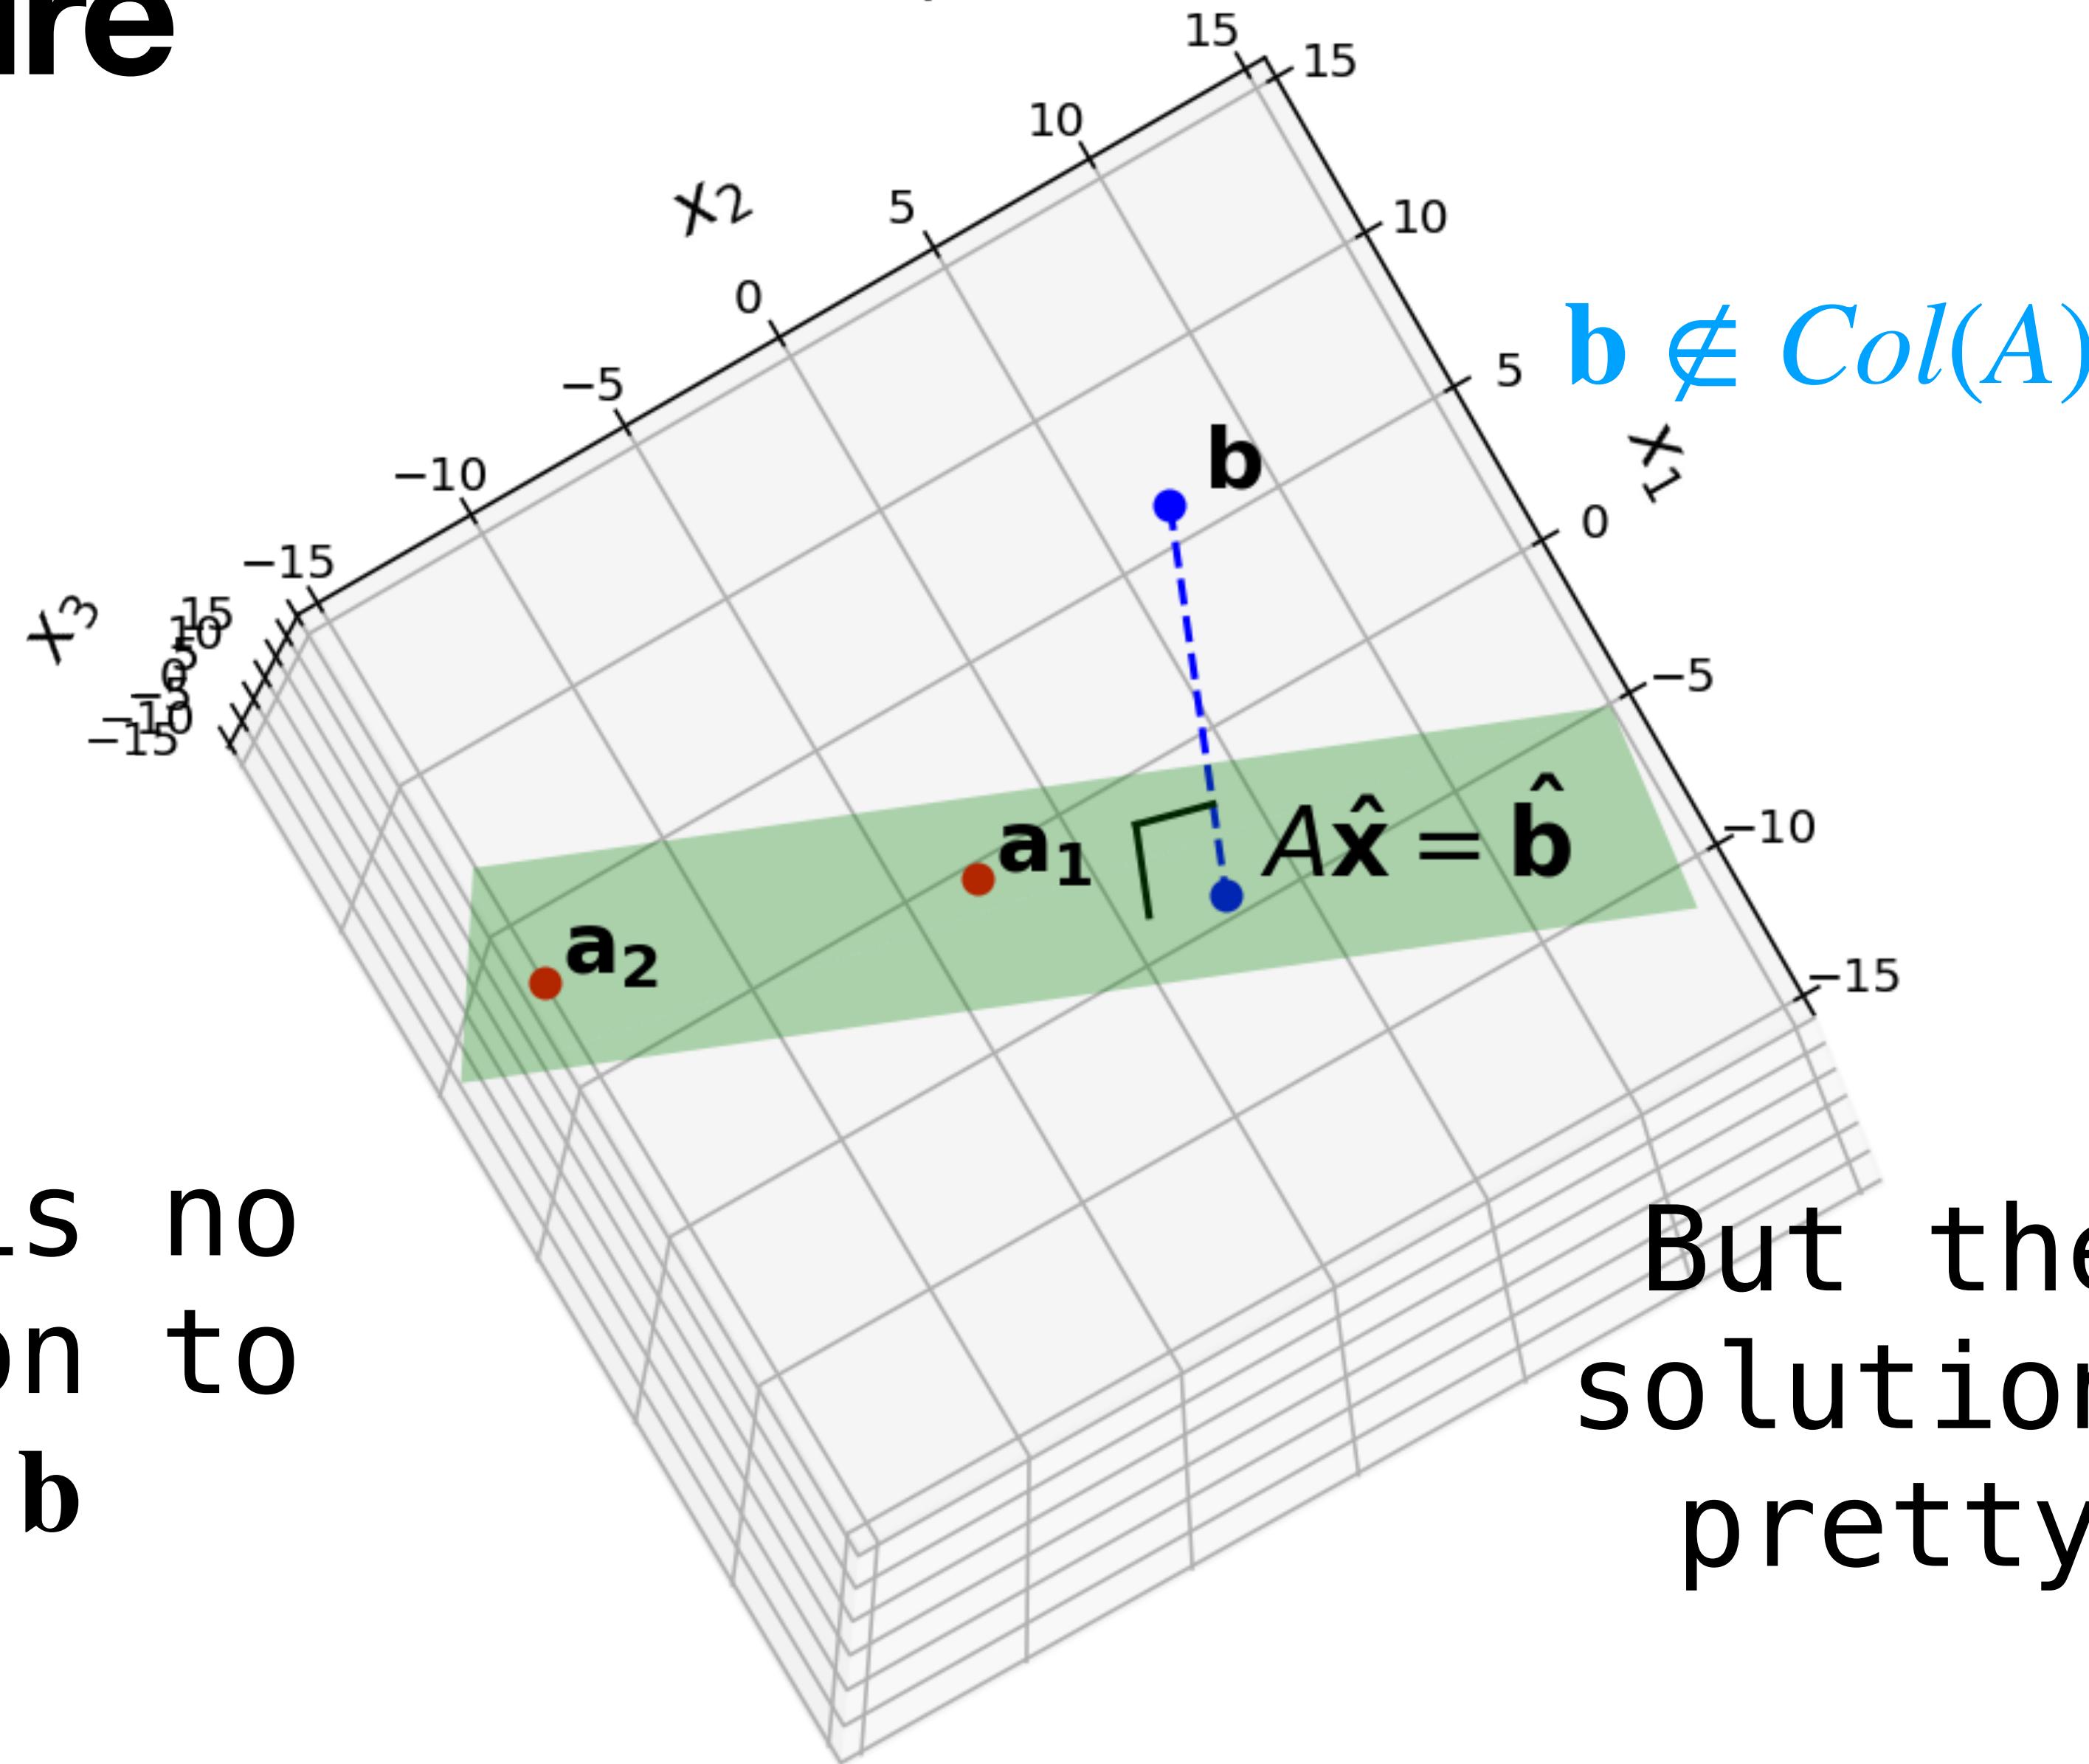


Figure 22.8

The Picture

$\hat{\mathbf{b}}$ is closest point in Col A to \mathbf{b}



There is no
solution to
 $Ax = b$

But there's a
solution that's
pretty close

Sum of Squares

$$\langle \tilde{A}\tilde{x} - \tilde{b}, \tilde{A}\tilde{x} - \tilde{b} \rangle = \|\tilde{A}\tilde{x} - \tilde{b}\|^2$$

\Downarrow

$$\|A\mathbf{x} - \mathbf{b}\|^2 = \sum_{i=1}^n ((A\mathbf{x})_i - \mathbf{b}_i)^2$$

Sum of Squares

$$\|A\mathbf{x} - \mathbf{b}\|^2 = \sum_{i=1}^n ((A\mathbf{x})_i - \mathbf{b}_i)^2$$

It is equivalent to minimize $\|A\mathbf{x} - \mathbf{b}\|^2$, which can be viewed as a **sum of squares**

Sum of Squares

$$\|A\mathbf{x} - \mathbf{b}\|^2 = \sum_{i=1}^n ((A\mathbf{x})_i - \mathbf{b}_i)^2$$

It is equivalent to minimize $\|A\mathbf{x} - \mathbf{b}\|^2$, which can be viewed as a **sum of squares**

These things come up everywhere

Sum of Squares

$$\|Ax - b\|^2 = \sum_{i=1}^n ((Ax)_i - b_i)^2$$

It is equivalent to minimize $\|Ax - b\|^2$, which can be viewed as a **sum of squares**

These things come up everywhere

(Advanced.) This error is everywhere differentiable, whereas $\sum_{i=1}^n |(Ax)_i - b_i|$ is not

L¹ error

Least Squares Solution

Definition. Given a $m \times n$ matrix A and a vector b in \mathbb{R}^m , a **least squares solution** of $Ax = b$ is a vector \hat{x} from \mathbb{R}^n such that

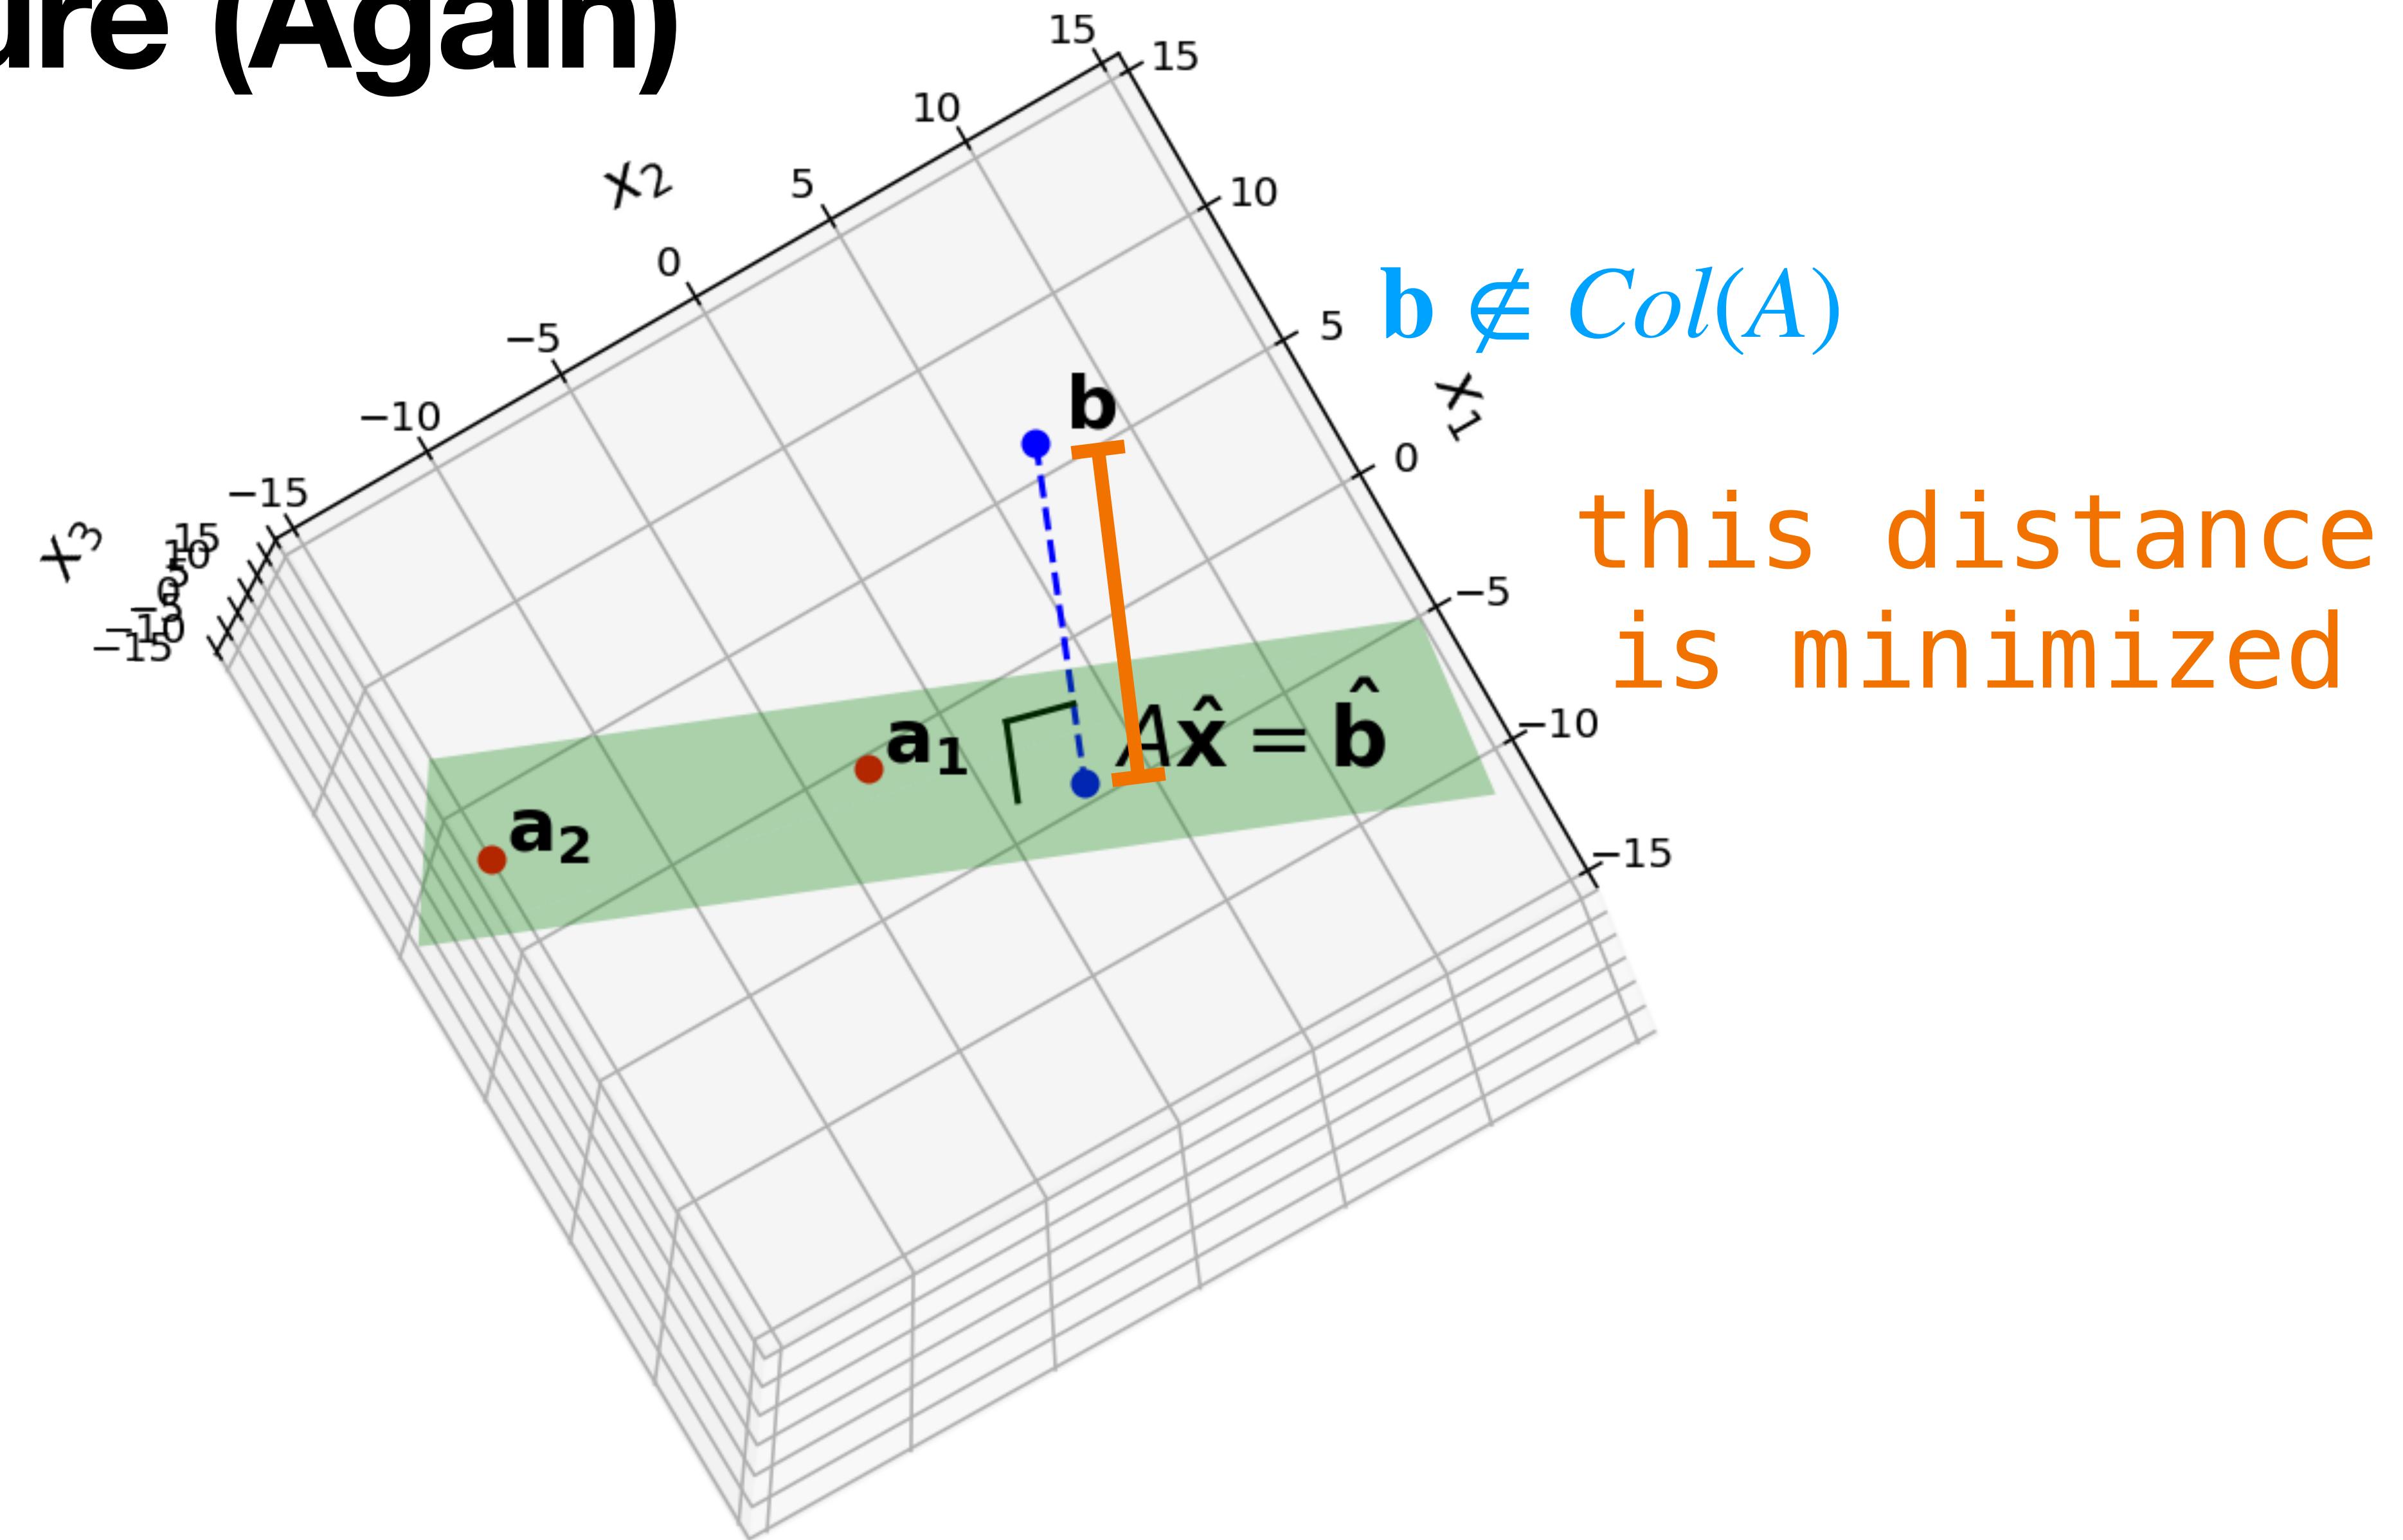
$$\underbrace{\|A\hat{x} - b\|}_{\text{for any } x \text{ in } \mathbb{R}^n} \leq \underbrace{\|Ax - b\|}$$

for any x in \mathbb{R}^n

Again, $\|A\hat{x} - b\|$ is as small as possible

Figure 22.8

The Picture (Again)



Argmin

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|$$

Argmin

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|$$

Another way of framing this is via $\arg \min$

Argmin

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|$$

Another way of framing this is via $\arg \min$

Defintion. $\arg \min_{x \in X} f(x) = \hat{x}$ where $f(\hat{x}) = \min_{x \in X} f(x)$

Argmin

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|$$

Another way of framing this is via $\arg \min$

Defintion. $\arg \min_{x \in X} f(x) = \hat{x}$ where $f(\hat{x}) = \min_{x \in X} f(x)$

\hat{x} is the *argument* that *minimizes* f

Argmin

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|$$

Another way of framing this is via $\arg \min$

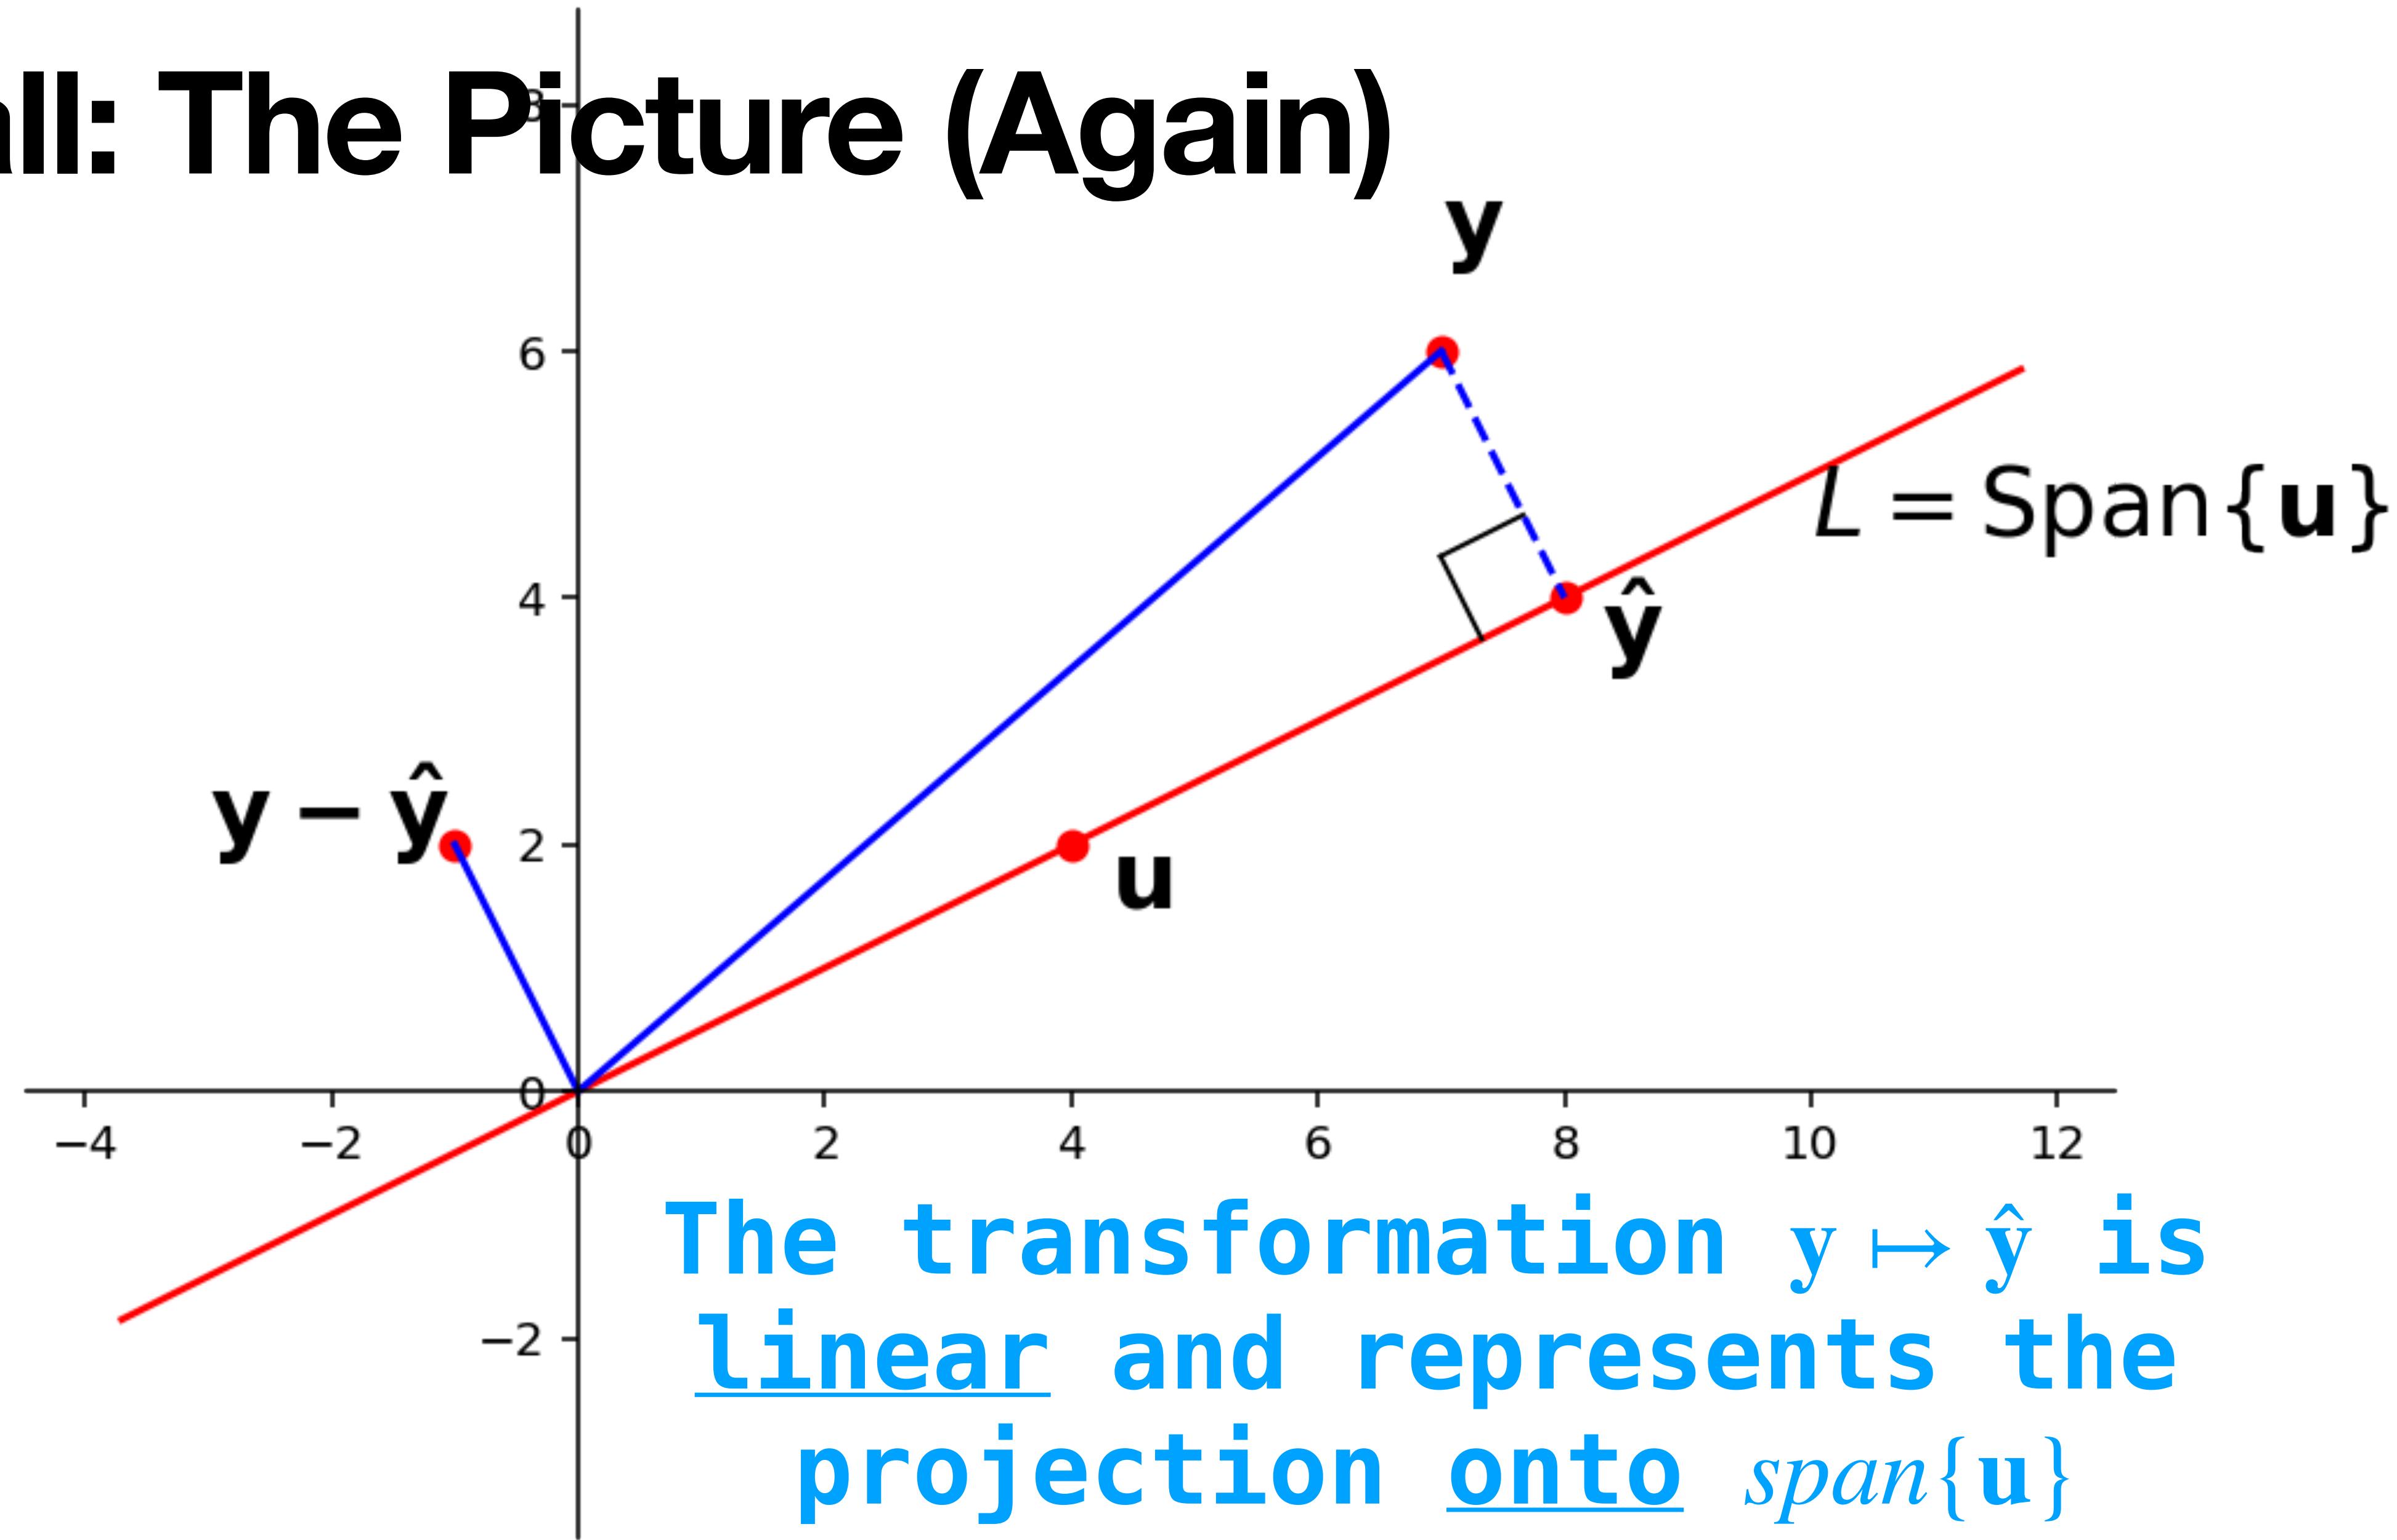
Defintion. $\arg \min_{x \in X} f(x) = \hat{x}$ where $f(\hat{x}) = \min_{x \in X} f(x)$

\hat{x} is the *argument* that *minimizes* f

This is now an optimization problem

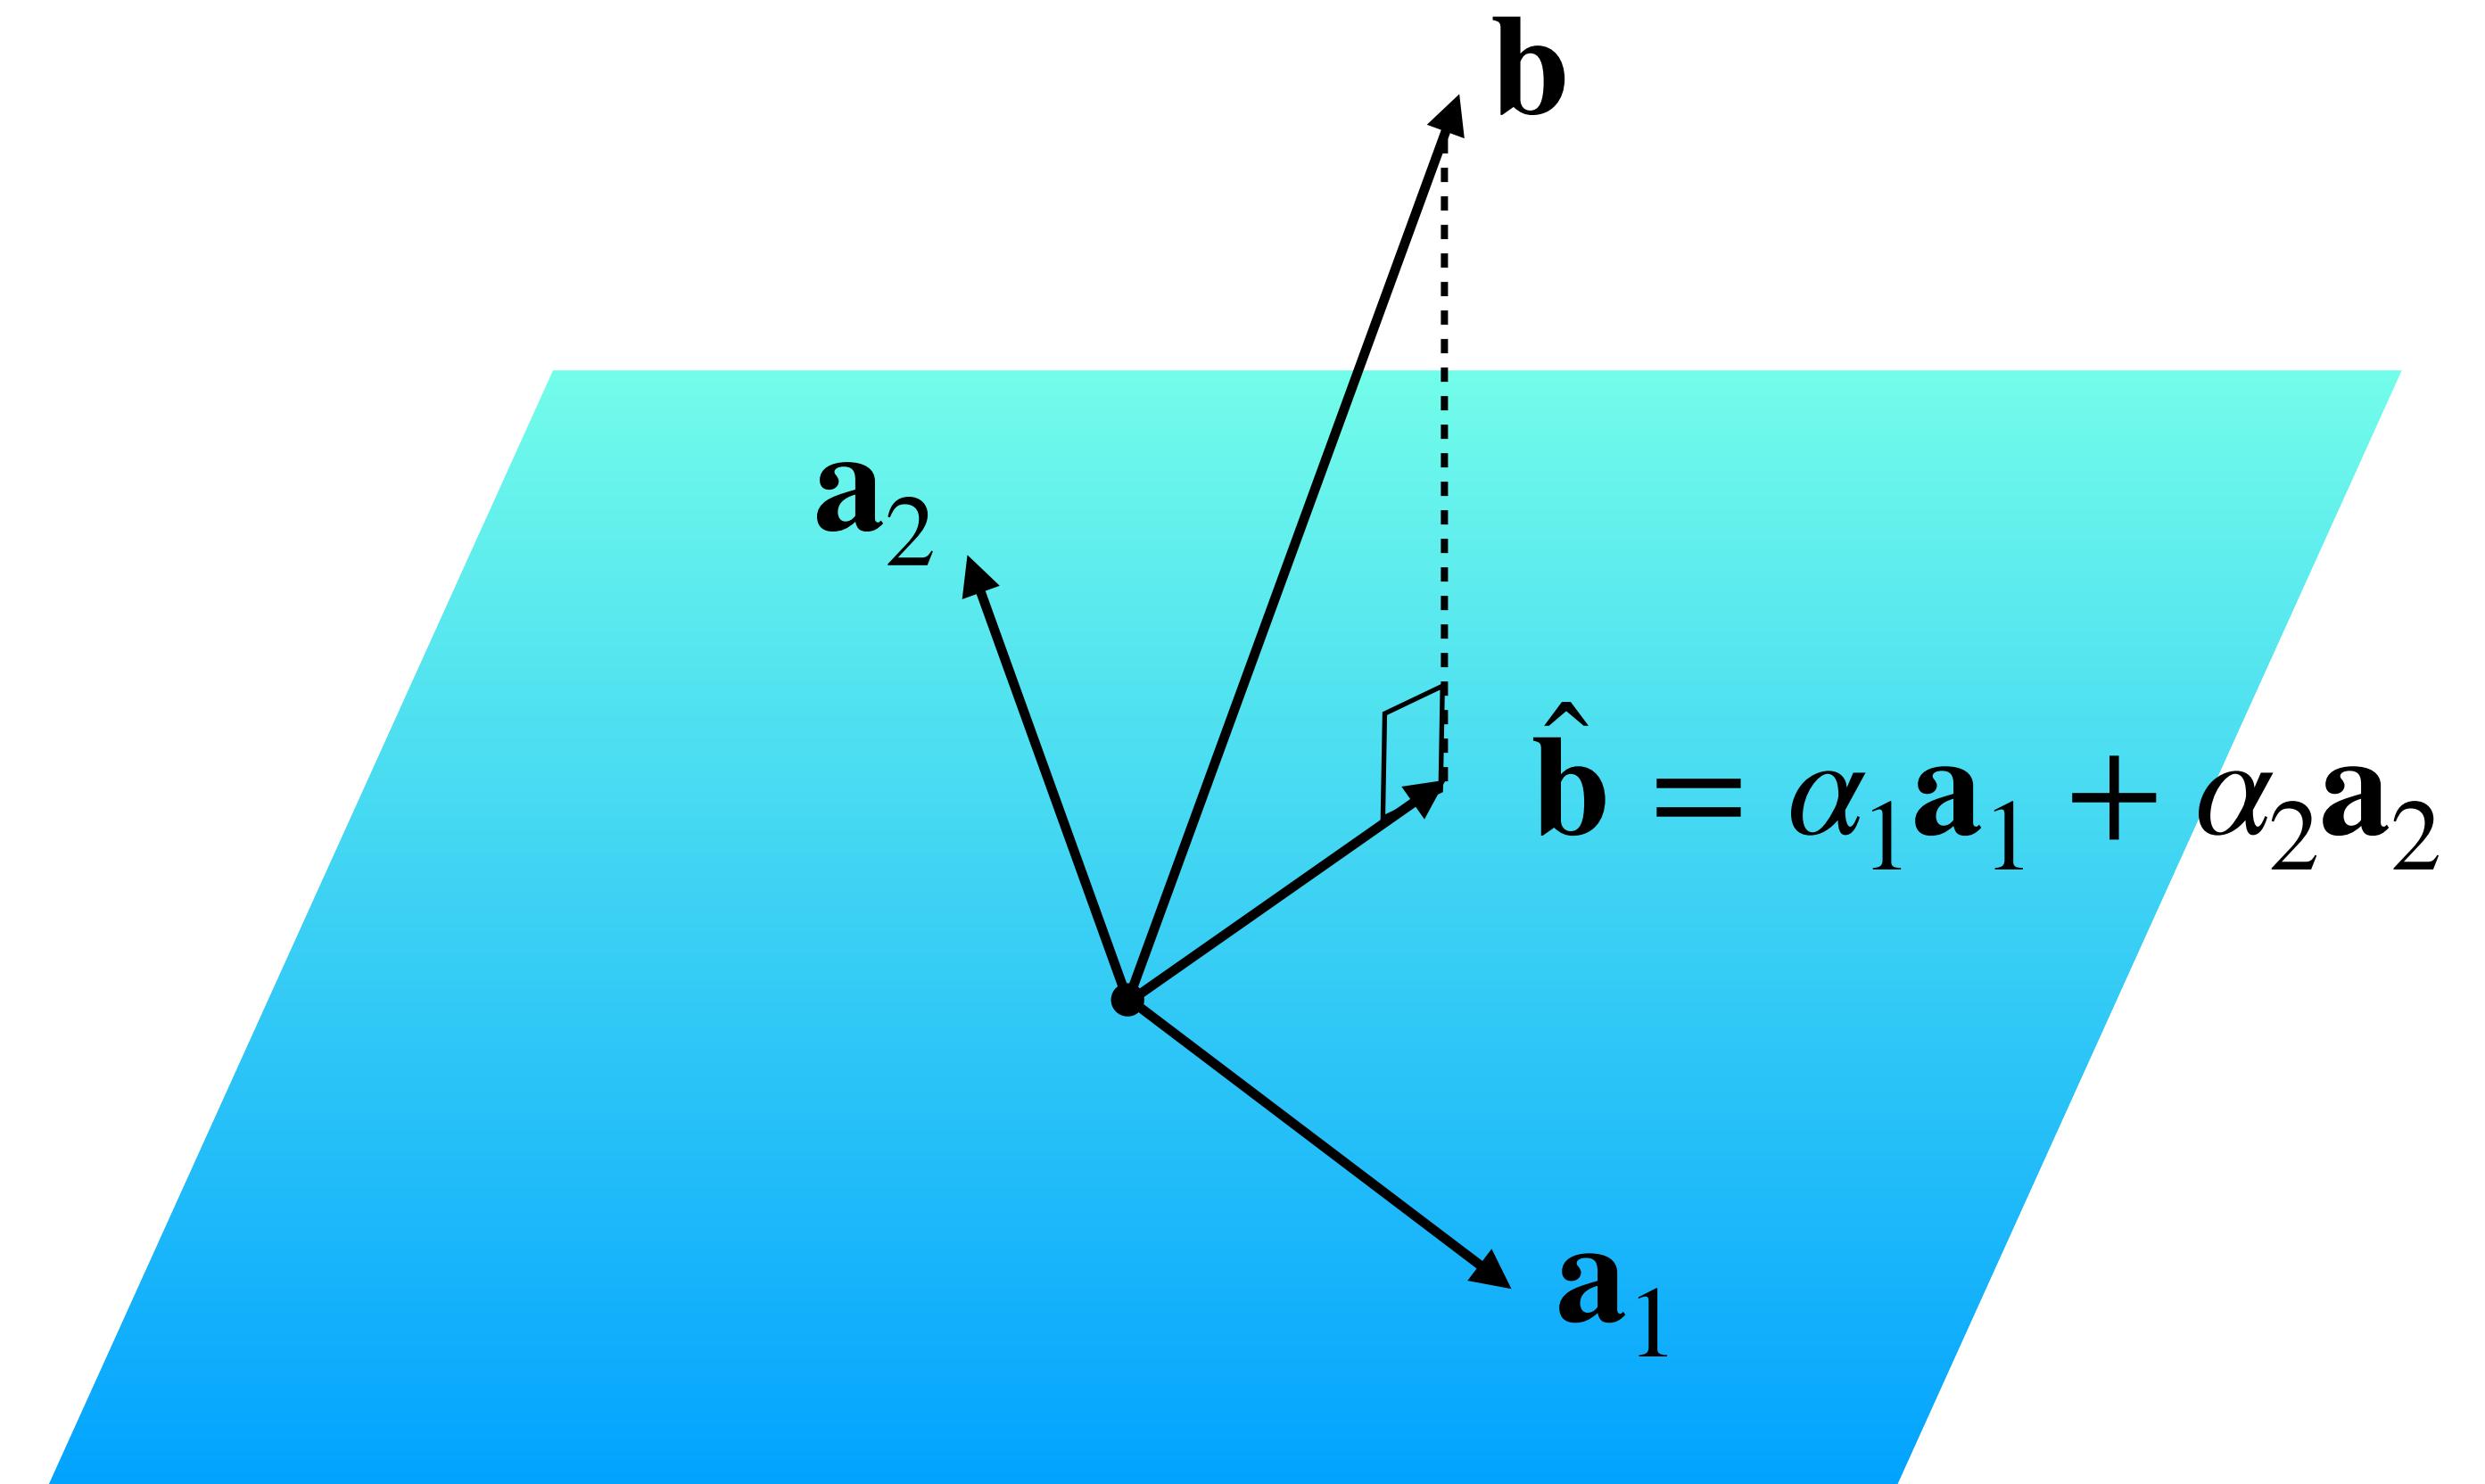
Solving the General Least Squares Problems

Recall: The Picture (Again)



Projects onto other Spans

The transformation
 $b \mapsto \hat{b}$ is the
projection of b
onto $\text{span}\{a_1, a_2\}$



The High Level Approach.

Question. Find a least squares solutions to

$$\tilde{A}\tilde{x} = \tilde{b}$$

Solution.

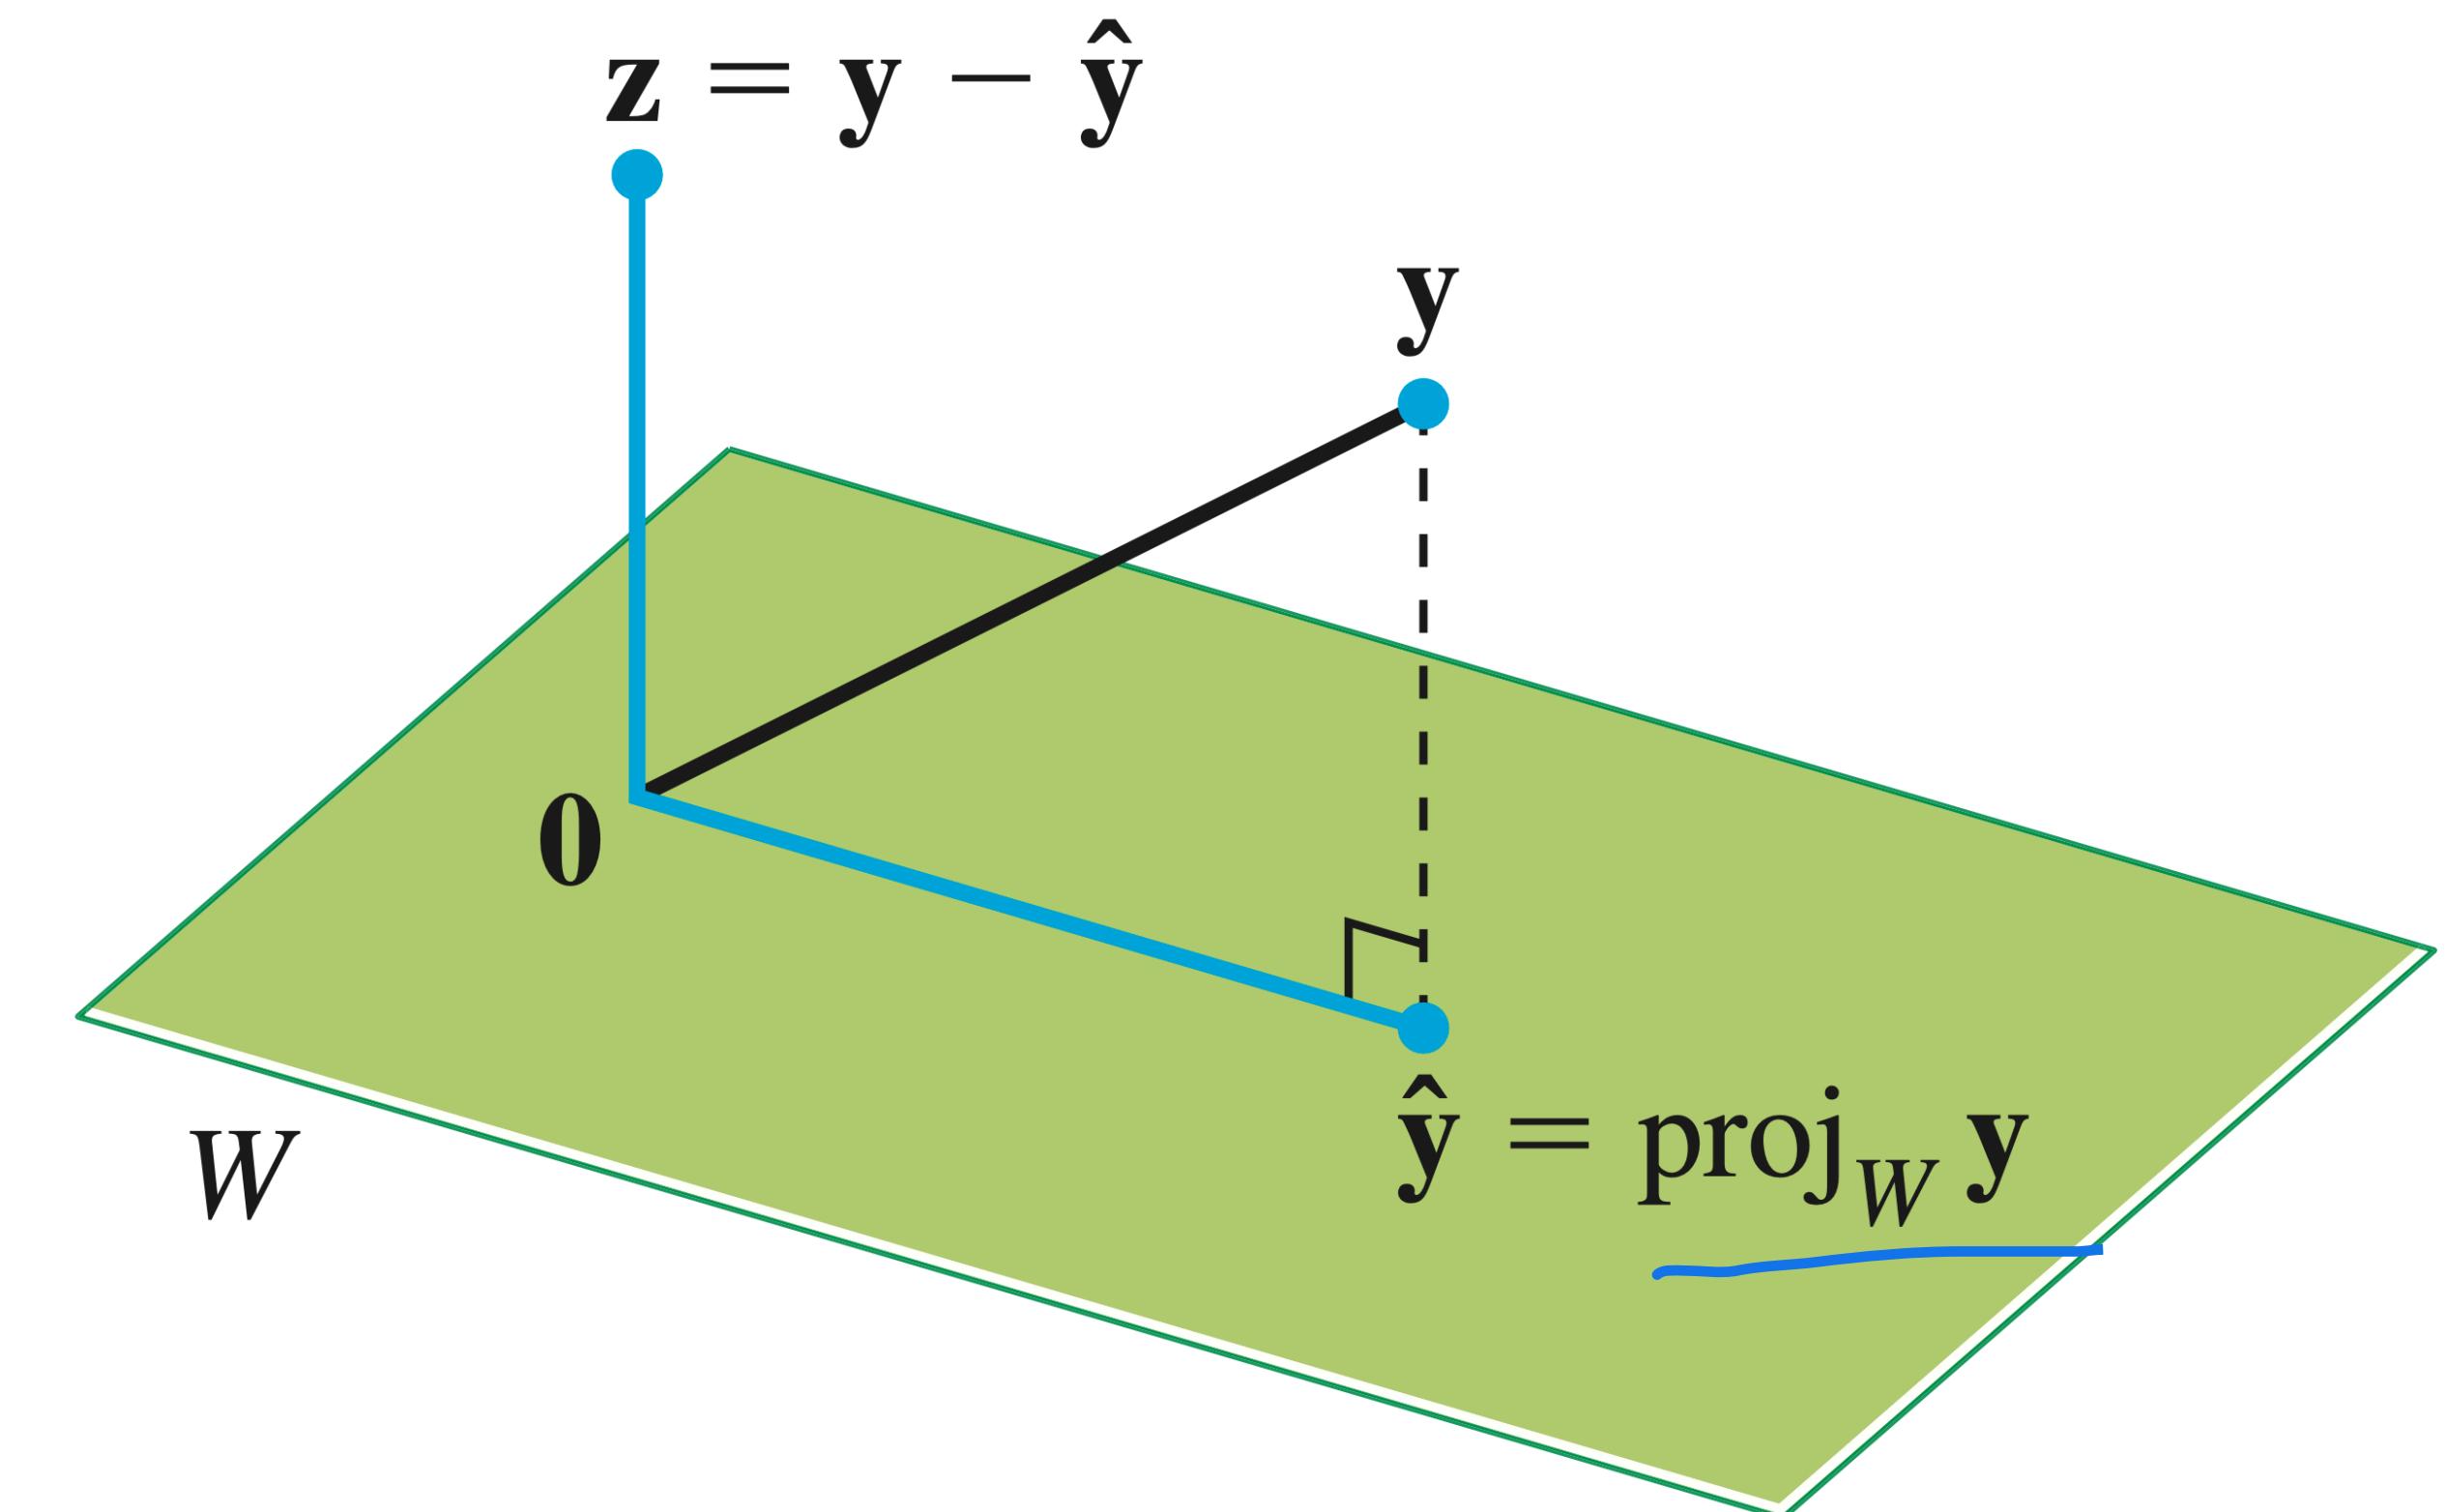
1. Find the closest point \hat{b} in $Col(A)$ to b
2. Solve the equation $Ax = \hat{b}$ instead

Orthogonal Decomposition Theorem

Theorem. Let W be a subspace of \mathbb{R}^n . Every vector y in \mathbb{R}^n can be written uniquely as

$$y = \hat{y} + z$$

where $\hat{y} \in W$ and z is orthogonal to every vector in W

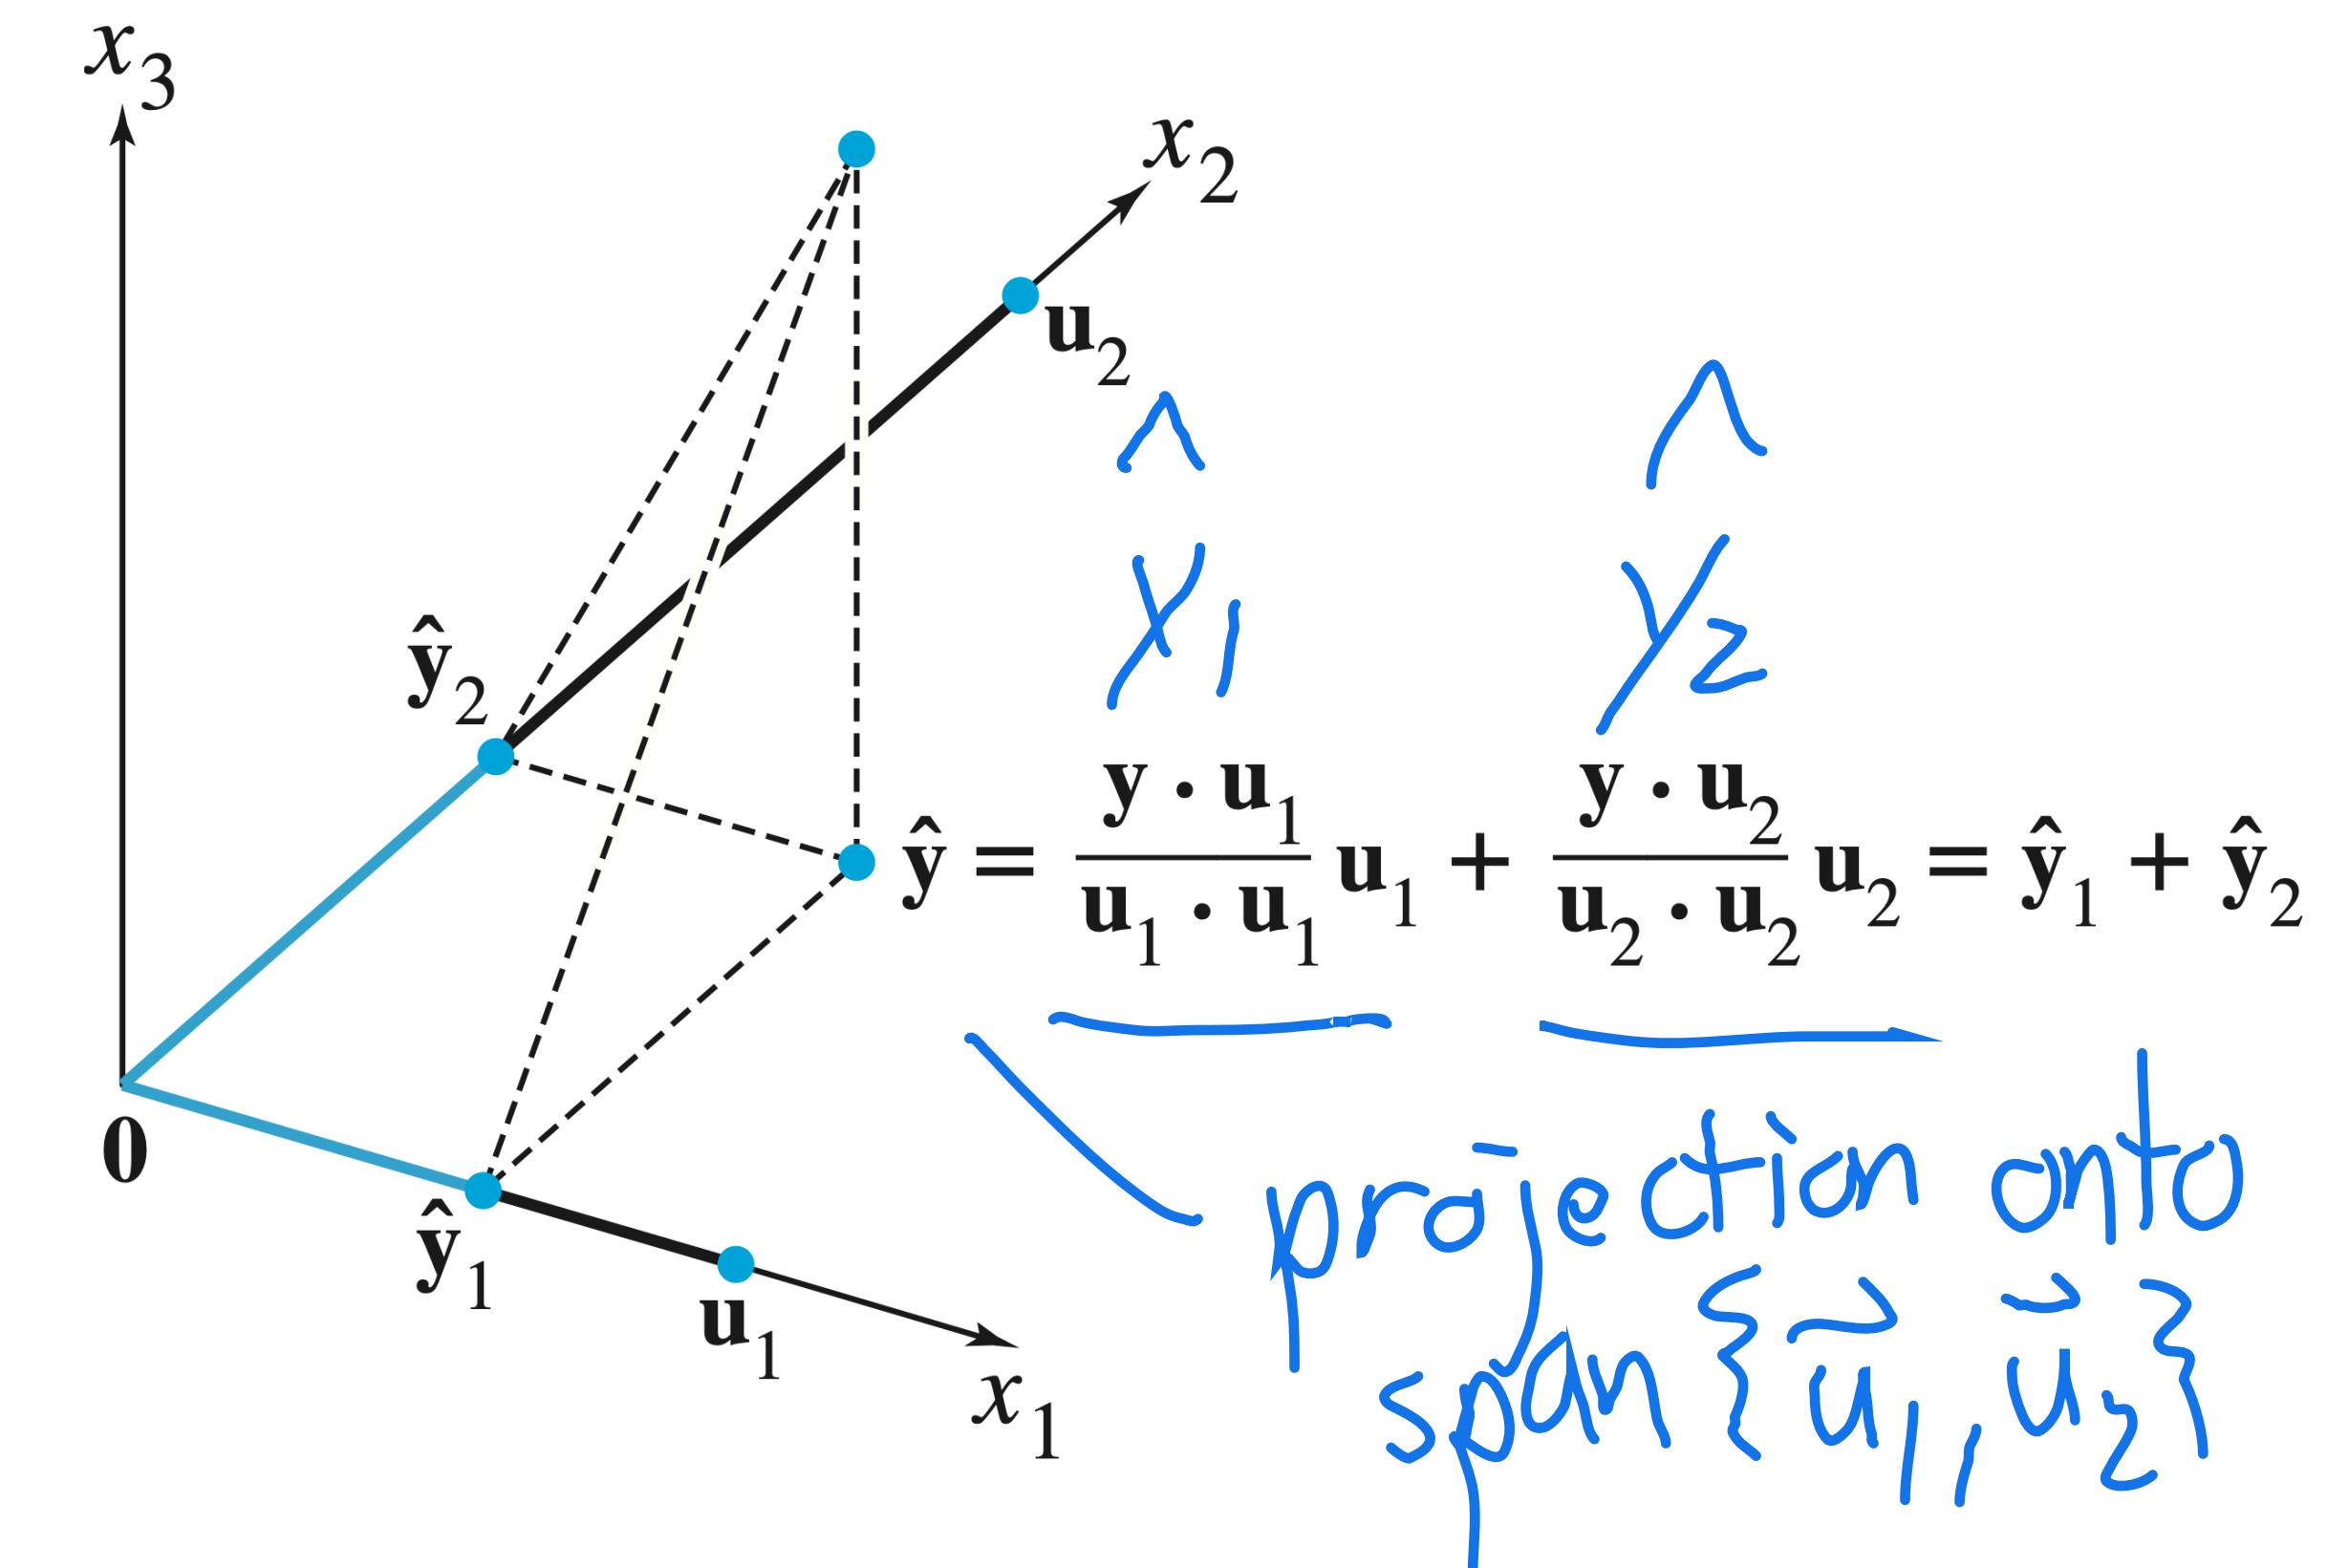


Projection via Orthogonal Bases

We can determine \hat{y} by projecting onto an orthogonal basis

Every subspace has an orthogonal basis (we won't prove this)

Gram-Schmidt



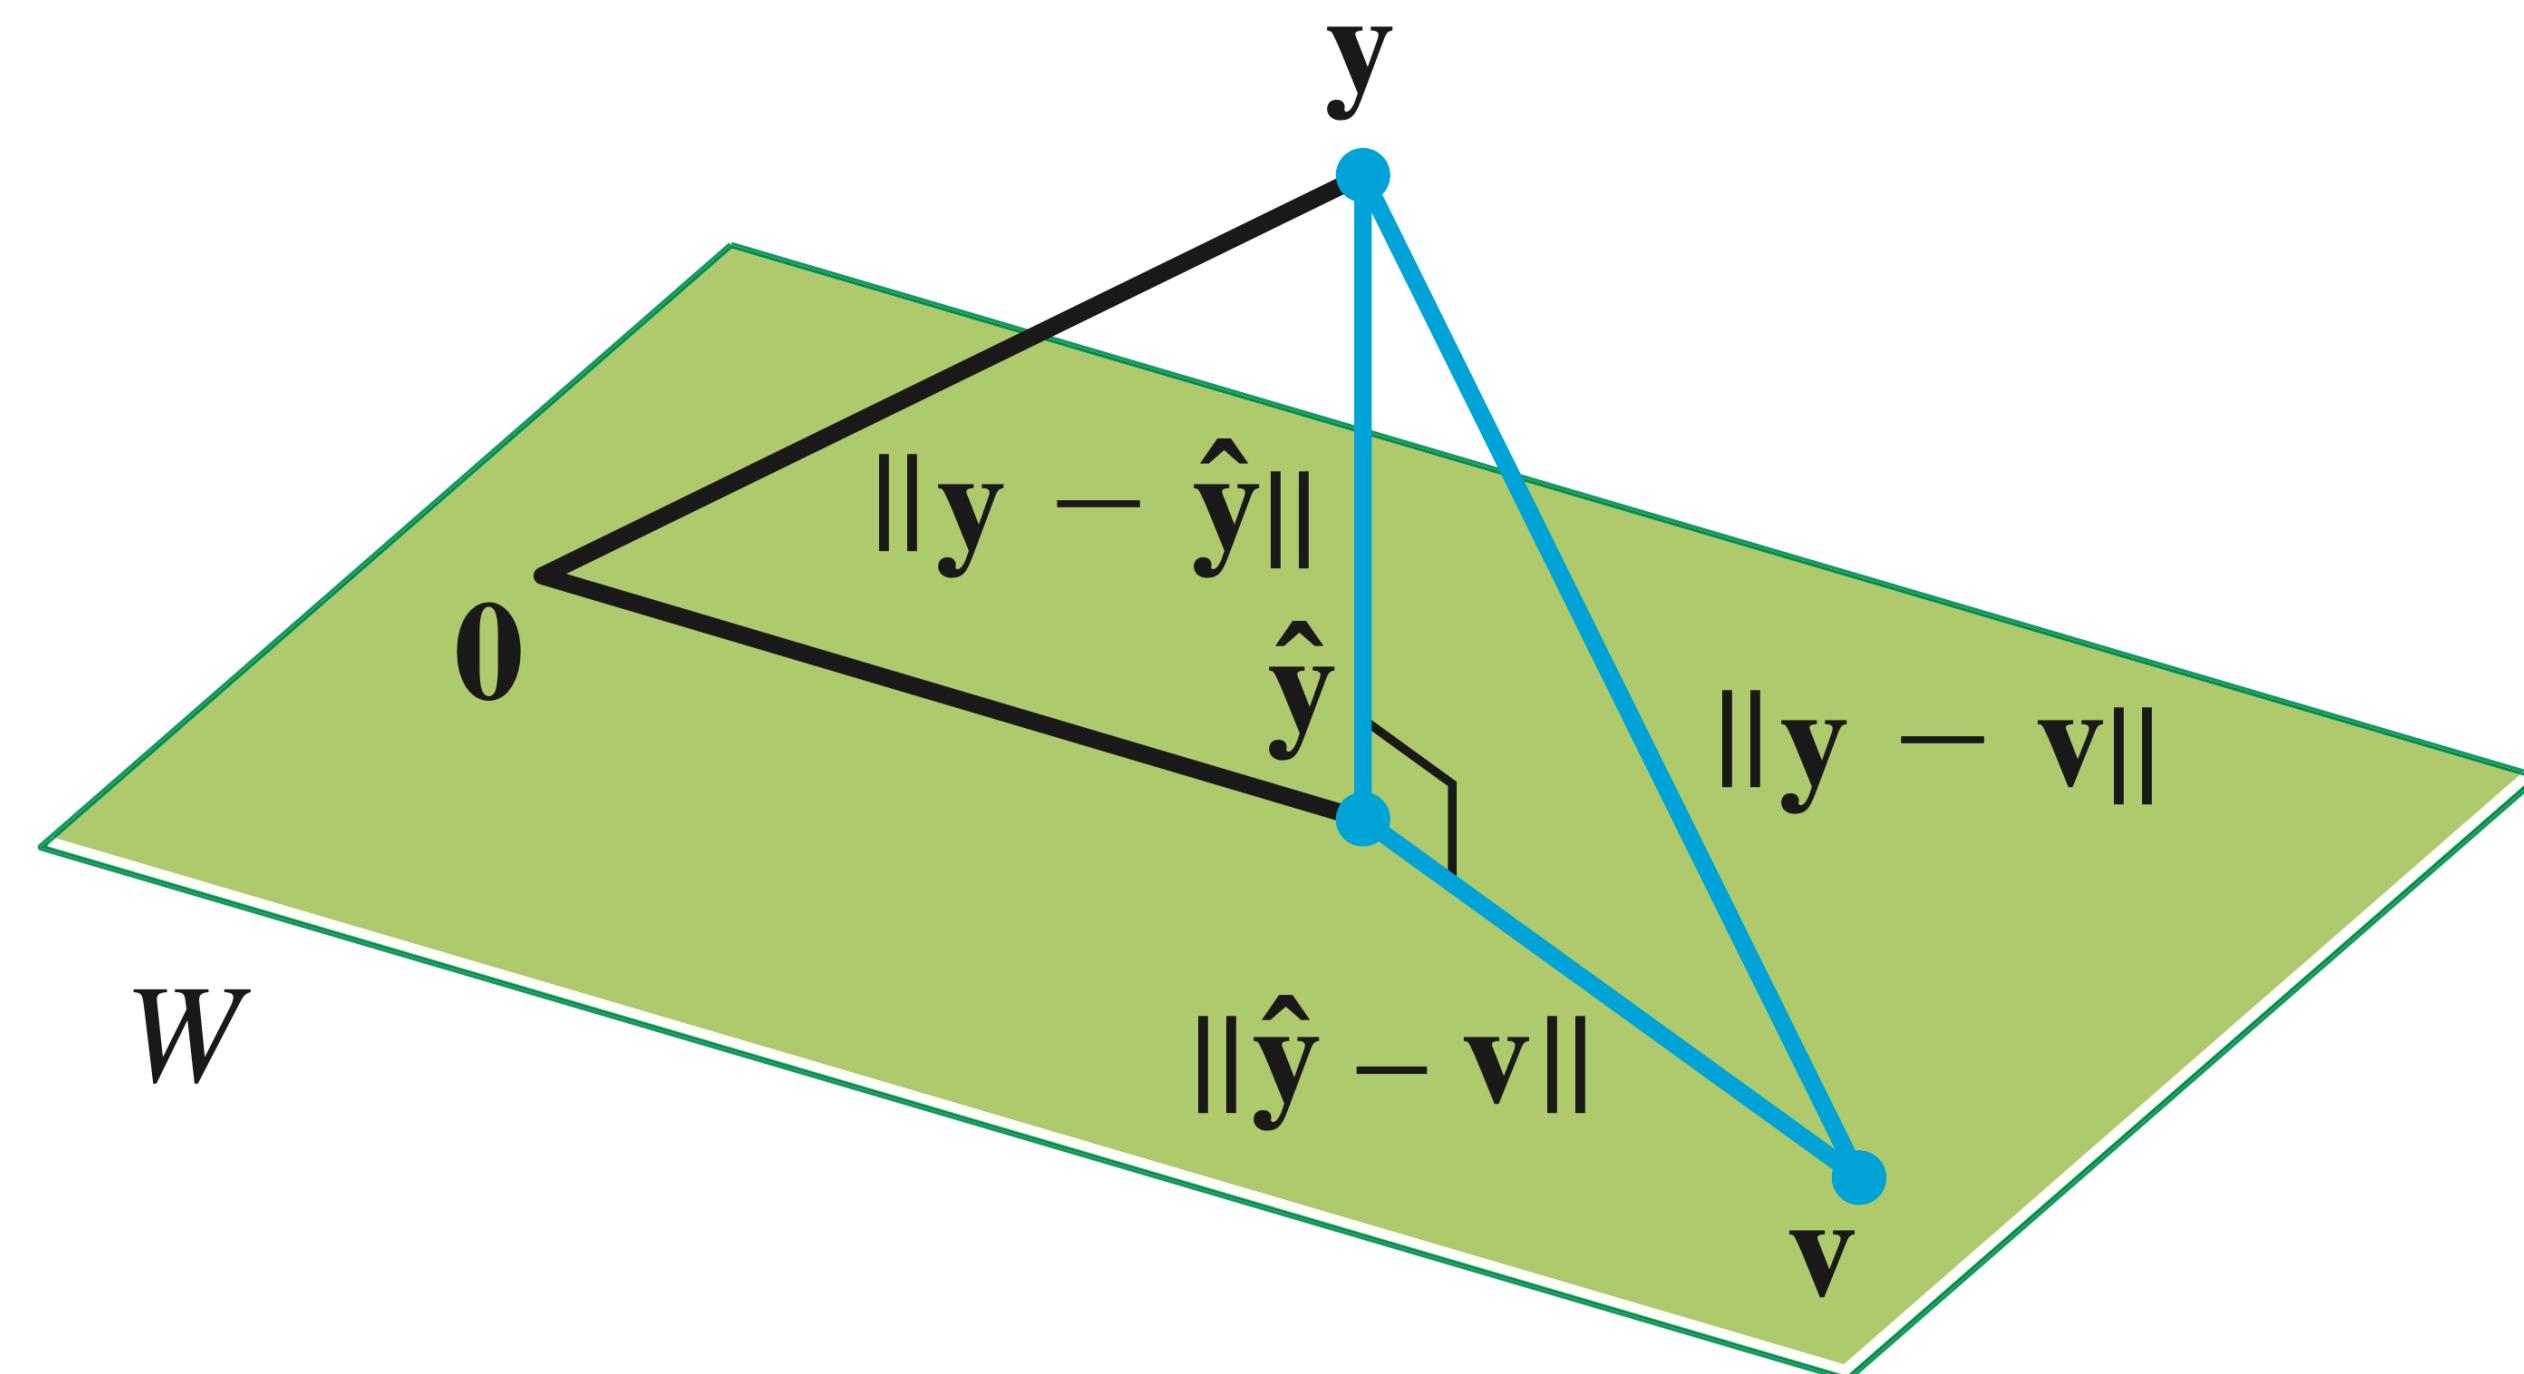
The Best-Approximation Theorem

Theorem. Let W be a subspace of \mathbb{R}^n , and let \hat{y} be the orthogonal projection of y onto W . Then

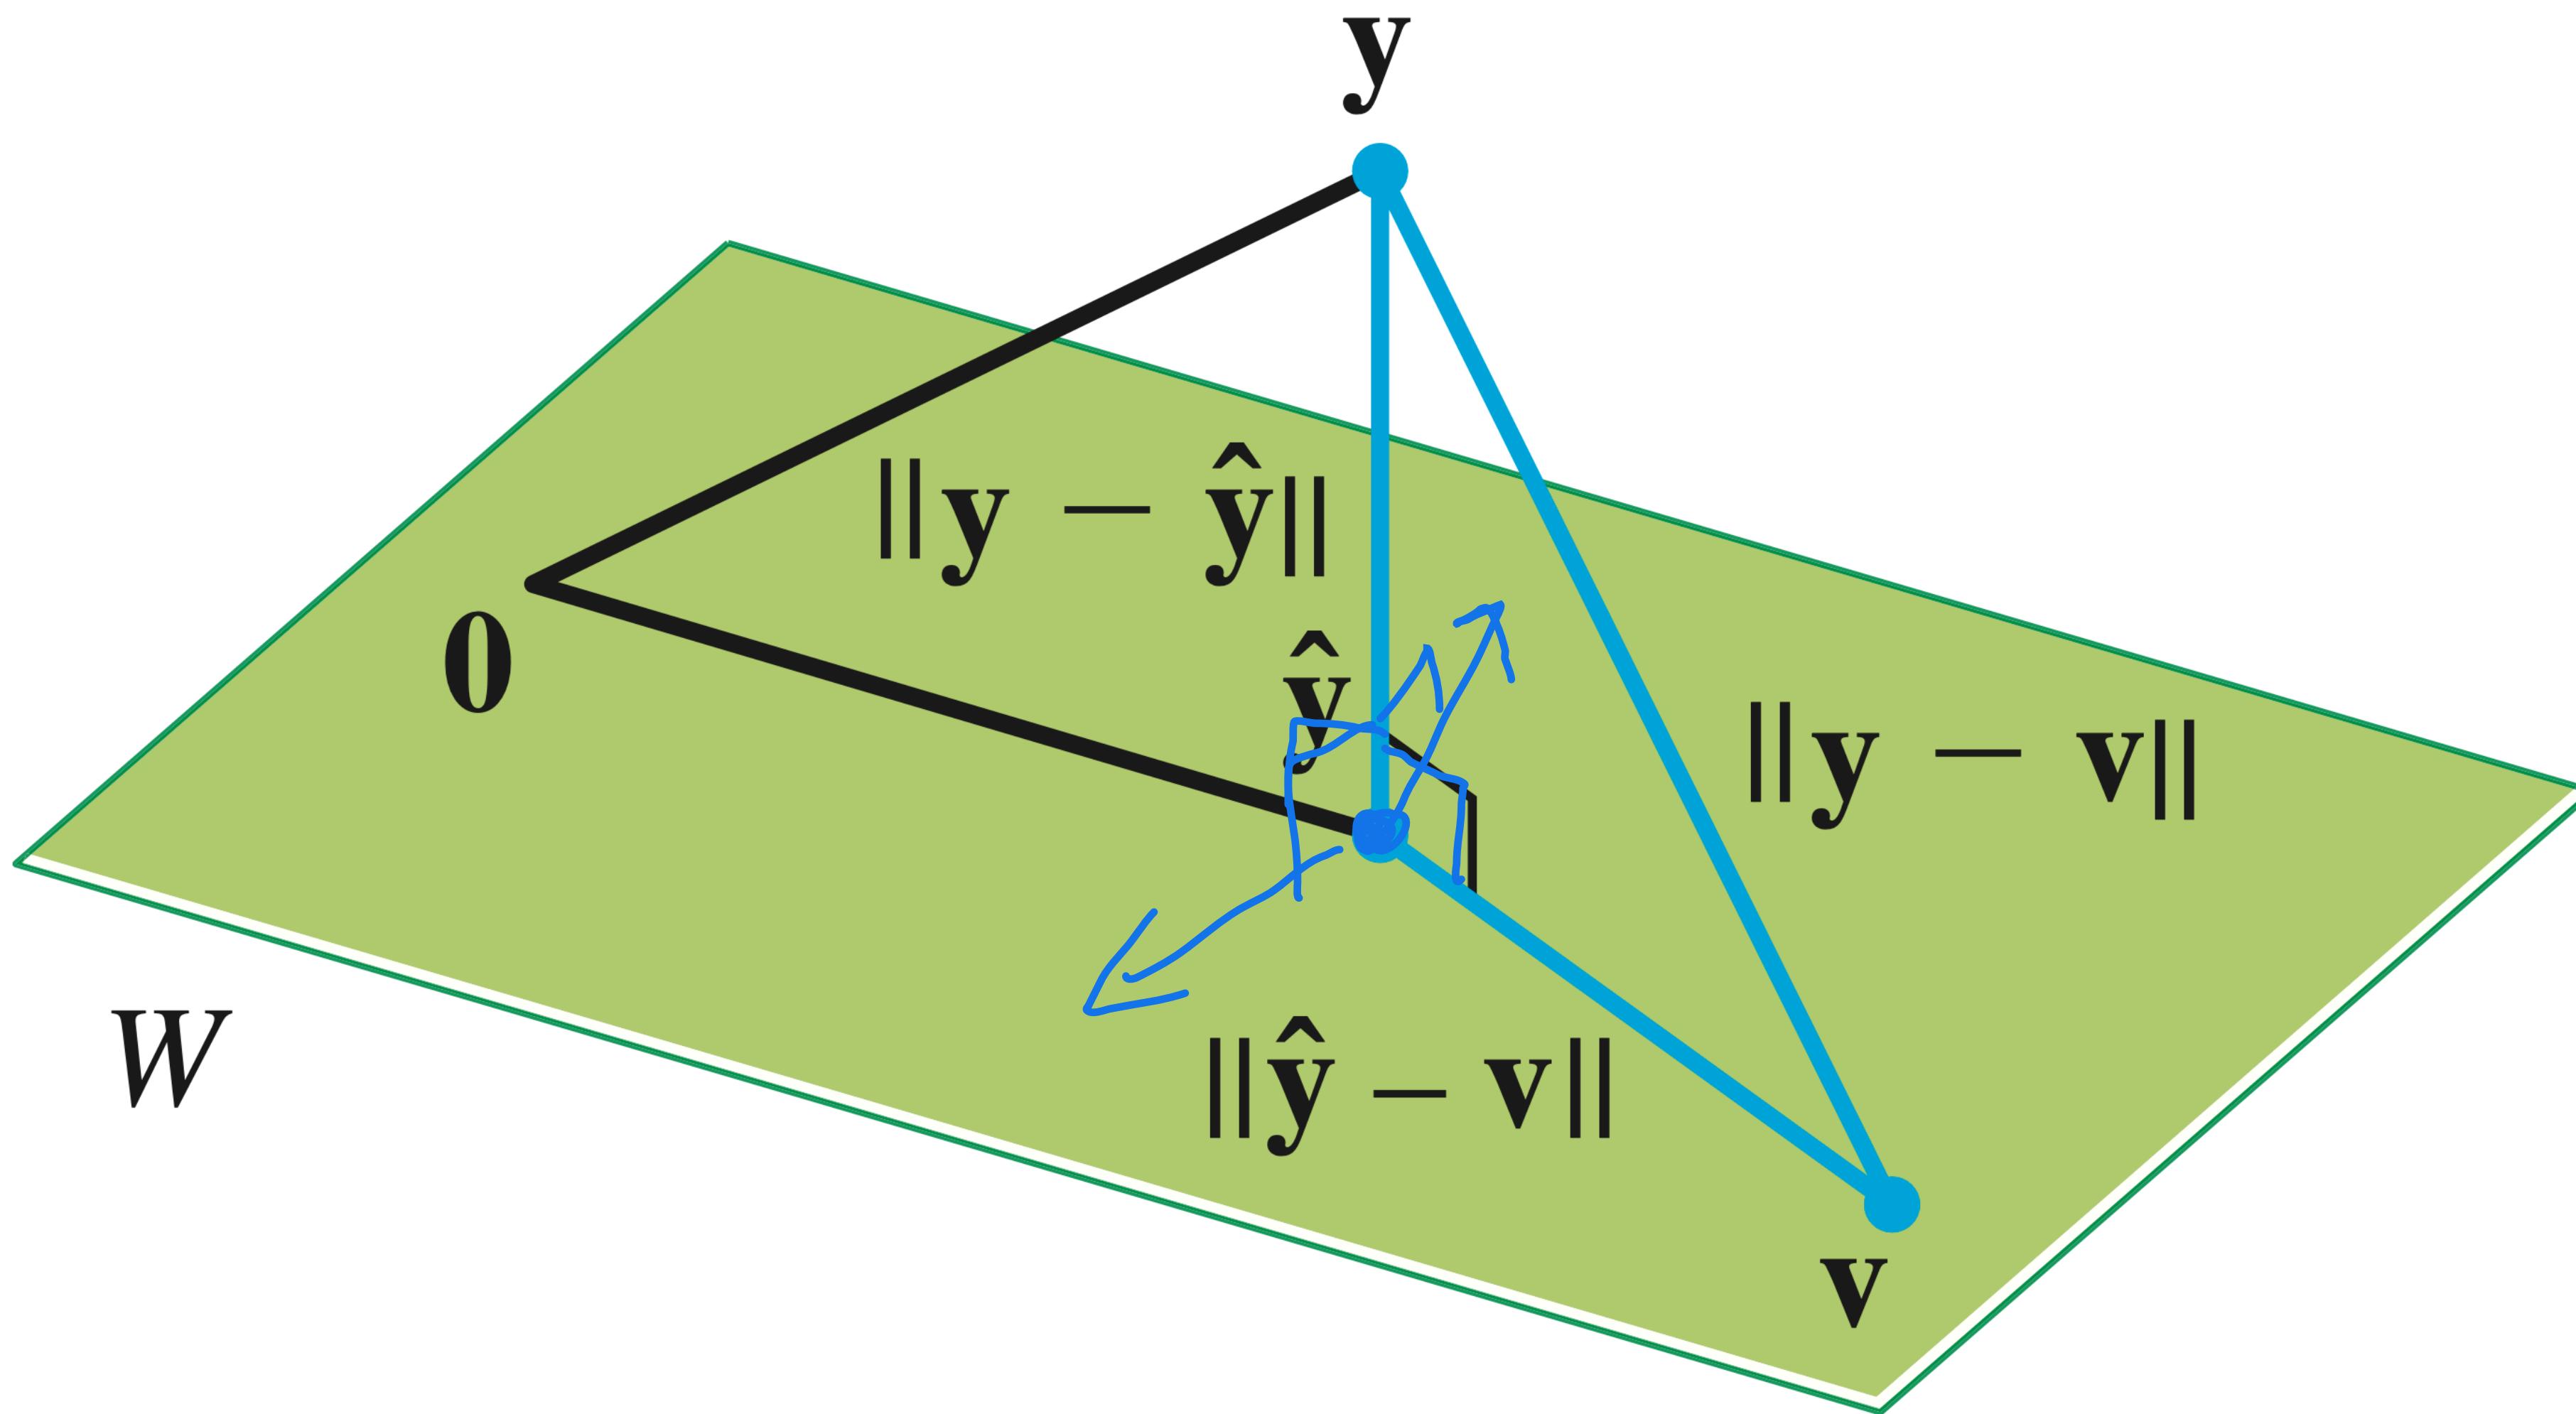
$$\|y - \hat{y}\| \leq \|y - w\|$$

for any vector w in W

\hat{y} is the closest point in W to y



Proof by Inspection

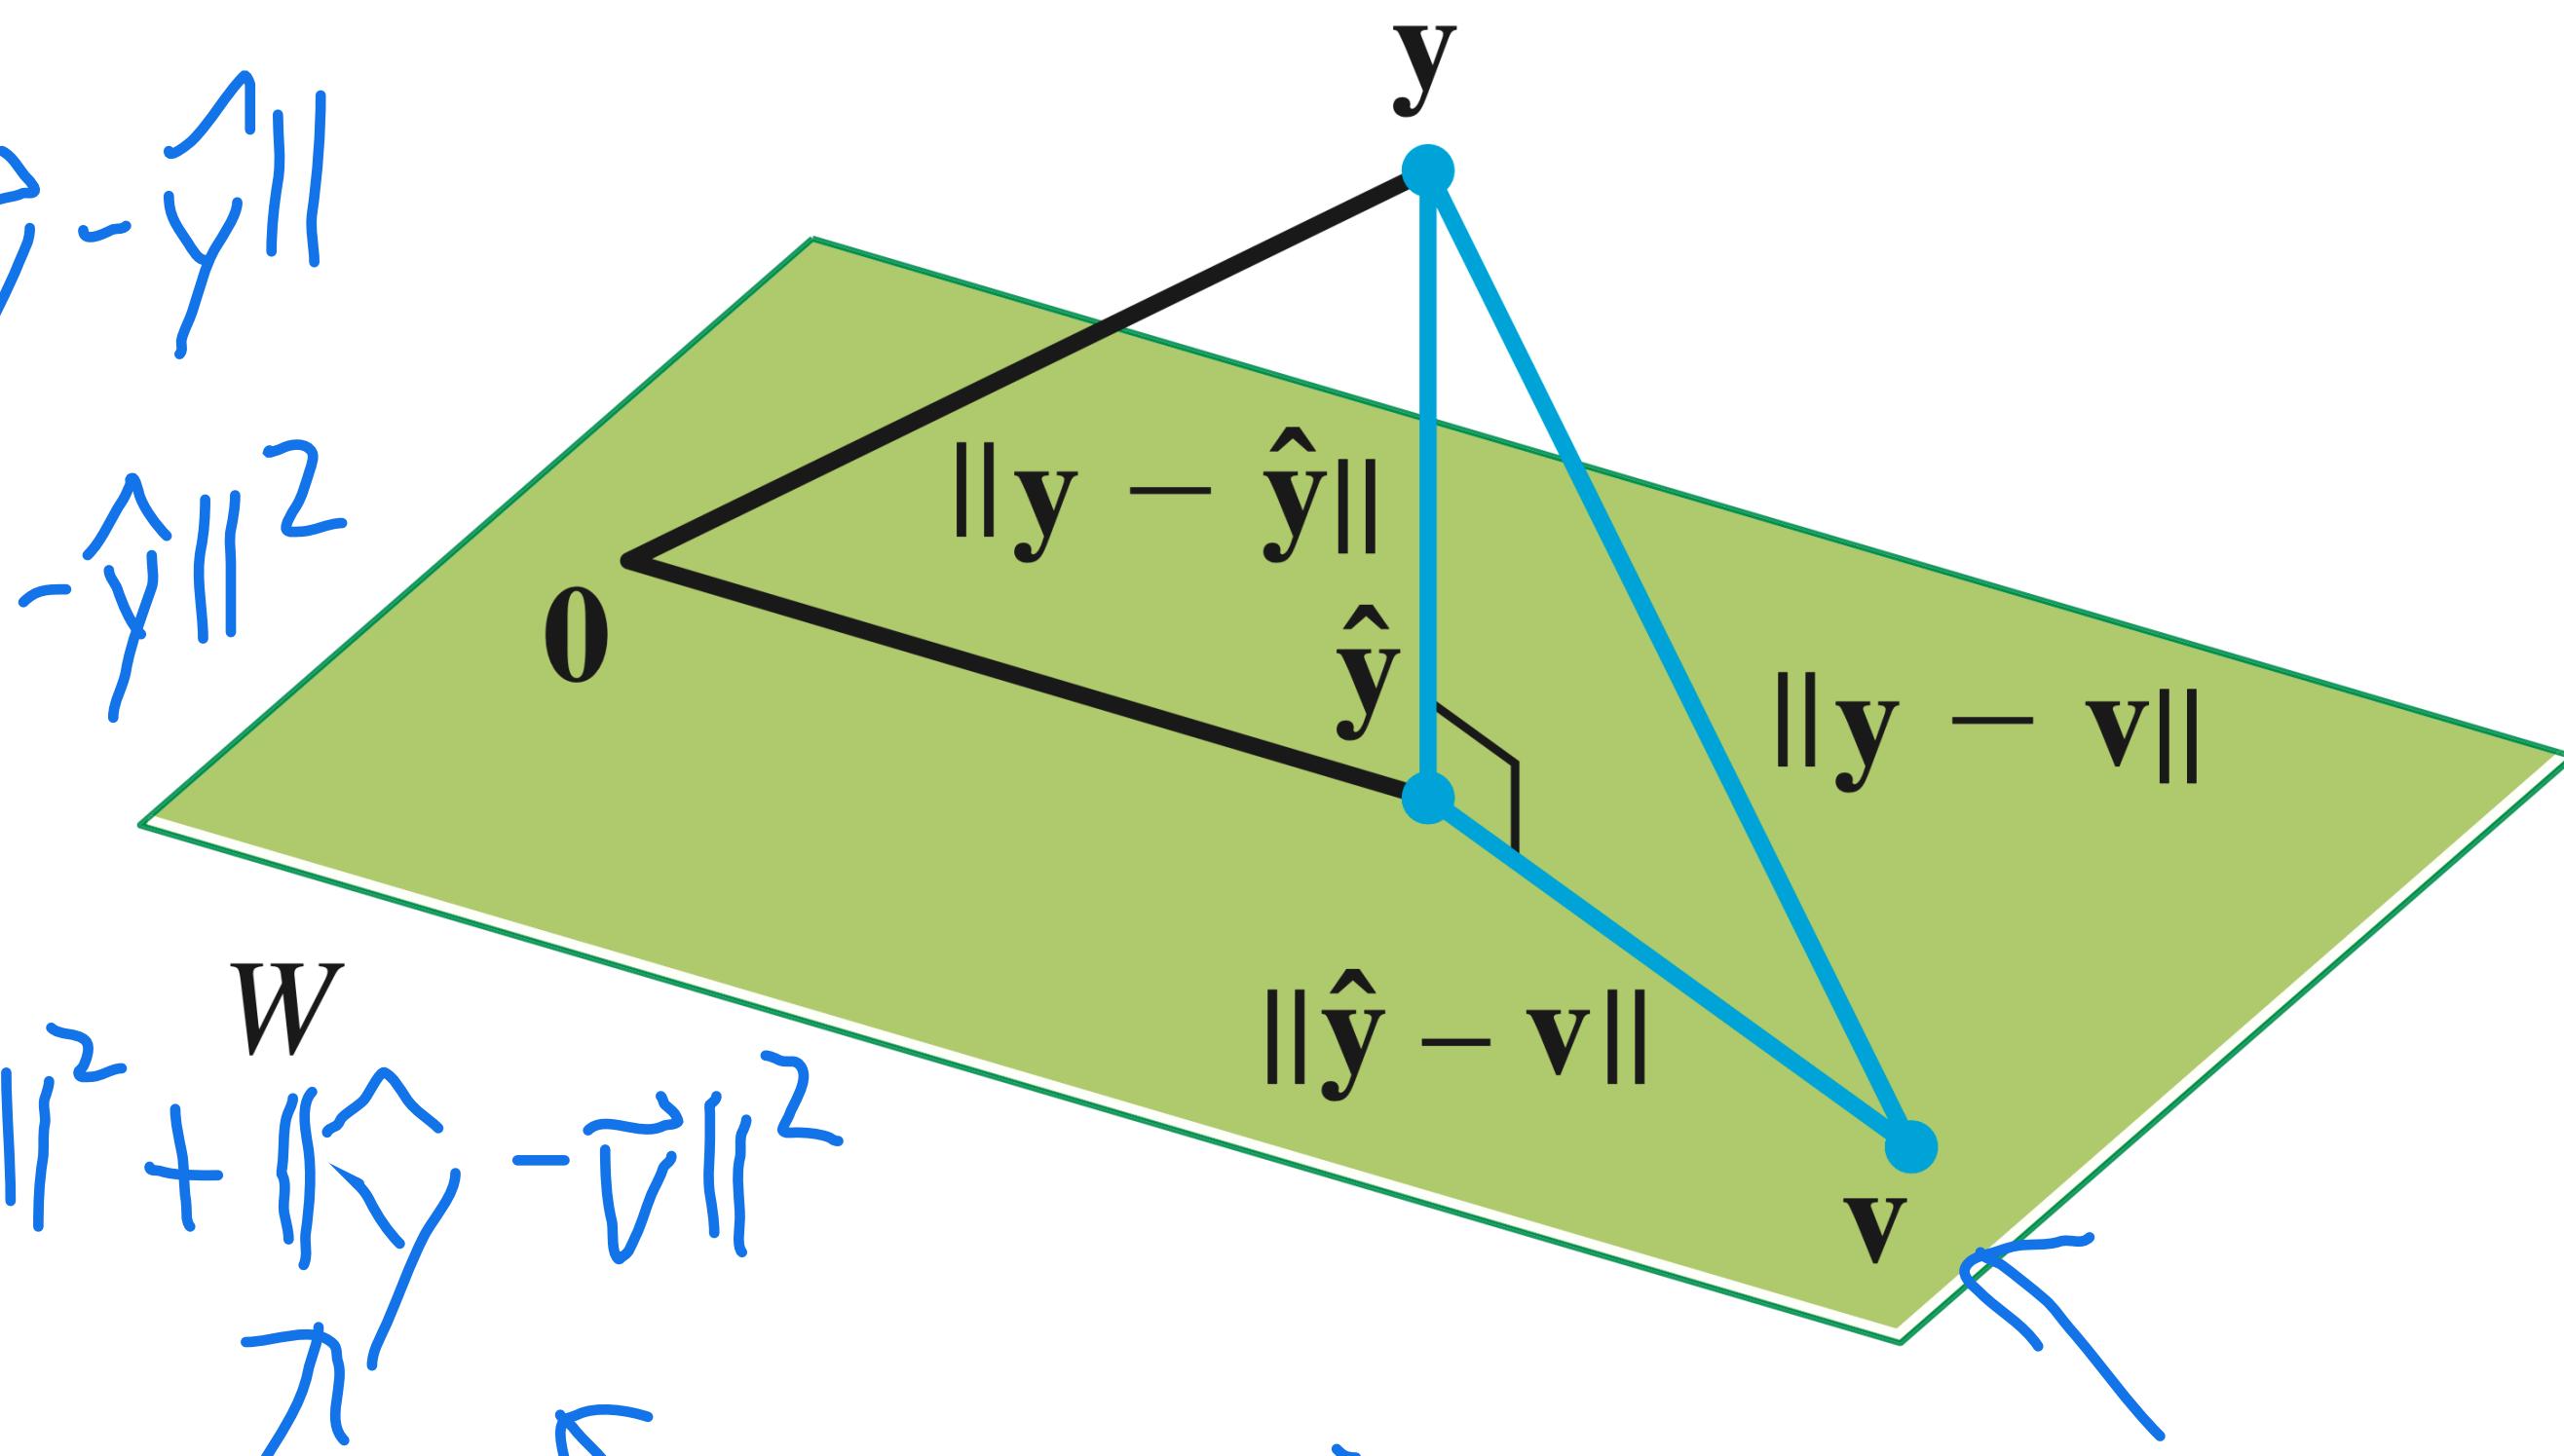


Proof by Algebra

Verify: For any \vec{v} in W
Goal: $\|\vec{y} - \vec{v}\| \geq \|\vec{y} - \hat{\vec{y}}\|$

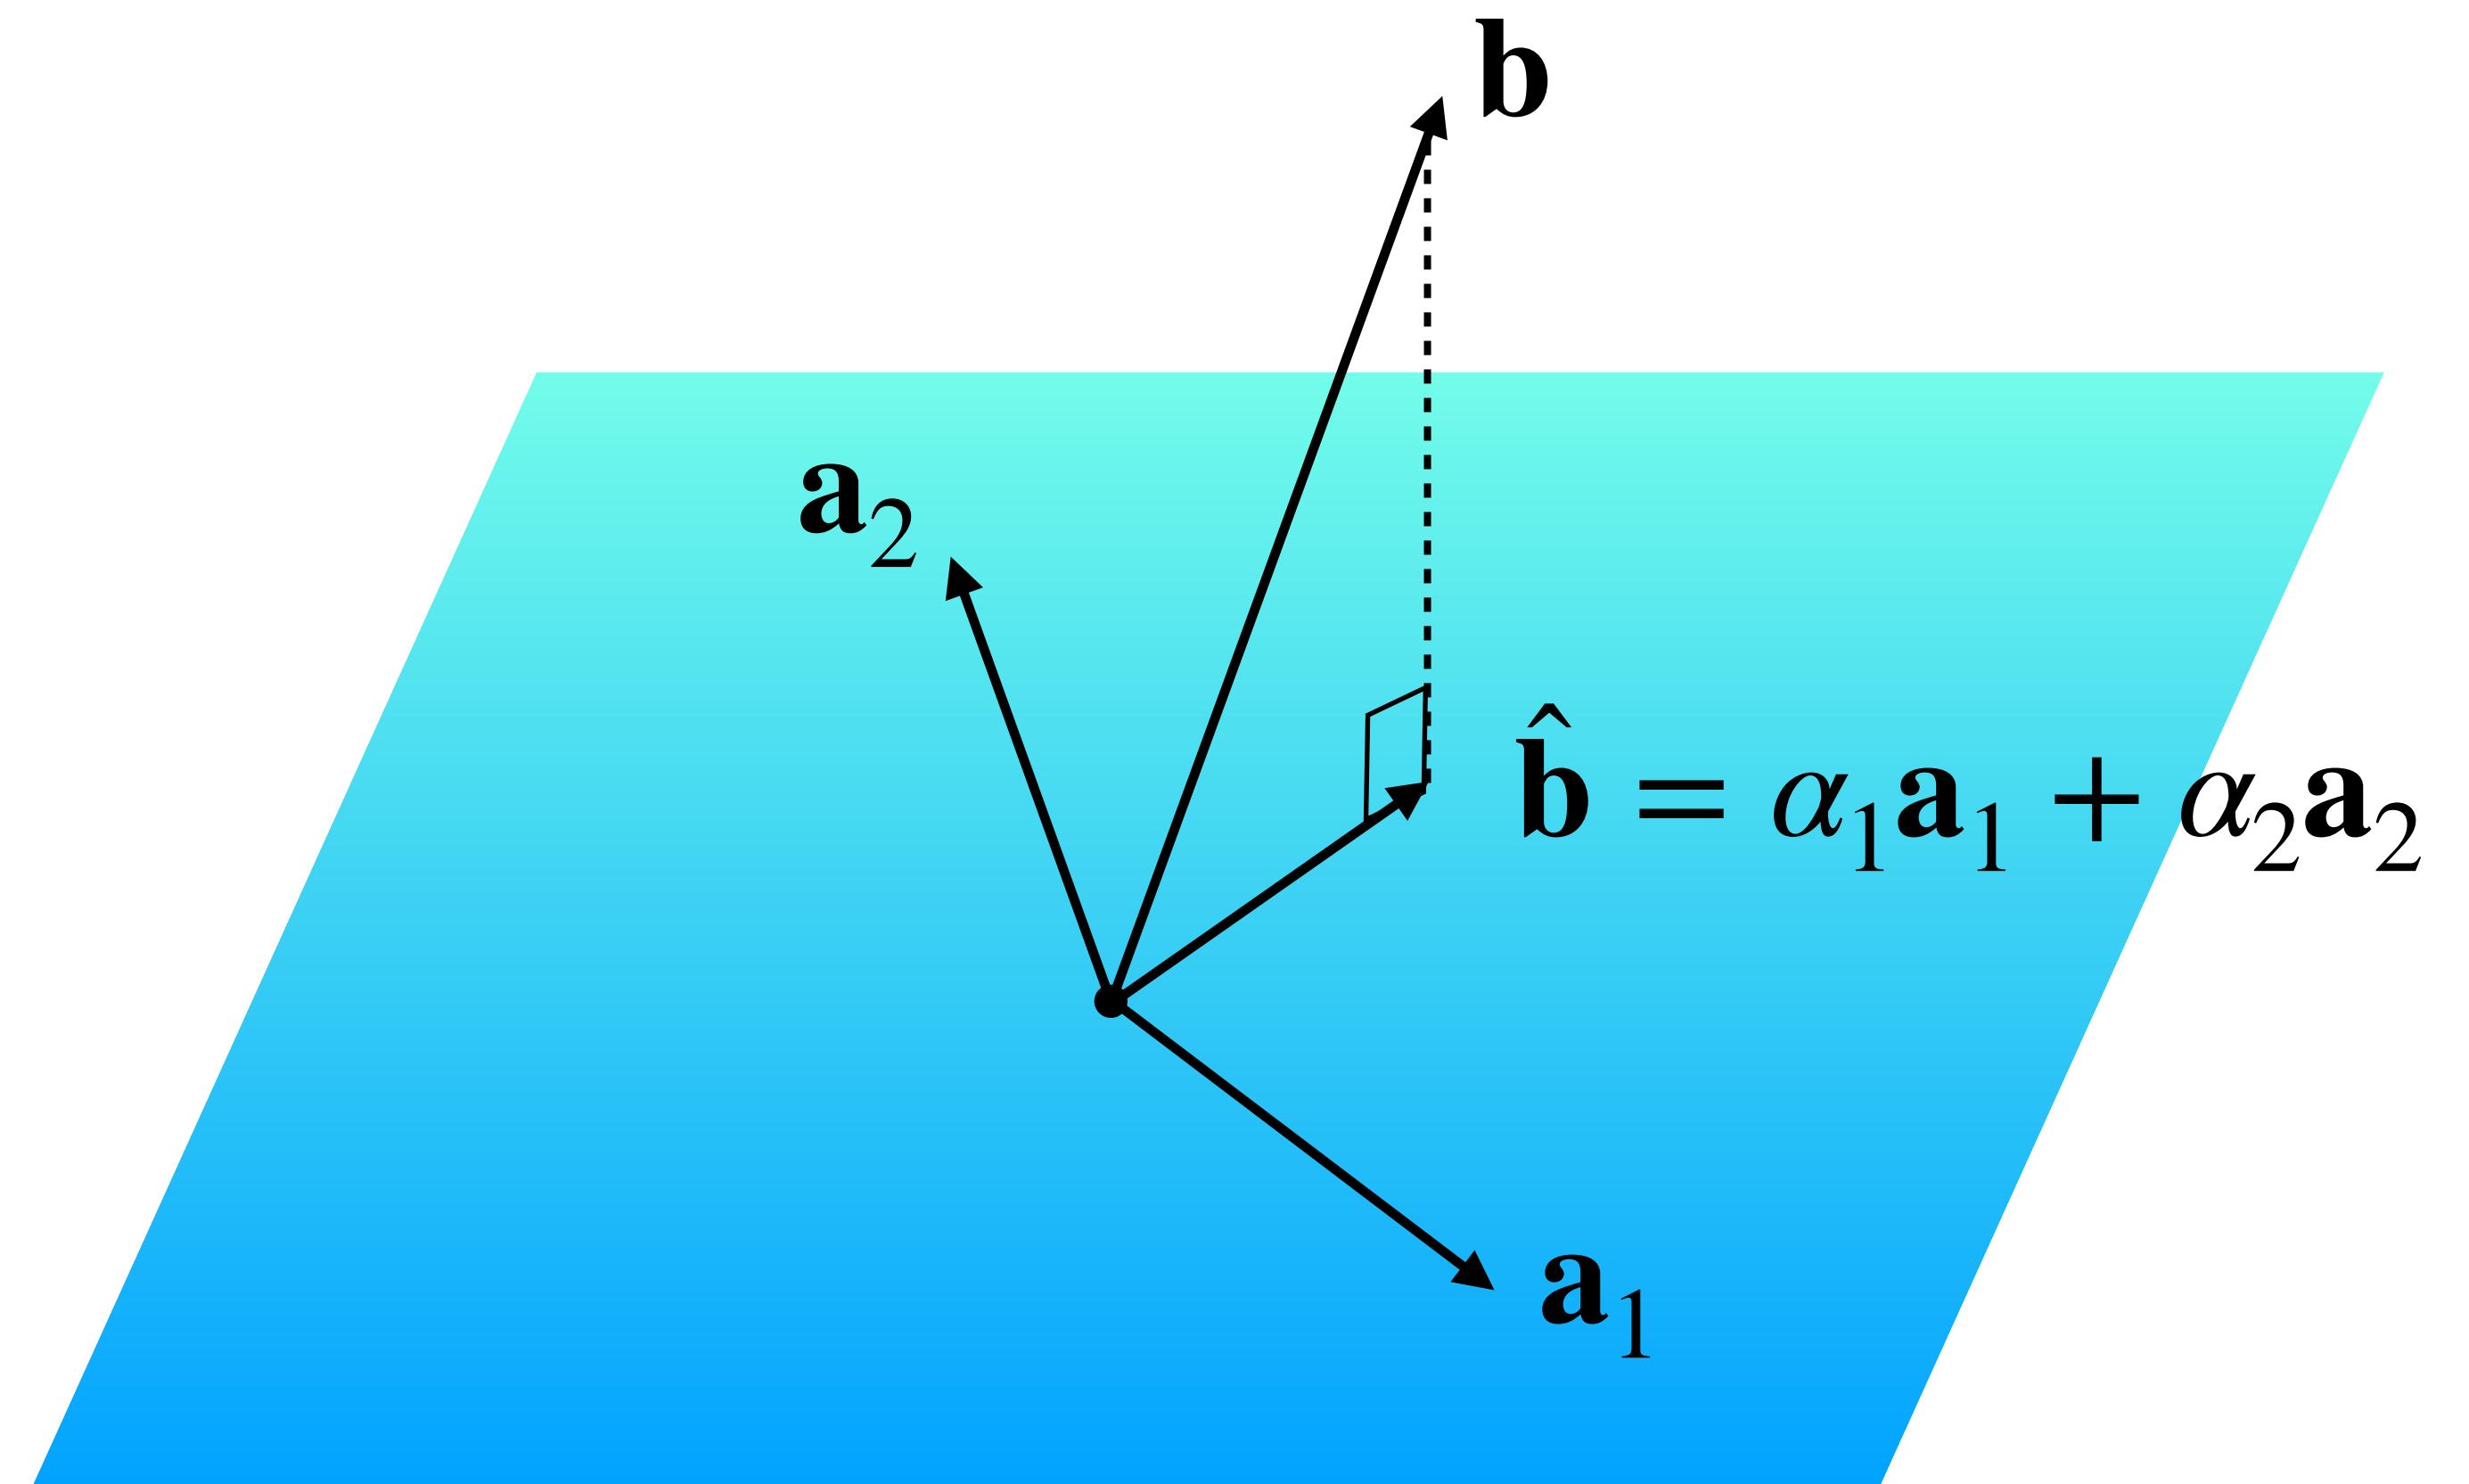
$$\|\vec{y} - \vec{v}\|^2 \geq \|\vec{y} - \hat{\vec{y}}\|^2$$

$$\begin{aligned} \|\vec{y} - \vec{v}\|^2 &= \|\vec{y} - \hat{\vec{y}}\|^2 + \|\hat{\vec{y}} - \vec{v}\|^2 \\ &\geq 0 \end{aligned}$$



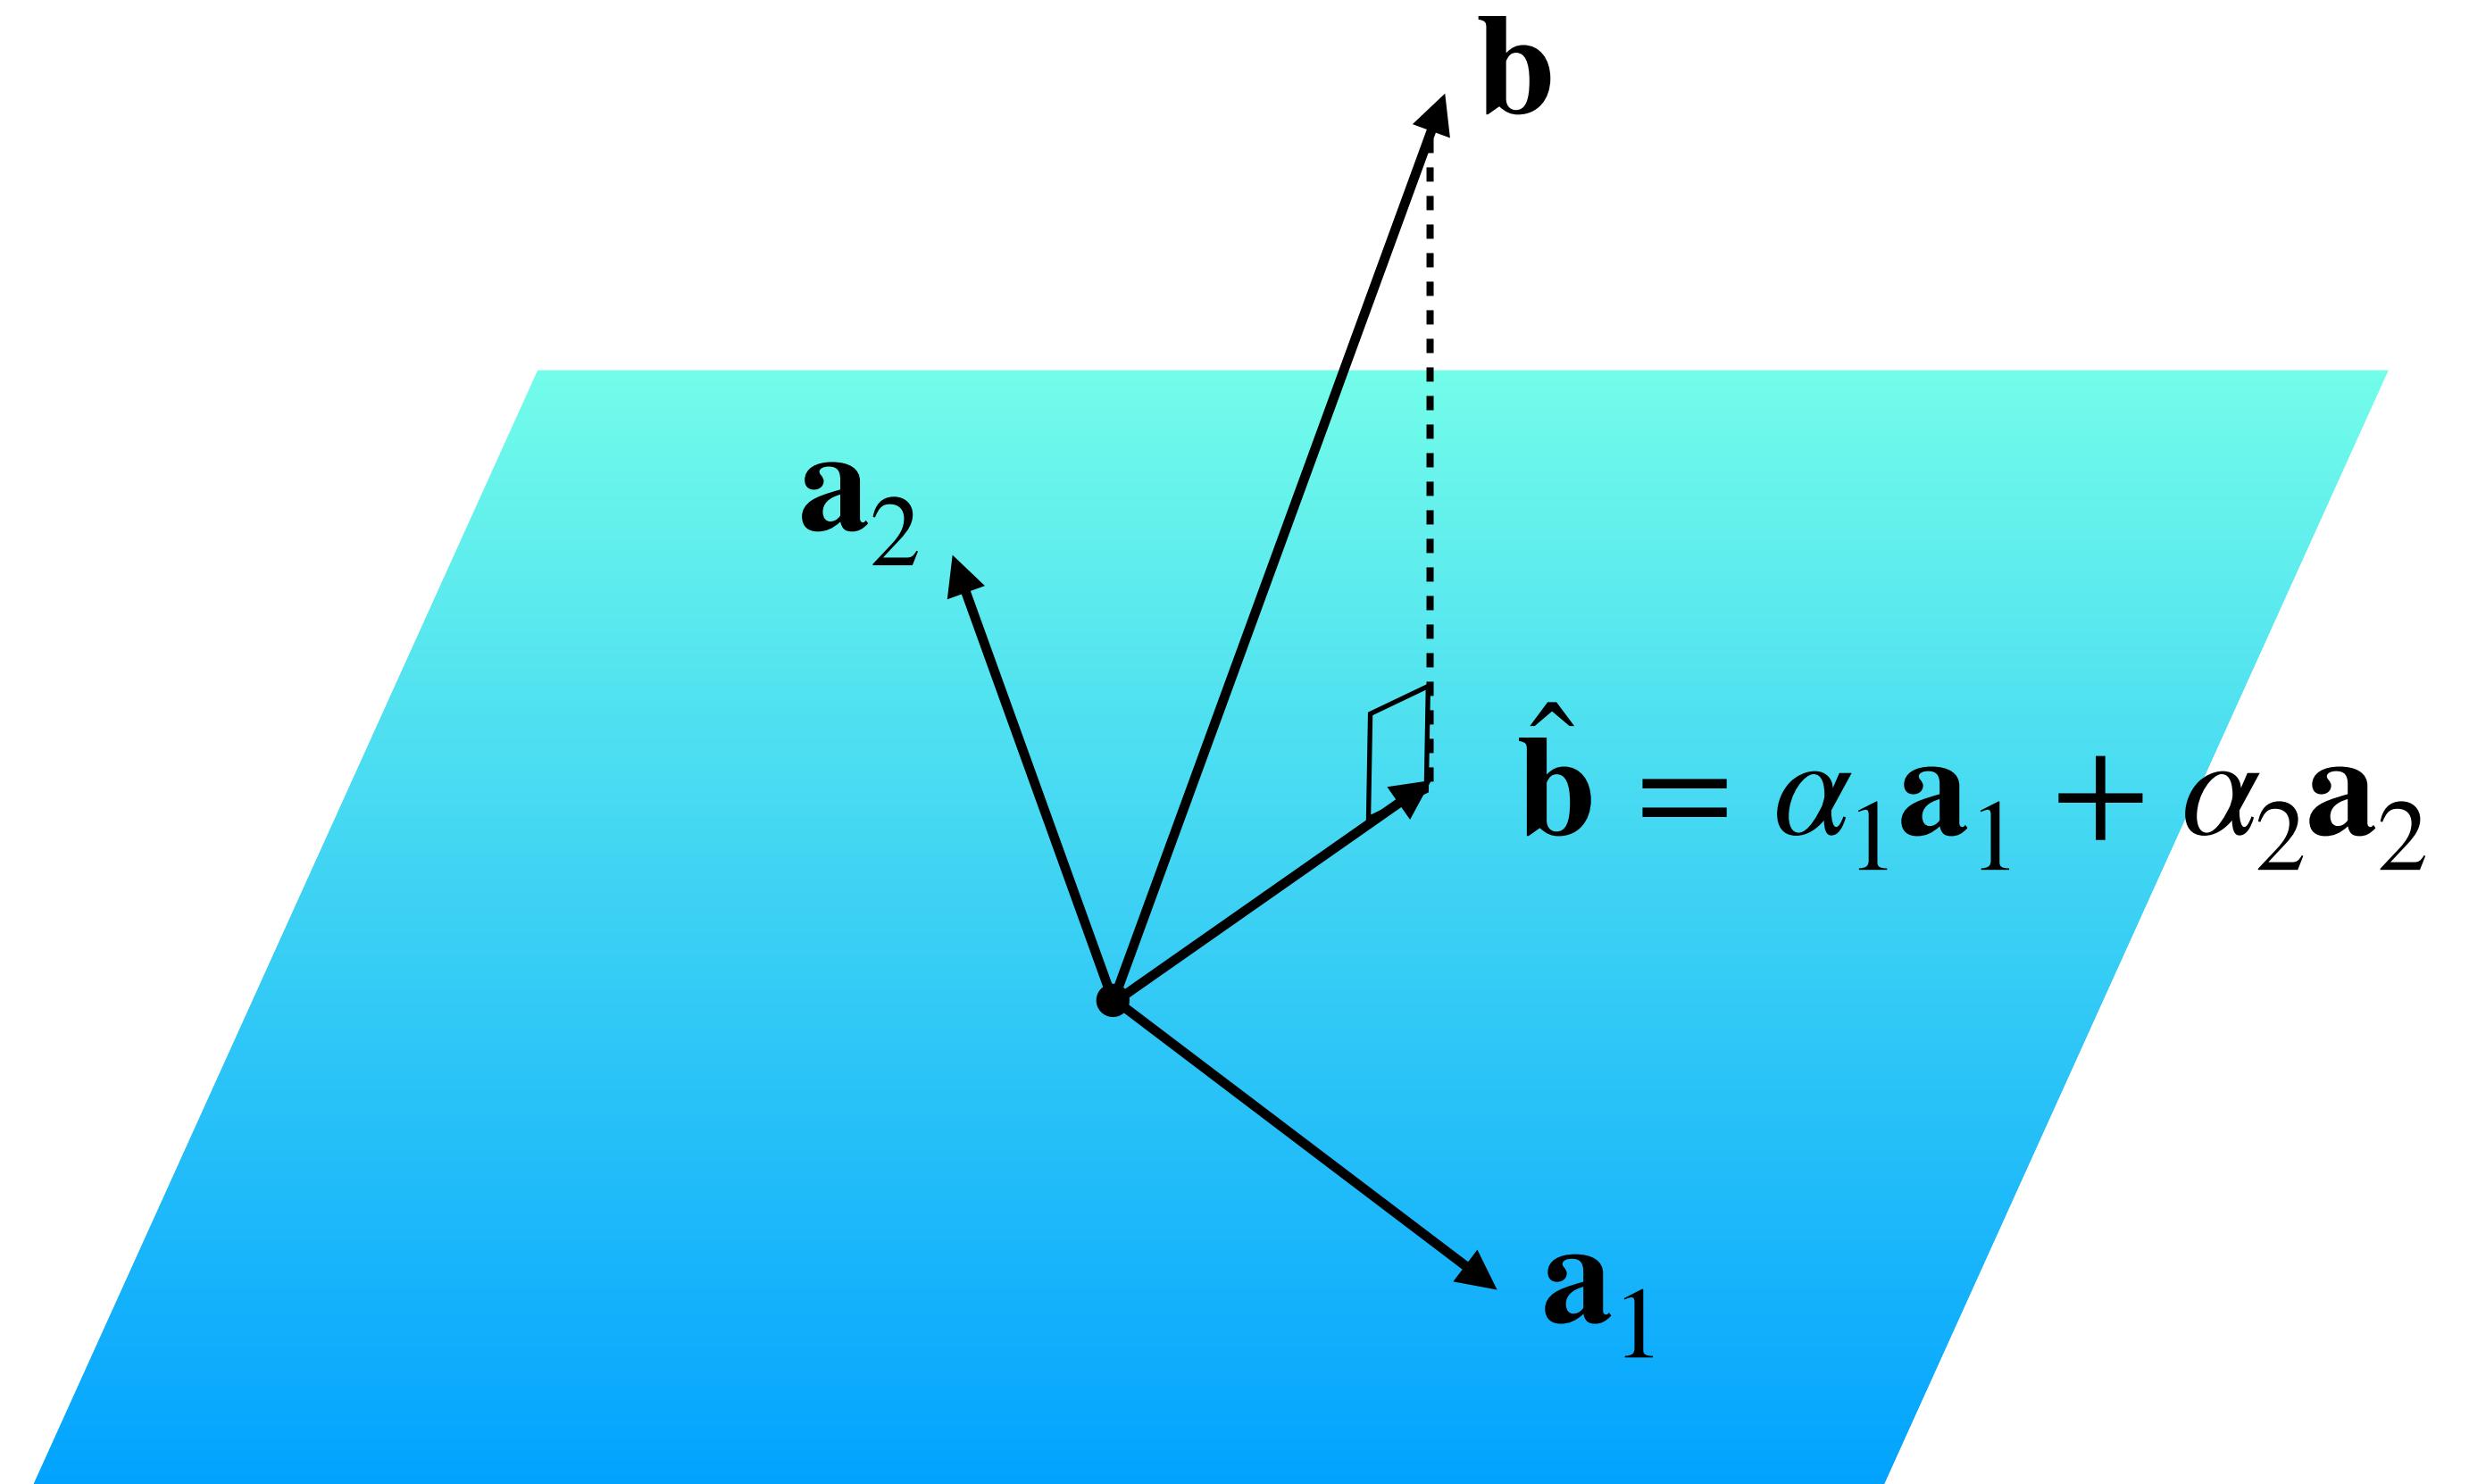
equality when $\vec{v} = \hat{\vec{y}}$

The Point



The Point

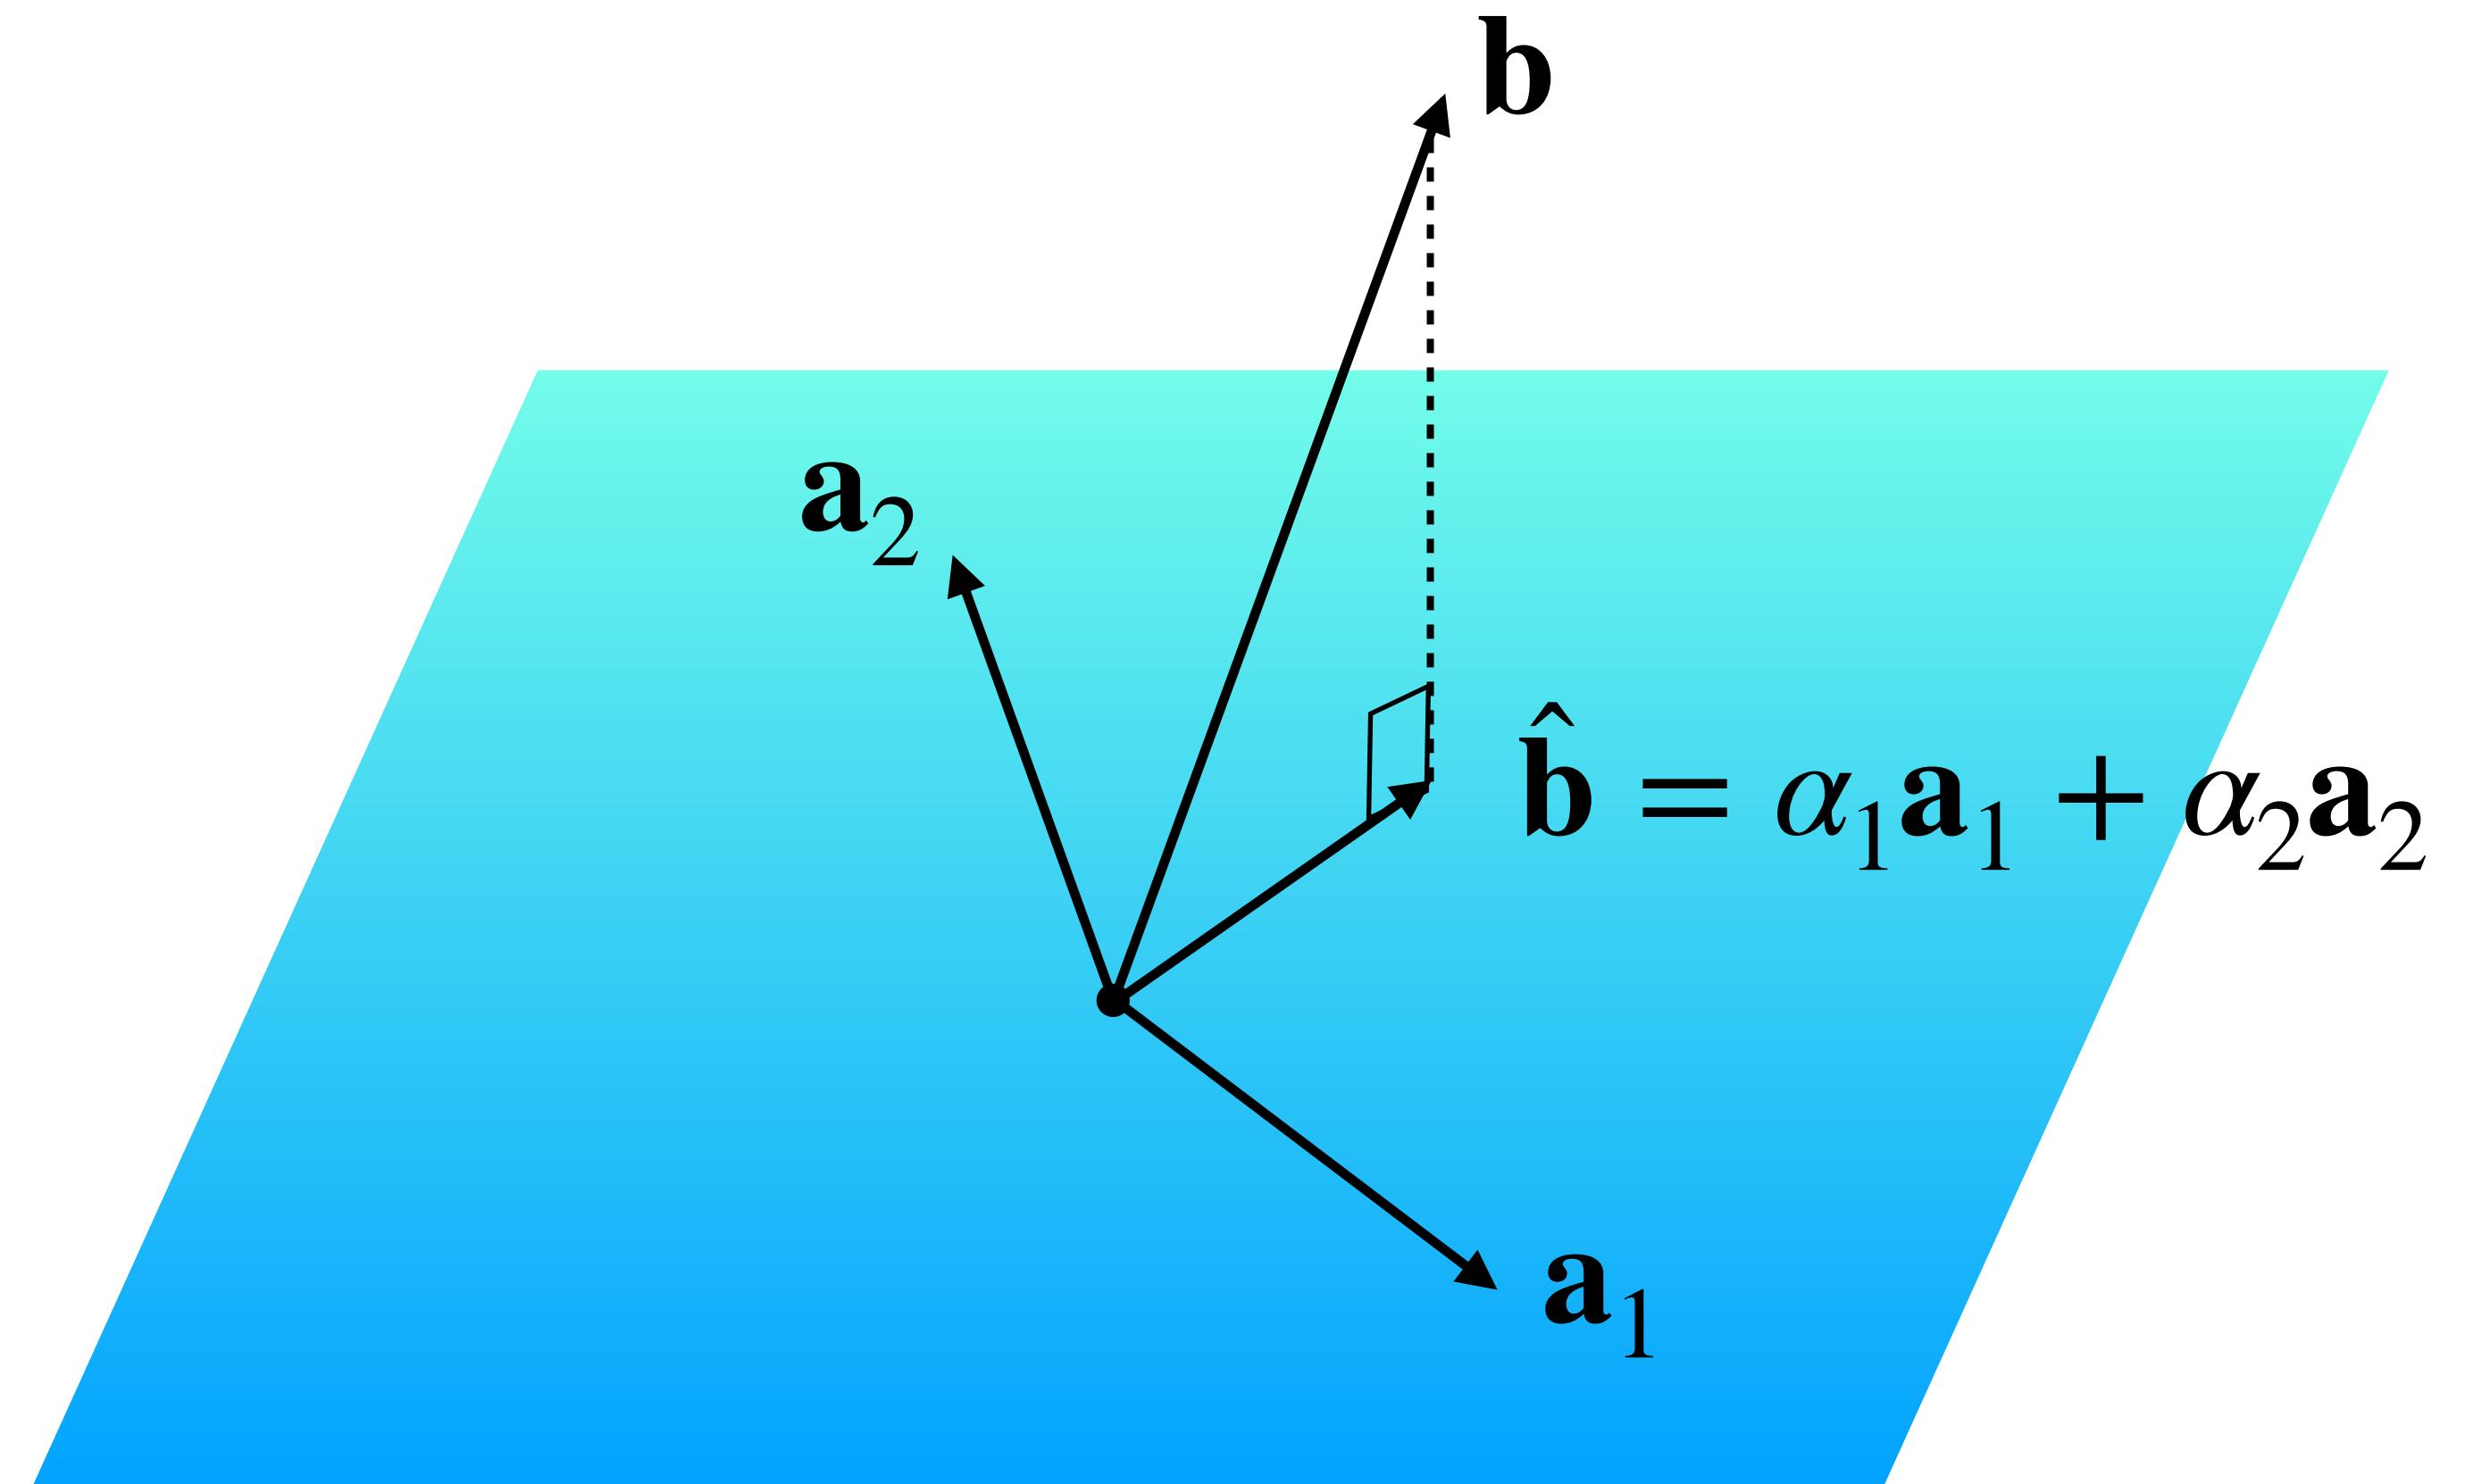
$\hat{\mathbf{b}}$ is in $Col(A)$ so $A\mathbf{x} = \hat{\mathbf{b}}$
has a solution



The Point

$\hat{\mathbf{b}}$ is in $Col(A)$ so $A\mathbf{x} = \hat{\mathbf{b}}$
has a solution

At this point, we could
call it a day:

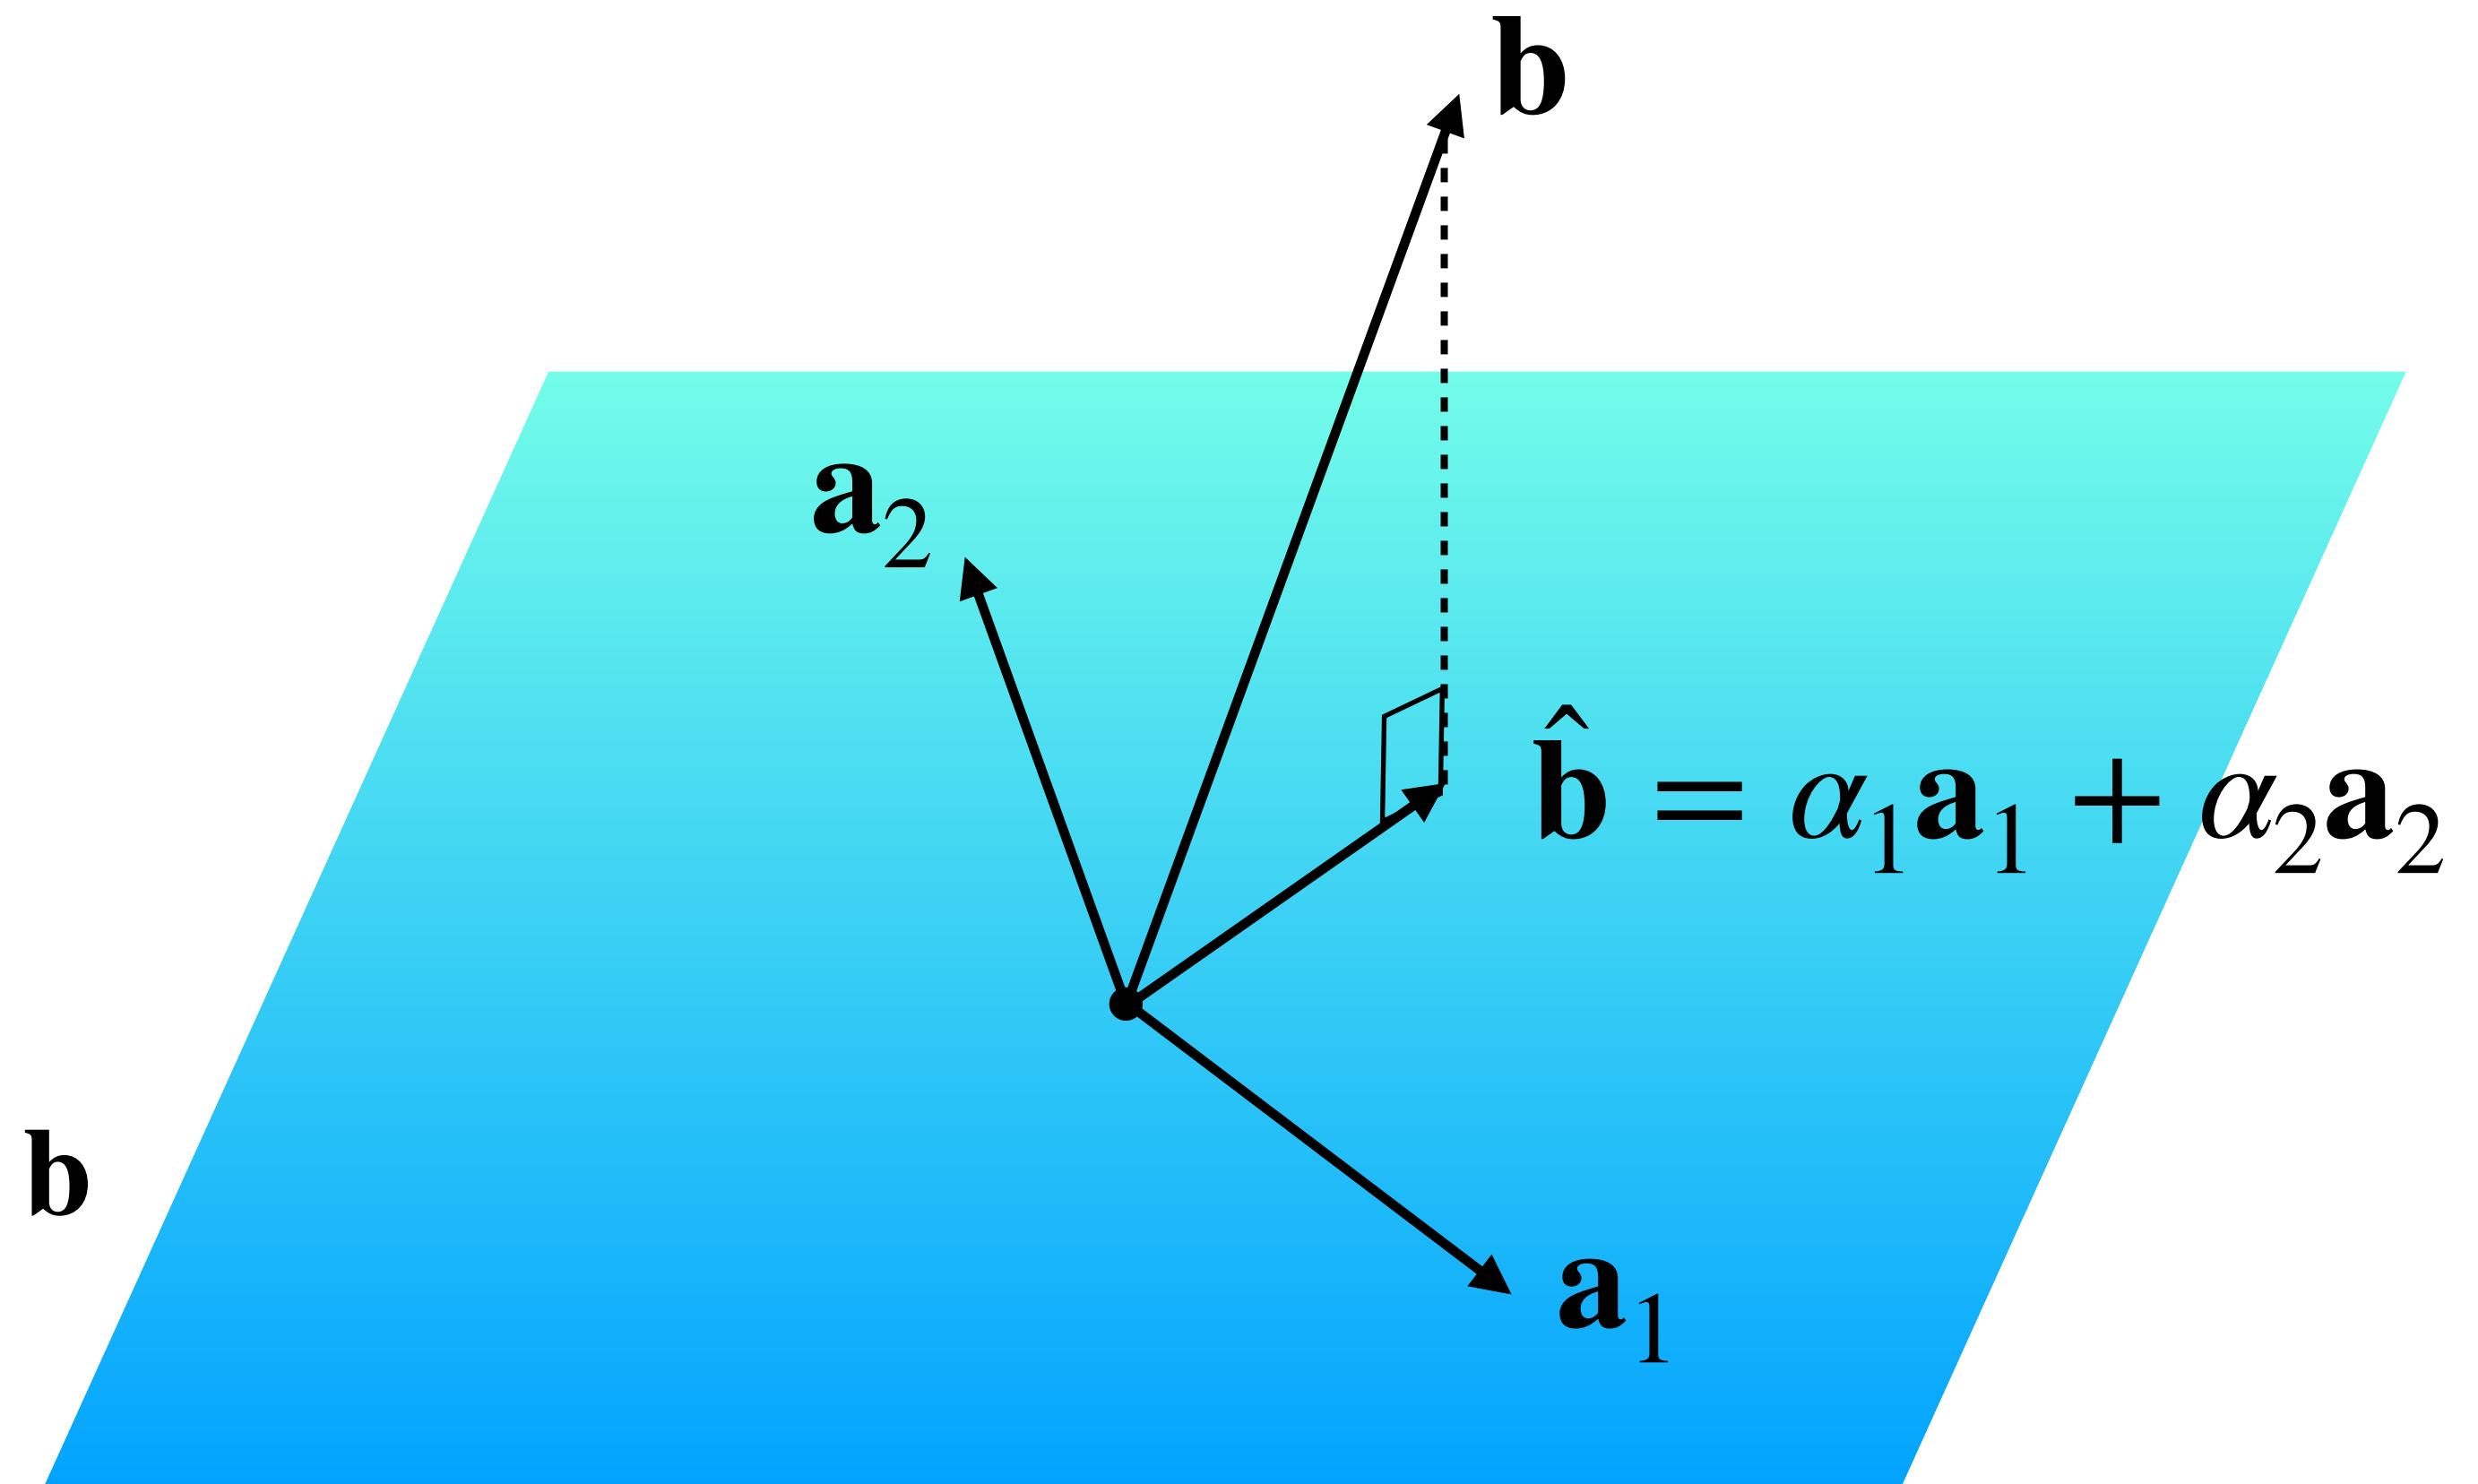


The Point

$\hat{\mathbf{b}}$ is in $Col(A)$ so $A\mathbf{x} = \hat{\mathbf{b}}$
has a solution

At this point, we could
call it a day:

Question. Find a least
squares solution to $A\mathbf{x} = \mathbf{b}$



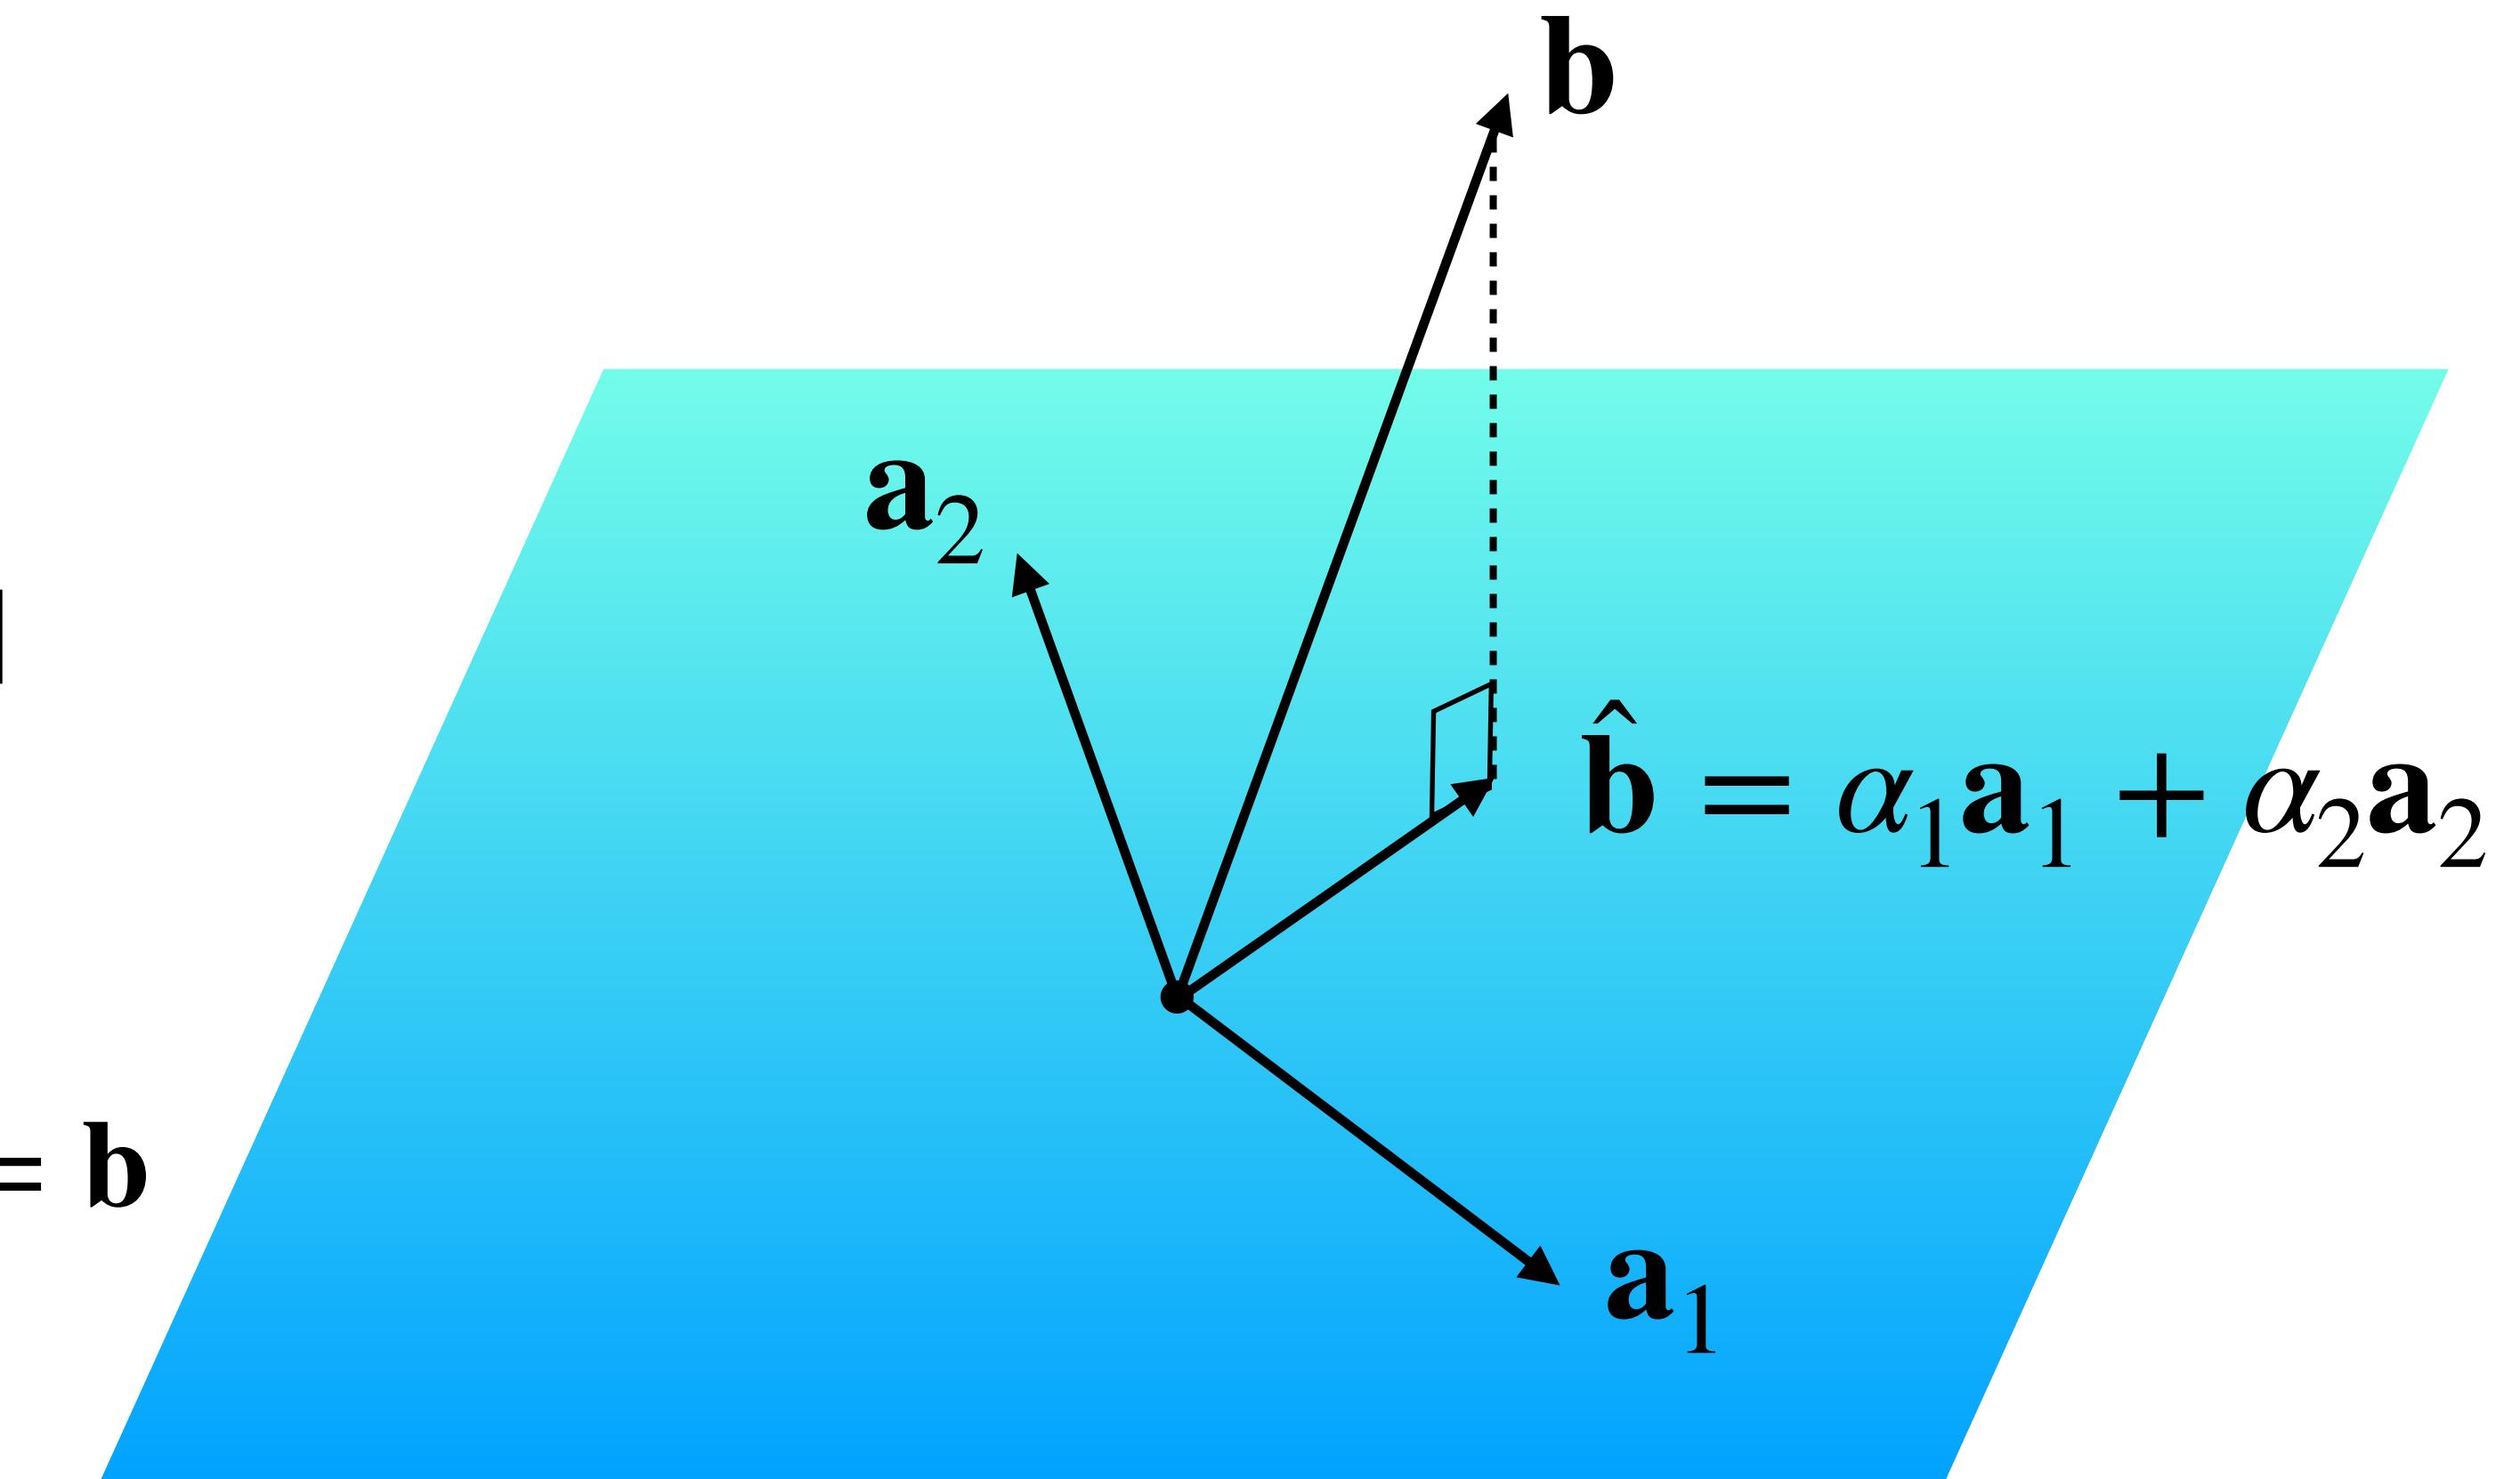
The Point

$\hat{\mathbf{b}}$ is in $Col(A)$ so $A\mathbf{x} = \hat{\mathbf{b}}$
has a solution

At this point, we could
call it a day:

Question. Find a least
squares solution to $A\mathbf{x} = \mathbf{b}$

Solution. Find $\hat{\mathbf{b}}$, then
solve $A\mathbf{x} = \hat{\mathbf{b}}$



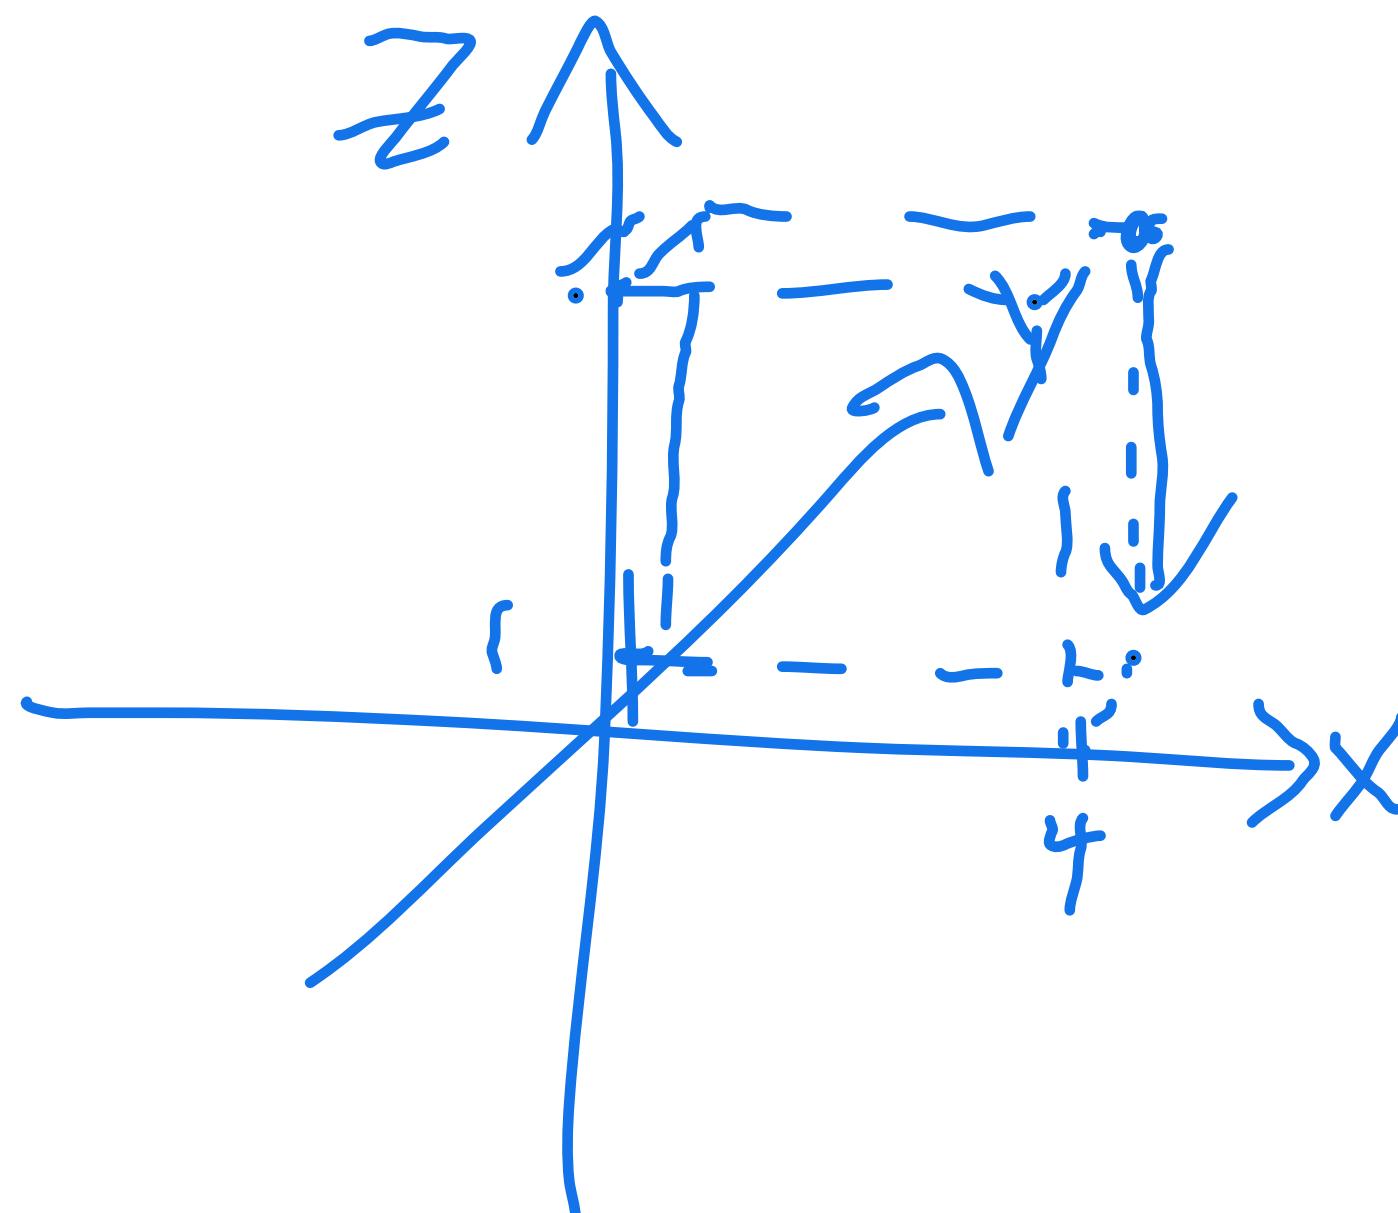
Example

$$\begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix} \begin{bmatrix} 2 \\ 3 \\ 0 \end{bmatrix} x = \begin{bmatrix} 4 \\ 1 \\ 4 \end{bmatrix} \leftarrow \vec{b}$$

$$GIA = \left\{ \begin{bmatrix} x \\ x \\ 0 \end{bmatrix} \mid x \in \mathbb{R} \right\}$$

\geq entries are all 0

Let's determine the least squares solution for the above system:



$$\hat{b} = \begin{bmatrix} 4 \\ 1 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 4 \\ -1 & 3 & 1 \end{bmatrix} \sim \begin{bmatrix} 1 & 2 & 4 \\ 0 & 5 & 5 \end{bmatrix}$$

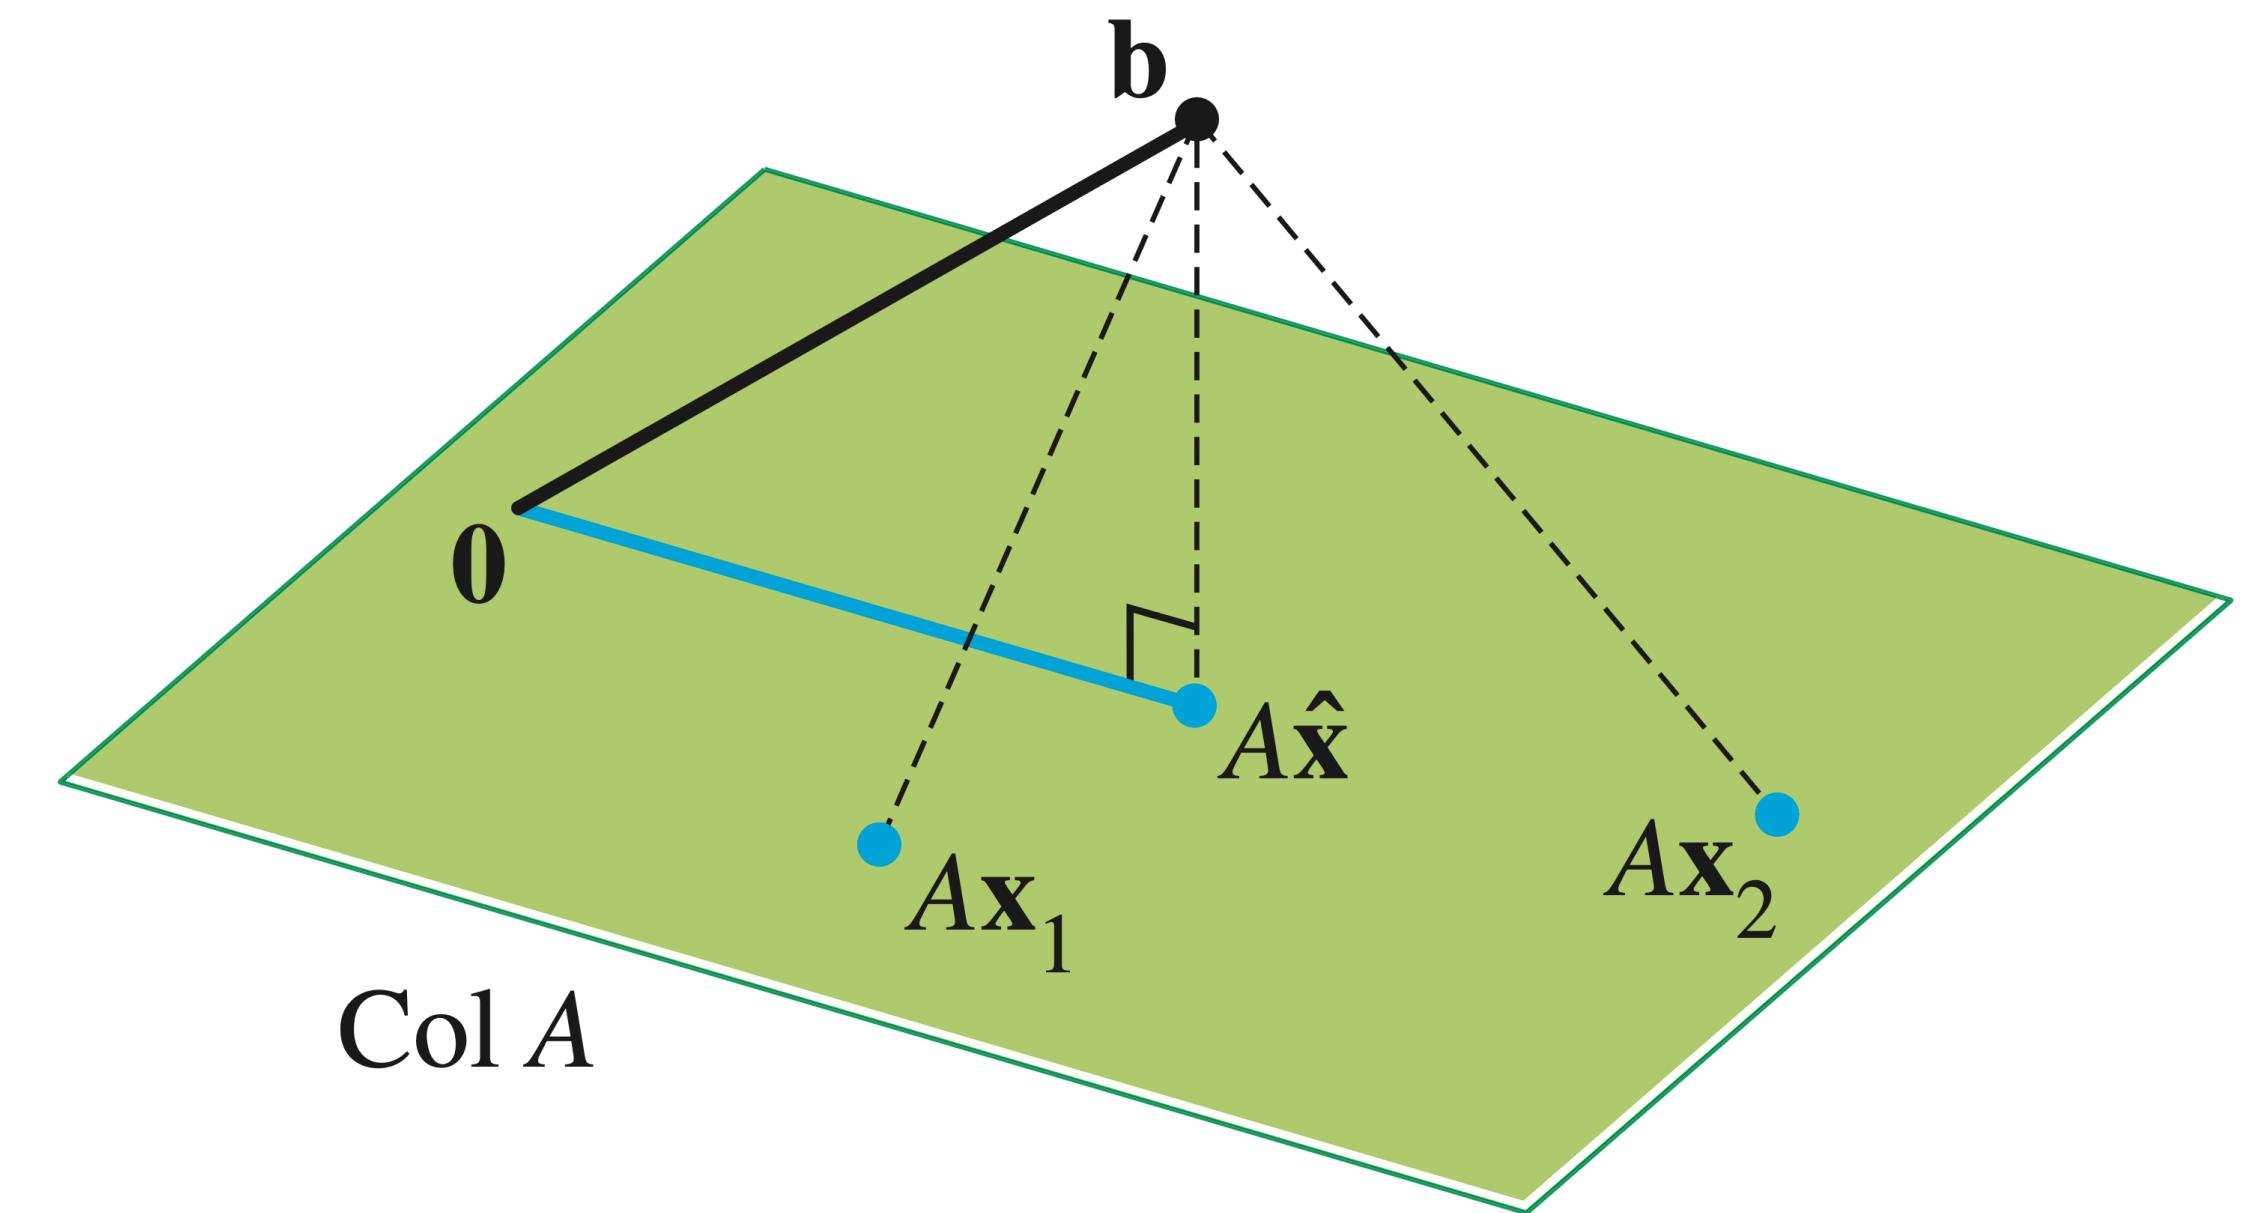
$$x = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 2 \\ 0 & 1 & 1 \end{bmatrix} \sim \begin{bmatrix} 1 & 2 & 4 \\ 0 & 1 & 1 \end{bmatrix}$$

least-squares sol'n

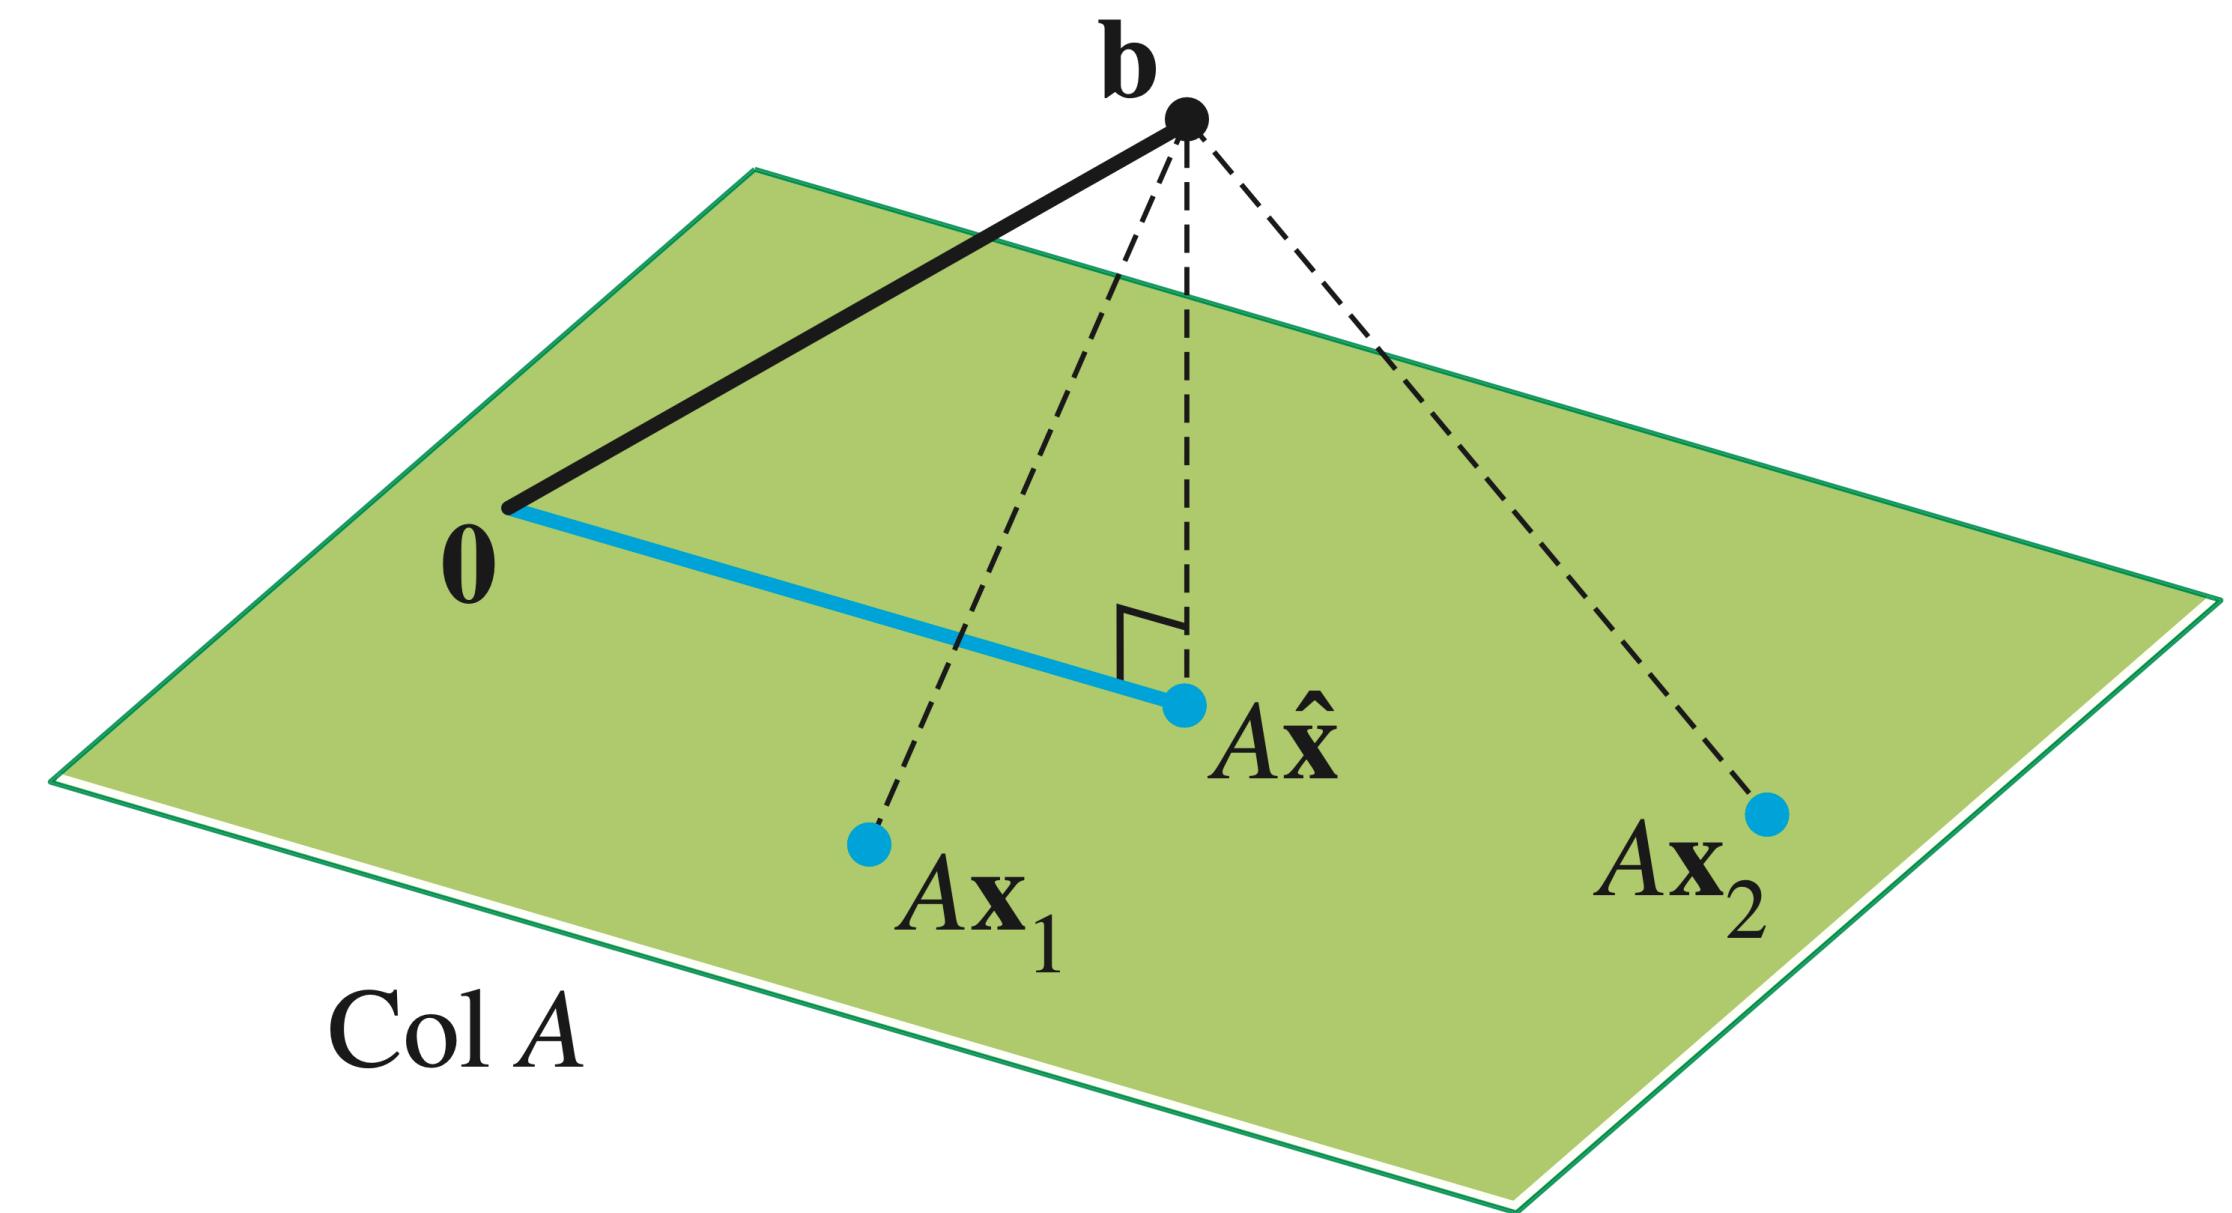
The Normal Equations

A Couple Observations



A Couple Observations

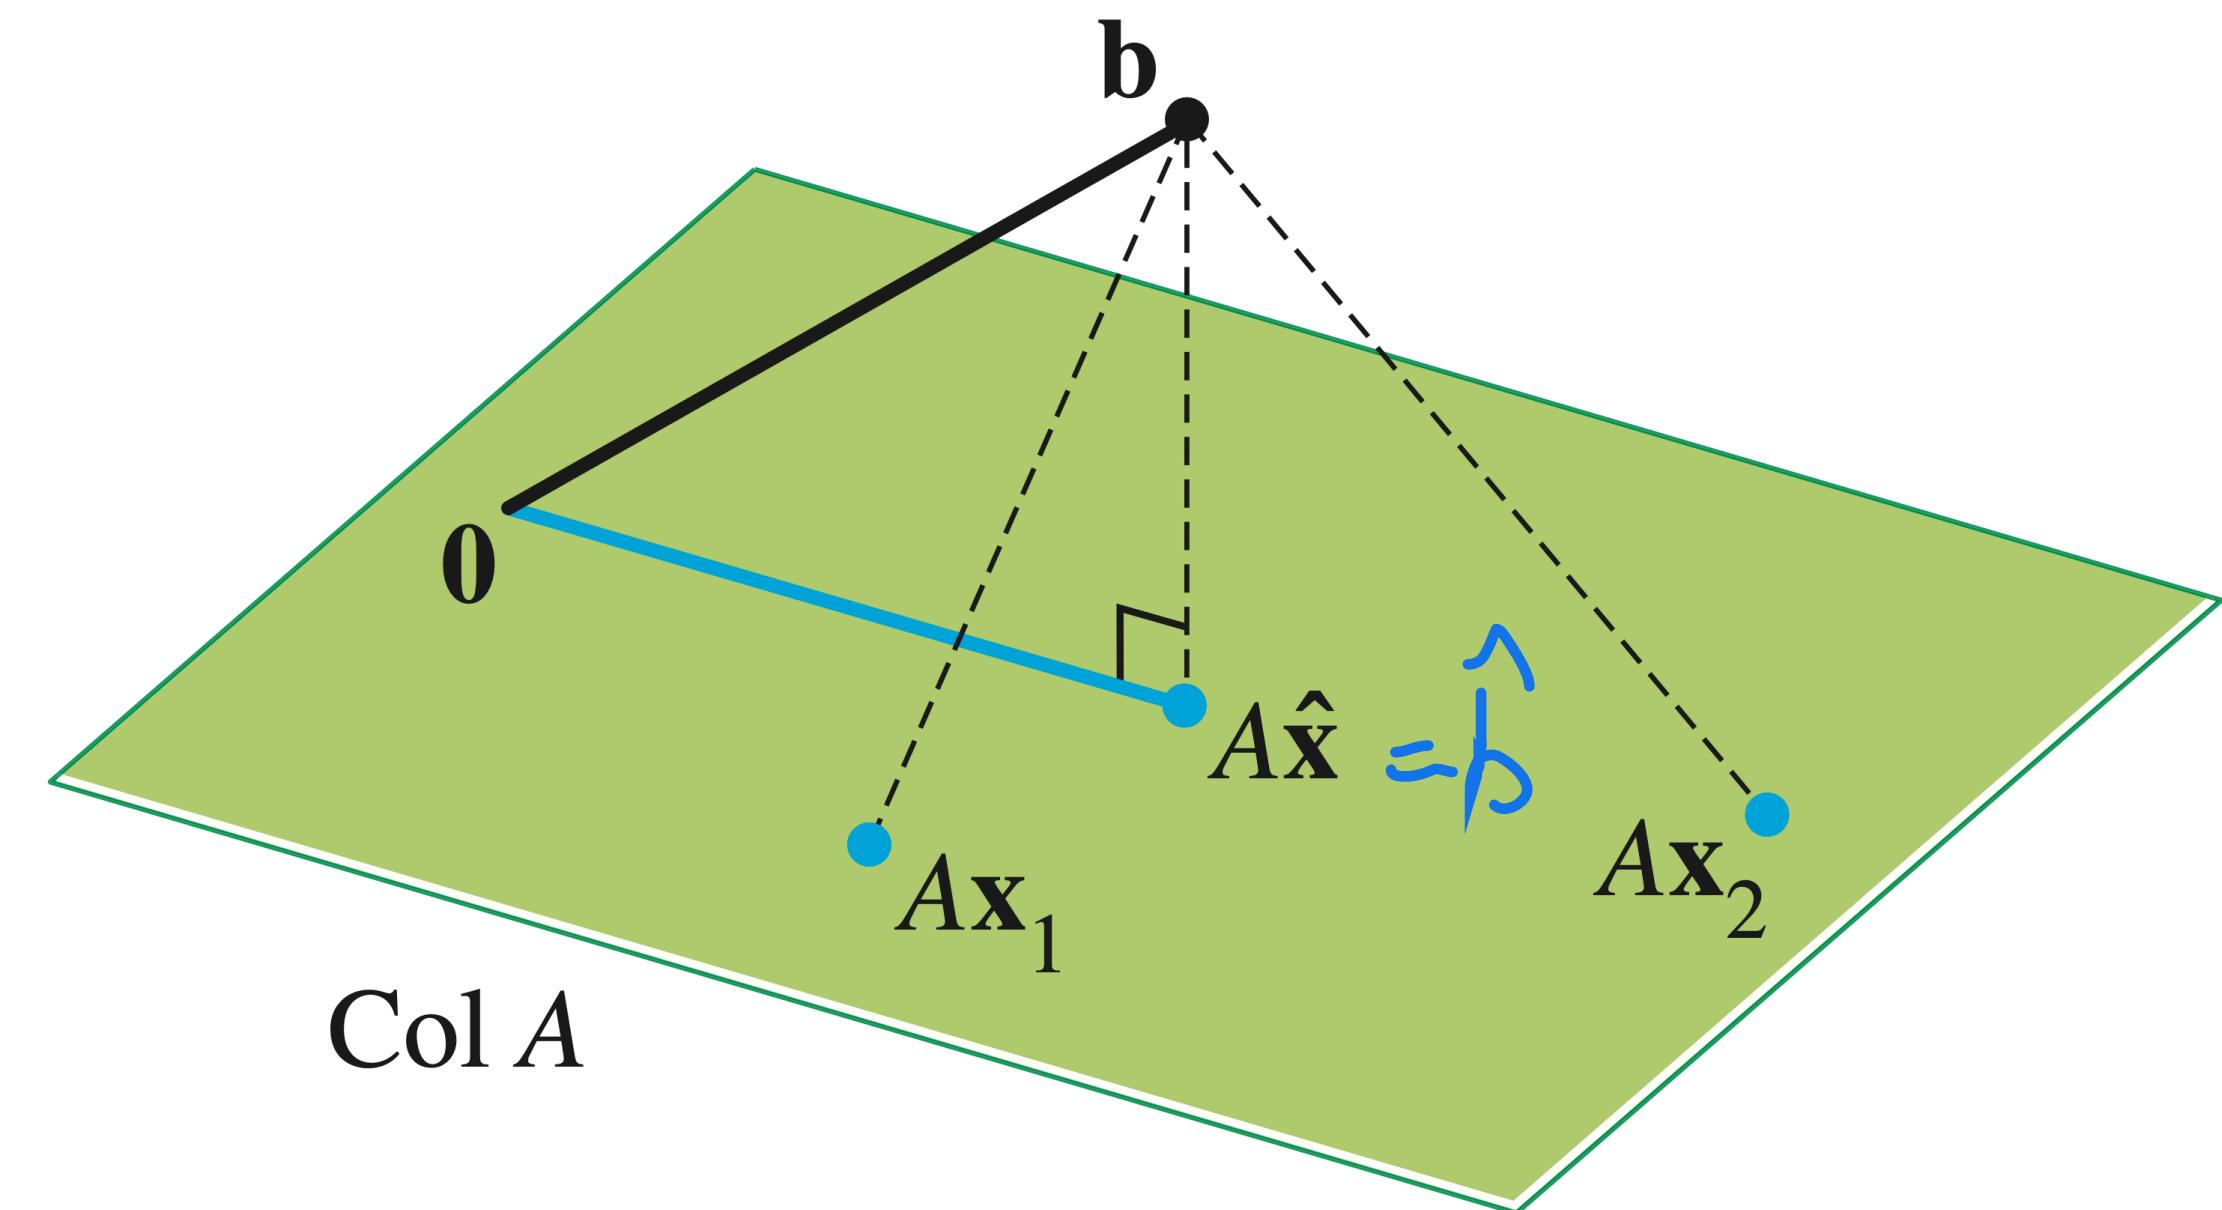
Suppose that \hat{x} is a least squares solution to A , so $A\hat{x} = \hat{b}$



A Couple Observations

Suppose that $\hat{\mathbf{x}}$ is a least squares solution to A , so $A\hat{\mathbf{x}} = \hat{\mathbf{b}}$

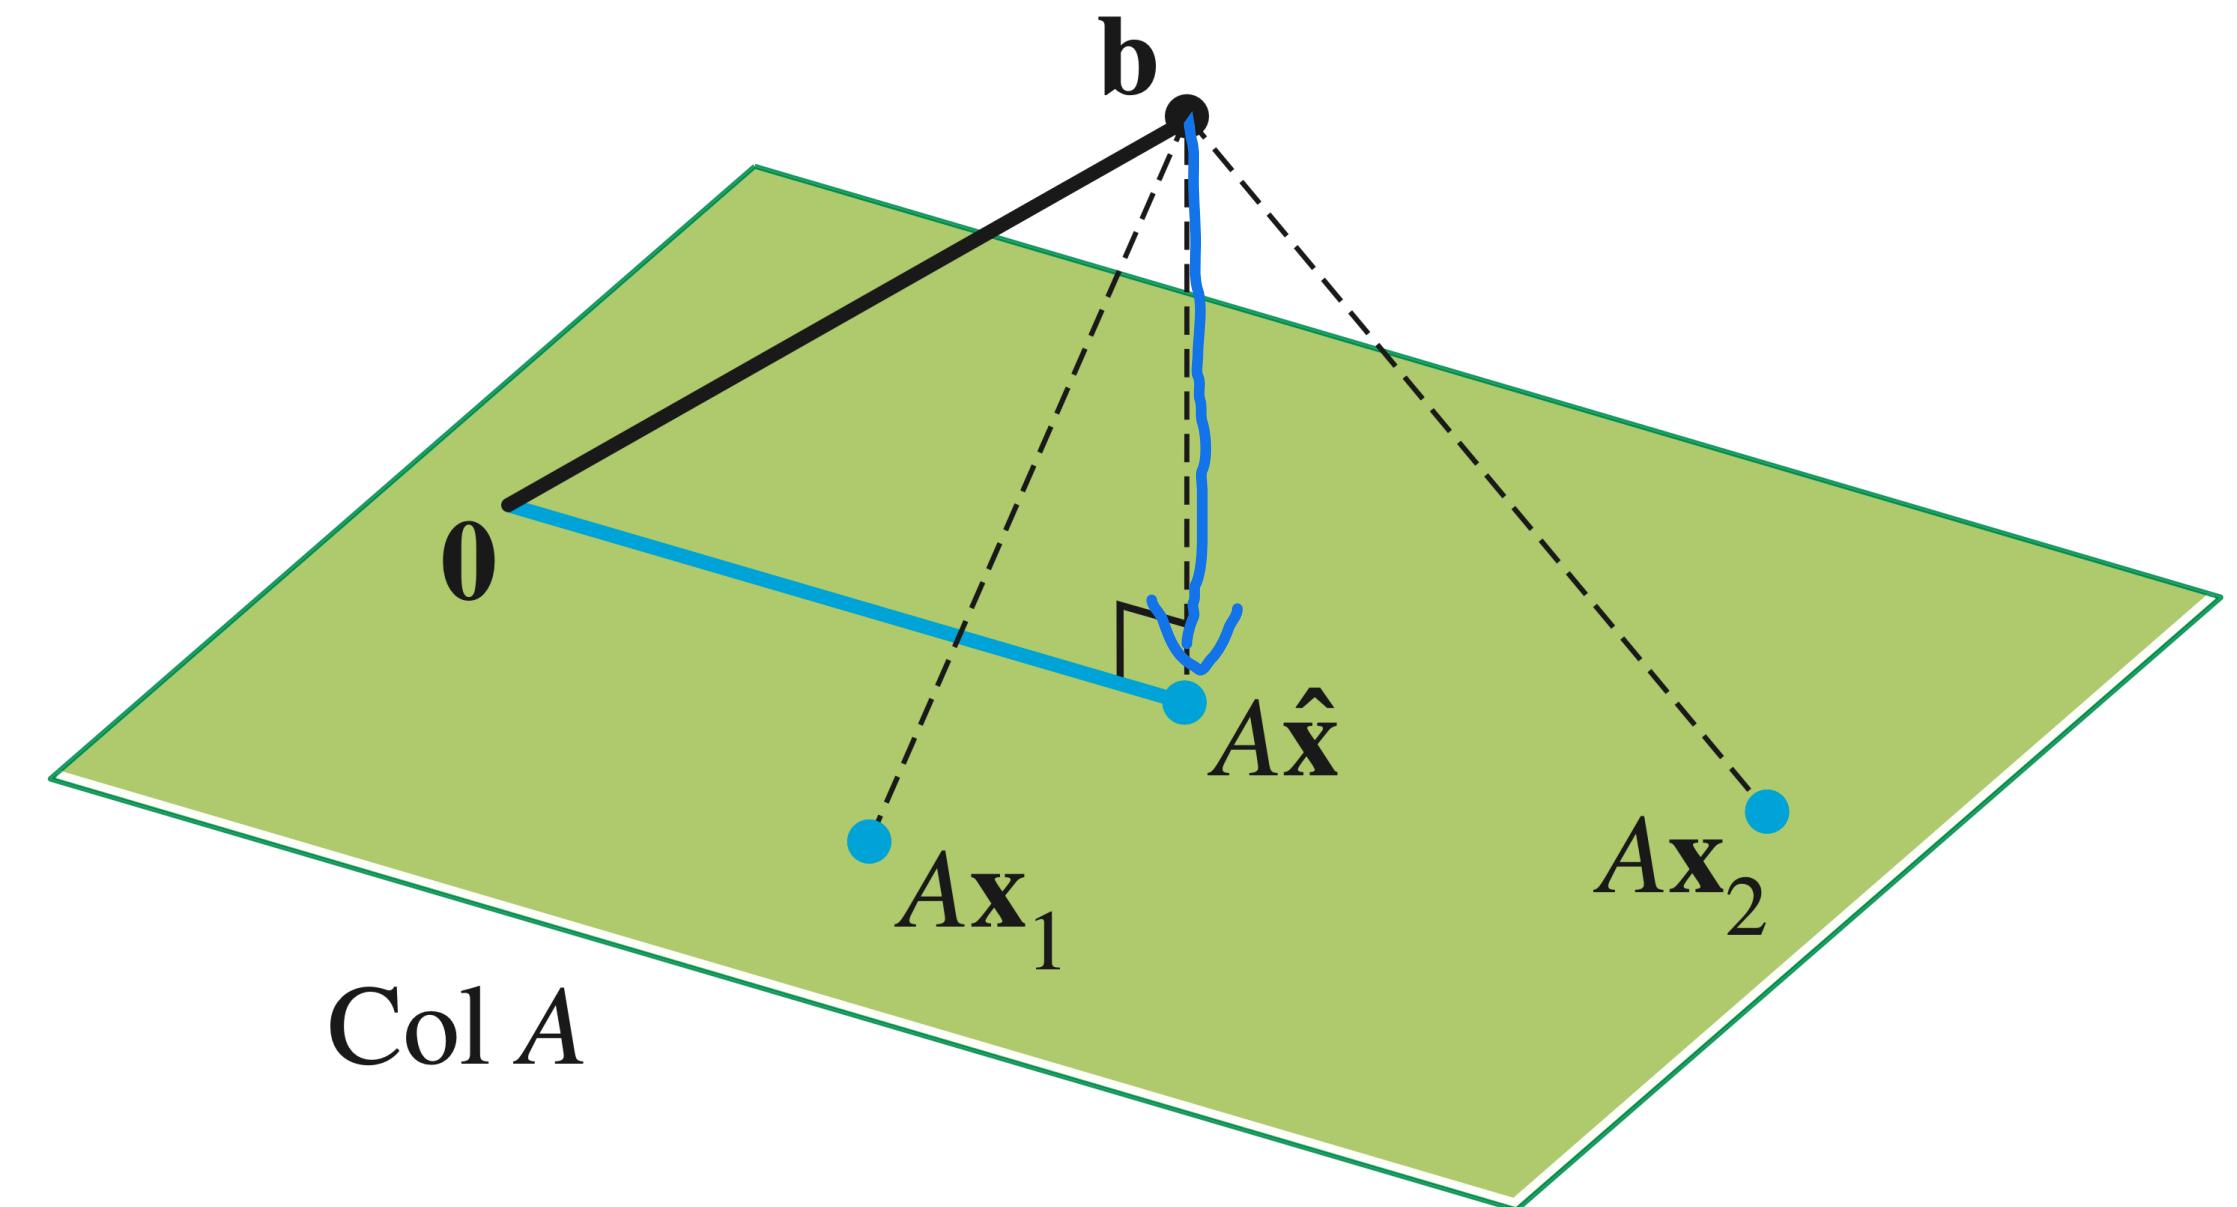
- $\hat{\mathbf{b}} - \mathbf{b}$ is orthogonal to $Col(A)$



A Couple Observations

Suppose that \hat{x} is a least squares solution to A , so $A\hat{x} = \hat{b}$

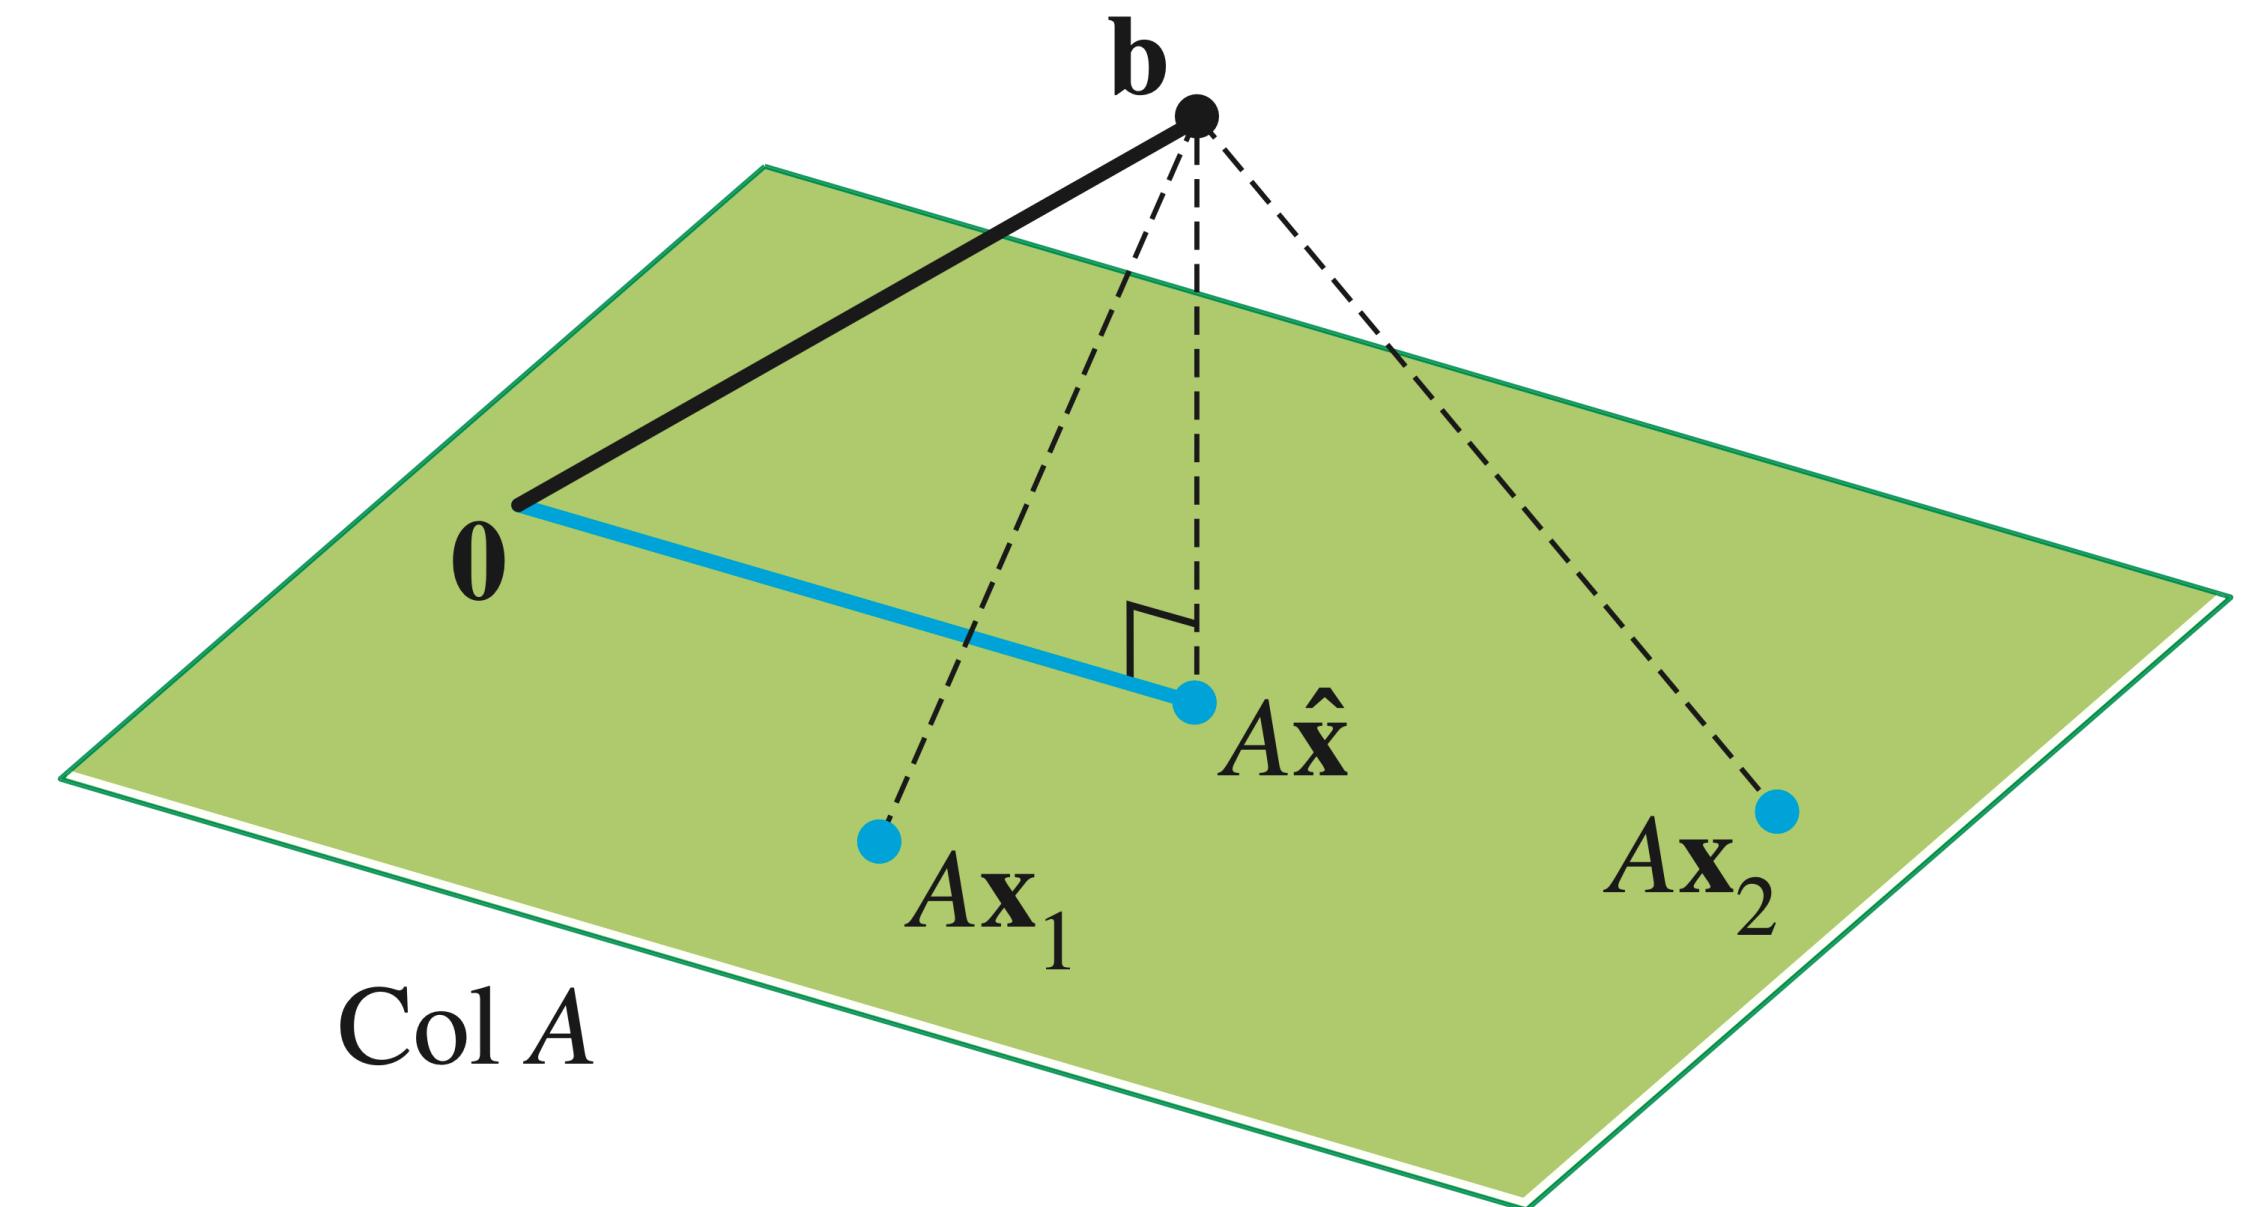
- $\hat{b} - b$ is orthogonal to $Col(A)$
- $A\hat{x} - b$ is orthogonal to $Col(A)$



A Couple Observations

Suppose that \hat{x} is a least squares solution to A , so $A\hat{x} = \hat{b}$

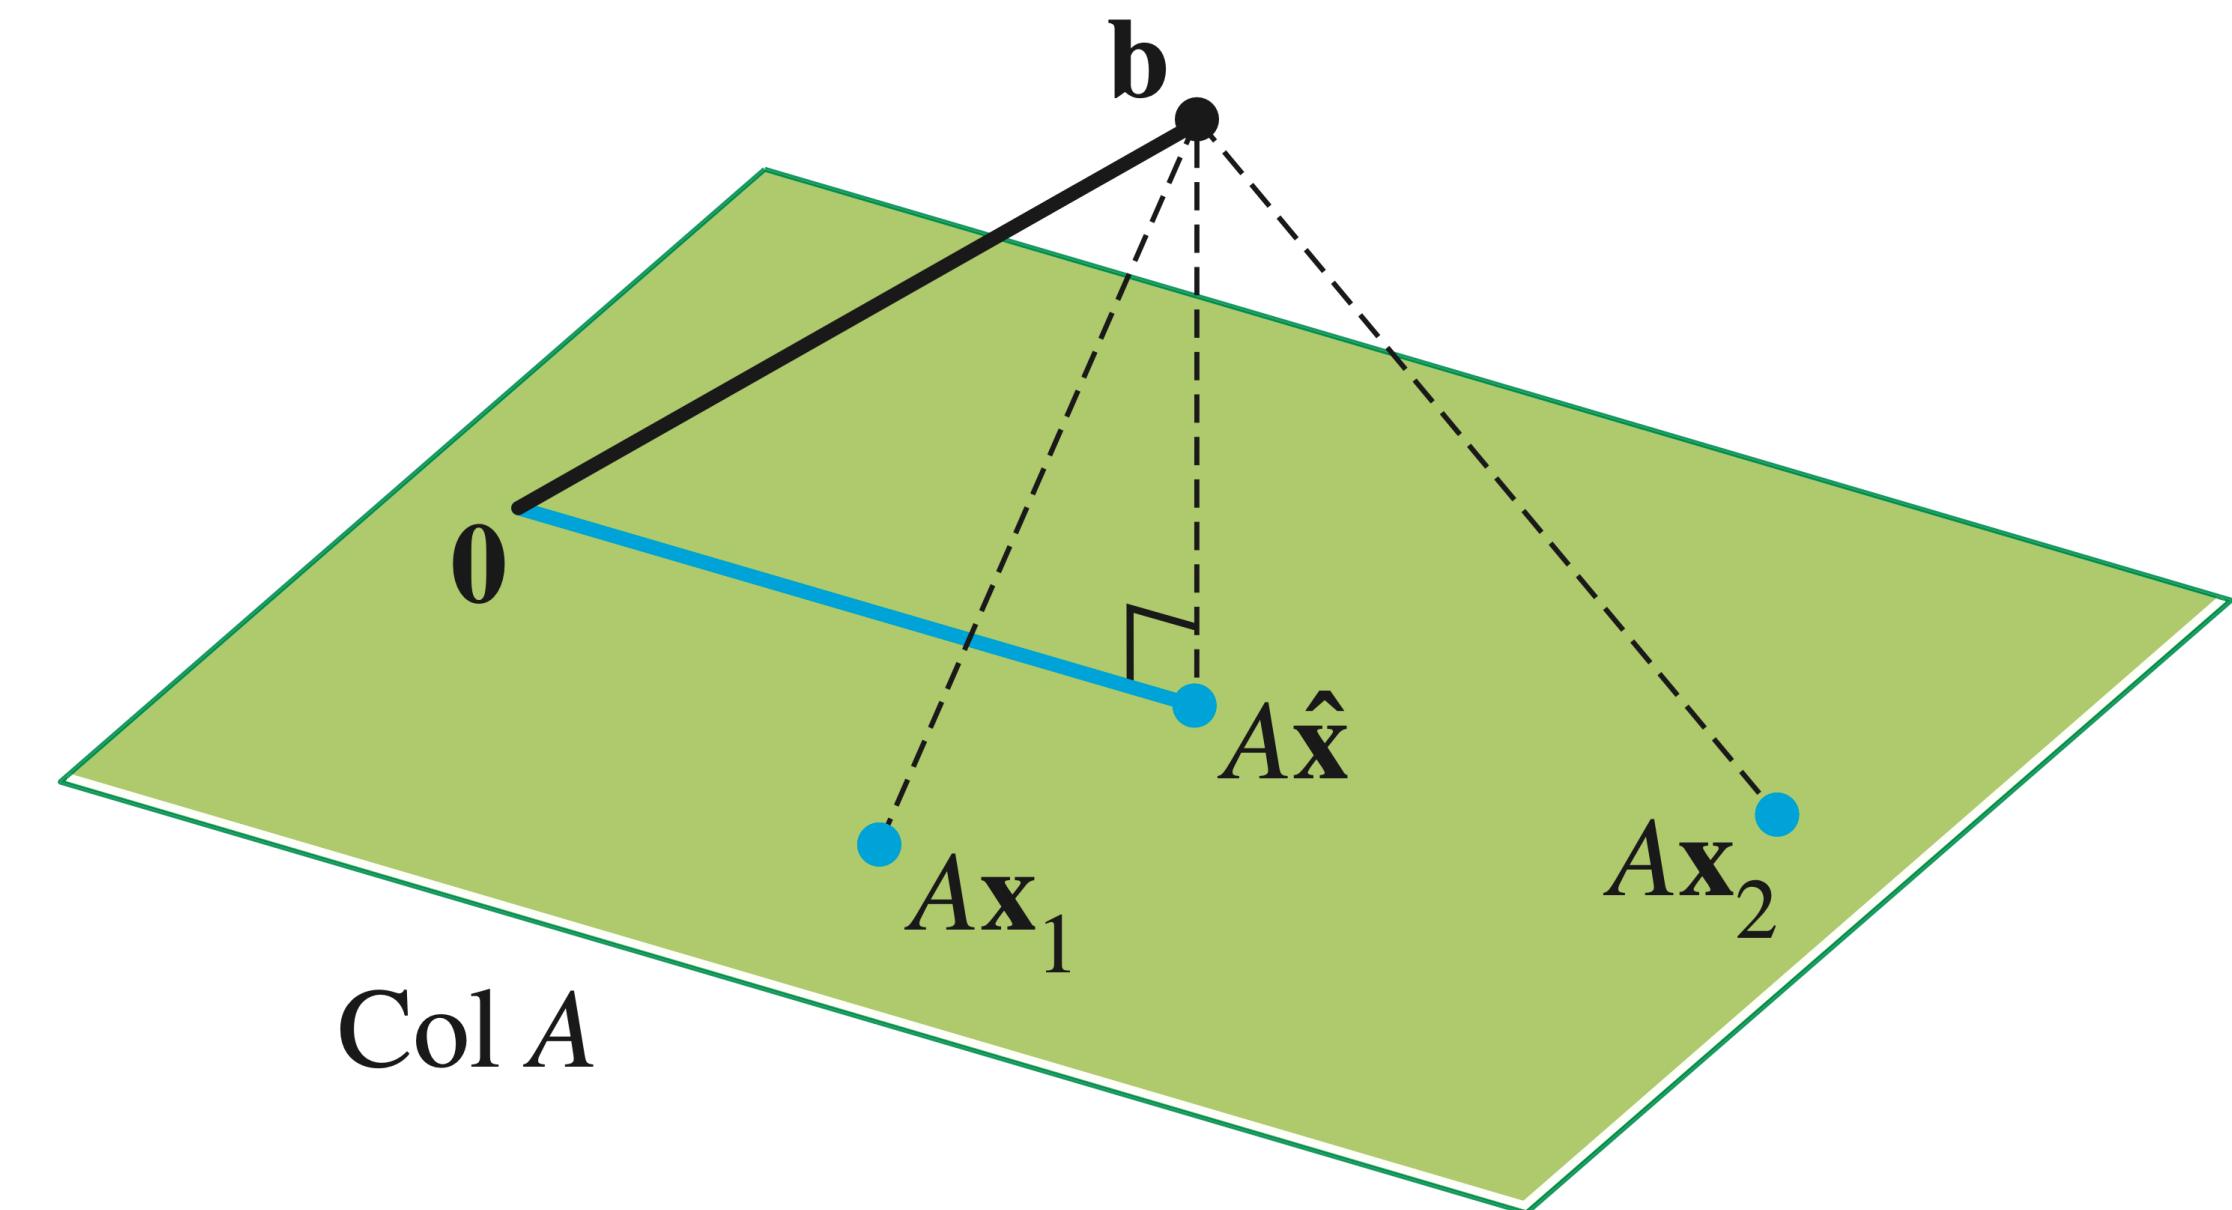
- $\hat{b} - b$ is orthogonal to $Col(A)$
- $A\hat{x} - b$ is orthogonal to $Col(A)$
- If $A = [a_1 \ a_2 \ \dots \ a_n]$ then $A\hat{x} - b$ is orthogonal to each a_1, a_2, \dots, a_n



A Couple Observations

Suppose that $\hat{\mathbf{x}}$ is a least squares solution to A , so $A\hat{\mathbf{x}} = \hat{\mathbf{b}}$

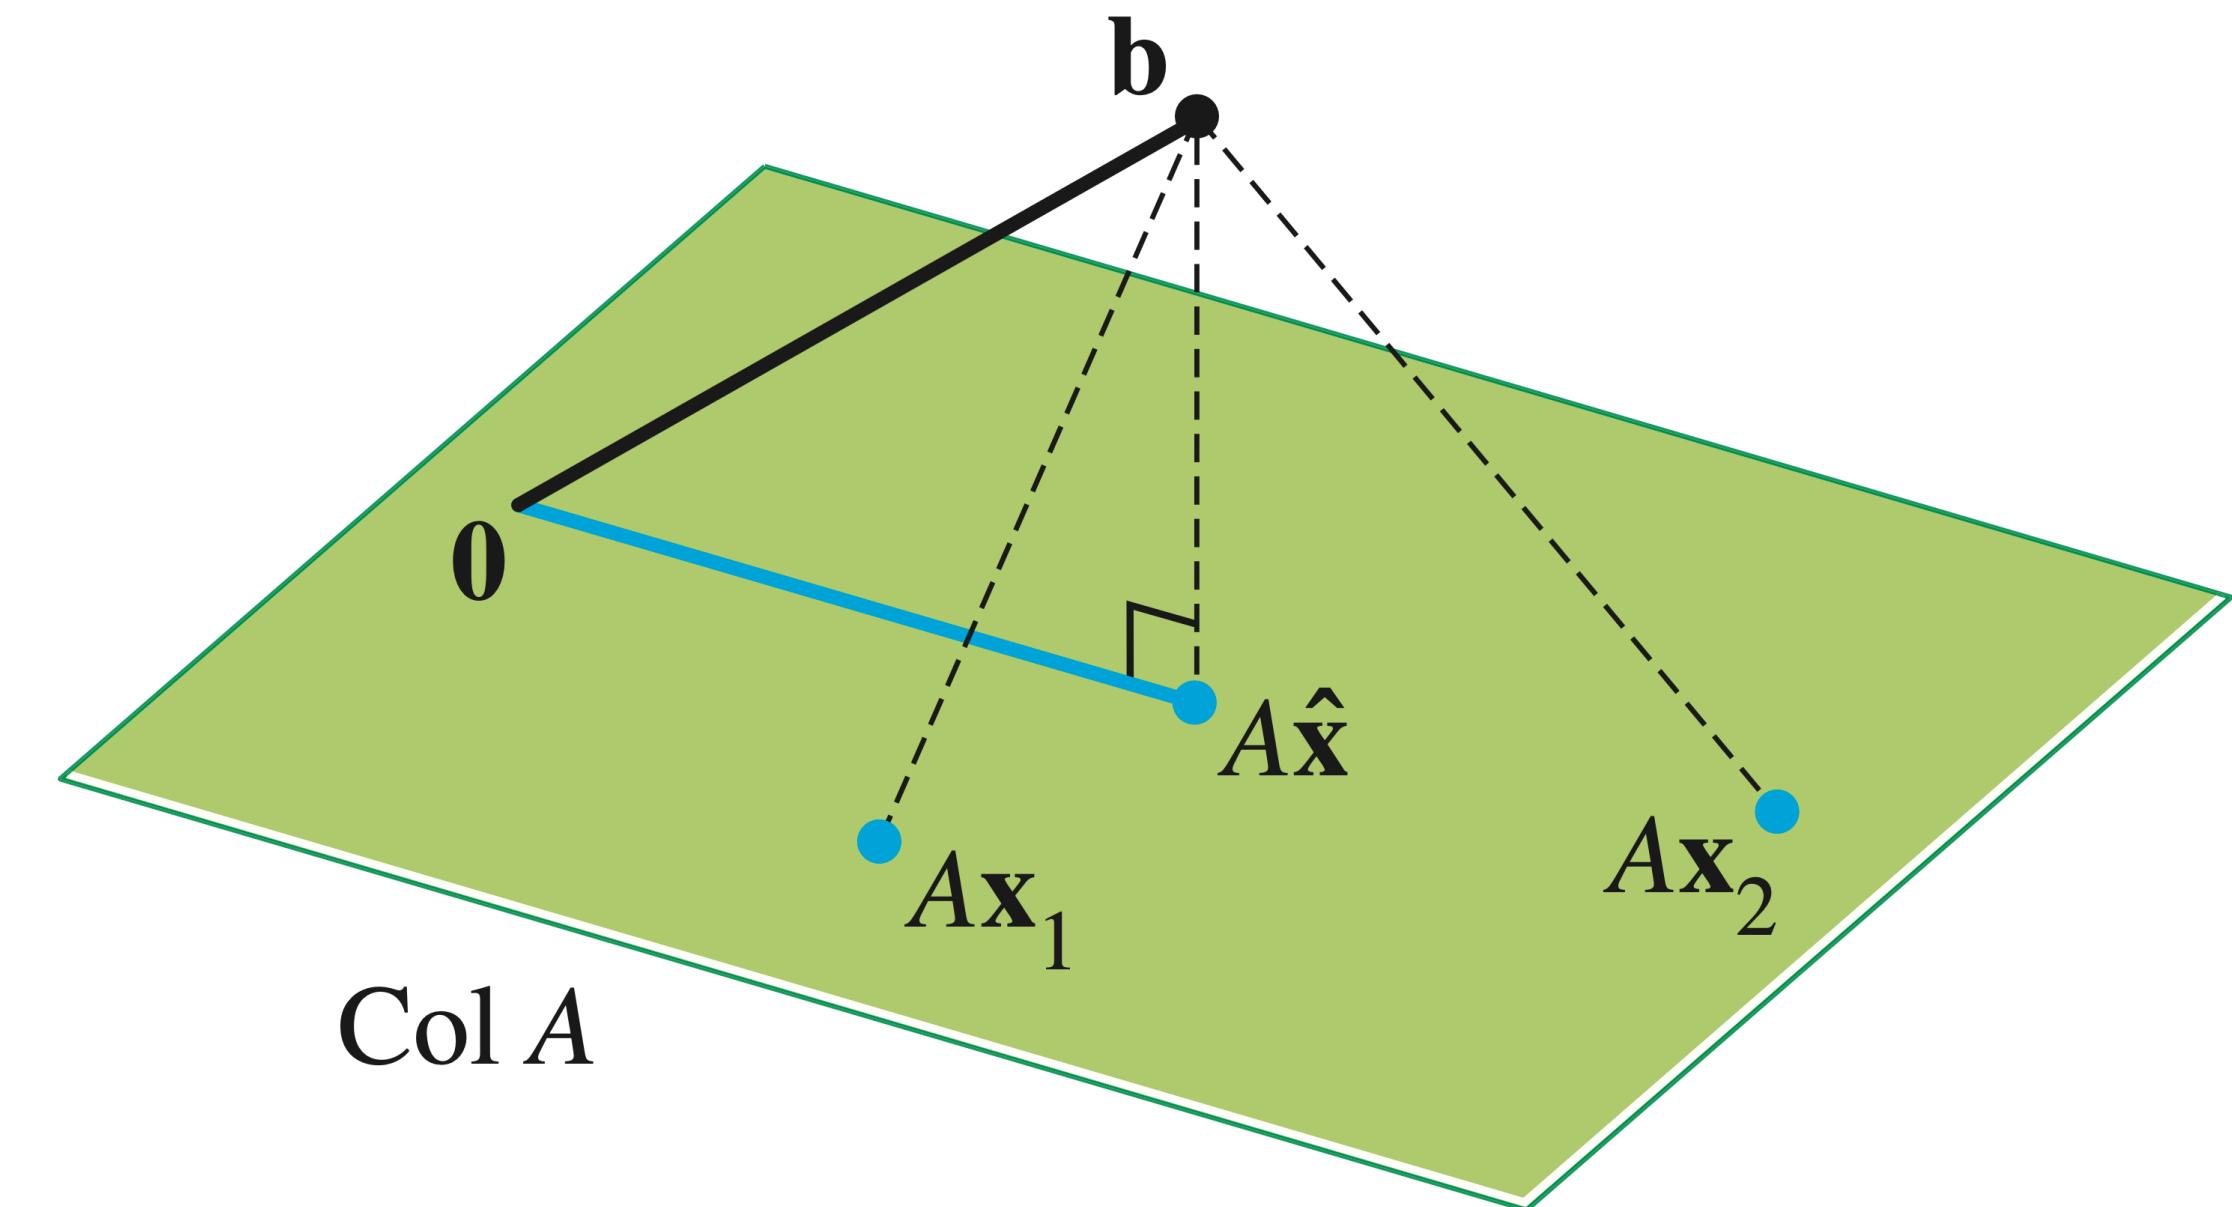
- $\hat{\mathbf{b}} - \mathbf{b}$ is orthogonal to $Col(A)$
- $A\hat{\mathbf{x}} - \mathbf{b}$ is orthogonal to $Col(A)$
- If $A = [\mathbf{a}_1 \ \mathbf{a}_2 \ \dots \ \mathbf{a}_n]$ then $A\hat{\mathbf{x}} - \mathbf{b}$ is orthogonal to each $\underline{\mathbf{a}_1}, \underline{\mathbf{a}_2}, \dots, \underline{\mathbf{a}_n}$
- $\mathbf{a}_i^T(A\hat{\mathbf{x}} - \mathbf{b}) = 0$



A Couple Observations

Suppose that \hat{x} is a least squares solution to A , so $A\hat{x} = \hat{b}$

- $\hat{b} - b$ is orthogonal to $Col(A)$
- $A\hat{x} - b$ is orthogonal to $Col(A)$
- If $A = [\mathbf{a}_1 \ \mathbf{a}_2 \ \dots \ \mathbf{a}_n]$ then $A\hat{x} - b$ is orthogonal to each $\mathbf{a}_1, \mathbf{a}_2, \dots, \mathbf{a}_n$
- $\mathbf{a}_i^T(A\hat{x} - b) = 0$
- $A^T(A\hat{x} - b) = 0$ *← matrix-vector equation*



A bit more magic

Let's simplify $A^T(A\hat{x} - b)$:

$$A^T A \hat{x} - A^T \vec{b} = 0$$
$$\underbrace{A^T A \hat{x}}_{=} = \underbrace{A^T \vec{b}}_{}$$

The Normal Equations

The Normal Equations

Theorem. The set of least-squares solutions of $Ax = b$ is the same as the set of solutions to

$$\underline{A^T A \hat{x} = A^T b}$$

The Normal Equations

Theorem. The set of least-squares solutions of $A\mathbf{x} = \mathbf{b}$ is the same as the set of solutions to

$$A^T A \mathbf{x} = A^T \mathbf{b}$$

In particular, this set of solutions is nonempty

The Normal Equations

Theorem. The set of least-squares solutions of $Ax = b$ is the same as the set of solutions to

$$A^T A x = A^T b$$

In particular, this set of solutions is nonempty

(We just showed that if \hat{x} is a least squares solution then $A^T A \hat{x} = A^T b$)

Example

$$A = \begin{bmatrix} 4 & 0 \\ 0 & 2 \\ 1 & 1 \end{bmatrix} \quad b = \begin{bmatrix} 2 \\ 0 \\ 11 \end{bmatrix}$$

Let's find the normal equations for $Ax = b$:

$$A^T A = \begin{bmatrix} 4 & 0 & 1 \\ 0 & 2 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 4 & 0 \\ 0 & 2 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 17 & 1 \\ 1 & 5 \end{bmatrix}$$

$$A^T b = \begin{bmatrix} 4 & 0 & 1 \\ 0 & 2 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 0 \\ 11 \end{bmatrix} = \begin{bmatrix} 19 \\ 11 \end{bmatrix}$$

$$\begin{bmatrix} 17 & 1 \\ 1 & 5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 19 \\ 11 \end{bmatrix}$$

Example

$$\begin{bmatrix} 17 & 1 \\ 1 & 5 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 19 \\ 11 \end{bmatrix}$$

$$\det(A^T A) = (17 \times 5) - 1 = 84$$

Let's solve the normal equations for $Ax = b$:

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 17 & 1 \\ 1 & 5 \end{bmatrix}^{-1} \begin{bmatrix} 19 \\ 11 \end{bmatrix} = \frac{1}{\det(A^T A)} \begin{bmatrix} 5 & -1 \\ -1 & 17 \end{bmatrix} \begin{bmatrix} 19 \\ 11 \end{bmatrix}$$

$\tilde{A}^T / \tilde{A} \tilde{x}$

$$= \frac{1}{84} \begin{bmatrix} 95 - 11 \\ -19 + 187 \end{bmatrix} = \frac{1}{84} \begin{bmatrix} 84 \\ 168 \end{bmatrix}$$

$$\hat{b} = \tilde{A} \tilde{x} = \begin{bmatrix} 4 & 0 \\ 0 & 2 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 4 \\ 4 \\ 3 \end{bmatrix}$$

Orthogonal proj. onto Col A

Example

$$\begin{bmatrix} 1 & 2 \\ -1 & 3 \\ 0 & 0 \end{bmatrix} \vec{A} \quad \vec{x} = \begin{bmatrix} 4 \\ 1 \\ 4 \end{bmatrix} \vec{b}$$

$$\hat{\vec{b}} = \begin{bmatrix} 4 \\ 1 \\ 0 \end{bmatrix} \stackrel{?}{=} \vec{A}\vec{x} = \begin{bmatrix} 2+2 \\ -2+3 \\ 0 \end{bmatrix}$$

Let's do it again...

$$\vec{A}^T \vec{A} = \begin{bmatrix} 2 & -1 \\ -1 & 13 \end{bmatrix}$$

$$\vec{A}^T \hat{\vec{b}} = \begin{bmatrix} 1 & -1 & 0 \\ 2 & 3 & 0 \end{bmatrix} \begin{bmatrix} 4 \\ 1 \\ 4 \end{bmatrix} = \begin{bmatrix} 3 \\ 11 \end{bmatrix}$$

$\hat{\vec{b}} \stackrel{?}{=} \vec{b}$ NO

$$\vec{x} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

least-squares sol'n

$$\begin{bmatrix} 1 & 6 & ; & 2 \\ 6 & 1 & ; & 1 \end{bmatrix} \xrightarrow{\sim}$$

$$\begin{bmatrix} 1 & -13 & ; & -11 \\ 0 & 1 & ; & 1 \end{bmatrix}$$

Unique Least Squares Solutions

Question (Conceptual)

Is a least squares solution unique?

Answer: No

Remember that if $\mathbf{b} \in Col(A)$ then $\hat{\mathbf{b}} = \mathbf{b}$ and then we're asking if $A\mathbf{x} = \mathbf{b}$ has a unique solution for any choice of A


may have many solutions

When is there a unique solution?

The least squares method gives us to find an approximate solution when there is no exact solution

But it doesn't help us choose a solution in the case that there are many

Practically Speaking

numpy.linalg.lstsq

`linalg.lstsq(a, b, rcond='warn')`

[\[source\]](#)

Return the least-squares solution to a linear matrix equation.

Computes the vector x that approximately solves the equation $a @ x = b$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the “exact” solution of the equation. Else, x minimizes the Euclidean 2-norm $\|b - ax\|$. If there are multiple minimizing solutions, the one with the smallest 2-norm $\|x\|$ is returned.

Parameters: `a` : *(M, N) array_like*

“Coefficient” matrix.

`b` : *{(M,), (M, K)} array_like*

Ordinate or “dependent variable” values. If b is two-dimensional, the least-squares solution is calculated for each of the K columns of b .

`rcond` : *float. optional*

Practically Speaking

numpy.linalg.lstsq

`linalg.lstsq(a, b, rcond='warn')`

[\[source\]](#)

Return the least-squares solution to a linear matrix equation.

Computes the vector x that approximately solves the equation $a @ x = b$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the “exact” solution of the equation. Else, x minimizes the Euclidean 2-norm $\|b - ax\|$. If there are multiple minimizing solutions, the one with the smallest 2-norm $\|x\|$ is returned.

NumPy chooses the shortest vector

Parameters: `a` : (M, N) *array_like*

“Coefficient” matrix.

`b` : $\{(M,), (M, K)\}$ *array_like*

Ordinate or “dependent variable” values. If b is two-dimensional, the least-squares solution is calculated for each of the K columns of b .

`rcond` : *float. optional*

Practically Speaking

numpy.linalg.lstsq

`linalg.lstsq(a, b, rcond='warn')`

[\[source\]](#)

Return the least-squares solution to a linear matrix equation.

Computes the vector x that approximately solves the equation $a @ x = b$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the “exact” solution of the equation. Else, x minimizes the Euclidean 2-norm $\|b - ax\|$. If there are multiple minimizing solutions, the one with the smallest 2-norm $\|x\|$ is returned.

NumPy chooses the shortest vector

Parameters: `a` : (M, N) *array_like*

“Coefficient” matrix.

(why?...)

`b` : $\{(M,), (M, K)\}$ *array_like*

Ordinate or “dependent variable” values. If b is two-dimensional, the least-squares solution is calculated for each of the K columns of b .

`rcond` : *float. optional*

Unique Least Squares Solutions

Theorem. For a $m \times n$ matrix A the following are equivalent:

- » $Ax = b$ has a unique least squares solution for any choice of b
- » The columns of A are linearly independent
- » $A^T A$ is invertible

Unique Least Squares Solutions

$$\hat{\underline{x}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$$

If \mathbf{A} has linearly independent columns, then its unique least squares solution is defined as above:

$$\begin{aligned} \mathbf{A}^T \mathbf{A} \hat{\underline{x}} &= \mathbf{A}^T \tilde{\mathbf{b}} \\ \hat{\underline{x}} &= (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \tilde{\mathbf{b}} \end{aligned}$$

Projecting onto a subspace

$$\hat{\mathbf{b}} = \underline{A\hat{\mathbf{x}}} = \underline{A(A^T A)^{-1} A^T \mathbf{b}}$$

If the columns of A are linearly independent,
then **they form a basis**

Said another way: if \mathcal{B} is a basis, then we can
construct a matrix A whose columns are the
vectors in \mathcal{B}

This means we can find arbitrary projections